

Affective Motivational Collaboration Theory

by

Mahni Shayganfar - mshayganfar@wpi.edu

A PhD Dissertation

Presented at

WORCESTER POLYTECHNIC INSTITUTE

in partial fulfillment of the requirements for the

DOCTOR OF PHILOSOPHY

in

Computer Science

November 2016

APPROVED

Professor Charles Rich, Thesis Advisor

Professor Candace L. Sidner, Thesis Co-Advisor

Professor John E. Laird, Thesis Committee Member

Professor Stacy Marsella, Thesis Committee Member

© Copyright by Mahni Shayganfar 2016
All Rights Reserved

ABSTRACT

Abstract Here!

ACKNOWLEDGMENTS

Acknowledgments Here!

Contents

Abstract	i
Acknowledgments	ii
1. Introduction	1
1.1 Motivation	1
1.2 Thesis Statement and Scope	3
1.3 Contributions	4
2. Background and Related Work	7
2.1 Computational Collaboration Theories	7
2.1.1 Shared-Plans Theory	9
2.1.2 Joint Intentions Theory	16
2.1.3 STEAM – A Hybrid Collaboration Approach	20
2.1.4 Other Approaches	23
2.1.5 Similarities and Differences	23
2.1.6 Applications of Collaboration Theories	26
2.2 Affective Computing	31
2.2.1 Affect and Emotions	31
2.2.2 Functions of Emotions	31
2.2.3 Motivation and Theory of Mind	31
2.3 Computational Models of Emotions	31
2.3.1 Appraisal Theory	31

2.3.2	Other Computational Models	31
2.3.3	Similarities and Differences	31
2.3.4	Applications in Autonomous Agents and Robots	31
2.4	Conclusion	31
3.	Affective Motivational Collaboration Theory	33
3.1	Introduction	33
3.1.1	Scenario	35
3.1.2	Example of a Collaborative Interaction	35
3.2	Design and Architecture	35
3.2.1	Mechanisms	35
3.2.2	Functions of Emotions	35
3.2.3	Mental States	35
3.2.4	Attributes of Mental States	35
4.	Computational Framework	36
4.1	Introduction	36
4.2	Collaboration Mechanism	37
4.3	Appraisal Mechanism and Underlying Processes	39
4.3.1	Relevance	40
4.3.2	Desirability	44
4.3.3	Expectedness	46
4.3.4	Controllability	47
4.4	Goal Management	50
4.5	Coping Mechanism and Strategies	56
4.5.1	Planning	56
4.5.2	Active Coping	57
4.5.3	Seeking Social Support for Instrumental Reasons	58
4.5.4	Acceptance	58
4.5.5	Mental Disengagement	58

4.5.6	Shifting Responsibility	59
4.5.7	Activation of Coping Strategies	59
4.6	Motivation Mechanism	60
4.6.1	Satisfaction Motive	62
4.6.2	Achievement Motive	64
4.6.3	External Motive	65
4.7	Theory of Mind	67
4.8	Perception and Action	67
4.9	Emotion Instances	68
5.	Evaluation	71
5.1	Evaluating Appraisal Algorithms (Crowd Sourcing)	71
5.1.1	Experimental Scenario	71
5.1.2	Hypothesis and Methodology	72
5.1.3	Results	74
5.1.4	Discussion	80
5.2	End-to-End System Evaluation	80
5.2.1	Implementation	81
5.2.2	Experimental Scenario	83
5.2.3	Hypotheses and Methodology	87
5.2.4	Results	90
5.2.5	Discussion	105
6.	Conclusion	107
6.1	Discussion	107
6.2	Future Work	107
Appendix A	118

List of Figures

1.1	A robotic arm collaborating with a human to achieve a shared goal using <i>Affective Motivational Collaboration Framework</i>	5
2.1	Plans for collaborative action [28].	13
4.1	Collaboration structure (shared plan).	37
4.2	Using Collaboration structure in Appraisal (mechanisms in our framework).	40
4.3	Using Appraisals' outcome to influence Collaboration structure (mechanisms in our framework).	52
4.4	Cost values indicated by tuples with (second number) and without (first number) the influence of emotions.	53
4.5	Conditions for selecting coping strategies	61
4.6	Three functions of satisfaction motive (blue: valence = 0, red: valence = positive, green: valence = negative). The x-axis indicates the satisfaction drive's delta value in [-1, +1], and the y-axis indicates the magnitude of satisfaction motive in [-1, +1].	63
4.7	Two functions of the achievement motive (blue: valence = +1, red: valence = 0, green: valence = -1, orange: valence = close to zero from negative side). The x-axis indicates the success probability value of achieving a goal which is in [0, +1], and the y-axis indicates the magnitude of achievement motive in [-1, +1].	64

4.8	Two functions of external motive (blue: valence = -1, red: valence = 0, green: valence = +1, orange: valence = close to zero from negative side). The x-axis indicates the success probability value of achieving a proposed goal which is in [0, +1], and the y-axis indicates the magnitude of the achievement motive in [-1, +1].	66
4.9	Appraisal values.	69
4.10	Conditions for selecting emotion instances	70
5.1	Collaboration task model for the evaluation.	72
5.2	Expectedness results.	75
5.3	Example expectedness question.	75
5.4	Controllability results.	76
5.5	Example controllability question.	77
5.6	Desirability results.	77
5.7	Example desirability question.	78
5.8	Relevance results.	79
5.9	Example relevance question.	80
5.10	Computational framework based on Affective Motivational Collaboration theory (arrows indicate primary influences between mechanisms and data flow).	82
5.11	Collaboration structure used as the task model.	82
5.12	The layout of the available spots for the human and the robot to place their pegs during the collaboration.	83
5.13	The Graphical User Interface (GUI) used during interaction.	84
5.14	Experimental setup.	88
5.15	The 31 Likert scale questions organized according to their groups.	91
5.16	Results of the Likert scale survey for Likability questions. The p-value for the difference between means is $\ll 0.001$ for all questions.	92

5.17 Results of the Likert scale survey for questions related to trust. The p-value for the difference between means is $\ll 0.001$ for all questions. . .	93
5.18 Results of the Likert scale survey for questions related to the robot's performance. The p-value for the difference between the means for questions 10, 11 and 12 are 0.001, 0.063 and $\ll 0.001$, respectively. . .	94
5.19 Results of the Likert scale survey for the questions related to the robot's understanding of human emotions. The p-value for the difference between the means is $\ll 0.001$ for all of the questions except Question 14, for which the p-value is 0.003.	95
5.20 Results of the Likert scale survey for questions related to the robot's understanding of goals. The p-value for the difference between the means for all questions is $\ll 0.001$, except Question 19, for which the p-value is 0.006.	96
5.21 Results of the Likert scale survey for questions related to the human's feeling about the collaboration. The p-value for the difference between the means is $\ll 0.001$ for questions 22, 25, 26, and 28. The p-value for Questions 23, 24 and 27 are 0.02, 0.008 and 0.001, respectively.	97
5.22 Results of the Likert scale survey for questions related to satisfaction with collaborative partner. The p-value for the difference between means is $\ll 0.001$ for all questions.	98
5.23 Open-ended questionnaire questions and results. (*Note: Because we are evaluating whether humans prefer an emotion-aware robot, these results are taken as negative test results when calculating the p-value using the binomial distribution. Only those participants who clearly indicated a preference for the emotion-aware robot are taken as positive test results.)	99
5.24 Impact of age on results of Likert scale questions related to likability.	102
5.25 Impact of age on results of Likert scale questions related to trust. . .	102

5.26 Impact of age on results of Likert scale questions related to performance.	103
5.27 Impact of age on results of Likert scale questions related to robot's understanding of human's emotions.	103
5.28 Impact of age on results of Likert scale questions related to robot's understanding of goals.	104
5.29 Impact of age on results of Likert scale questions related to human's feeling about collaboration.	104
5.30 Impact of age on results of Likert scale questions related to satisfac- tion with collaborative partner.	105

List of Tables

5.1 Number of participants	74
--------------------------------------	----

CHAPTER 1

INTRODUCTION

1.1 Motivation

The idea of robots or other intelligent agents living in a human environment has been a persistent dream from science fiction books to artificial intelligence and robotic laboratories. Collaborative robots are expected to become an integral part of humans' environment to accomplish their industrial and household tasks. In these environments, humans will be involved in robots' operations and decision-making processes. The involvement of humans influences the efficiency of robots' interaction and performance, and makes the robots sensitive to humans' cognitive abilities and behaviors.

A key aspect of the sociability of robots is their ability to collaborate with humans in the same environment. Collaboration is a coordinated activity in which the participants work jointly to satisfy a shared goal [34]. There are many challenges in achieving a successful collaboration between robots and humans. To meet these challenges, it is crucial to understand what makes a collaboration not only successful, but also efficient. Existing computational models of collaboration explain some of the important concepts underlying collaboration; such as the presence of a reason for collaborators' commitment, and the necessity of communicating about mental states in order to maintain progress over the course of a collaboration. The most prominent collaboration theories are based on plans and intentions [16] [34] [51], and are derived from Bratman's BDI architecture [4]. Two theories, Joint Intentions [16]

and SharedPlans [31, 32, 34], have been used to support teamwork and collaboration between humans and robots or virtual agents [8] [60] [80] [89]. However, these theories explain only the structure of a collaboration. For instance, in SharedPlans theory collaborators build a shared plan containing a collection of beliefs and intentions about the actions in the plan. Collaborators communicate these beliefs and intentions via utterances about actions that contribute to the shared plan. This communication leads to the incremental construction of a shared plan, and ultimately successful completion of the collaboration. In contrast, in Joint Intentions theory, the notion of joint intention is viewed as a persistent commitment of the team members to a shared goal. In this theory, once an agent enters into a joint commitment with other agents, it should communicate its private beliefs to other team members.

Although existing collaboration theories explain the important elements of a collaboration structure, the underlying processes required to dynamically create, use, and maintain the elements of this structure are largely unexplained. For instance, a general mechanism has yet to be developed that allows an agent to effectively integrate the influence of its collaborator's perceived or anticipated emotions into its own cognitive mechanisms to prevent shared task failures while maintaining collaborative behavior. Therefore, a process view of collaboration must include certain key elements. It should inherently involve social interactions since all collaborations occur between social agents, and it should essentially constitute a means of modifying the content of social interaction as the collaboration unfolds. The underlying processes of emotions possess these two properties, and social functions of emotions explain some aspects of the underlying processes in collaboration. This thesis makes the case for emotion-driven processes within collaboration and demonstrates how it furthers collaboration between humans and robots.

1.2 Thesis Statement and Scope

In this thesis, we develop and validate a framework based on *Affective Motivational Collaboration Theory* which can improve the effectiveness of collaboration between agents/robots and humans. This thesis is established based on the reciprocal influence of collaboration structure and the appraisal processes in a dyadic collaboration. We focus only on two-participant collaboration; teamwork collaboration is out of our scope. Furthermore, this work focuses on a) the influence of emotion-regulated processes on the collaboration structure, and b) prediction of the observable behaviors of the other during a collaborative interaction.

We describe the cognitive processes involved in a collaboration in the context of a cognitive architecture. There are several well-developed cognitive architectures, e.g., Soar [48] and ACT-R [44], each with different approaches to defining the basic cognitive and perceptual operations. There have also been efforts to integrate affect into these architectures [20, 55]. In general, however, these cognitive architectures do not focus on processes to specifically produce emotion-regulated goal-driven collaborative behaviors. At the same time, existing collaboration theories, e.g., Shared-Plans [34] theory, focus on describing the structure of a collaboration in terms of fundamental mental states, e.g., mutual beliefs or joint intentions. However, they do not describe the associated processes, their relationships, and influences on each other. *Affective Motivational Collaboration Theory* deals with some of the major affect-driven processes having an impact on the collaboration structure. This theory is informed by research in psychology and artificial intelligence which is reviewed in Chapter 2. Our contribution, generally speaking, is to synthesize prior work on appraisal and collaboration, and motivation to provide a new theory which describes some of the prominent emotion-regulated goal-driven phenomena in a dyadic collaboration.

1.3 Contributions

Throughout this work we aim to show how a robot can leverage emotion-driven processes using appraisal algorithms to improve collaboration with humans. As such, in this thesis work, we introduce a novel framework, called Affective Motivational Collaboration (AMC) framework, which allows a robotic agent to collaborate with a human while incorporating the underlying emotion-driven processes and the expressed emotion of the human collaborator. Such a framework is built based on computational models of collaboration and appraisal allowing for task-driven interaction with robots or other agents. The theoretical foundation, computational models and algorithms as well as the overall framework, and the end-to-end evaluation of the framework make the following contributions:

1. Introducing *Affective Motivational Collaboration Theory*:

(Chapter 3) As mentioned earlier, since the theoretical foundation of AMC framework is built on the combination of SharedPlans theory of collaboration [34] and cognitive appraisal theory of emotions [57] [76], one of the contributions of our work is to introduce theoretical concepts incorporating key notions of both theories in a dyadic collaboration context. Applying cognitive appraisal theory in the collaboration context is novel. Other models of the appraisal theory have not paid attention to the dynamics of the collaboration.

2. Developing new computational models and algorithms for *Affective Motivational Collaboration Framework*:

(Chapter 4) Another contribution of our work is to create computational models and algorithms to compute the value of appraisal variables in a dyadic collaboration. We use the collaboration structure to compute appraisal variables. Reciprocally, we use the evaluative nature of the appraisal to make changes to the collaboration structure as required. We have also developed

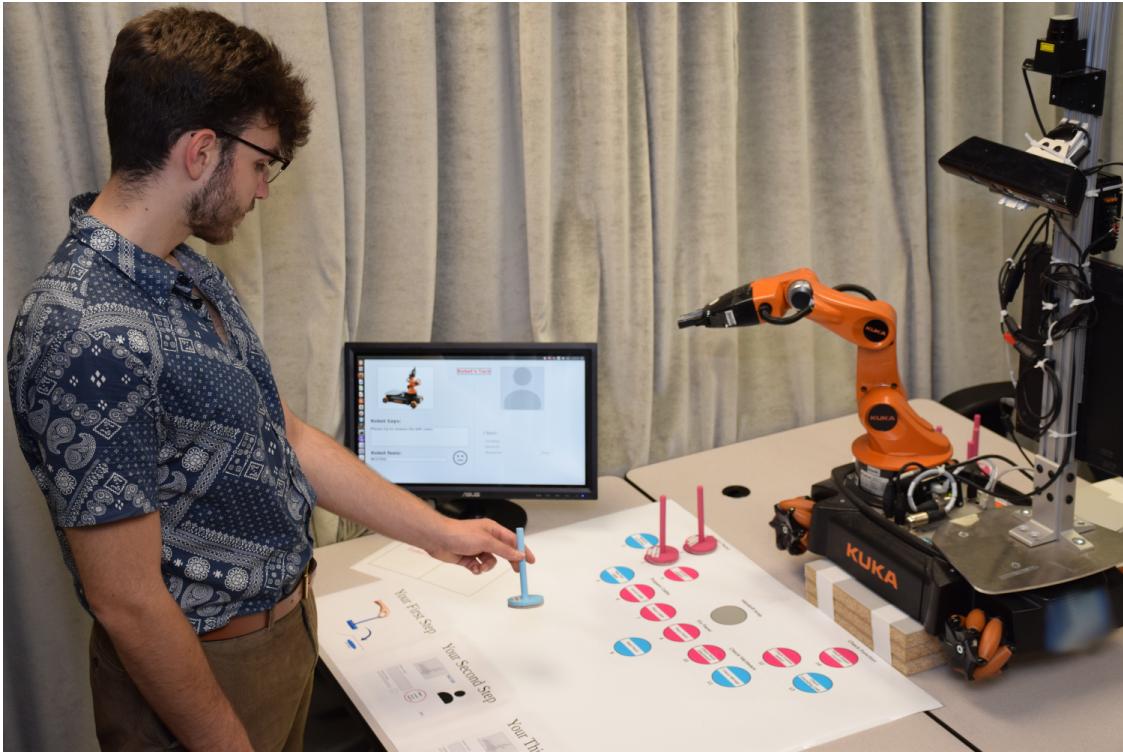


Figure 1.1: A robotic arm collaborating with a human to achieve a shared goal using *Affective Motivational Collaboration Framework*.

a new algorithm for emotion-driven goal management in the context of collaboration. Goal management is one of the important functions of emotions during collaboration. Existing models and implementations of emotions focus only on how emotions regulate and control internal processes and sometimes behaviors. This part of our work shows how appraisal components of the self and the human collaborator contributes to goal management as an emotion function.

3. Developing a computational framework based on *Affective Motivational Collaboration Theory*:

(Chapter ??) In order to evaluate our computational models and algorithms within an interaction with human collaborators, we have developed a computational framework based on our theoretical foundations in *Affective Motiva-*

tional Collaboration Theory. Our computational framework implements the key concepts related to *Affective Motivational Collaboration Theory* as well as minimal implementation of other processes which are required for validation of the model but are not part of this thesis' contributions. The emphasis of the model is on the underlying cognitive processes of collaboration and appraisal concepts, rather than the Perception and the Action mechanisms.

4. Validating *Affective Motivational Collaboration Theory*:

(Chapters 4 and 5) We have conducted two user studies a) to validate our appraisal algorithms before further development of our framework, and b) to investigate the overall functionality of our framework within an end-to-end system evaluation with participants and a robot. The second user study was also conducted to evaluate the benefit of using our computational framework in human-robot collaboration. In the first user study, we crowd sourced our questionnaires to test our hypothesis that humans and our algorithms will provide similar answers to questions related to different factors within our appraisal algorithms. In the second user study, we investigated the importance of emotional awareness in human-robot collaboration, and the overall functionality of the AMC framework with the participants in our study environment.

CHAPTER 2

BACKGROUND AND RELATED WORK

2.1 Computational Collaboration Theories

The construction of computer systems and robots that are intelligent, collaborative problem-solving partners is important in Artificial Intelligence (AI) and its applications. It has always been important for us to make computer systems better at helping us to do whatever they are designed for. To build collaborative systems, we need to identify the capabilities that must be added to individual agents so that they can work with us or other agents. As Grosz says, collaboration must be designed into systems from the start; it cannot be patched on [28].

Collaboration is a special type of coordinated activity in which the participants work jointly, together performing a task or carrying out the activities needed to satisfy a shared goal [32]. Collaboration involves several key properties both in structural and functional levels. For instance, most collaborative situations involve participants who have different beliefs and capabilities; most of the time collaborators only have partial knowledge of the process of accomplishing the collaborative activities; collaborative plans are more than the sum of individual plans; collaborators are required to maintain mutual beliefs about their shared goal throughout the collaboration; they need to be able to communicate with others effectively; they need to commit to the group activities and to their role in it; collaborators need to commit to the success of others; they need to reconcile between commitments to the existing collaboration and their other activities; and they need to interpret

others' actions and utterances in the collaboration context [29]. These collaboration properties are captured by the existing computational collaboration theories.

As we mentioned, to be collaborative, partners, e.g., a robot and a human, need to meet the specifications stipulated by collaboration theories. These theories argue for an essential distinction between a collaboration and a simple interaction or even a coordination in terms of commitments [27, 52]. This document briefly provides descriptions of major computational collaboration theories, their similarities and differences, and their application in AI and robotics. It primarily focuses on Joint Intention, SharedPlans and hybrid theories of collaboration. In this document, we do not present the theories in formal language, but simply describe their features in general terms.

The prominent collaboration theories are mostly based on plans and joint intentions [16] [34] [51], and they were derived from the BDI paradigm developed by Bratman [4] which is fundamentally reliant on folk psychology [69]. The two theories, Joint Intentions [16] and SharedPlans [34], have been extensively used to examine and describe teamwork and collaboration.

The SharedPlans theory is based on the theories of Bratman and Pollack [7, 65, 66], who outline a mental-state view of plans in which having a plan is not just knowing how to do an action, but also having the intention to do the actions entailed. Bratman's views of intention goes back to the philosophical views of Anscombe [2] and *Castañeda* [11] about intention. Also, as Grosz and Sidner mention in [34] the natural segmentation of discourse reflects intentional behaviors in each segment. These intentions are designated as Discourse Segment Purposes (DSPs) which are the basic reasons for engaging in different segments of discourse. DSPs are a natural extension of Gricean intentions at the utterance level [62].

Cohen and Levesque also mention that in Joint Intentions theory their view of intention is primarily future-directed [17] which makes their view similar to Bratman's theory of intention [5], contra Searle [78].

Commitment – One of the most important concepts of teamwork and collaboration is the concept of commitment. Collaboration theories are required to meet the notion of commitment, otherwise the participants are just doing some coordinated works. Since the prominent computational collaboration theories, reviewed in this paper, are based on Bratman’s view of intention, we briefly provide his view of commitment here before describing these theories. Bratman defines certain prerequisites for an activity to be considered shared and cooperative [6]. He stresses the importance of:

- a) **Mutual commitment to joint activity** – which can be achieved by agreement on the joint activity, and prevention of abandoning the activity without involving teammates;
- b) **Mutual support** – which can be achieved by team members if they actively try to help teammate activity;
- c) **Mutual responsiveness** – which means team members should take over tasks from teammates if necessary.

In the following sections, we are also going to see how each collaboration theory addresses the notion of commitment.

2.1.1 Shared-Plans Theory

The SharedPlans model of collaborative action, presented by Grosz and Sidner [31, 32, 34], aims to provide the theoretical foundations needed for building collaborative robots/agents [28]. SharedPlans is a general theory of collaborative planning that requires no notion of joint intentions (see Section 2.1.2), accommodates multi-level action decomposition hierarchies and allows the process of expanding and elaborating partial plans into full plans (see Section 2.1.1). SharedPlans theory explains how a group of agents can incrementally form and execute a shared plan that then guides and coordinates their activity towards the accomplishment of a

shared goal. SharedPlans is rooted in the observation that collaborative plans are not simply a collection of individual plans, but rather a tight interleaving of mutual beliefs and intentions of different team members. In [32] Grosz and Kraus use first-order logic to present the formalization of SharedPlans.

Grosz and Sidner in [34] present a model of plans to account for how agents with partial knowledge collaborate in the construction of a domain plan. They are interested in the type of plans that underlie discourse in which the agents are collaborating in order to achieve a shared goal. They propose that agents are building a shared plan (see Section ??) in which participants have a collection of beliefs and intentions about the actions in the plan. Agents have a library of how to do their actions, i.e. recipes (see Section 2.1.1). These recipes might be partially specified as to how an action is executed, or contributes to a goal (see Section 2.1.1). Then, each agent communicates their beliefs and intentions by making utterances about what actions they can contribute to the shared plan. This communication leads to the construction of a shared plan, and ultimately termination of the collaboration with each agent mutually believing that there exists one agent who is going to execute an action in the plan, and the fact that that agent has intention to perform the action, and that each action in the plan contributes to the goal [34] [53].

Later in Section ??, we are going to see that to successfully complete a plan the collaborators must mutually believe that they have a common goal and have agreed on a sequence of actions for achieving that goal. They should believe that they are both capable of performing their own actions and intend to perform those actions while they are committed to the success of their plans.

Recipes

The SharedPlans theory differentiates between knowing how to accomplish a goal (a recipe) and having a plan, which includes intentions. The SharedPlans definition of mutual beliefs states that when agents have a shared plan for doing some action, they must hold mutual beliefs about the way in which they should perform that action

[32, 34]. Following Pollack [66], the term recipe refers to what collaborators know when they know a way of doing an action. Recipes are specified at a particular level of detail. Although the agents need to have mutual beliefs about actions specified in the recipe, they do not need to have mutual beliefs about all levels of performing actions. Therefore, having mutual beliefs of the recipe means that the collaborators hold the same beliefs about the way in which an action should be accomplished. Consequently, the collaborators need to agree on how to execute an action. Recipes are aggregations of action-types and relations among them. Action-types, rather than actions, are the main elements in recipes. Grosz and Sidner in their earlier work [34] have considered only simple recipes in which each recipe consisted of only a single action-type relation [53]. Recipes can be partial, meaning they can expand and be modified over time.

Grosz and Sidner propose that collaboration must have the following three elements, which also indicates the importance of the shared plan:

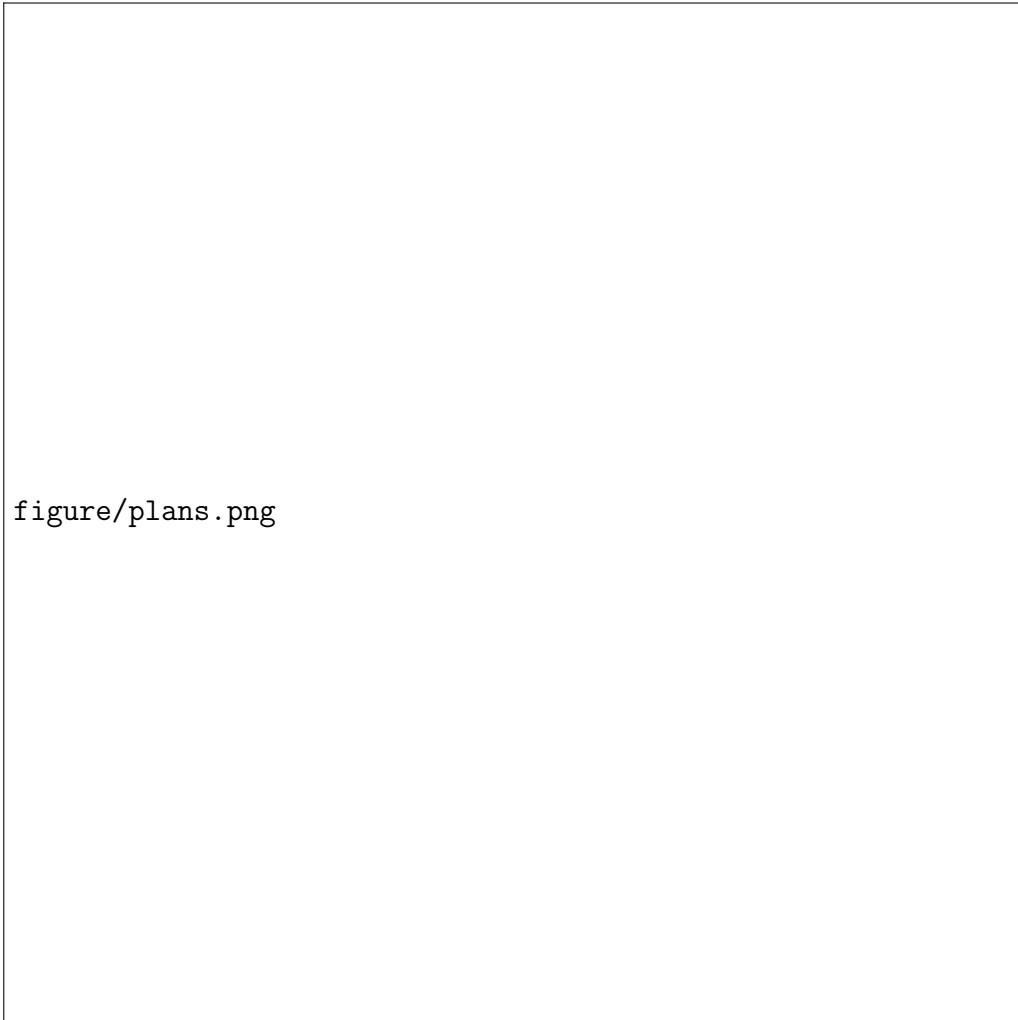
1. the participants must have commitment to the shared activity;
2. there must be a process for reaching an agreement on a recipe for the group action;
3. there must be commitment to the constituent actions.

Shared plan is an essential concept in the collaboration context. The definition of the shared plan is derived from the definition of plans Pollack introduced in [65, 66] since it rests on a detailed treatment of the relations among actions and it distinguishes the intentions and beliefs of an agent about those actions. However, since Pollack's plan model is just a simple plan of a single agent, Grosz and Sidner extended that to plans of two or more collaborative agents. The concept of the shared plan provides a framework in which to further evaluate and explore the roles that particular beliefs and intentions play in collaborative activity [53]. However, this formulation of shared plans (a) could only deal with activities that directly decomposed

into single-agent actions, (b) did not address the requirement for the commitment of the agents to their joint activities, and (c) did not adequately deal with agents having partial recipes [32]. Grosz and Kraus in [32], reformulate Pollack’s definition of the individual plans [66], and also revise and expand the SharedPlans to address these shortcomings.

Figure 2.1 shows what we need to add to individual plans in order to have plans for group actions. The top of the figure lists the main components for individual plans. First, an individual agent needs to know the recipe for an action, whereas agents in a group need to have a mutual belief of a recipe for an action (bottom of the figure). In the case of a group plan, having a mutual belief of a recipe, leads the agents to agree on how they are going to execute the action. Then, similar to individual agents that need to have the ability to perform the constituent actions in an individual plan and must have intentions to perform them, the participants in a group activity need to have individual or group plans for each of the constituent actions in the mutually agreed recipe [28, 34].

As shown in Figure 2.1 (bottom), plans for group actions include two essential constituents that do not have correlates in the individual plan. First, the agents need to have a commitment to the group activity; All the agents need to intend that (see Section 2.1.1) the group will do the action. For instance, a robot and an astronaut need to have intentions that they install solar panels together. Among other things, these intentions will keep them both working on the panels until the panels are installed. Second, the participants need to have some commitment to the other agents to succeed in their own their actions. For instance, the robot must have an intention that the astronaut be able to measure the quality of installation successfully. This intention will prevent the robot from interrupting the astronaut’s measurement action or prevent the robot from using the astronaut’s measurement tool [28, 34].



figure/plans.png

Figure 2.1: Plans for collaborative action [28].

Full Vs. Partial Shared Plan

The SharedPlans formalization distinguishes complete plans and partial plans. A shared plan can be either a *Full Shared Plan (FSP)* or a *Partial Shared Plan (PSP)*. An *FSP* is a complete plan in which agents have fully determined how they will perform an action. A *PSP* definition provides a specification of the minimal mental state requirements for collaboration to exist and gives criteria governing the process of completing the plan.

An *FSP* to do α represents a situation where every aspect of a joint activity α is

fully determined. This includes mutual belief and agreement in the complete recipe to do α . A recipe is a specification of a set of actions A_i , which constitutes the performance of α when executed under specified constraints. $FSP(\mathbf{P}, \Theta, \alpha, T_p, T_\alpha, \mathbf{R}_\alpha)$ denotes a group Θ 's plan \mathbf{P} at time T_p to do action α at time T_α using recipe \mathbf{R}_α . In short, FSP holds if and only if the following conditions are satisfied:

1. All members of group Θ mutually believe that they intend to do α .
2. All members of group Θ mutually believe that \mathbf{R}_α is the recipe for α .
3. For each step A_i in recipe \mathbf{R}_α :
 - A subgroup Θ_j has an FSP for A_i , using recipe \mathbf{R}_{A_i} .
 - Other members of group Θ believe that there exists a recipe such that subgroup Θ_j can bring about A_i and have an FSP for A_i .
 - Other members of group Θ intend that subgroup Θ_j can bring about A_i using some recipe.

Most of the times a team and its members do not possess an FSP to achieve their shared goal. In this case, the concept of FSP puts limits on the SharedPlans theory. However, SharedPlans uses the concept of PSP as a snapshot of the team's mental states in different situations, which further leads to communication and planning to fulfill the conditions of an FSP . The idea behind PSP is enabling the agents to modify the shared plan over the course of planning without impairing the achievement of the shared goals. Notice that for the same reason recipes also can be partial [32, 34].

Communicating Intentions

In SharedPlans theory Grosz and Sidner are interested in the type of plans that underlie a discourse in which the agents collaborate to achieve a shared goal. Here

we present their view of discourse structure, since it is directly related to the intentions behind collaborators' actions. In [34], Grosz and Sidner argue that the SharedPlans theory recognises three interrelated levels of discourse structure, and the components of the discourse structure are a trichotomy of linguistic structure, intentions structure and the attention state. In their work, the linguistic structure of a discourse is a sequence of utterances aggregating into discourse segments just as the words in a single sentence form constituent phrases. They also discuss the idea of the discourse purpose as the intention that underlies engagement in the particular discourse. They believe this intention is the reason behind performing a discourse rather than some other actions, and also the reason behind conveying a particular content of the discourse rather than some other contents. They describe mechanisms for plan analysis looking at Discourse Segment Purposes (DSPs). In fact, the DSPs specify how the discourse segments contribute to achieving the overall discourse purpose. Finally, the third component in their theory, the attentional state, provides an abstraction of the agent's focus of attention as the discourse unfolds. The focusing structure contains DSPs and the stacking of focus spaces reflects the relative salience of the entities in each space during the discourse. In short, the focusing structure is the central repository for the contextual content required for processing utterances during the discourse [34]. Using discourse plans can help to encode the knowledge about conversation.

Intention-to and Intention-that

In Grosz and Sidner's SharedPlans theory [34], two intentional attitudes are employed: *intending to* (do an action) and *intending that* (a proposition will hold). The notion of *intention to*, as an individual-oriented intention, models the intention of an agent to do any single-agent action while the agent not only believes that it is able to execute that action, but it also commits to doing so. In short, it is an intention to perform an action, similar to Bratman's view of intention. In contrast with *intention to*, an *intention that*, as an intention directed toward group activity, does

not directly imply an action. In fact, an individual agent’s *intention that* is directed towards its collaborators’ action or towards a group’s joint action. *Intention that* guides an agent to take actions (including communication), that enable or facilitate other collaborators to perform assigned tasks. This leads an agent to behave collaboratively. Therefore, agents will adopt intentions to communicate about the plan [32]. As another difference, *Intention to* commits an agent to means-end reasoning and acting [4] while *Intention that* does not necessarily entail this commitment. The key point about *Intention to* and *intention that* is that both commit an agent not to adopt conflicting intentions, and constrain replanning in case of failure. Further, an agent can *intention that* another agent achieve the specified proposition.

2.1.2 Joint Intentions Theory

Following Bratman’s guidelines, Cohen and Levesque propose a formal approach to building artificial collaborative agents. The Joint Intentions theory of Cohen and Levesque [16, 17, 18, 19, 50] represents one of the first attempts to establish a formal theory of collaboration, and due to its clarity and expression, is a widely used teamwork theory.

The basic idea of Joint Intentions theory is based on individual and joint intentions (as well as commitments) to act as a team member. Their notion of joint intention is viewed not only as a persistent commitment of the team to a shared goal, but also implies a commitment on part of all its members to a mutual belief about the state of the goal. In other words, Joint Intentions theory describes how a team of agents can jointly act together by sharing mental states about their actions while an intention is viewed as a commitment to perform an action. A joint intention is a shared commitment to perform an action while in a group mental state [17].

In [16] Cohen and Levesque establish that joint intention cannot be defined simply as individual intention with the team regarded as an individual. The reason is that after the initial formation of an intention, team members may diverge in their beliefs and their attitudes towards the intention. Instead, Cohen and Levesque gen-

eralize their own definition of intention. First, they present a definition of individual persistent goal (see Section 2.1.2) and individual intention (see Section 2.1.2). Then, they define analogues of these concepts by presenting mutual belief in place of individual belief. The definition of joint persistent goal (see Section 2.1.2) requires team members to commit to informing other members, if it comes to believe that the shared goal is in its terminal status. As a result, in Cohen and Levesque’s theory, a team with a joint intention is a group that shares a common objective and a certain shared mental state [39].

In this theory, once an agent entered into a joint commitment with other agents, the agent should communicate its private beliefs with other team members if the agent believes that the joint goal is in its terminal status, i.e., either the joint goal is achieved, or it is unachievable, or irrelevant [88]. Thus, as we mentioned above, team members are committed to inform other team members when they reach the conclusion that a goal is achievable, impossible, or irrelevant. For instance, if a robot and an astronaut are collaborating to install a solar panel, and the robot reaches the conclusion that the welding tool has deficiency, it is essential for the robot to have an intention to communicate with the astronaut and make this knowledge common. Therefore, according to this theory, in a collaboration, agents can count on the commitment of other members, first to the goal and then to the mutual belief of the status of the goal.

Individual Commitment

As we mentioned earlier, intentions and commitments are the basic ideas of Joint Intentions theory. Here, we provide the definition of “individual commitment” (also called *persistent goal*) by Cohen et. al. in [15]. According to their definition an agent has a persistent goal relative to q to achieve p only when:

1. agent believes that p is currently false;
2. agent wants p to be true;

3. it is true (and agent knows it) that (2) will continue to hold until the agent comes to believe either that p is true, or that it will never be true, or that q is false.

Note that the condition q is an “escape” clause, which can be omitted for brevity, or it can be used as a reason for the agent to drop a commitment, even though it could be quite vague.

Individual Intention

As we mentioned above, Joint Intention theory adopts Bratman’s view of future-directed properties of intention. In this theory, an intention is defined to be a commitment to act in a certain mental state. In other words, an agent intends relative to some condition to do an action when it has a persistent goal or commitment (relative to that condition) of having done the action and, moreover, believing throughout that it is doing that action [16].

Intention inherits all the properties of commitment (e.g., consistency with mental states). Typically, an agent uses an intention as a decision within a subgoal-supergoal hierarchy to do a particular action. For instance, initially, the agent commits to p becoming true without having any concern about who or how p is going to be accomplished. Then, the agent commits to x or y as a mean to accomplish p. Lastly, the agent selects one of the actions (e.g., x) and forms an intention to do it. This intention will be given up when for whatever reason p is accomplished.

Joint Commitment

Before talking about joint commitment, we provide the definition of the *Weak Achievement Goal* (WAG) concept in Joint Intentions theory which shows the state of a team member nominally working on a goal. The concept of WAG is used to provide the definition of the Joint Commitment in this theory.

An agent has a WAG relative to q and with respect to a team to bring about p if either of the following conditions holds:

- The agent has a normal achievement goal to bring about p ; that is, the agent does not yet believe that p is true and wants p to be true as a goal.
- The agent believes that p is true, will never be true, or is irrelevant, but has as a goal that the status of p be mutually believed by all the team members.

Joint commitment – A joint intention of a team Θ is based on its joint commitment, which is defined as a *Joint Persistent Goal* (JPG). A JPG to achieve a team action p , denoted $\text{JPG}(\Theta, p)$ requires all team members to mutually believe that p is currently false and want p to eventually be true. A JPG guarantees that team members cannot decommit until p is mutually known to be *achieved*, *unachievable* or *irrelevant*. Basically, $\text{JPG}(\Theta, p)$ requires team members to each hold p as a *Weak Achievement Goal* (WAG), where $\text{WAG}(\mu, p, \Theta)$ in which μ is a team member in Θ , requires μ to achieve p if it is false. However, if μ privately believes that p is either achieved, unachievable or irrelevant, $\text{JPG}(\Theta, p)$ is dissolved, but p is left with a commitment to have this belief become Θ 's mutual belief. Such a commitment is required to establish mutual belief in Θ ; this commitment typically makes an agent communicate with its teammates [16].

An important consequence of achieving joint commitment in a team is that it predicts future communication which is critical within the course of a collaboration. Thus, this communication leads team members to attain mutual beliefs which is a fundamental concept in teamwork activities. Notice that the minimum mutual belief for team members to attain is the achievement or failure of the shared goal which terminates collaboration.

Joint Intention

Joint intention is defined to be a joint commitment to the team members trying to do a joint action. Based on Cohen and Levesque's definition of joint intention, a

team of agents jointly intends (relative to some escape condition) to do an action if and only if the members have a JPG (relative to that condition) of them having the action completed, and having it completed mutually believing throughout that they are doing it (knowingly) [16].

Teamwork & Communication

In summary, according to Joint Intentions theory, the notion of teamwork is characterized by joint commitment, also known as joint persistent goal (see Section 2.1.2). The definition of JPG states that the agents mutually believe they have the appropriate goal, and that they mutually believe a persistent weak achievement goal (which represents the one-way commitment of one agent directed towards another) to achieve it persists until the agents mutually believe that the goal has either been achieved, or become impossible, or irrelevant.

Joint Intentions theory claims that an efficient collaboration requires communication. Sharing information through communication is critical given that collaborators have different capabilities, and each individual often has only partial knowledge relevant to solving the problem, and sometimes diverging beliefs about the state of the collaborative activity. Communication is important in coordinating team members' roles and actions to accomplish their goal. For instance, it can help team members to establish and maintain a set of mutual beliefs regarding the current state of the collaboration, and the respective roles and capabilities of each member.

2.1.3 STEAM – A Hybrid Collaboration Approach

Tambe in [86] argues that teamwork in complex, dynamic, multi-agent domains requires the agents to obtain flexibility and reusability by using integrated capabilities. Tambe created STEAM (simply, a Shell **TEAM**work) based on this idea. STEAM's operationalization in complex, real-world domains is the key in its development to addressing important teamwork issues, some of which are discussed in Section 2.1.6.

STEAM is founded on the Joint Intentions theory and it uses joint intentions as the basic building block of teamwork while it is informed by key concepts from SharedPlans theory.

Building on the well developed theory of joint intentions [16] and sharedPlans [32, 34], the STEAM teamwork model [86] was operationalized as a set of domain-independent rules that describe how teams should work together. According to Tambe’s claim, several advantages accrue due to this use of Joint Intentions theory, such as achieving a principled framework for reasoning about coordination and communication in a team, which the joint intention can provide. Another advantage is the guidance for monitoring and maintenance of a team activity which the joint commitment in joint intention again provides. And lastly, Tambe believes the joint intention in a team can facilitate reasoning about team activity and team members’ contribution to that activity.

However, he also believes that for a high level team goal, one single joint intention is not sufficient to achieve all these advantages. Thus, STEAM borrows some of the concepts of SharedPlans theory. First, STEAM uses the concept of “intention that” (see Section 2.1.1) towards an activity as well as the fact that SharedPlans theory mandates team members’ mutual belief in a common recipe and shared plans for individual steps in the common recipe. Thus, in this case, SharedPlans helps STEAM to achieve coherency within the teamwork. Besides, STEAM uses joint intentions to ensure the teamwork coherency to build the mental attitudes of team members. In other words, as the recipe evolves, STEAM requires all team members to agree on the execution of a step and form joint intentions to execute it while other joint intentions are formed, leading to a hierarchy. A second concept STEAM borrows from SharedPlans is the amount of information that a team member needs to know to perform an action. According to SharedPlans, team members require to know only that a recipe exists to enable them to perform actions (not recipe details – see Section 2.1.1). Similarly in STEAM, team members only track the responsible subteam or individual team member to perform a specific step while this tracking

does not need detailed plan recognition. The third issue is parallel to what is called an unreconciled case in SharedPlans theory, which in STEAM is handled by replanning and communication between team members assigning the unassigned or unachieved task. The last issue is communication between team members which also borrows the concept of “intention that” from SharedPlans theory, to help the generalization of STEAM’s communication capabilities beside what Joint Intentions theory offers.

In summary, STEAM builds on both Joint Intention theory and SharedPlans theory and tries to overcome their shortcomings. Based on joint intentions, STEAM builds up hierarchical structures that parallel the SharedPlans theory. Hence, STEAM formalizes commitments by building and maintaining Joint Intentions, and uses SharedPlans to formulate the team’s attitudes in complex tasks.

In [86] Tambe argues that the novel aspects of STEAM relate to its teamwork capabilities. The key novelty in STEAM has team operators beside individual team member operators. In STEAM when agents select a team operator for execution, they instantiate a team’s joint intentions. Team operators explicitly express a team’s joint activities, unlike the regular individual operators which express an agent’s own activities. Hence, STEAM agents maintain their own private (to apply individual operators) and team states, e.g., mutual belief about the world (to apply team operators).

However, Tambe added more practical concepts into STEAM’s architecture. For instance, STEAM has a team synchronization protocol to establish joint intention (see JPG in Section 2.1.2), or it has constructs for monitoring joint intentions which helps the agent to be able to monitor team performance. STEAM facilitates this monitoring by exploiting its explicit representation of team goals and plans. In particular, STEAM allows an explicit specification of monitoring conditions to determine achievement, unachievability or irrelevancy conditions of team operators. Finally, in STEAM, communication is driven by commitments embodied in the Joint Intentions theory, i.e., team members may communicate to obtain mutual be-

lief while building and disbanding joint intentions. Thus, joint intentions provide STEAM with a principled framework for reasoning about communication. Also, STEAM addresses some practical issues, not addressed in other teamwork theories. One of these issues is STEAM’s detailed attention to communication overheads and risks, which can be significant [85]. Furthermore, operationalization of STEAM is based on enhancements to the Soar architecture [48], plus a set of about 300 domain-independent Soar rules.

2.1.4 Other Approaches

There are other frameworks, approaches, and models focusing on teamwork and collaborative agents. For instance, Jennings provides the Joint Responsibility framework which is specified formally using modal, temporal logic. Joint Responsibility stresses the role of joint intentions (based on Joint Intentions theory) specifying how both individuals and teams should behave whilst engaged in collaborative problem solving [40, 41, 42, 43]. Jennings has developed *Generic Rules and Agent model Testbed Environment* (GRATE) as a prototype system based on the Joint Responsibility framework. In [46] Kinny et. al. elaborate the concept of Planned Team Activity and introduce a language for representing joint plans for teams of agents and describe how agents can organize the formation of a skilled team to achieve a joint goal. They use joint intentions to capture the mental properties which characterize team activity.

2.1.5 Similarities and Differences

There are some similarities between SharedPlans and Joint Intentions theories. Here, we specify some of these similarities:

1. Similar to SharedPlans theory, Joint Intentions theory specifies what it means for agents to execute actions as a team [83].

2. Both theories follow Bratman’s basic ideas about intention’s roles in relational actions which prevent the collaborative agents from adopting conflicting intentions. Besides, these two theories are also agreed and follow Bratman’s BDI model.
3. Just as SharedPlans theory, Joint Intentions theory also states that a joint action could not be seen as a collection of individual actions but that agents working together need to share beliefs.
4. Both theories in their latest articles show that the agents are required to communicate to maintain collaboration. SharedPlans theory requires collaborators to communicate to establish and maintain the shared plan which is crucial especially when collaborators only have partial shared plan. Similarly in Joint Intentions theory, communication is an explicit requirement of collaborative agents until the shared goal is achieved, unachievable or irrelevant.
5. Both Joint Intentions and SharedPlans theories are concerned about commitment to the joint activity. Although, these two theories use different concepts to fulfill the requirements of commitment during collaboration.

There are also differences between SharedPlans and Joint Intention theories; we address some of them here in this section:

1. Although the crucial components of the SharedPlans theory (see Section 2.1.1) lack the notion of a joint intention, which is the most significant notion within the Joint Intentions theory, Grosz and Sidner do not believe that such a phenomenon (joint intention) exists in a collaboration. They believe their notion of “intention that” and mutual beliefs about states of the collaboration can provide similar functionalities as described in Joint Intentions theory (see Section 2.1.2).
2. In SharedPlans theory teammates agree on the shared plan, whereas in Joint Intentions theory teammates agree on intentions.

3. In contrast to Joint Intentions, the SharedPlans theory employs hierarchical structures over intentions, thus it overcomes the shortcoming of a single joint intention for complex team tasks.
4. The SharedPlans theory describes the way to achieve a common goal through the hierarchy of plans, whereas the Joint Intentions theory describes only this common goal [81].
5. Joint Intentions theory assumes that knowledge about the teammates is always available, whereas SharedPlans theory uses the concept of partial plan/recipe to make the process of dynamically achieving information possible throughout the collaboration.
6. Communication requirements are derived from “intention that” in SharedPlans theory, as opposed to being “hard-wired” in Joint Intentions theory.

A critique to Joint Intention theory – Castelfranchi criticizes the necessary and sufficient conditions (see Section 2.1.2) for the joint persistent goal which plays a crucial role in the Joint Intentions theory. According to his example, if a French scientist and an American scientist are both working on an AIDS vaccine and both have the final goal of p “vaccine anti-AIDS be found” relative to the belief q that “if vaccine is found, AIDS is wiped out”, they both share the mental attitudes described in Joint Intentions theory. It means that they mutually believe that p is currently false, and they mutually know they both want p to be true, and it is true that until they come to believe either that p is true, that p will never be true, or that q is false, they will continue to mutually believe that they each have a weak achievement goal (see Section 2.1.2) relative to q and with respect to the team (i.e., the WAG with respect to the team has been defined as “a goal that the status of p be mutually believed by all the team members”). The problem is that we can not claim the French and American professors are working as a team. In fact, given their personal

goals of finding the vaccine, they might come to strongly compete with each other [12].

2.1.6 Applications of Collaboration Theories

There are many research focusing on different aspects of collaboration based on different collaboration theories, i.e., SharedPlans, Joint Intentions, and hybrid theories of collaboration. In this section, we provide some examples of homogeneous and heterogeneous agent/robot and human collaborations.

There are some works focusing on the concepts of robot assistants [14], or teamwork and its challenges in cognitive and behavioral levels [63, 74]. Some researchers have an overall look at a collaboration concept at the architectural level. In [23] authors present a collaborative architecture, COCHI, to support the concept of emotional awareness. In [22] authors present the integration of emotional competence into a cognitive architecture which runs on a robot, MEXI. In [82] authors discuss the challenges of integrating natural language, gesture understanding and spatial reasoning of a collaborative humanoid robot situated in space. The importance of communication during collaboration has also been considered by some researchers from human-computer interaction and human-robot collaboration [13, 58, 72] to theories describing collaborative negotiation, and discourse planning and structures [1, 33, 79]. There are other concepts such as joint actions and commitments [30], dynamics of intentions during collaboration [50], and task-based planning providing more depth in the context of collaboration [10, 70]. The concept of collaboration has also received attention in the industry and in research in robotic laboratories [26].

Applications of SharedPlans Theory – COLLAGEN [71, 72] is the first implemented system based on the SharedPlans theory. It incorporates certain algorithms for discourse generation and interpretation, and is able to maintain a segmented interaction history, which facilitates the discourse between the human user and the intelligent agent. The model includes two main parts: (1) a representation

of a discourse state and (2) a discourse interpretation algorithm for the utterances of the user and agent [73]. In [35] Heeman presents a computational model of how a conversational participant collaborates in order to make a referring action successful. The model is based on the view of language as goal-directed behaviour, and in his work, he refers to SharedPlans as part of the planning and conversation literature. In [53], Lochbaum and Sidner modify and expand the SharedPlan model of collaborative behavior [34]. They present an algorithm for updating an agents beliefs about a partial shared plan and describe an initial implementation of this algorithm in the domain of network management. Lochbaum, also in [52], provides a computational model (based on the collaborative planning framework of SharedPlans [32]) for recognizing intentional structure and utilizing it in discourse processing. In short, she presents a SharedPlans model for recognizing Discourse Segment Purposes (DSPs) [34] [79] and their interrelationships. CAST (Collaborative Agents for Simulating Teamwork) [89] [90] is a teamwork framework based on the SharedPlans theory. CAST focuses on flexibility in dynamic environments and on proactive information exchange enabled by anticipating what information team members will need. Petri Nets are used to represent both the team structure and the teamwork process, i.e., the plans to be executed. Researchers in [37] discuss developing an ontology of microsocial concepts for use in an instructional system for teaching cross-cultural communication. They believe being acquainted with one another is not a strong enough relationship from which to create a society. Hence, there is a need for commitment and shared plans (as the basis of social life) to achieve a shared goal. In this work, Grosz and Sidner's SharedPlans theory [34] is used to explain the concept of shared plans within the interpersonal relationships of societies in an industrial environment. In [38] Hunsberger and Grosz discuss the idea of whether the rational, utility-maximizing agents should determine commitment to a group activity when there is an opportunity to collaborate. They call this problem the “initial-commitment decision problem” (ICDP) and provide a mechanism that agents can use to solve the ICDP. They use the representation of action, act-types

and recipes in the SharedPlans theory. In [91] an integrated agent-based model for Group Decision Support Systems is proposed and discussed. The decisional model that authors outline in this paper is based on the SharedPlans theory. Rauenbusch and Grosz in [68] formally define a search problem with search operators that correspond to the team planning decisions. They provide an algorithm for making the three types of interrelated decisions by recasting the problem as a search problem. Their model respects the constraints on mental states specified by the SharedPlans theory of collaboration. Babaian et. al. in [3] describe Writer’s Aid, a system that deploys AI planning techniques to enable it to serve as an author’s collaborative assistant. While an author writes a document, Writer’s Aid helps in identifying and inserting citation keys and by autonomously finding and caching potentially relevant papers and their associated bibliographic information from various on-line sources. They believe the underlying concepts of SharedPlans is relevant since in collaborative interfaces like Writers Aid, the users establish shared goals with the system and user and the system both take initiative in satisfying them. In [60] researchers address high-level robot planning issues for an interactive cognitive robot that acts in the presence of or in collaboration with a human partner. They describe a Human Aware Task Planner (HATP) which is designed to provide socially acceptable plans to achieve collaborative tasks. They use notions of plans based on SharedPlans theory. In [80] Sidner and Dzikovska argue that robots, in order to participate in conversations with humans, need to make use of conventions of conversation and the means to be connected to their human counterparts. They provide an initial research on engagement in human-human interaction and applications to stationary robots in hosting activities. They believe hosting activities are collaborative because neither party completely determines the goals to be undertaken nor the means of reaching the goal. To build a robot host, they rely on an agent built using COLLAGEN which is implemented based on the SharedPlans theory.

Applications of Joint Intentions Theory – In [46] authors introduce a lan-

guage for representing joint plans for teams of agents. They describe how agents can organize the formation of a suitably skilled team to achieve a joint goal, and they explain how such a team can execute these plans to generate complex, synchronized team activity. In this paper, authors adopt the underlying concepts of the Joint Intentions theory as the structure of their collaborative agents. Breazeal et. al. in [8] present an overview of their work towards building socially intelligent, cooperative humanoid robots, Leonardo, that can collaborate and learn in partnership with humans. They employ the Joint Intentions theory of collaboration to implement the collaborative behaviors while performing a task in collaboration with humans. In [83] the researchers' goal is to develop an architecture (based on the concepts of Joint Intentions theory) that can guide an agent during collaborative teamwork. They describe how a joint intention interpreter that is integrated with a reasoner over beliefs and communicative acts can form the core of a dialogue engine. Ultimately, the system engages in dialogue through the planning and execution of communicative acts necessary to attain the collaborative task at hand. Mutlu et. al. in [61] discuss key mechanisms for effective coordination toward informing the design of communication and coordination mechanisms for robots. They present two illustrative studies that explore how robot behavior might be designed to employ these mechanisms (particularly joint attention and action observation) to improve measure of task performance in human-robot collaboration. Their work uses Joint Intentions theory to develop shared task representations and strategies for task decomposition. The system GRATE* by Jennings [42] is based on the Joint Intention theory. GRATE* provides a rule-based modelling approach to cooperation using the notion of Joint Responsibilities, which in turn is based on Joint Intentions. GRATE* is geared towards industrial settings in which both agents and the communication between them can be considered to be reliable.

Applications of Hybrid Theories – This domain independent teamwork model, STEAM, has been successfully applied to a variety of domains. From com-

bat air missions [36] to robot soccer [47] to teams supporting human organizations [67] to rescue response [75], applying the same set of STEAM rules has resulted in successful coordination between heterogeneous agents. The successful use of the same teamwork model in a wide variety of diverse domains provides compelling evidence that it is the principles of team-work, rather than exploitation of specific domain phenomena, that underlies the success of teamwork based approaches. In [56] authors provide their RoboCup (robotics soccer testbed) in which their focus is on teamwork and learning challenges. Their research investigation in RobotCup is based on ISI Synthetic, a team of synthetic soccer-players. They also investigate the use of STEAM as their model of teamwork which is influenced by the Joint Intentions and SharedPlans theories. In [45] researchers propose a behavioral architecture C²BDI that allows the enhancement of the knowledge sharing using natural language communication between team members. They define collaborative conversation protocols that provide proactive behavior to agents for the coordination between team members. Their agent architecture provides deliberative and conversational behaviors for collaboration, and it is based on both of the SharedPlans and Joint Intentions theories.

2.2 Affective Computing

2.2.1 Affect and Emotions

2.2.2 Functions of Emotions

2.2.3 Motivation and Theory of Mind

2.3 Computational Models of Emotions

2.3.1 Appraisal Theory

2.3.2 Other Computational Models

2.3.3 Similarities and Differences

2.3.4 Applications in Autonomous Agents and Robots

2.4 Conclusion

In this response, we started by defining the concept of collaboration based on Grosz and Sidner’s work [34], and listed a number of collaboration properties. Then, we provided the background of two prominent collaboration theories which helped develop a better understanding of the actual theories and how they relate to each other. Next, we presented the SharedPlans theory and its major properties, e.g., partial shared plan, recipe, and two notions of intention. Afterwards, we delivered key concepts of the Joint Intentions theory including joint commitment and joint intention. Then, we continued with the hybrid approach of modeling collaboration and provided one of the most well-established models, STEAM. We also briefly mentioned some other approaches. Later, we presented two different lists to compare similarities and differences between SharedPlans and Joint Intentions collaboration theories. We ended this document with different categories of applications of these theories in agent/robot and human collaboration areas.

We believe the SharedPlans and Joint Intentions collaboration theories are the most well-defined and well-established theories in computer science. We found SharedPlans theory more convincing than the other major and subordinate approaches, with respect to its inclusive explanation of the collaboration structure and its association to discourse analysis which directly improves the communicative aspects of a collaboration theory. We also understand the value of Joint Intentions theory due to its clarity and closeness to the foundations of collaboration concepts. These specifications of the Joint Intentions theory can make it applicable in multi-agent system designs and human-robot collaboration. We also consider hybrid approaches valuable, such as STEAM, if they clearly understand drawbacks with existing theories and successfully achieve better collaborative agents by infusing different concepts from different theories. Although all these theories are well-defined and properly introduce collaboration concepts, they mostly explain the structure of a collaboration and they lack the underlying domain-independent processes with which collaborative procedures could be defined more systematically and effectively in different applications.

CHAPTER 3

AFFECTIVE MOTIVATIONAL COLLABORATION THEORY

3.1 Introduction

Current computational theories used for human-robot collaboration specify the structure of collaborative activities, but are weak on the underlying processes that generate and maintain these structures. We argue that emotions are crucial to these underlying processes and we have developed a new computational theory, called Affective Motivational Collaboration Theory, that combines emotion-based processes, such as appraisal and coping, with collaboration processes, such as planning, in a single unified framework. This work is implemented as part of a larger effort to build robots capable of generating and recognizing emotions in order to be better collaborators. We have investigated the mutual influences of affective and collaborative processes in a cognitive theory to support interaction between humans and robots or virtual agents. We build primarily on the *cognitive appraisal* theory of emotions and the *SharedPlans* theory of collaboration to investigate the structure, fundamental processes and functions of emotions in a collaboration. We have developed new algorithms for appraisal processes as part of a new overall computational model. We have evaluated our implemented appraisal algorithms by conducting an online user study.

Although existing collaboration theories explain the important elements of a col-

laboration structure, the underlying processes required to dynamically create, use, and maintain the elements of this structure are largely unexplained. For instance, a general mechanism has yet to be developed that allows an agent to effectively integrate the influence of its collaborator’s perceived or anticipated emotions into its own cognitive mechanisms to prevent shared task failures while maintaining collaborative behavior. Therefore, a process view of collaboration must include certain key elements. It should inherently involve social interactions since all collaborations occur between social agents, and it should essentially constitute a means of modifying the content of social interaction as the collaboration unfolds. The underlying processes of emotions possess these two properties, and social functions of emotions explain some aspects of the underlying processes in collaboration. This work is implemented as part of a larger effort to build robots capable of generating and recognizing emotions in order to be better collaborators.

There is also a communicative aspect of emotions. For instance, emotions are often intended to convey information to others [24]. Emotions are also involved in verbal behaviors. For instance, an utterance can include both content and relational meaning. An emotion might appear to be elicited by the content of the utterance, but in fact be an individual’s response to the relational meaning [64]. The interpretation of these relational meanings are handled by the appraisal of events. Appraisal processes give us a way to view emotion as social [87]. Meaning is created by an individual’s social experiences in the social world, and individuals communicate these meanings through utterances. Consequently, the meaning of these utterances and the emotional communication change the dynamic of social interactions. A successful and effective emotional communication necessitates ongoing reciprocal adjustments between interactants that can happen based on interpretation of each other’s behaviors [54]. This adjustment procedure requires a baseline and an assessment procedure. While the components of the collaboration structure, e.g., shared plan, provide the baseline, emotion-related processes (e.g., appraisal) provide the assessment procedure.

3.1.1 Scenario

3.1.2 Example of a Collaborative Interaction

3.2 Design and Architecture

3.2.1 Mechanisms

3.2.2 Functions of Emotions

3.2.3 Mental States

3.2.4 Attributes of Mental States

CHAPTER 4

COMPUTATIONAL FRAMEWORK

4.1 Introduction

There are several appraisal models (e.g., EMA [57]) contributing in different applications such as social sciences, virtual agents, and robotics. However, none of these models have focused on the appraisal processes during collaboration. We believe appraisal plays a key role in collaboration due to its regulatory and evaluative nature. Also, collaboration induces some changes to appraisal processes due to its unique nature. For instance, although the appraisal models mostly use utility to compute the relevance of an event, we have found new cognitive components involved in determining utility because of the influence of the collaboration. These components, such as the recurrence of a belief by the human collaborator or the influence of the human collaborator’s perceived emotion on the robot’s decisions emphasize the fact that collaboration requires different procedures in appraisal processes. One of our contributions is to ground general appraisal concepts in the specific context and structure of collaboration.

Furthermore, we believe collaboration and appraisal have reciprocal influences on each other. In this chapter, we also talk about the influence of appraisal on collaboration through the goal management process. Also, we discuss our coping mechanism and strategies within the collaboration context. Then, we provide our computational model of three different motives used in our framework. Finally, we briefly discuss other mechanisms in our framework.

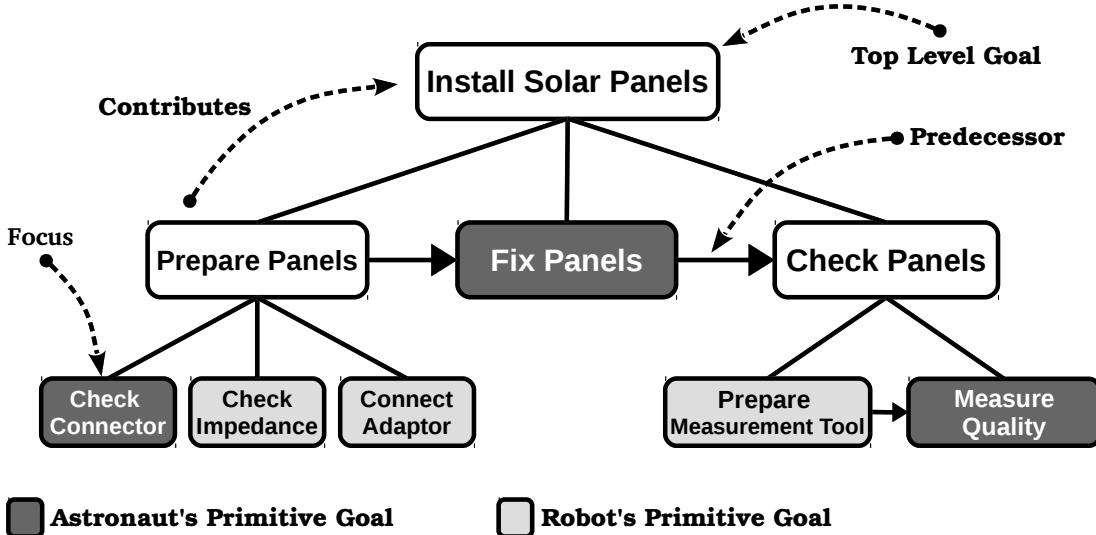


Figure 4.1: Collaboration structure (shared plan).

4.2 Collaboration Mechanism

The Collaboration and Appraisal mechanisms have reciprocal influences on each other. In this section, we focus on information about the collaboration structure which will be incorporated in appraisal processes in Section 4.3. We describe some of the methods in our Collaboration mechanism which are used to retrieve information about the collaboration structure.

The Collaboration mechanism constructs a hierarchy of goals associated with tasks in the form of a hierarchical task network (see Figure 4.1), and also manages and maintains the constraints and other required details of the collaboration including the inputs and outputs of individual tasks, the *preconditions* (specifying whether it is appropriate to perform a task), and the *postconditions* (specifying whether a just-completed task was successful). Collaboration also keeps track of the focus of attention, which determines the salient objects, properties and relations at each point, and shifts the focus of attention during the interaction.

Here, we describe the methods which retrieve information about the collaboration structure, and are used in our algorithms to compute the values of appraisal variables. In these methods, ε_t is the event corresponding to time t , and g_t is a given goal at time t .

- $\text{recognizeGoal}(\varepsilon_t)$ returns the unique goal to which the given event (action, utterance, or emotional expression) directly contributes; it is only one goal since the robot can only do one primitive action at a time in our collaboration model, i.e, in the goal tree, a given primitive action can only directly contribute to one parent goal. The method returns *ambiguous* if it does not recognize a goal in the plan¹.
- $\text{getGoalStatus}(g_t)$ returns whether g_t 's status is ACHIEVED, FAILED, BLOCKED, INAPPLICABLE, PENDING, or IN PROGRESS. In our example, “Check Connector” is the current (focused) goal and it is PENDING, and the “Prepare Panels” and “Install Solar Panels” are IN PROGRESS. The focused goal is the goal that the robot currently pursues.
- $\text{getTopLevelGoal}(g_t)$ returns g_t 's top level goal.
- $\text{precondStatus}(g_t)$ returns the status of the precondition for the given goal whether it is SATISFIED, UNSATISFIED or UNKNOWN. For instance, the precondition for fixing a panel is whether the panel is appropriately located on its frame.
- $\text{isLive}(g_t)$ returns *true* if all the predecessors of g_t are ACHIEVED and all the preconditions are SATISFIED, i.e., PENDING or IN PROGRESS goals; otherwise returns *false*.
- $\text{isFocusShift}(g_t)$ returns *true* if the given goal is not the previous focus (top of the stack); otherwise returns *false*.
- $\text{isNecessaryFocusShift}(g_t)$ returns *true* if the status of the previous focus was ACHIEVED; otherwise returns *false* [49].
- $\text{isPath}(g_1, g_2)$ returns *true* if there is a path between g_1 and g_2 in a plan tree structure; otherwise returns *false*.

¹Ambiguity introduces some extra complexities which are beyond scope of this thesis.

- $getContributingGoals(g_t)$ returns g_t 's children.
- $getPredecessors(g_t)$ returns g_t 's predecessors.
- $getInputs(g_t)$ returns all required inputs for g_t . For example, the goal “Fix Panels” requires inputs such as *welding tool* and *panel*.
- $isAvailable(g_t)$ returns whether the given input is available. For instance, whether the *welding tool* is available for the goal “Fix Panels”.
- $isFocused(g_t)$ returns whether the focus is on g_t .
- $getResponsible(g_t)$ returns responsible agent(s) for g_t . In a dyadic collaboration, both of the agents (jointly) can be partly responsible for a nonprimitive goal, while each (self or other) is responsible for one or more primitive goals. For instance, both the Robot and the Astronaut are responsible for the non-primitive goal of “Install Solar Panels”, whereas it is only the Robot who is responsible for the primitive goal of “Prepare Measurement Tool”.

4.3 Appraisal Mechanism and Underlying Processes

In this section, we focus on the specific problem of appraising the *Relevance* (since other appraisals are only computed for relevant events), *Desirability* (since it discriminates facilitating and inhibitory events towards the collaboration progress), *Expectedness* (since it underlies a collaborative robot's attention), and *Controllability* (since it is associated with the agent's coping ability) of events within a collaborative interaction. There are other appraisal variables introduced in psychological [76] and computational literature [25]. We believe most of these variables can be straightforwardly added to our appraisal mechanism whenever they are required. All of the algorithms in this section use mental states of the robot (discussed in Section 3.2.3) which are formed based on the collaboration structure (Figure 4.2). These algorithms use the corresponding recognized goal of the most recent event at

each turn.

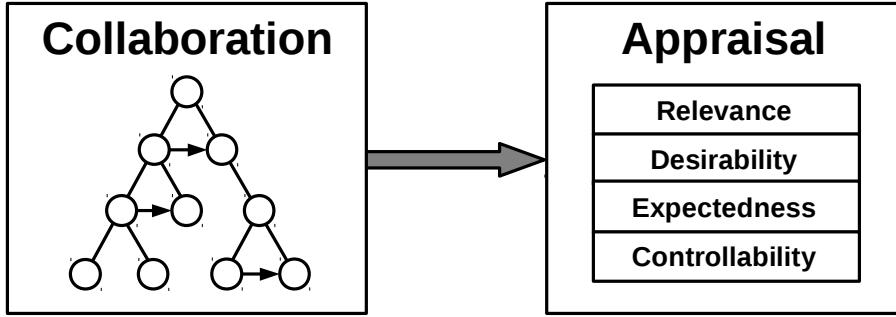


Figure 4.2: Using Collaboration structure in Appraisal (mechanisms in our framework).

4.3.1 Relevance

Relevance is an important appraisal variable since the other appraisal variables are meaningful only for relevant events. Relevance as an appraisal variable measures the significance of an event for the self. An event can be evaluated to be relevant if it has a non-zero utility [57]. However, the utility of an event is also influenced by the other collaborator's emotional expressions as the reflection of the other collaborator's mental state with respect to the status of the collaborative environment. Other appraisal models only consider the utility of an event based on the self's goal and plan.

Algorithm 1 determines the relevance of the given event with respect to the current mental state. The relevance of the event depends on the significance of the event with respect to the collaboration status, which is determined based on the utility of the event as presented in [25, 57]. Our algorithm for computing the relevance of an event during collaboration involves other factors that other appraisal models do not consider. For instance, the human's perceived emotion, recurrence of a belief, or occurrence of a belief about an unrelated goal by the human play important roles by influencing the utility of an event during collaboration. As a result, evaluating the relevance of events can cause a collaborative robot to respond effectively which can

positively impact the status of the shared goal, without dedicating all its resources to every event.

After perceiving an event, the belief about that event represents the event in the robot’s mental state. *recognizeGoal* returns the goal to which the current event contributes, unless it is *ambiguous*; g_t represents the shared goal at time (turn) t within the shared plan. We compute the utility ($-1 \leq \mathcal{U} \leq 1$) of the event using the values of the attributes associated with the existing beliefs, and the attributes of the motive associated with the recognized goal (see details below). We use three belief attributes (see Section 3.2.3) to compute the belief-related part of the utility:

Algorithm 1 (Relevance)

```

1: function ISEVENTRELEVANT(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{recognizeGoal}(\varepsilon_t)$ 
3:    $\mathcal{U} \leftarrow \text{GETEVENTUTILITY}(g_t)$ 
4:    $\tau_t \leftarrow \text{GETEMOTIONALTHRESHOLD}(g_t)$ 
5:   if  $(\tau_t \leq |\mathcal{U}|)$  then
6:     return RELEVANT
7:   else
8:     return IRRELEVANT

```

- *Strength*: The extent to which the preconditions (α), postconditions (β), predecessors (λ), and contributing goals (μ) of a goal are known (SATISFIED or UNSATISFIED) makes beliefs about the goal stronger. An UNKNOWN pre and postcondition status of a goal and its predecessors and contributing goals forms weaker beliefs. For instance, if one knows all predecessors of a pursued goal (e.g., “Check Panels”) are SATISFIED (i.e., “Fix Panels” and “Prepare Panels”), failure of the pursued goal will elicit one’s negative emotion (due to the strong beliefs related to the goal); whereas not knowing the status of the goal-related factors (e.g., whether the Astronaut could find the tool to fix a panel) causes one to form weaker beliefs about the goal.

- *Saliency (S)*: Beliefs related to the focused goal are more salient than beliefs related to any other goal in the plan; according to Figure 4.1, if one of the collaborators is preparing a solar panel, beliefs related to all of the other *live* (PENDING or IN PROGRESS) goals (e.g. “Connect Adaptor”) will be less salient than beliefs related to the focused goal, i.e., “Check Connector”. Beliefs’ saliency decreases according to their corresponding *live* goal’s distance from the focused goal in the shared plan. *Non-live* goals will not be salient.
- *Persistence (P)*: The recurrence of a belief over time (turns) increases the persistence of the belief. Beliefs occurring only once have the lowest value of persistence. For instance, if the Astronaut repeatedly says that she can not find the measurement tool to check the connector, the Robot could pursue a new goal outside of the shared plan to acknowledge Astronaut’s concern.

We also use two motive attributes discussed in Section 3.2.3 to compute the motive related part of the utility (\mathcal{U}):

- *Urgency (γ)*: There are two factors impacting the urgency of a motive: a) whether the goal directing the given motive is the predecessor of another goal for which the other collaborator is responsible, and b) whether achieving the goal directing the given motive can mitigate the other collaborator’s negative valenced emotion. For instance, if the Robot has a private goal to fetch another panel while the Astronaut is waiting for the Robot to connect the adaptor, connecting the adaptor will be more urgent than Robot’s private goal.
- *Importance (η)*: A motive is important if failure of the directing goal causes an impasse in the shared plan (i.e., no further goal is available to achieve), or achievement of the directing goal removes an existing impasse. For example, if the Robot cannot find the adaptor (an impasse to connect the adaptor), and the Astronaut provides another adaptor (external motive), the new motive becomes important to remove the impasse in the shared plan.

We provide the utility function (\mathcal{U}) in Equation 4.1. This function uses: saliency (S) and persistence (P) of the belief related to the recognized goal, the recognized goal's status (v), and the aggregation of belief and motive attributes (Ψ) according to Equation 4.1.

$$\mathcal{U}(\varepsilon_t) = \begin{cases} vP \cdot S^\Psi & \Psi > 0 \\ 0 & \Psi = 0 \end{cases} \quad (4.1)$$

Intuitively, we use v to generate positive and negative utility values. The v 's value becomes +1 if the status of the corresponding goal is ACHIEVED, PENDING, or IN PROGRESS, and v 's value becomes -1 if the status of the corresponding goal is FAILED, BLOCKED, or INAPPLICABLE. The P influences the value of utility only as a coefficient since recurrent beliefs are not formed frequently during collaboration. The Ψ value indicates the magnitude of the influence of beliefs and motives using their attributes. Hence, the Ψ value impacts the saliency value of beliefs exponentially, helping to differentiate between beliefs.

In equation 4.2, the subscript k refers to the *known* goal-related factors (SATISFIED or UNSATISFIED); whereas the subscript *all* includes both *known* and *unknown* goal-related factors. In this equation, both urgency (γ) and importance (η) attributes of motives can impact the outcome of the goal-related belief attributes' ratio, and ultimately the Ψ value.

$$\begin{aligned} \Psi &= \frac{\alpha_k + \beta_k + \lambda_k + \mu_k}{\alpha_{all} + \beta_{all} + \lambda_{all} + \mu_{all}} + \eta + \gamma & (4.2) \\ \eta, \gamma &\in \mathbb{N}, & \eta, \gamma \geq 0 \\ \alpha_k, \beta_k, \lambda_k, \mu_k &\in \mathbb{N}, & \alpha_k, \beta_k, \lambda_k, \mu_k \geq 0 \\ \alpha_{all}, \lambda_{all}, \mu_{all} &\in \mathbb{N}, & \alpha_{all}, \lambda_{all}, \mu_{all} \geq 0 \\ \beta_{all} &\in \mathbb{N}, & \beta_{all} \geq 1 \end{aligned}$$

The significance of an event in a collaborative environment is based on the utility of the event and the human's perceived emotion. The human's perceived emotion

influences the relevance of the event in the form of a threshold value τ_t . In Equation 4.3, we use the valence of the perceived emotion (\mathcal{V}_{e_h}) to compute τ_t .

$$\tau_t = \begin{cases} 1 - \mathcal{V}_{e_h} & \mathcal{V}_{e_h} > 0 \\ |\mathcal{V}_{e_h}| & \mathcal{V}_{e_h} \leq 0 \end{cases} \quad (4.3)$$

$$\mathcal{V}_{e_h} \in \mathbb{R}, \quad -1 \leq \mathcal{V}_{e_h} \leq 1$$

Hence, perceiving human's positive emotion (e.g., happiness) reduces the threshold value which makes the robot find an event RELEVANT with even a slightly positive utility. Similarly, an event can be considered IRRELEVANT even though the utility has a relatively positive value, because of perceiving the human's negative emotion.

4.3.2 Desirability

Desirability characterizes the value of an event to the robot in terms of whether the event facilitates or thwarts the collaboration goal. Desirability captures the valence of an event with respect to the robot's preferences [25]. In a collaborative robot, preferences are biased towards those events facilitating progress in the collaboration. Desirability plays an important role in the overall architecture; it makes the processes involved in the other mechanisms (e.g., Motivation and Theory of Mind) and consequently the robot's mental state, congruent with the collaboration status which is a collaborative robot's desire. Therefore, it causes the robot to dismiss events causing inconsistencies in the robot's collaborative behavior. Moreover, desirability is also crucial from the collaboration's point of view.

Algorithm 2 provides a process in which the desirability of an event is computed with regard to the status of the shared goal; i.e., it operates based on whether and how the event changes the status of the current shared goal. It distinguishes between the top level goal and the current goal because the top level goal's change of status attains a higher positive or negative value of desirability. For instance, failure of

Algorithm 2 (Desirability)

```
1: function ISEVENTDESIRABLE(Event  $\varepsilon_t$ )  
2:    $g_t \leftarrow \text{recognizeGoal}(\varepsilon_t)$   
3:    $g_{top} \leftarrow \text{getTopLevelGoal}(g_t)$   
4:   if ( $\text{getGoalStatus}(g_{top}) = \text{ACHIEVED}$ ) then  
5:     return MOST-DESIRABLE  
6:   else if ( $\text{getGoalStatus}(g_{top}) = \text{FAILED}$ ) then  
7:     return MOST-UNDESIRABLE  
8:   else if ( $\text{getGoalStatus}(g_{top}) = \text{BLOCKED}$ )  
9:     ( $\text{getGoalStatus}(g_{top}) = \text{INAPPLICABLE}$ ) then  
10:    return UNDESIRABLE  
11:   else if ( $\text{getGoalStatus}(g_{top}) = \text{PENDING}$ )  
12:     ( $\text{getGoalStatus}(g_{top}) = \text{INPROGRESS}$ ) then  
13:       if ( $\text{getGoalStatus}(g_t) = \text{ACHIEVED}$ ) then  
14:         return DESIRABLE  
15:       else if ( $\text{getGoalStatus}(g_t) = \text{FAILED}$ ) then  
16:         return MOST-UNDESIRABLE  
17:       else if ( $\text{getGoalStatus}(g_t) = \text{BLOCKED}$ )  
18:         ( $\text{getGoalStatus}(g_t) = \text{INAPPLICABLE}$ ) then  
19:           return UNDESIRABLE  
20:         else if ( $\text{getGoalStatus}(g_t) = \text{PENDING}$ )  
21:           ( $\text{getGoalStatus}(g_t) = \text{INPROGRESS}$ ) then  
22:             return NEUTRAL
```

the top level goal (e.g., installing solar panel) is more undesirable than failure of a primitive goal (e.g., measuring the quality of the installed panel).

A top level goal' status must be ACHIEVED (i.e., SATISFIED postcondition) to consider the event MOST-DESIRABLE. When the goal's status is FAILED (i.e., UNSATISFIED postcondition) or BLOCKED, the associated event has the MOST-UNDESIRABLE or UNDESIRABLE values respectively. A goal is BLOCKED if any of the required goals or goals recursively through the parent goal are not ACHIEVED. An INAPPLICABLE goal is also considered as UNDESIRABLE. A goal is INAPPLICABLE if any of its predecessors are not ACHIEVED, and/or its preconditions are not SATISFIED. For PENDING and INPROGRESS top level goals, the status of the current goal associated

with the top level goal determines the status of the event ε_t . Only a non-primitive goal can have INPROGRESS status, if it has been started but is not yet completed. A goal can be PENDING if it is live, or if it is a non-primitive goal that has not been started yet. ACHIEVED current goals mark an event (ε_t) as DESIRABLE, while FAILED or BLOCKED current goals render the event associated with them as MOST-UNDESIRABLE and UNDESIRABLE respectively. PENDING or INPROGRESS current goals mark their associated events as NEUTRAL.

4.3.3 Expectedness

Expectedness is the extent to which the truth value of a state could have been predicted from causal interpretation of an event. In the collaboration context the expectedness of an event evaluates the congruency of the event with respect to the existing knowledge about the shared goal. Thus, expectedness underlies a collaborative robot’s attention. The collaboration mechanism uses expectedness to maintain the robot’s attention and subsequently its mental state with respect to the shared goal. Reciprocally, the appraisal mechanism uses the underlying information of the collaboration structure to evaluate the expectedness of an event [84].

In Algorithm 3 we provide the process of computing the expectedness based on the shared plan and status of the shared goal. The key point in this algorithm is the status of the current shared goal (g_t), which is associated with the event ε_t and its relationship with the top level goal (g_{top}).

The intuition captured here is that one expects the current goal to be finished before undertaking another activity, but the goals that can be the next focus of attention are also to be expected. Therefore, if the goal is live, the algorithm checks whether the goal has not changed, or whether the interpretation of the last event results in a necessary focus shift. Shifting the focus to a new goal is necessary when the former goal is achieved and a new goal is required. Consequently the new event is the MOST-EXPECTED one. However, even if the focus shift is not necessary, the new event can be considered as EXPECTED, since the corresponding goal is already

Algorithm 3 (Expectedness)

```
1: function ISEVENTEXPECTED(Event  $\varepsilon_t$ )  
2:    $g_t \leftarrow \text{recognizeGoal}(\varepsilon_t)$   
3:    $g_{top} \leftarrow \text{getTopLevelGoal}(g_t)$   
4:   if ( $\text{isLive}(g_t)$ ) then  
5:     if ( $\neg \text{isFocusShift}(g_t)$   
6:        $\text{isNecessaryFocusShift}(g_t)$ ) then  
7:         return MOST-EXPECTED  
8:       else  
9:         return EXPECTED  
10:    else  
11:      if ( $\text{isPath}(g_t, g_{top})$ ) then  
12:        return UNEXPECTED  
13:      else  
14:        return MOST-UNEXPECTED
```

live. For goals that have not yet been started (that is, are not live), the algorithm must determine how unexpected it would be to pursue one now; if the goal is at least in the plan, i.e., on the path to the top level goal, it is just UNEXPECTED while any others are MOST-UNEXPECTED.

4.3.4 Controllability

Controllability is the extent to which an event can be influenced; it is associated with a robot’s ability to cope with an event [25]. Thus, a robot can determine whether an event’s outcome can be altered by actions under either of the collaborators’ control. In other words, controllability is a measure of a robot’s ability to maintain or change a particular state as a consequence of an event.

Controllability is important for the overall architecture. For instance, the robot can choose to ask or negotiate about a collaborative task which is not controllable, or form a new motive to establish an alternative goal for the current uncontrollable event. In general, other mechanisms in the architecture use the controllability output in their decision making processes; meanwhile controllability uses information from

Algorithm 4 (Controllability)

```
1: function ISEVENTCONTROLLABLE(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{recognizeGoal}(\varepsilon_t)$ 
3:    $\mathcal{M} \leftarrow \text{GETAGENCYRATIO}(g_t)$ 
4:    $\mathcal{R} \leftarrow \text{GETAUTONOMYRATIO}(g_t)$ 
5:    $\mathcal{P} \leftarrow \text{GETSUCCPREDECESSORSRATIO}(g_t)$ 
6:    $\mathcal{I} \leftarrow \text{GETAVAILABLEINPUTS}(g_t)$ 
7:    $\mathcal{V}_{e_h} \leftarrow \text{GETEMOTIONVALENCE}(g_t)$ 
8:    $\omega \leftarrow \text{GETWEIGHTS}(g_t)$ 
9:    $\mathcal{X} \leftarrow \frac{\omega_0 \cdot \mathcal{M} + \omega_1 \cdot \mathcal{R} + \omega_2 \cdot \mathcal{P} + \omega_3 \cdot \mathcal{I}}{\omega_0 + \omega_1 + \omega_2 + \omega_3} + \mathcal{V}_{e_h}$ 
10:  if ( $\mathcal{X} > 0$ ) then
11:    return CONTROLLABLE
12:  else
13:    return UNCONTROLLABLE
```

the collaboration structure, e.g., predecessors of a goal.

An important determinant of one’s emotional response is the sense of control over occurring events. This sense of subjective control is based on one’s reasoning about self’s power. For instance, the robustness of one’s plan for executing actions can increase one’s sense of power and subsequently the sense of control. In the collaboration context, we have translated the sense of control into a combination of four different factors including a) *agency* and b) *autonomy* of the robot, as well as the ratios of c) *successful predecessors*, and d) the *available inputs* of a given goal (i.e., g_t) in the shared plan.

In Algorithm 4, we partially compute the controllability of an event based on the above four factors. We use weighted averaging of these factors to determine their impact on the controllability of an event (line 9). The value of all these weights are set to *1.0* for the purpose of simplicity (**getWeights**). These weights can be adjusted after further investigating the influence of these factors, and implementing

other mechanisms in the overall architecture. We believe that the human’s perceived emotion also impacts the controllability of an event (**getEmotionValence**). The ($-1.0 \leq \mathcal{V}_{e_h} \leq 1.0$) is the valence value of the human’s perceived emotion. Positive emotions, e.g., happiness, possess positive values, and negative emotions, e.g., anger, have negative values. The magnitude of this value can change with respect to the intensity of the perceived emotion. Thus, a positive controllability value indicates that an event is **CONTROLLABLE**; otherwise **UNCONTROLLABLE**.

GetAgencyRatio: *Agency* is the capacity of an individual to act independently in a given environment. In a collaborative environment collaborators are sometimes required to act independently of each other. Hence, they need to have some internal motives that are formed based on their own mental states rather than motives that are reinforced by the other. These internal motives will lead the collaborators to acquire new intentions when required. If the robot’s mental state possesses only an internal motive supporting the recognized goal, we consider a maximum agency value denoted as \mathcal{M} in Algorithm 4 (i.e., $\mathcal{M} = 1.0$); otherwise we consider the minimum agency value (i.e., $\mathcal{M} = 0.0$).

GetAutonomyRatio: *Autonomy* is the ability to make decisions without the influence of others, and implies acting on one’s own and being responsible for that. In a collaborative environment, tasks are delegated to the collaborators based on their capabilities. Therefore, each collaborator is responsible for the delegated task and the corresponding goal. In Algorithm 4, \mathcal{R} denotes the value of autonomy with regard to the goal g_t . This value ($0.0 \leq \mathcal{R} \leq 1.0$) is the ratio of the number of goals contributing to g_t for which the robot is responsible over the total number of contributing goals, if the goal associated with the current event is a nonprimitive goal. However, if the associated goal of the current event corresponds to a primitive goal the value of \mathcal{M} would be 0.0 or 1.0. In general, higher autonomy leads to a more positive value of controllability.

GetSuccPredecessorsRatio: The structure of a shared plan contains the order of the required *predecessors* of a goal. Predecessors of a goal, g_t , are goals that the

collaborators should achieve before trying to achieve goal g_t . We use the ratio of successfully achieved predecessors of the recognized goal over the total number of predecessors of the same goal. If all of the predecessors of the given goal are achieved, then $\mathcal{P} = 1.0$ which is the maximum value for \mathcal{P} . On the contrary, failure of all of the predecessors will lead to $\mathcal{P} = 0.0$. Therefore, a higher \mathcal{P} value positively impacts the value of controllability for the current event.

GetAvailableInputs: Finally, *inputs* of a task are the required elements that the collaborators use to achieve the specified goal of the task. These inputs are also part of the structure of a shared plan. We compute the ratio of the available required inputs over the total required inputs of the goal associated with the current event. This value (denoted as \mathcal{I} in Algorithm 4) will be bound between 0.0 and 1.0. Similar to the other factors in the controllability process, the closer the value of \mathcal{I} gets to 1.0, the more positive impact it has on the overall controllability value of the event.

In summary, the output of these four appraisal processes serves as critical input for the other mechanisms of the Affective Motivational Collaboration Framework, shown in Chapter 3. By providing adequate interpretation of events in the collaborative environment, the appraisal mechanism enables the robot to carry out proper collaborative behaviors.

4.4 Goal Management

A collaborative robot needs to be able to regulate and manage shared goals during collaboration. Emotion has a crucial influence on this goal management process. In this section, we provide a cost function that we use to choose the goal in the shared plan with the lowest cost value out of a set of alternative goals. This cost function is a) based on the goal attributes, b) with respect to the reverse appraisal of the perceived emotion, and c) the appraisal of the collaborative environment.

Goals represent an important part of the context during collaboration. However, not all goals are appropriate to pursue at the moment, depending on conditions. In fact, it can be destructive for a collaboration to pursue a good goal in a poor context. Therefore, a collaborative robot must be able to manage shared goals during collaboration. The goal management process has a critical influence on a collaborative robot’s behavior by maintaining or shifting the focus of attention to an appropriate goal based on the collaboration status.

Changes in a collaboration environment alter the balance of alternative goals. These changes can reflect the collaborators’ internal changes and the influence of their actions. In a collaboration environment, emotions represent the outcome of underlying mental processes of the collaborators. Emotions have many different functions [77] including goal management. Goal-oriented emotions such as anger, frustration and worry regulate the mental processes influenced by one’s internal goals. In our example in this section, a robot and an astronaut are collaborating to install solar panels. When one of the astronaut’s goals is blocked, the robot must manage the shared goals in order to prevent failure of the collaboration. By using reverse appraisal [21] of the astronaut’s emotion and its own appraisal of individual goals, the robot is able to successfully shift the focus of attention from the blocked goal (eliciting worry in the astronaut) to an appropriate one to maintain the collaboration. A similar example is provided our conducted user study, which is explained in Chapter 5.

Here, we describe the goal management process in our framework using an astronaut-robot collaboration example. We introduce the goal management process based on a cost function including the influence of affective appraisal and reverse appraisal processes. Goal management is a crucial part of our investigation of the reciprocal influence of appraisal on a collaboration structure (see Figure 4.3).

As we mentioned earlier, we use four appraisal variables including: relevance, desirability, expectedness and controllability. The outcome of each appraisal process is a specific value for the corresponding appraisal variable. The vector containing

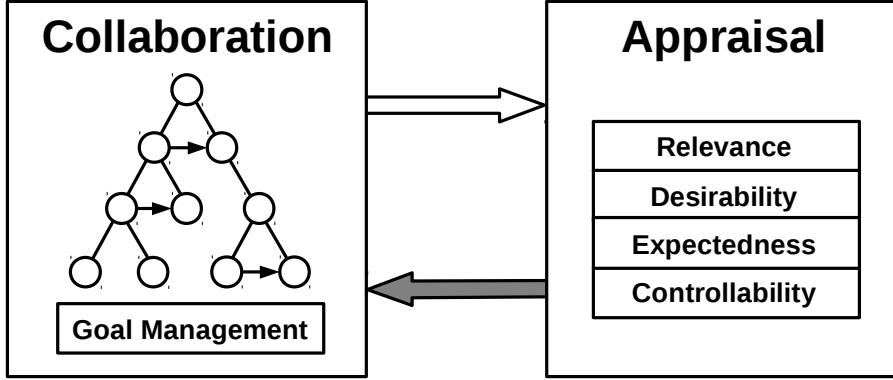


Figure 4.3: Using Appraisals’ outcome to influence Collaboration structure (mechanisms in our framework).

these appraisal variables can be mapped to a particular emotion instance at each point in time when required. Moreover, the functions of emotions, such as goal management, in a social setting and the meaning of the collaborator’s perceived emotion in collaboration context are also important.

A collaboration structure provides a hierarchy and constraints of the shared goals in the form of a shared plan (Figure 5.1) which contains both the robot and the human collaborator’s goals. The robot pursues the goals for which the robot is responsible in the shared plan. However, there can be several live goals available for the robot to pursue at each point in time during collaboration. A goal is live if all of its predecessors are achieved and all of its preconditions are satisfied. Therefore, a collaborative robot requires a mechanism to choose between a set of live goals. We believe appraisal processes are crucial to choose between the available live goals, since the appraisals are the immediate outcome of the robot’s assessment of the collaboration environment.

For instance, Figure 4.4 shows a non-primitive “Prepare Panels” goal decomposed into three primitive goals. Therefore, if “Prepare Panels” is live, its primitive goals can be pursued by the responsible agent. In our example, the astronaut is responsible for the “Check Connector” goal; the robot is responsible for the remaining two primitive goals. According to the collaboration mechanism in our overall

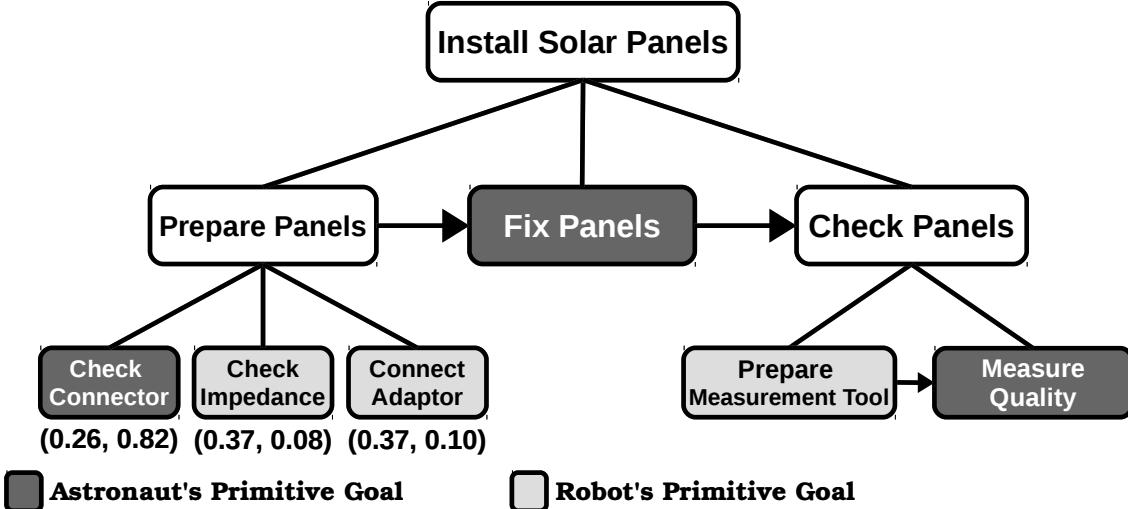


Figure 4.4: Cost values indicated by tuples with (second number) and without (first number) the influence of emotions.

framework, “Check Connector” is in focus, with the astronaut pursuing this goal. Suddenly, however the astronaut tells the robot that she can not find the connector and she is *worried* about failure of this goal. The robot’s response to this situation will be explored below as we discuss details of our cost function.

Equation 4.4 shows the function to calculate the cost of each live goal. The base in the equation calculates the cost of pursuing any given goal. The three functions used to calculate the cost are: *proximity* $P(g)$, *difficulty* $D(g)$, and *specificity* $S(g)$ (see equations 4.6 to 4.8).

$$Cost(g) = (\omega_0.P(g) + \omega_1.D(g) + \omega_2 \cdot \frac{1}{S(g) + 1})^\Gamma \quad (4.4)$$

For simplicity in this example, we assume equal values for the weights: $\omega_i=1$.

$$\Gamma = -C[(R_r + 1)D_r + \alpha(R_h + 1)D_h] \quad (4.5)$$

The exponent part of our cost function (Equation 4.5) captures a) the influence of the human’s perceived emotional instance, and b) the influence of self appraisal of the given goal. $R_h \in [0, 1]$ and $D_h \in [-1, 1]$ are the relevance and desirability

values respectively, which are based on the *reverse* appraisal of the human’s perceived emotion. For instance, if the astronaut is *worried*, D_h is negative, e.g., -0.8 (depending on how undesirable the event is according to reverse appraisal), and R_h will be 1 for the active goal and its value descends to 0 for other live goals depending on their distance to the active goal in the shared plan (e.g., 0.1).

$R_r \in [0, 1]$ and $D_r \in [-1, 1]$ are relevance and desirability values, provided by the *self* appraisal functions for all of the live goals. For instance, for the active goal for which the astronaut was *worried*, D_r can be positive, e.g., 0.8 (depending on the self’s desirability appraisal function); R_r can be 1, since the active goal is relevant for the robot. These values will change for the other live goals depending on how relevant they are with respect to the collaboration status (e.g., 0.9 and 0.8). Finally, $C \in [1, \infty)$ is a constant (e.g., 2) used to control the influence of affect on cost value. It is negative since undesirability (negative values) should increase the cost. $\alpha \in [1, \infty)$ is another constant (e.g., 3) used to control the importance of reverse appraisal relative to self appraisal.

The *proximity* of a goal indicates how far the goal is from the current active goal in the shared plan. It is calculated by the distance function (Equation 4.6) which returns the number of edges between the current active goal g_{act} , and the given goal g in the shared plan. In our example, $P(g)$ is 2 for both “Check Impedance” and “Connect Adaptor” goals.

$$P(g) = \max\{1, \text{distance}(g_{act}, g)\} \quad (4.6)$$

The *difficulty* of a goal is a function of three parameters (Equation 4.7) which consider the difficulty based on a) topology of the shared plan tree (domain independent), and b) the amount of effort required to pursue a given goal (domain dependent). The $\sum pred_e(g)$ is the sum of efforts that all the *predecessors* of a given goal g require. The $\sum desc_e(g)$ is the sum of efforts that all the *descendants* of a given goal g require. The effort values represent the amount of effort for the goals

with respect to the domain. In our example, we assume the values of all the goal efforts are 1 for simplicity. The $H(g)$ is the height of the given goal g . The heights of all primitives under “Prepare Panel” goal are 0 in our example.

$$D(g) = (H(g) + 1) \times \left[\sum_{m=0}^M pred_e(g) + \sum_{n=0}^N desc_e(g) \right] \quad (4.7)$$

The *specificity* of a goal is the function of *depth* (distance from the root) and *degree* (number of children in the graph) of a given goal g . The first non-primitive goal (root) is the least specific goal, and the primitives (leaves) are the most specific goals. As calculated based on Figure 4.4, the values of $S(g)$ for the three primitives under the “Prepare Panels” are 2.

$$S(g) = \frac{depth(g)}{degree(g) + 1} \quad (4.8)$$

The tuples below the goals in Fig. 4.4 indicate the cost value of each goal. The first number in each tuple is the normalized cost value without the influence of the affective part of the cost function, i.e., the exponent is equal to 1 in Equation 4.4. The second number of each tuple indicates the normalized value of the cost including the influence of affective appraisal and the astronaut’s perceived emotion.

Based on our cost function, the cost of completing the primitive goal “Check Connector” is 0.82 (see Figure 4.4). As shown, when affect is not considered the cost is 0.26; the negative emotion of the astronaut (worry) significantly increases the cost of the current goal, and also impacts the other two primitive live goals under the same parent. Therefore, instead of insisting on pursuing the same blocked goal which has caused the astronaut’s negative emotion, the robot can mitigate the astronaut’s emotions by adapting to her worry. The robot shifts the focus of attention to “Check Impedance” to maintain progress and prevent failure of the collaboration.

We use our proposed cost function in our goal management algorithm to integrate affective appraisal into the collaboration mechanism in our framework. A similar situation is used in our conducted user study (see Chapter 5) to evaluate the human’s

perception of the robot’s behavior as a result of the goal management process.

4.5 Coping Mechanism and Strategies

We have implemented the Coping mechanism to determine how the agent would respond to events using our framework. Our Coping mechanism includes a set of coping strategies that can be triggered based on different conditions (see Figure 4.5). All of these coping strategies are known in the literature. Some of our coping strategies, i.e., *planning*, *active coping* and *seeking social support for instrumental reasons*, are categorized as problem-focused and others, i.e., *acceptance*, *mental disengagement*, and *shifting responsibility*, are categorized as emotion-focused startegies as described in [25]. We selected these six coping strategies since they let our agent demonstrate distinct behaviors with respect to the output of the appraisal mechanism and the agent’s mental state in our framework. The behaviors and underlying processes associated with these coping strategies are described as follows.

4.5.1 Planning

The *planning* coping strategy works based on the shared plan and the task structure introduced as an input to our framework. The task structure includes the hierarchy and ordering of the tasks, the required inputs of each task as well as the preconditions and postconditions of individual tasks. We use this task structure to create our shared plan which includes the primitive and non-primitive goals that our agent and its collaborator want to achieve throughout their collaboration. Therefore, our agent executes actions related to its own goals based on this shared plan, and uses the same shared plan to associate goals and their status with the human collaborator. To achieve a goal the agent is required to execute an action, and to execute an action the agent needs to have the right intention. In our framework, whenever this coping strategy is activated the Coping mechanism provides the selected intention to the Action mechanism. The Action mechanism executes an action based on the

given intention to achieve the corresponding goal in the shared plan.

4.5.2 Active Coping

The *active* coping strategy can provide one or all of the following three different intentions with respect to whether this coping strategy is activated and the required conditions are provided. Firstly, this coping strategy can provide an intention to *acknowledge* the human’s emotions. For instance, if the human expresses an emotion with negative valence, the agent can acknowledge human’s negative emotion accordingly. Secondly, the active coping strategy can provide an intention to *respond* to the human if the human asks a question. Currently, in our framework, the agent can respond to the human if the human asks the agent: a) what input is required to achieve a goal, b) how to do a task to achieve a goal, c) to achieve a goal, d) who is responsible to achieve a given goal. For instance, if the human asks the agent to achieve a goal, the active coping strategy forms an intention to either accept the human’s proposal (if achieving the given goal is controllable for the agent), or reject the human’s proposal (if it is not controllable for the agent). Thirdly, the active coping strategy can form an intention to *delegate* a task to the human collaborator. The intention for task delegation can be formed if the agent fails to achieve its own goal, and the human’s perceived emotion is not negative. As mentioned earlier, any or all of these intentions can be formed if active coping is selected. The agent acts accordingly by passing these intentions to the Action mechanism. For instance, if the human is frustrated about a failure that occurred when using a tool to perform its own task and asks the agent whether the agent can provide its own tool, the active coping strategy forms a new intention to acknowledge the human’s frustration and responds to the human by providing the right tool (input) to use and fulfill the task. In this example, there will be no new intention to delegate a new goal to the human since the agent perceives the human’s negative emotion.

4.5.3 Seeking Social Support for Instrumental Reasons

The *seeking social support for instrumental reasons* strategy forms new intentions for the agent whenever the agent needs the human's help and needs to ask questions from the human collaborator to make progress in collaboration. The questions that our agent can ask are the reciprocal of those questions that the human can ask and the human can respond as we mentioned above. Therefore, our agent can ask a) what input is required to achieve a goal, b) how to do a task to achieve a goal, c) the human to achieve a goal, d) who is responsible to achieve a given goal. Reciprocally, again, the agent expects the human collaborator to accept or reject the agent's proposals. In our framework, whenever this strategy is activated the agent considers human's perceived emotion. For instance, if the human is worried about the outcome of a task failure, the agent does not form an intention to ask questions about any of the above cases and consequently prevents asking for more help.

4.5.4 Acceptance

The *acceptance* coping strategy forms an intention to drop the intention of pursuing a goal. In our framework, if this strategy becomes activated, the intention to pursue the current goal will be dropped; see Figure 4.5. The acceptance strategy also forms an intention to inform the human collaborator about the agent's decision on not pursuing the current goal.

4.5.5 Mental Disengagement

The *mental disengagement* coping strategy forms new intention to lower the negative emotional charge associated with a goal in the event of a failure or an impasse. We use our goal management algorithm (see section 4.4) as the underlying process used as the result of selecting this strategy to dissociate the collaboration process and subsequently disengage the collaborator from a negative event (e.g., failure to achieve

a goal). This disengagement helps the agent to lower the utility of an unsuccessful goal achievement attempt and focus on other achievable goals with respect to their costs to facilitate progress of collaboration. In our framework, this coping strategy forms an intention to run the goal management process. As the result of mental disengagement activation, the mechanism also forms another intention to inform the human about the outcome of the goal management process, i.e., whether the agent proposes switching to pursue another goal with lower cost, or if there is not much the agent can do since there is no other goal with a lower cost to pursue. The process and example of choosing another goal with a lower cost are shown in section 4.4.

4.5.6 Shifting Responsibility

The *shifting responsibility* strategy forms new intention to shift the blame from the agent to another entity. In our framework, we use this strategy to mitigate the influence of negative events causing negative emotions in the agent or the human collaborator. For instance, if this strategy becomes activated as a result of a failure, a new intention will be formed to blame the third person who provided the input (if the task needed a tool as an input) or the other collaborator. It can also form an intention to give the credit to the human collaborator to mitigate human's negative emotions.

4.5.7 Activation of Coping Strategies

In our Coping mechanism, there are three components involved as the activation criteria for each coping strategy. The first criterion is the conjunction of emotion valences of the self and the other collaborator (see Emotion Valence column in Figure 4.5). For instance, if the valence of the human collaborator's emotion is *negative* and the valence of the agent's emotion is also *negative*, the active coping, the acceptance, and the mental disengagement coping strategies are the coping strategy

candidates that have potential to become activated if the other activation criteria also exist for any of them. For example, if the human collaborator is frustrated and the agent’s elicited emotion is guilt, the three above mentioned coping strategies become potential candidates to be selected as the agent’s active coping strategy. The second criterion is the need for the agent to cope with an event. The values of our three different motives (i.e., *satisfaction*, *achievement*, and *external*) are involved in the decision of whether there is a need for a particular coping strategy to become activated. We use conjunction of satisfaction motive’s value with the disjunction of achievement and external motives. For instance, if we have highly negative values for all three motives for the potential candidates of coping strategies based on the example we mentioned above, the acceptance coping strategy will be selected as the strategy with the highest need for the agent. For example, this kind of condition can occur when the agent fails doing its own task and pursuing the current goal (negative satisfaction motive), and can not find another goal to overcome the impasse (negative achievement motive). The details about how the motive values are computed is presented in Section 4.6. Finally, the ability to cope with an event is the third criterion that impacts the decision of whether the selected coping startegy can be activated. The controllability of an event represents whether the agent is able to control the situation occurring with the given event. In our example, if the agent finds the event uncontrollable, the acceptance coping strategy becomes activated (see Figure 4.5).

4.6 Motivation Mechanism

As we discussed in Chapters 2 and 3, motives are goal-driven emotion-regulated constructs indicating an urge related to their goal. There are several motives in psychological and computational literatues as we reviewed in Chapter 2. However, none of these computational models have particularly focused on the application of motives in the collaboration context. In our framework, we have implemented

Coping Strategy	Emotions (AND)		Need [a AND (b OR c)]			Ability
	Other	Self	Satisfaction Motive (a)	Achievement Motive (b)	External Motive (c)	Controllability
Planning	Neutral Positive	Any	-/+	high +	high +	High
Active Coping	Any	Neutral Negative	-/+	med +	med +	High
Seeking Social Support for Instrumental Reasons	Neutral Positive	Any	-/+	low +	low +	Low
Acceptance	Negative	Negative	high -	high -	high -	No
Mental Disengagement	Neutral Negative	Neutral Negative	low/med -	low/med -	low/med -	No
Shifting Responsibility	Neutral Positive	Negative	high -	-/+	-/+	No Low

Figure 4.5: Conditions for selecting coping strategies

three computational models of motives for *satisfaction*, *achievement*, and *external* motives. We use the values of these three motives in other mechanisms including the Coping mechanism as we described in Section 4.5 and show in Figure 4.5.

4.6.1 Satisfaction Motive

The satisfaction motive indicates the satisfaction level with the collaboration for the agent and its human collaborator. The satisfaction motive process maintains the value of *satisfaction drive* throughout the collaboration. The satisfaction drive is the quantitative weighted accumulation of desirability values between -1 and +1 over time. For instance, if the desirability values of the agent's appraisal over three consecutive turns are $\{0.75, 0, -0.25\}$, and their corresponding weights are $\{0.25, 0.5, 1.0\}$, the satisfaction drive value will be $(0.25)(0.75) + (0.5)(0) + (1.0)(-0.25)$ which is -0.0625. Notice that the latest desirability values get higher weights. Intuitively, it is because older desirable events have less influence on overall desirability and consequently the satisfaction level of the collaboration. The same process computes the satisfaction drive values for the agent and the human collaborator. Only the sources of desirability values are different, i.e., appraisal for the agent and reverse appraisal for the human collaborator. Then, the satisfaction motive process computes the difference between the current and the previous satisfaction drives, called delta of satisfaction drive value, δ_{sat} . As shown in equation 4.9, we use the δ_{sat} value in all three functions to compute the overall satisfaciton motive's value \mathcal{M}_{sat} . We also use three different functions with respect to the valence value of the the human collaborator's perceived emotion. Our satisfaction motive's model has three user defined parameters $\mathcal{S}_{sat} \in [0, 1.5]$, i.e. strength of motive, $\mathcal{B}^{\mathcal{L}}$ where \mathcal{B} is the base parameter of the function in $(1, \infty)$ and \mathcal{L} is the exponential parameter of the same function in $(0, \infty)$; together \mathcal{B} and \mathcal{L} define *unsatisfiability* value. In our framework, we set the \mathcal{S}_{sat} value to 1.5, the \mathcal{B} to 3.0, and the \mathcal{L} to 2.0.

$$\mathcal{M}_{sat}(\varepsilon_t) = \begin{cases} \arctan(\mathcal{S}_{sat} \times \delta_{sat}) & valence = 0 \\ \mathcal{B}^{\mathcal{L} \times (\delta_{sat}-1)} & valence > 0 \\ -\mathcal{B}^{-\mathcal{L} \times (\delta_{sat}+1)} & valence < 0 \end{cases} \quad (4.9)$$

Intuitively, if the human collaborator does not express any emotion, the satisfaction motive's value can vary between -1 and +1 (blue curve in Figure 4.6). However, if the agent perceives positive emotion, there will be no negative satisfaction value since the other collaborator is in positive state of mind (red curve in Figure 4.6), and in contrast, if the agent perceives negative emotion, the satisfaction motive value only changes between -1 and 0 (green curve in Figure 4.6) with respect to how satisfied the agent is according to the status of its own goals during collaboration.

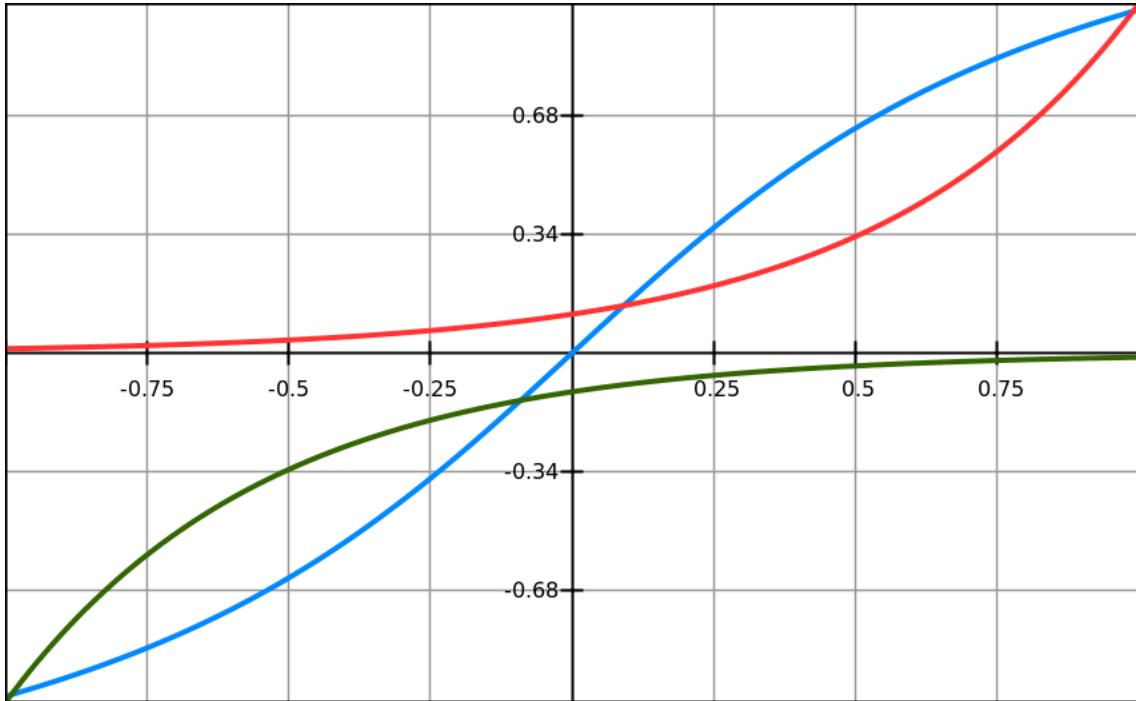


Figure 4.6: Three functions of satisfaction motive (blue: valence = 0, red: valence = positive, green: valence = negative). The x-axis indicates the satisfaction drive's delta value in [-1, +1], and the y-axis indicates the magnitude of satisfaction motive in [-1, +1].

4.6.2 Achievement Motive

The achievement motive drives the agent’s need to achieve a goal during the collaboration. According to the literature, e.g. [59], the achievement motive is based on the estimation of success probability and the difficulty of achieving a goal. In our framework, we compute the probability of success as the multiplication of the *controllability* and *expectedness* appraisal values. Intuitively, the more controllable and expected the events are, the probability of successful achievement of their related goal is higher.

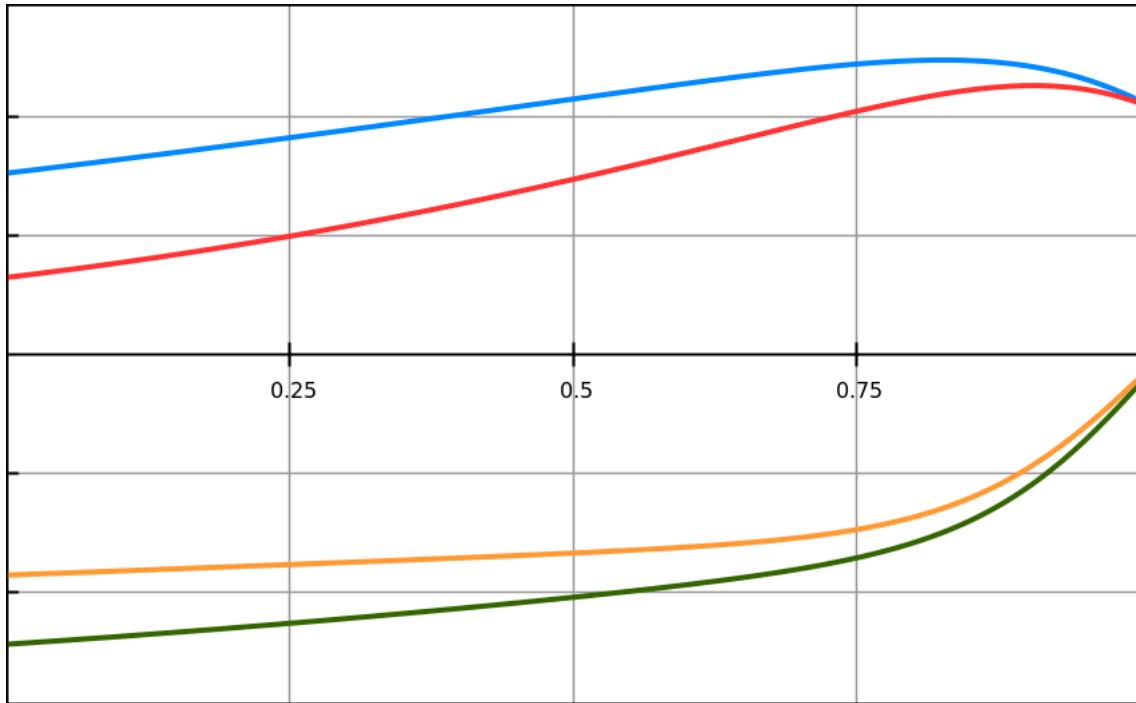


Figure 4.7: Two functions of the achievement motive (blue: valence = +1, red: valence = 0, green: valence = -1, orange: valence = close to zero from negative side). The x-axis indicates the success probability value of achieving a goal which is in $[0, +1]$, and the y-axis indicates the magnitude of achievement motive in $[-1, +1]$.

In our framework we use two sigmoid-based functions to compute the achievement motive’s value. These functions values change based on the probability of success and valence of the human collaborator’s emotion. We use Equation 4.10

when the perceived emotion of the human has positive or zero valence value, and we use Equation 4.11 when the perceived emotion of the human has a negative valence value. As shown in Figure 4.7, when the value of the valence changes between 0 and +1, the output of \mathcal{M}^{+ach} function changes between the red and the blue lines respectively. Conversely, when the value of the valence changes between -1 and a small negative number (close to zero), the output of \mathcal{M}^{-ach} function changes between the green and the orange lines.

$$\mathcal{M}^{+ach}(\varepsilon_t) = \frac{2.0}{1 + e^{(2.0-valence) \times (1.05-p(success))}} - \frac{1.0}{1 + e^{(12.0-valence) \times (1.2-p(success))}} \quad (4.10)$$

$$\mathcal{M}^{-ach}(\varepsilon_t) = \frac{1.0}{1 + e^{(0.5+valence) \times (1.05)-p(success)}} - \frac{1.0}{1 + e^{(12.0+valence) \times (p(success)-1.02)}} \quad (4.11)$$

By intuition, as the probability of success increases the agent is more motivated to achieve a goal and this motive gets higher when the human's emotion is positive or at least neutral. The human's negative emotions cause lower values of achievement motive since taking care of and acknowledging the human's negative emotion should have higher priority for a collaborative agent than achieving a goal.

4.6.3 External Motive

The external motive drives the agent's need to achieve a proposed goal by the human collaborator during the collaboration. In our framework, the external motive is also based on the estimation of success probability and the difficulty of achieving a goal, but this goal is proposed by the human collaborator. The probability of success for the external motive is computed the same way as the achievement motive's probability of success, i.e. the multiplication of *controllability* and *expectedness* appraisal values.

The only difference is that we use Equations 4.10 and 4.11 in reverse order for the external motive; i.e. we use Equation 4.11 when the valence of human's perceived emotion is positive, and Equation 4.10 when the valence of the human's perceived emotion is negative or zero.

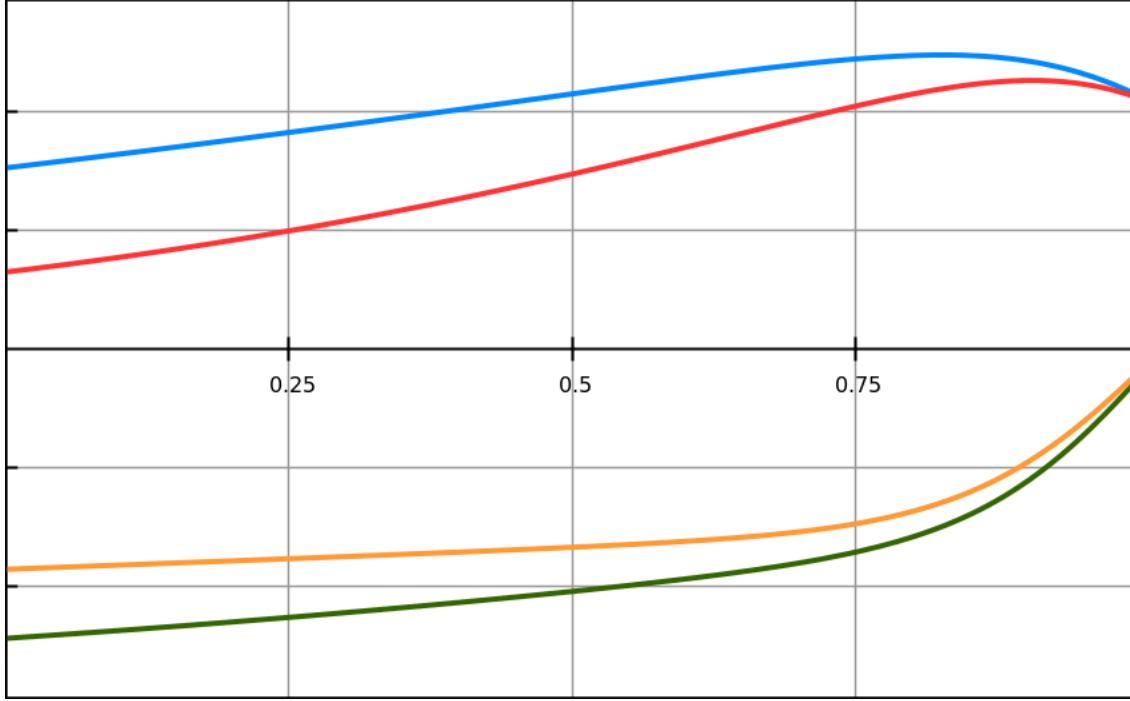


Figure 4.8: Two functions of external motive (blue: valence = -1, red: valence = 0, green: valence = +1, orange: valence = close to zero from negative side). The x-axis indicates the success probability value of achieving a proposed goal which is in $[0, +1]$, and the y-axis indicates the magnitude of the achievement motive in $[-1, +1]$.

Intuitively, when the human proposes a new goal while expressing a negative emotion the agent should be more motivated to acknowledge human's proposal and pursue the proposed goal to mitigate human's negative emotion and maintain the collaboration.

4.7 Theory of Mind

The Theory of Mind mechanism uses the same collaboration structure and functions as well as appraisal processes to form anticipated beliefs about the human's mental and emotional states. The agent uses the collaboration structure during the human's turn to compute appraisal values with respect to the human's current emotional state and the current goal in the shared goal structure. The outcome of the reverse appraisal forms beliefs about the anticipated mental and emotional state of the human collaborator.

We use the same *relevance*, *expectedness* and *controllability* algorithms for the reverse appraisal as those algorithms we described in Section 4.3. In these three algorithms the Theory of Mind mechanism substitutes the agent's required goal and its corresponding constraints and information with the human's goal and its corresponding information which is provided to the agent within the shared plan structure. In other words, since our agent knows about the human's goals (as part of the shared plan), it can apply the human goals to the same algorithms during the human's turn of the collaboration. However, only for the reverse appraisal of *desirability* we chose to simply use the valence value of the human's perceived emotion and interpret negative, neutral and positive valence values as undesirable, neutral and desirable values respectively. In this way, our agent could directly infer whether the occurrence of the current event and its corresponding goal is desirable for the human. The outcome of all of these processes is a vector of reverse appraisal values that could be used by other mechanisms in our framework.

4.8 Perception and Action

As described in Chapter 3, the Perception and Action mechanisms are not part of our theoretical work. Therefore, we only implemented these mechanisms to the extent to which they could help us to run and test our framework. The Perception

mechanism only redirects the input values from the system's users to the framework. For instance, in our conducted user study described in Chapter 5, the Perception mechanism only receives the valence of human's emotion from the input and provides it to the framework. On the other hand, the Action mechanism executes some functions based on the intentions formed and provided by the Coping mechanism described in this section. We group all of these functions into three categories in our framework. The first group of functions includes all of the functions capable of executing some actions with respect to the domain. The second category includes all of the functions involved in revealing the agent's utterances by writing on the screen or conveying through the agent's voice and text to speech systems. The last category includes all of the functions to express the agent's emotion. The emotions can be expressed through colors, emoticons, voice and text. For example, in the user study described in Chapter 5, we expressed the agent's emotions by using emoticons and utterances through the text on the screen as well as the agent's voice.

4.9 Emotion Instances

We have also implemented 10 different emotion instances that can be elicited by the agent or anticipated from the human during collaboration in our framework (see Figure 4.10). These emotion instances have meanings in social context and more specifically in collaboration. There are two components involved in selecting a particular emotion: appraisal variables and collaboration context.

We use the outcome of the four appraisal processes discussed in section 4.3 to determine the potential emotion instance to be elicited (if the agent wants to express an emotion), or to anticipate a potential emotion from the human collaborator (if the human response is anticipated). The outcome of appraisal processes can be one of the values presented in Figure 4.9 with respect to the corresponding process.

We also use the collaboration context as our second determinant to select a particular emotion. We define the collaboration context based on: *goal achievement*,

Appraisal Variable	Relevance	Desirability	Expectedness	Controllability
Values	RELEVANT	HIGH_DESIRABLE	MOST_EXPECTED	HIGH_CONTROLLABLE
		DESIRABLE	EXPECTED	LOW_CONTROLLABLE
	IRRELEVANT	NEUTRAL	UNEXPECTED	
		UNDESIRABLE	MOST_UNEXPECTED	UNCONTROLLABLE
		HIGH_UNDESIRABLE		

Figure 4.9: Appraisal values.

goal failure, proposal of a goal, acceptance of the proposed goal, and rejection of the proposed goal. All of these situations can occur by either of the collaborators, i.e., agent or human (see Figure 4.10). There is only one exception and it is when the desirability value is neutral the associated emotion to the event is always neutral without considering the collaboration context and the values of other appraisal variables.

As an example, if the agent finds an event *uncontrollable, unexpected, undesirable* and *relevant* as the result of the human’s proposal of a new goal to the agent (in the agent’s turn), the elicited emotion instance will be *worry* which can be expressed by the agent to indicate the agent’s concern. Similarly, the agent will anticipate *worry* for the human if the same appraisal values are computed while for instance the agent rejects the human’s proposal of the new goal (in the human’s turn).

#	Emotion Instance	Context	Relevance	Desirability	Expectedness	Controllability
1	Neutral			NEUTRAL		
2	Joy	human agent	<i>HUMAN_ACHIEVED</i> <i>AGENT_ACHIEVED</i>	<i>RELEVANT</i>	<i>DESIRABLE</i> <i>HIGH_DESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>
3	Sadness	human agent	<i>HUMAN_FAILED</i> <i>AGENT_FAILED</i>	<i>RELEVANT</i>	<i>UNDESIRABLE</i> <i>HIGH_UNDESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>
4	Gratitude	human agent	<i>AGENT_ACCEPTED</i> <i>AGENT_ACHIEVED</i> <i>HUMAN_ACCEPTED</i> <i>HUMAN_ACHIEVED</i>	<i>RELEVANT</i>	<i>DESIRABLE</i> <i>HIGH_DESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>
5	Positive Surprise	human agent	<i>AGENT_PROPOSED</i> <i>AGENT_ACCEPTED</i> <i>AGENT_ACHIEVED</i> <i>HUMAN_PROPOSED</i> <i>HUMAN_ACCEPTED</i> <i>HUMAN_ACHIEVED</i>	<i>RELEVANT</i>	<i>DESIRABLE</i> <i>HIGH_DESIRABLE</i>	<i>MOST_UNEXPECTED</i>
6	Negative Surprise	human agent	<i>AGENT_PROPOSED</i> <i>AGENT_REJECTED</i> <i>AGENT_FAILED</i> <i>HUMAN_PROPOSED</i> <i>HUMAN_REJECTED</i> <i>HUMAN_FAILED</i>	<i>RELEVANT</i>	<i>UNDESIRABLE</i> <i>HIGH_UNDESIRABLE</i>	<i>MOST_UNEXPECTED</i>
7	Anger	human agent	<i>AGENT_REJECTED</i> <i>AGENT_FAILED</i> <i>HUMAN_REJECTED</i> <i>HUMAN_FAILED</i>	<i>RELEVANT</i>	<i>HIGH_UNDESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>
8	Worry	human agent	<i>AGENT_PROPOSED</i> <i>AGENT_REJECTED</i> <i>HUMAN_FAILED</i> <i>AGENT_FAILED</i> <i>HUMAN_PROPOSED</i> <i>HUMAN_REJECTED</i> <i>HUMAN_FAILED</i> <i>AGENT_FAILED</i>	<i>RELEVANT</i>	<i>UNDESIRABLE</i> <i>HIGH_UNDESIRABLE</i>	<i>UNEXPECTED</i>
9	Frustration	human agent	<i>AGENT_PROPOSED</i> <i>AGENT_FAILED</i> <i>HUMAN_PROPOSED</i> <i>HUMAN_FAILED</i>	<i>RELEVANT</i>	<i>UNDESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>
10	Guilt	human agent	<i>HUMAN_FAILED</i> <i>AGENT_FAILED</i>	<i>RELEVANT</i>	<i>UNDESIRABLE</i> <i>HIGH_UNDESIRABLE</i>	<i>EXPECTED</i> <i>MOST_EXPECTED</i>

Figure 4.10: Conditions for selecting emotion instances

CHAPTER 5

EVALUATION

In this chapter, we provide the explanation and results of two different user studies. The first user study (see Section 5.1) was conducted online to evaluate our appraisal algorithms. The goal of this study was to validate the effectiveness of the factors involved in our appraisal algorithms. We prepared online questionnaires and asked participants to tell us what their decision would be in the simple situations provided. The participants' answers to our questionnaires were compared with the results of our algorithms for the given situations. The results are provided in Section 5.1. The second user study (see Section 5.2) was conducted in the Human-Robot Interaction lab. The goal of this user study was to provide an end-to-end system evaluation using our overall framework. We provided several pre- and post-study questionnaires as well as an open-ended questionnaire to study the humans' evaluation of a robot collaborating using our framework. The results are provided in Section 5.2.

5.1 Evaluating Appraisal Algorithms (Crowd Sourcing)

In this section, we present a crowd-sourcing user-study and the results, which we conducted to validate the components of our appraisal processes.

5.1.1 Experimental Scenario

We developed an experimental scenario in which participants were asked to carry out a sequence of hypothetical collaborative tasks between themselves and an imaginary

friend, Mary, in order to accomplish their shared goal. To minimize the background knowledge necessary for our test subjects, we used a simple domestic example of preparing a peanut butter and jelly sandwich, and a hard boiled egg sandwich for a hiking trip. The tasks did not require the participants to solve problems; rather, the tasks were part of simple daily activities that should be familiar to all participants.

5.1.2 Hypothesis and Methodology

Hypothesis

We conducted this user study to test our hypothesis that humans and our algorithms will provide similar answers to questions related to different factors used to compute four appraisal variables: relevance, desirability, expectedness, and controllability.

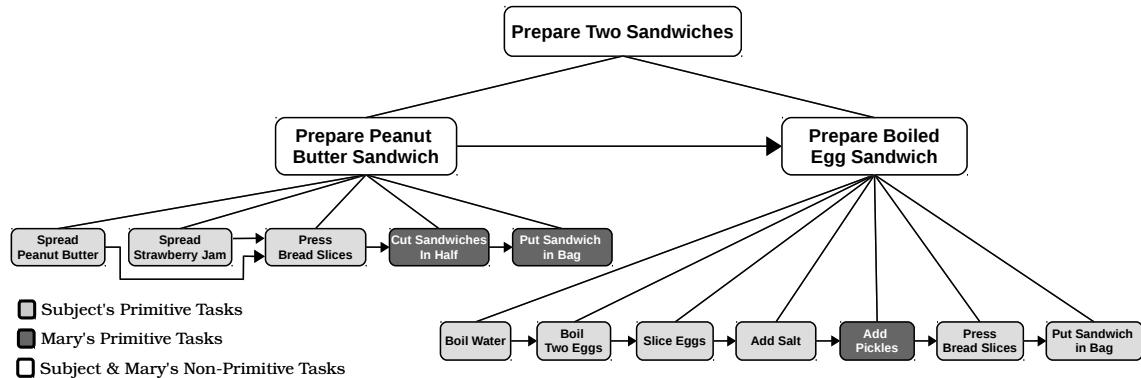


Figure 5.1: Collaboration task model for the evaluation.

Procedure

We conducted a between-subject user study using an online crowdsourcing website – CrowdFlower¹. We had a questionnaire for each appraisal variable. There were 12 questions (including 2 test questions) in the controllability and expectedness questionnaires, 14 questions (including 2 test questions) in the desirability questionnaire, and 22 questions (including 3 test questions) in the relevance questionnaire.

¹<http://www.crowdflower.com>

We provided textual and graphical instructions for all questionnaires; Figure 5.1 shows the corresponding task model. The instructions, provided in the Appendix, presented a sequence of hypothetical collaborative tasks to be carried out by the test subject and an imaginary friend, Mary, in order to accomplish their goal of preparing two sandwiches. We also provided a simple definition and an example of each appraisal variable. The collaboration structure and the instructions were the same for all questionnaires. The questions introduced specific situations related to the shared plan, which included blocked tasks and failure or achievement of a shared goal. Each question provided three answers which were counterbalanced in the questionnaire. We provided an option like C in all questions (see Figure 5.3), because we did not want to force participants to choose between two options when they did not have a good reason. There were two questions designed based on each factor that we use in our algorithms (see Section 4.3). The questions were randomly placed in the questionnaire. Figure 5.3 shows an example question from the relevance questionnaire which was designed to test whether participants perceive saliency as a factor in relevance. The input for our algorithms was the task model depicted in Figure 5.1.

Participants

Each participant group originally had 40 participants. We limited the participant pools to those with the highest confidence level on the crowdsourcing website in the United States, Britain, and Australia. Test questions were included to check the sanity of the answers. We eliminated participants providing wrong answers to our sanity questions, and participants with answering times less than 2 minutes. The final number of accepted participants in each group is provided in Table 5.1.

Table 5.1: Number of participants

appraisal variables	# of participants
Relevance	29
Desirability	35
Expectedness	33
Controllability	33

5.1.3 Results

Each question in our questionnaires was designed based on different factors that we use in our algorithms (see Section 4.3). For each of the four questionnaires we provide an example question, and describe how each question relates to a specific factor within the corresponding algorithm. The input for our algorithms was the task model depicted in Figure 5.1. The complete list of questions is provided in the Appendix. Additionally, we provide the p-value for each question, using a binomial distribution, with a probability of success of 0.33, which is the probability of selecting the right answer if the participant is simply guessing.

Expectedness

Figure 5.3 shows an example question from the expectedness questionnaire. In this example, with respect to Algorithm 3 (line 6), option A is more expected because the task related to this option provides the next available task in the focus stack (see the task model in Figure 5.1). Although the task in option B is part of the existing task model, it is considered as unexpected by our algorithm, since it is not live in the plan. We provided option C to determine whether the participants will similarly differentiate between these two options. This question was presented to the participants to determine whether their decision for the expectedness of this event is similar to the output of the expectedness algorithm. For this question, the human decision was 97% similar to the algorithm’s output.

Results for the expectedness questionnaire are presented in Figure 5.2 (Refer to Expectedness summary table). As shown in this table, there is strong evidence that

Question	Factor	Equivalent Condition	Number of Matching Answers	p-Value
1	Live goal vs. Necessary focus shift	No	31	« 0.001
2	Live goal vs. Not part of shared plan	No	32	« 0.001
3	Live goal vs. Not part of current branch	No	27	« 0.001
4	Necessary focus shift vs. Not part of shared plan	No	33	« 0.001
5	Necessary focus shift vs. Not part of current branch	No	32	« 0.001
6	Not part of shared plan vs. Not part of current branch	No	24	« 0.001
7	Live goal	Yes	14	0.093
8	Not part of current branch	Yes	14	0.093
9	Necessary focus shift	Yes	22	« 0.001
10	Not part of shared plan	Yes	29	« 0.001

Figure 5.2: Expectedness results.

the results are not random; in fact, for questions 1-6 and 9-10, human participants showed between 67 and 100 % agreement with our algorithms, with p-values of «0.001 when compared with a random population. Questions 7 and 8 were the only two questions that did not show a statistically significant p-value. It should be noted that these questions are comparing equally expected or equally unexpected situations, none of which our algorithms would consider most-expected or most-unexpected.

Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Which of the following two actions is **more expected**?

- A. Mary puts the given sandwich into a zip lock bag after cutting it in half.
- B. Mary puts some pickles on another slice of bread.
- C. Equally expected.

Figure 5.3: Example expectedness question.

Controllability

Figure 5.5 shows an example question from the controllability questionnaire. The algorithm's output is option B, and is determined by Algorithm 4 (line 3), similarly to the expectedness example above. In this example, option B is more controllable than option A, because the self over total ratio of the responsibility of the predecessors of the given task (see *Autonomy* in Section 4.3.4) is higher than the ratio in option A; i.e., self is responsible to spread peanut butter on one slice of bread and strawberry jam on another slice of bread. In this question, the humans decision was 90% in agreement with the algorithm's output.

Question	Factor	Equivalent Condition	Number of Matching Answers	p-Value
1	Agency	No	28	« 0.001
2	Autonomy (contributors)	No	17	0.009
3	Autonomy (predecessors)	No	30	« 0.001
4	Succeeded predecessors ratio	No	19	0.001
5	Available inputs	No	30	« 0.001
6	Agency	Yes	30	« 0.001
7	Autonomy (contributors)	Yes	24	« 0.001
8	Autonomy (predecessors)	Yes	18	0.003
9	Succeeded predecessors ratio	Yes	23	« 0.001
10	Available inputs	Yes	25	« 0.001

Figure 5.4: Controllability results.

Results for the controllability questionnaire are presented in Figure 5.4 (insert controllability summary table). As shown in the table, the p-value is < 0.01 for each of the ten questions. The two questions with the lowest human agreement with the algorithms both relate to autonomy of the participants with 52% and 55%.

Desirability

Figure 5.7 shows an example question from the desirability questionnaire. The output based on the Algorithm 2 (line 14) is option C, since in both option A and

Imagine you want to make a peanut butter sandwich. Which of the following two actions is **more controllable**?

- A. You can spread the peanut butter on one slice of bread and you need Mary to spread strawberry jam on the second slice of bread.
- B. You can spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread.
- C. Equally controllable.

Figure 5.5: Example controllability question.

option B, the focus goal has been achieved successfully. Therefore, in this example, both options A and B are desirable. The humans' decision was 77% in agreement with the algorithm's output in this question.

Question	Factor	Equivalent Condition	Number of Matching Answers	p-Value
1	Top level goal is failed	No	35	« 0.001
2	Top level goal is achieved	No	29	« 0.001
3	Predecessors or preconditions of the top level goal	No	35	« 0.001
4	Focus is achieved	No	34	« 0.001
5	Focus is failed	No	35	« 0.001
6	Predecessors or preconditions of the focus	No	35	« 0.001
7	Pending or in-progress focus	Yes	16	0.040
8	Top level goal is failed	Yes	23	« 0.001
9	Predecessors or preconditions of the top level goal	Yes	19	0.003
10	Focus is achieved	Yes	20	0.001
11	Focus is failed	Yes	21	« 0.001
12	Predecessors or preconditions of the focus	Yes	27	« 0.001

Figure 5.6: Desirability results.

The results of the desirability questionnaire are presented in 5.6 (insert desirability summary table). As shown in the results table, the p-value is less than 0.05 for all of the desirability questions. However, an interesting trend is that human

participants had a level of agreement of 83%-100% when the algorithm's output selected one alternate as more desirable than another alternate. When the algorithm's output chose option C (i.e. rating two situations as equally desirable), the human participants only showed 46%-77% agreement. This may indicated that a higher level of granularity is required in the algorithm when evaluating options with similar levels of desirability.

Which of the following two actions is more desirable ?
A. Imagine you pressed two slices of bread together with peanut butter and strawberry jam on them, and passed them to Mary. Mary cuts the peanut butter sandwich in half and puts them in the zip lock bag.
B. Imagine you want to make the egg sandwich. You have sliced the eggs, put them on one slice of bread, salted them, and waiting for Mary to put some pickles on your eggs. Mary puts some pickles on your eggs.
C. Equally desirable.

Figure 5.7: Example desirability question.

Relevance

In the example shown in Figure 5.9, with respect to Algorithm 1, option A is relevant because of Mary's perceived negative emotion (see Equation 4.1). Although option B is relevant (since it achieves the next goal in the shared plan), 83% of participants consider it as less relevant than option A; we believe this is due to the effect of Mary's perceived negative emotion which also generates a higher utility value in our relevance algorithm. Another question also tested belief saliency. However, the options provided only related to the shared plan (i.e., no human emotions in the options). In this case 87% of participants chose the option that accomplished the next goal in the shared plan. Interestingly, when confronted with a negative emotion from their collaborator, human participants deviated from the shared plan and found their collaborator's emotion more relevant than the original plan. It is noteworthy that in both the absence and the presence of emotions the participants

chose the more salient option with respect to our definition of saliency, which was not referenced or provided in the questionnaire.

Question	Factor	Equivalent Condition	Number of Matching Answers	p-Value
1	Belief Saliency	No	25	« 0.001
2	Belief Strength	No	13	0.063
3	Belief Recency	No	28	« 0.001
4	Motive Insistence	No	25	« 0.001
5	Motive Urgency	No	19	« 0.001
6	Motive Intensity	No	21	« 0.001
7	Goal Proximity	No	20	« 0.001
8	Goal Specificity	No	23	« 0.001
9	Belief Saliency	Yes	26	« 0.001
10	Belief Strength	Yes	22	« 0.001
11	Belief Recency	Yes	21	« 0.001
12	Motive Insistence	No	26	« 0.001
13	Motive Urgency	Yes	29	« 0.001
14	Motive Intensity	Yes	29	« 0.001
15	Goal Proximity	Yes	24	« 0.001
16	Goal Specificity	Yes	26	« 0.001
17	Belief Saliency	No	17	« 0.001
18	Motive Insistence	No	3	0.995
19	Goal Proximity	No	4	0.982

Figure 5.8: Relevance results.

The complete summary of results for the relevance questionnaire is provided in 5.8 (Insert summary table for relevance). As shown in the table, all questions show 59%-100% agreement with our algorithms and statistically significant p-values except for questions 2, 18 and 19. Question 2 addresses belief strength. Questions 18 and 19 address motive insistence and goal proximity, respectively; both of these questions present situations in which participants must choose whether an intense emotional circumstance, or adherence to the collaboration plan is more relevant (refer to the questionnaire provided in the Appendix). Our algorithms choose that the strong emotional circumstance will be more relevant; however, human participants

generally selected adherence to the collaboration plan to be more relevant.

Imagine you have made the peanut butter sandwich and passed it to Mary to cut it in half. Which of the following two actions is **more relevant**?

- A. Mary starts crying since she cut her finger with a knife.
- B. You begin to boil the water to boil the eggs for your second sandwich.
- C. Equally relevant.

Figure 5.9: Example relevance question.

5.1.4 Discussion

As shown in the results tables in sections 5.1.3 through 5.1.3, the human participants agreed 100% on some questions, while on some other questions there was a much lower level of agreement. Our results indicate that people largely performed as our hypothesis predicted. The p -values obtained based on a binomial distribution show the probability of human participants' answers being generated from a random set. The very small p -values indicate that the data set is not random; in fact, the high percentage of similarity confirms our hypothesis and shows that the algorithms can help us to model appraisal in a collaboration. The very low level of agreement on a handful of questions may indicate algorithm components that require further refinement before implementation.

5.2 End-to-End System Evaluation

As mentioned earlier, collaborative robots need to take into account humans' internal states while making decisions during collaboration. Humans express emotions to reveal their internal states in social contexts including collaboration [9]. Due to the existence of such expressions robots' emotional-awareness can improve the quality of collaboration in terms of humans' perception of performance and preferences. Hence, collaborative robots need to include affect-driven mechanisms in

their decision-making processes to be able to interpret and generate appropriate responses and behaviors. Our aim in this setup was to study the importance of emotional awareness and the underlying affect-driven processes in human-robot collaboration. We examined how emotional-awareness impacts different aspects of humans' preferences by comparing the results from our participants collaborating with an emotion-aware and an emotion-ignorant robot.

5.2.1 Implementation

The implementation of this user-study included three separate parts. The first part incorporated the Affective Motivational Collaboration Framework consisting of all Mental Processes (see left-side of Figure 5.10) as we described in Chapter ???. The second part was implemented to receive action commands from the framework and forward them to the robot to control joints and actuators (see right-side of Figure 5.10). A wizard was the third part of this setting. The wizard did nothing but inform the robot/framework whether the current task performed by either the robot or the participant was achieved successfully. The wizard was completely invisible to the participants, and the wizard had no impact on the robot's decision other than providing input regarding tasks' failure or success.

Framework

The framework includes all of the mechanisms depicted as mental processes in Figure 5.10 along with the mental states. The mental states shown in Figure 5.10 comprise the knowledge base required for all of the mechanisms in the overall model. The details about these mental processes and mental states are described in Chapters 3 and ???. In this user-study, the Collaboration mechanism uses a hierarchy of goals associated with tasks in the hierarchical task network structure depicted in Figure 5.11.

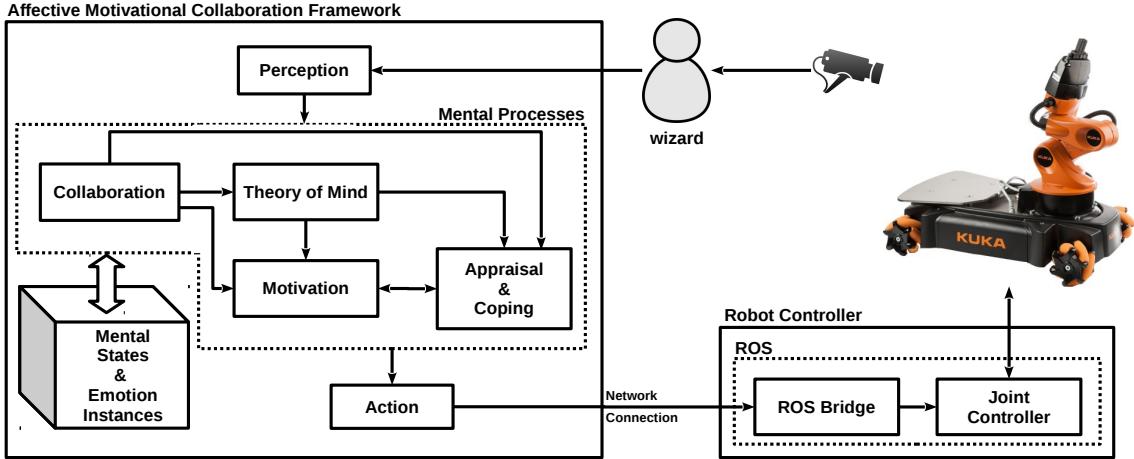


Figure 5.10: Computational framework based on Affective Motivational Collaboration theory (arrows indicate primary influences between mechanisms and data flow).

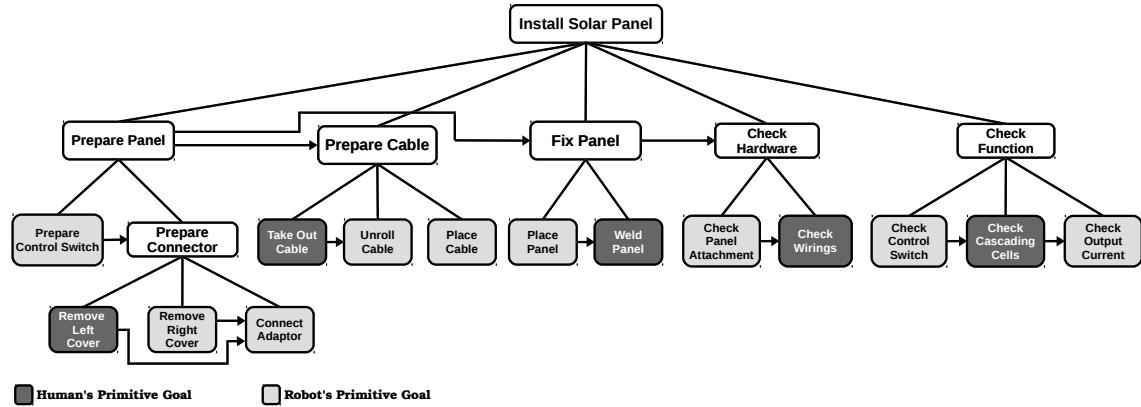


Figure 5.11: Collaboration structure used as the task model.

Robot Controller

The robot controller is comprised of two major components: 1) ROS-bridge and 2) joint controller (see Figure 5.10). ROS-bridge¹ provides an API to ROS functionality for non-ROS programs which enables us to send action commands from our framework (implemented in JAVA) to the robot's joint controller. The joint controller receives action commands and translates them into actual joint and actuator commands and sends them to the robot.

¹http://wiki.ros.org/rosbridge_suite

5.2.2 Experimental Scenario

Our scenario was based on a table top turn-taking game that we designed to simulate the installation of a solar panel. Participants collaborated one-on-one with our robot to complete all the given tasks required to install the solar panel. All of the tasks consisted of picking up and placing collaborators' available pegs on predefined spots on the board (see Figure 5.12). Each pick-and-place was associated with the robot's or the participant's task. The robot and the participants had their own unique primitive tasks that they had to accomplish in their own turns. The final goal of installing a solar panel required the robot and the participants to accomplish their own individual tasks. Failure of any task could create an impasse during the collaboration.

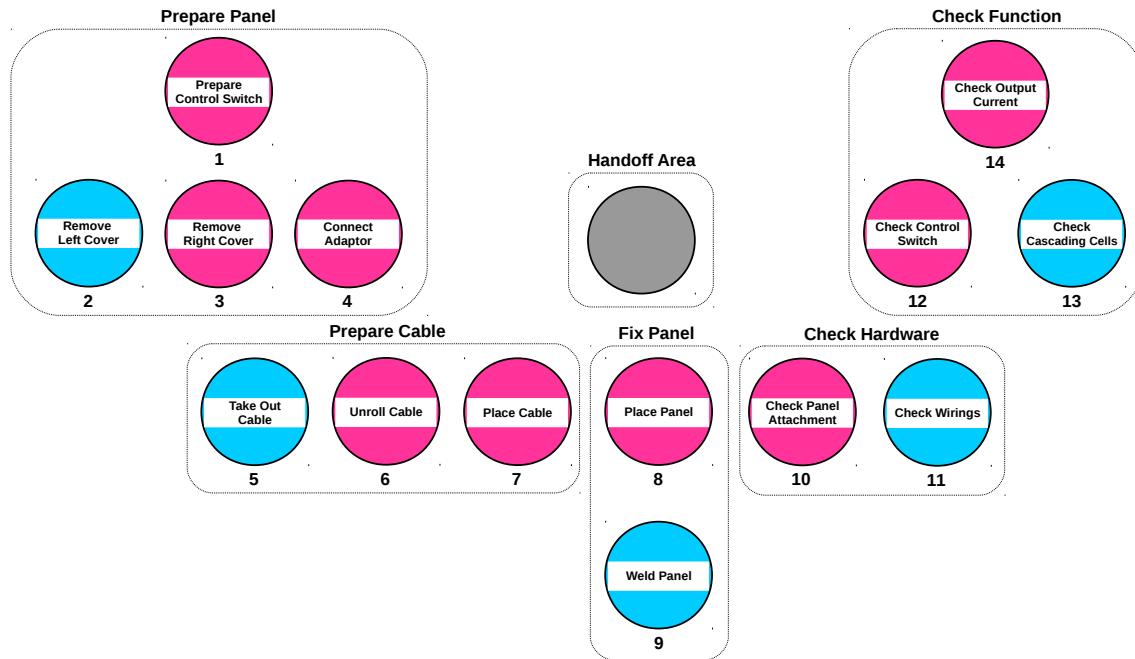


Figure 5.12: The layout of the available spots for the human and the robot to place their pegs during the collaboration.

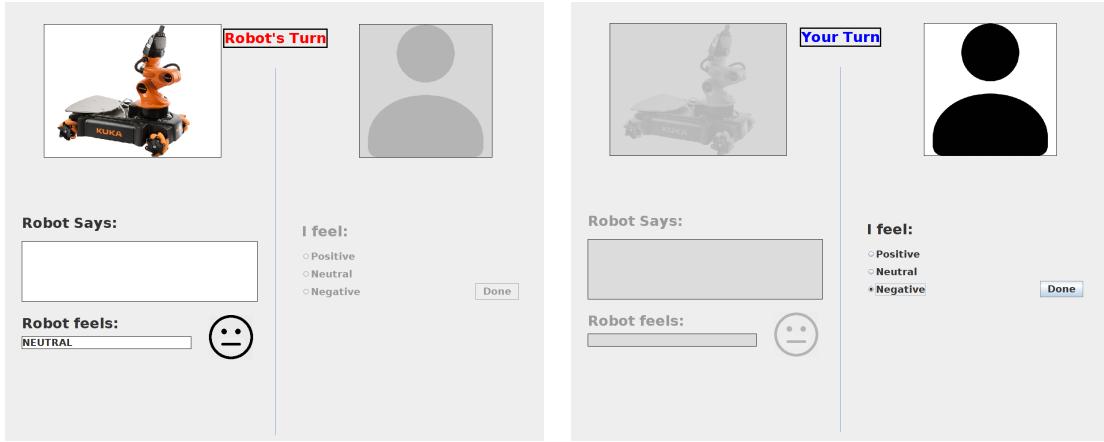


Figure 5.13: The Graphical User Interface (GUI) used during interaction.

The Robot

We conducted our experiment based on a KUKA Youbot (see Figure 5.14). The robot was stationary on top of a desk and was able to pick up and place available pegs corresponding to the robot’s task. The robot was operated based on Robot Operating System (ROS – indigo) and was receiving commands through the ROS-bridge from our Affective Motivational Collaboration framework (see Figure 5.10). We provided a simple GUI using a touch-screen monitor (see Figure 5.13 and) to a) express the robot’s positive, negative or neutral emotion through an emoticon, b) display the robot’s utterances, c) control turn-taking process of the collaboration, and d) let the participants express (report) their positive, negative or neutral emotion for each turn. The robot used MaryTTS an open-source, multilingual Text-to-Speech Synthesis platform to provide corresponding speech for its utterances in English.

Interaction Paradigms

At the beginning of each collaboration the robot asked each participant to achieve the overall shared goal, i.e., “installing the solar panel”. Then, before working towards a new goal, the robot informed the participant about the higher level non-

primitive goal (e.g. Prepare Panel – see Figure 5.11) of which the primitives were going to be working towards. The same procedure was used by the robot if there was a decision to switch to another nonprimitive due to the failure of a task in achieving the current goal. After achieving a new primitive goal, the robot either informed the human that it would pursue the next goal, or it informed and passed the turn to the human to execute the next task with respect to the human’s goal. In case of the human’s turn, the robot waited for the human to do a task, then the wizard let the robot know whether the human’s goal was achieved or not. Afterwards the robot made a decision about which goal to pursue and informed the human accordingly. The same procedure was applied to both conditions.

The robot interacted via a) speech, b) the corresponding utterance on the screen, c) negative, positive and neutral expression of emotion through an emoticon on the screen. There were two conditions of the robot: 1) emotion-aware and 2) emotion ignorant. The robot used neutral expression in the case of emotion-ignorance. The interaction was controlled autonomously by the framework we discussed in Section 5.2.1 in both the emotion-ignorant and the emotion-aware cases. The reasoning about which task should be done and controlling the robot was entirely autonomous. Only the perception of the task failure or achievement by the robot or by the participant was done by a wizard monitoring the collaboration outside of the test area. The interaction was structured based on the exact same goals in the same HTN for both conditions. The robot was using the same utterances in both conditions. In the emotion-aware condition the robot used a different behavior in comparison with the emotion-ignorant condition only if the participant was expressing a negative emotion in the event of a failure; i.e., the robot’s utterances were identical in emotion-ignorant and emotion-aware cases if in the latter the participant reported (expressed) a positive or a neutral emotion.

Three different behaviors could be generated only in the emotion-aware condition. These three behaviors were 1) mitigating the human’s negative emotion and postponing its own task to help the human, 2) goal-management to switch to another

goal which has lower cost with respect to the human’s negative emotion, and 3) task delegation to the human to overcome the impasse. In each run, the human had two pre-coordinated task failures, and the robot had one. If the human expressed negative emotion after the first human task-failure, the robot responded by mitigating the human’s negative emotion by saying “It was not your fault. I can help you with this task” and helping the human by providing a peg to fulfill the human’s task. If the human expressed negative emotion after the second human task-failure, the robot informed the human that they could proceed with another task to save time while simultaneously requesting a new peg (i.e. help) from the supervisor. If the human expressed negative emotion as a result of the robot’s task failure, the robot requested help from the human (who had the correct peg). In the event that the human expressed positive or neutral emotion during these three failures, the robot behaved identically in the emotion-ignorant and the emotion-aware cases, by asking the supervisor for help.

Environment and Tasks

The environment was set up in the Human-Robot Interaction lab and included the robot, the collaboration board on top of a desk, and the participant standing in front of the robot on the other side of the board (see Figure 5.14). One of the experimenters monitored the interactions using a live stream of a camera in a different room. The experimenter provided only the required perception, i.e., decision on success or failure of the tasks for the robot, through the entire time of the collaboration (see Section 5.2.2).

The tasks were defined based on the HTN structure shown in Figure 5.11 and were executed in a turn-taking fashion by either of the collaborators. For each task either the robot or the participant was responsible to pick up one of the corresponding pegs from their own inventory and place it on the right spot which was colored and tagged the same as the associated peg. Some pegs and corresponding spots on the board had hidden magnets which prevented the pegs from standing upright.

Any peg that fell over was considered a failed task.

5.2.3 Hypotheses and Methodology

Hypothesis

The non/social functions of emotions impact a collaboration process. Human collaborators prefer to collaborate with others whose behaviors are influenced by these functions of emotions depending on the context. We developed seven hypotheses on positive influence of emotion-awareness and usefulness of emotion function during collaboration:

Hypothesis 1. Participants will feel closer to the emotion-aware robot rather than the emotion-ignorant robot.

Hypothesis 2. Participants will find the emotion-aware robot to be more trustworthy than the emotion-ignorant robot.

Hypothesis 3. Participants will find the emotion-aware robot to have better performance in collaboration than the emotion-ignorant robot.

Hypothesis 4. Participants will find the emotion-aware robot to be more understanding of their feelings than the emotion-ignorant robot.

Hypothesis 5. Participants will find the emotion-aware robot to be more understanding of their goals than the emotion-ignorant robot.

Hypothesis 6. Participants will feel more satisfied about the collaboration when working with the emotion-aware robot rather than emotion-ignorant robot.

Hypothesis 7. Participants will perceive higher level of mutual satisfaction with the emotion-aware robot than emotion-ignorant robot.

Procedure

Participants were first given a brief description of the purpose of the experiment. After the short introduction, they were asked to review and sign a consent form. Participants were then provided with a written instruction of their task and the rules

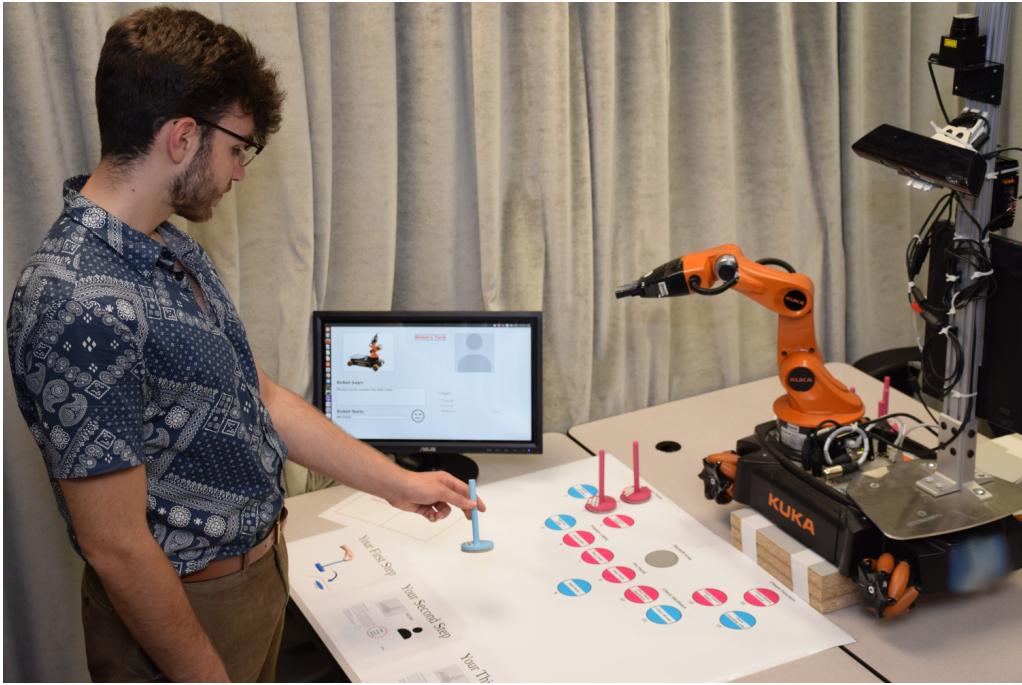


Figure 5.14: Experimental setup.

for collaborating with the robot. Then, one of the experimenters lead them into the experiment room and asked the participants to answer pre-experiment questionnaires. Afterwards, the experimenter went through all the details of the instructions with the participants standing in front of the collaboration board and the robot. The experimenter confirmed participants' correct understanding of the tasks and informed them of types of task failures that might occur during the collaboration. Participants were told that researchers were developing a collaborative robot and would like their help in evaluating their design. Participants were provided with identical instructions and randomly assigned to complete either the emotion-aware or the emotion ignorant condition first. They were told that, after their collaboration with the robot, they would be asked to answer a questionnaire on their experience. After completing the first round of collaboration, participants answered a post-experiment questionnaire that measured their perceptions of the robot, the task, and the collaboration procedure. After answering the first post-experiment questionnaire, participants were told that they were going to collaborate with the

robot one more time and the robot might not necessarily have the same collaborative behavior. After completing the second round of collaboration, participants were asked to answer the second post-experiment questionnaire which consisted of the same questions as the first post-experiment questionnaire. After all, participants were asked to answer an open-ended questionnaire which measured their perception of difference between two runs, their preference of collaborative robot between two runs, and their reasons of preference.

Measurements

In our study two basic conditions of the robot were tested: a) the emotion-ignorant condition, b) the emotion-aware condition. We measured participants' recall of the collaborative behaviors presented by the robot using an open-ended post-experiment questionnaire. We also specifically asked the participants what behavior of the robot they liked during their collaboration. We also evaluated participants' levels of satisfaction, trust, goal achievement, mutual understanding of goals, mutual understanding of feelings, mutual agreement, and also participants' beliefs about the efficiency of collaboration and their feeling of robot's collaborative behaviors. Seven-point Likert scales were used in these questionnaire items.

Participants

A total of 37 participants participated in the experiment in 74 trials. Participants were recruited from Worcester Polytechnic Institute's students and staffs as well as other civilians recruited from outside of the campus. The ages of the participants varied between 19 and 74 with an average of 34.2 years before our screening of 4 participants based on our sanity check questions. After this screening the ages of the participants varied between 19 and 54 with an average of 30.8 years old. Of the 33 participants, 21 were female and 12 were male. Each participant participated in 2 trials. In one trial the robot was aware of human's emotion and in the second

trial the robot was ignoring human's emotion. The order of these two trials were randomly assigned to each participant. In general we used emotion-ignorant robot first in 16 experiments, and emotion-aware robot first in 17 experiments.

5.2.4 Results

As discussed in Section 5.2.3, results of the user study were gathered through a 31-question Likert-scale survey that was given to each participant after each run with the robot, and through a 5-question open-ended summary questionnaire at the end of the experiment.

7-Point Likert Scale Survey Results

As mentioned previously, the 7-point Likert scale survey was administered at the end of the emotion-ignorant run and at the end of the emotion-aware run for each participant. The 31 questions are generally categorized to evaluate the humans' perceptions of the following seven categories, with 3-7 questions per group: (1) the likability of the robot (2) the level of trust the human feels in the robot (3) the human's perception of the robot's performance (4) the human's perception of the robot's understanding of human's emotions (5) the human's perception of the robot's understanding of human's and collaboration's goals and objectives (6) the human's feeling about the collaboration and (7) the human's perception of the human's and robot's mutual satisfaction with each other as collaborative partners. The questions presented are provided in Figure 5.15.

Question Category	Question	Question Number
Likability	I felt close to the robot.	Q1
	I would like to continue working with the robot.	Q2
	I like the robot.	Q3
	The robot was interesting.	Q4
Trust	I trust the robot.	Q5
	It was easy to express myself to the robot.	Q6
	I trust the robot to perform appropriately in our collaboration.	Q7
	I am confident in the robot's ability to help me.	Q8
	I trust the robot to assess my feelings appropriately in our collaboration.	Q9
Robot's Performance	The robot was repetitive.	Q10
	The robot made efficient decisions.	Q11
	The robot's decisions improved my performance during the collaboration.	Q12
Robot's Understanding of Human's Emotions	The robot understood my emotions.	Q13
	The robot is sometimes confused about what I feel about our activities.	Q14
	I feel that the robot, in its own unique ways, is genuinely concerned about me.	Q15
	The robot understands some of my feelings and takes them into account in our collaboration.	Q16
	The robot does not understand how I feel during our collaboration.	Q17
Robot's Understanding of Goals	The robot does not understand what we are trying to accomplish.	Q18
	The robot does not understand what I am trying to accomplish.	Q19
	The robot perceives accurately what my objectives are.	Q20
	The robot was committed to the collaboration.	Q21
Human Feeling about Collaboration	I find what the robot and I are doing is unrelated to my goals.	Q22
	I find what I am doing with the robot confusing.	Q23
	The robot and I are working towards mutually agreed-upon goals.	Q24
	The robot and I collaborate on setting goals for us to work on.	Q25
	The robot and I agree on what is important for us to work on.	Q26
	I believe that the robot and I achieved the goals we set.	Q27
	I am satisfied with the outcome of our collaboration.	Q28
Satisfaction of Collaborative Partner	The robot was satisfied with my collaborative behavior.	Q29
	I was satisfied with the robot.	Q30
	I understand the robot, and I think it understands me, at least in the best way it can.	Q31

Figure 5.15: The 31 Likert scale questions organized according to their groups.

The results were analyzed using a two-tailed paired t-test to analyze the difference of means between the emotion ignorant and the emotion-aware condition. Refer to Figures 5.18 - 5.22 for the results. As mentioned in Section 5.2.3, participants were randomly assigned to complete either the emotion-ignorant or the emotion-aware run first; analysis of the results revealed no statistically significant difference or consistent pattern based on which run the participant completed first.

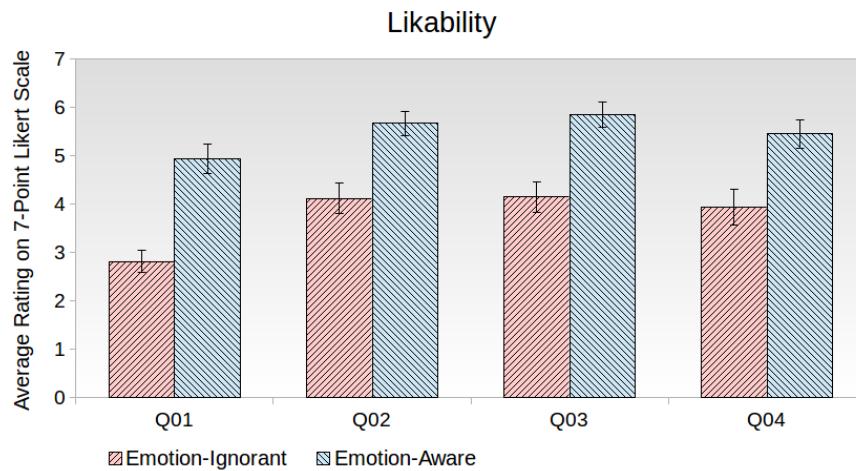


Figure 5.16: Results of the Likert scale survey for Likability questions. The p-value for the difference between means is $\ll 0.001$ for all questions.

Likability of the Robot

Questions 1 through 4 addressed the likability of the robot. As shown in Figure 5.16, participants rated the emotion-aware robot 1.5-2.1 points higher than the emotion-ignorant robot. These results indicate that participants felt closer with and preferred working with the emotion-aware robot; these results support Hypothesis 1, which stated that humans would prefer to work with the emotion-aware robot over the emotion-ignorant robot.

Human Trust in the Robot

Questions 5-9 were designed to measure the degree of trust that the human participants felt in the robot. As shown in Figure 5.17, participants trusted the emotion-

aware robot, on average, a minimum of 1.4 points more than the emotion-ignorant robot, both in general and in terms of collaboration performance. In Question 5, participants rated a general statement of trust 1.5 points higher in the emotion-aware case. Additionally, in Question 7, participants rated their trust in the emotion-aware robot to perform appropriately during collaboration an average of 5.9 on a 7-point Likert scale, where 7.0 would indicate maximum trust; this indicates an acceptable level of trust in the robot’s collaborative abilities. These results support Hypothesis 2, that posits that human participants would find the emotion-aware robot to be more trustworthy than the emotion-ignorant robot.

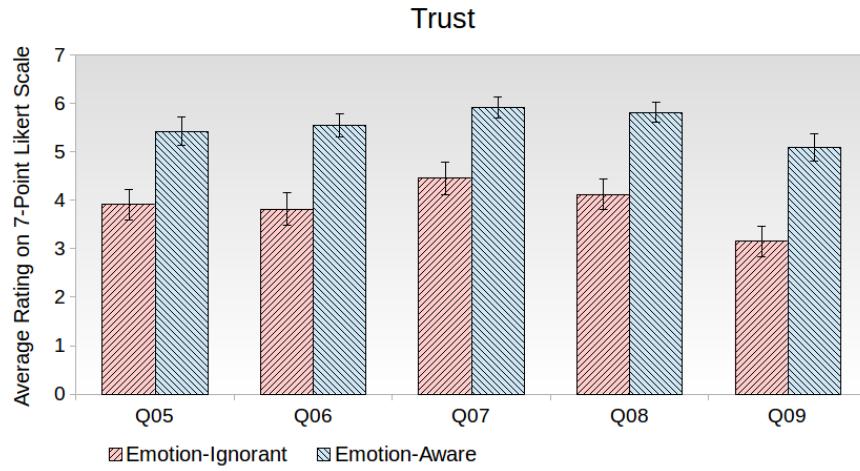


Figure 5.17: Results of the Likert scale survey for questions related to trust. The p-value for the difference between means is $\ll 0.001$ for all questions.

Perception of the Robot’s Performance

Question 10 (which is reverse-scored) measures the participant’s perception of repetitiveness in the robot during the collaboration. In both conditions, participants rated the robot as moderately repetitive, with the emotion-ignorant robot’s average response being about 1.1 points higher than the emotion-aware. This result correlates with several of the open-ended responses which described the emotion-aware robot’s behaviors as “cute” and “interesting”, refer to Section 5.2.4. Question 11, which asks

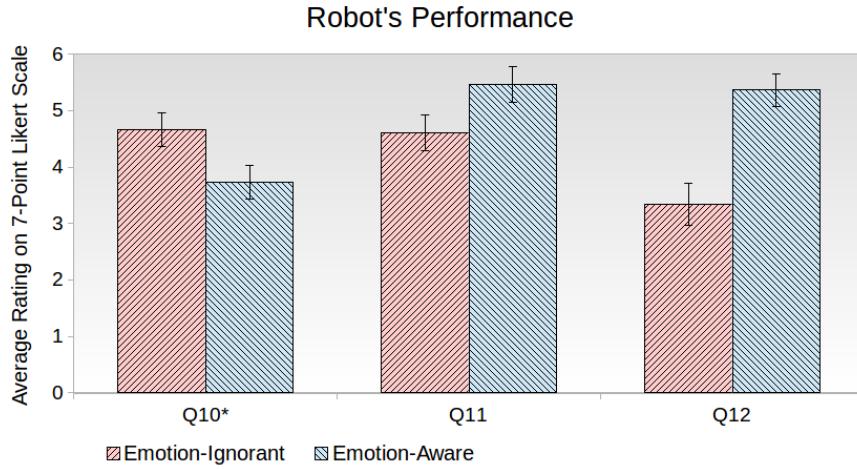


Figure 5.18: Results of the Likert scale survey for questions related to the robot’s performance. The p-value for the difference between the means for questions 10, 11 and 12 are 0.001, 0.063 and $\ll 0.001$, respectively.

about the efficiency of the robot’s decisions is the only question of the 31 questions that did not have a statistically significant difference between the emotion-aware and the emotion-ignorant case. This correlates with the result of the open-ended question asking which condition of the robot exhibited behaviors that could prevent human error (refer to 5.2.4); in response to this question, several respondents stated that it may be quicker or simpler to call the supervisor in the event of a task failure, rather than changing the order of the tasks. According to the results from Question 12, the participants felt that the emotion-aware robot’s decisions during collaboration improved their own performance, with an average rating of 5.4, while the emotion-ignorant robot only received an average rating of 3.3, indicating that participants felt it was not able to interact in such a way as to increase the human’s performance; refer to results from Question 6. These results support Hypothesis 3, which posited that humans will perceive the emotion-aware robot as being more capable than the emotion-ignorant robot.

Robot’s Understanding of Human Emotions

In Questions 13 through 17, participants evaluate the robot’s understanding of hu-

mans' emotions. In questions 13, 15, and 16, participants rated the emotion-aware robot, on average, a minimum of 1.8 points higher than the emotion-ignorant robot. In response to questions 14 and 17, which are reverse-scored, participants ranked the emotion-ignorant robot 1.2 and 2.0 points higher, respectively, than then emotion-aware robot. The results of all five questions in this category support Hypothesis 4.

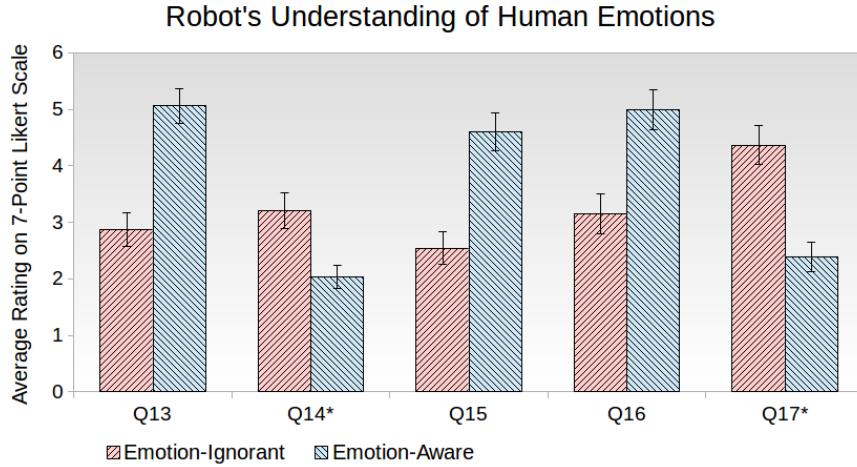


Figure 5.19: Results of the Likert scale survey for the questions related to the robot's understanding of human emotions. The p-value for the difference between the means is $\ll 0.001$ for all of the questions except Question 14, for which the p-value is 0.003.

Robot's Understanding of Human and Collaboration Goals

Questions 18 and 19 were reverse-scored questions intended to determine whether the humans felt that the robot understood the shared collaboration goal and the human's personal goal, respectively. For both conditions of the robot, the average scores were lower than 3.5, indicating that the human's perceived the robot as having some understanding of the goals. For both questions, the emotion-ignorant robot's average score was significantly higher than the emotion-aware robot's score. Similarly, Question 20 was a measure of whether the human perceived that the robot correctly perceived the human's goal. On average, participants provided an average rating for the emotion-aware robot that was 1.5 points higher than that for the emotion-ignorant robot. Question 21 measured the human perception of the robot's

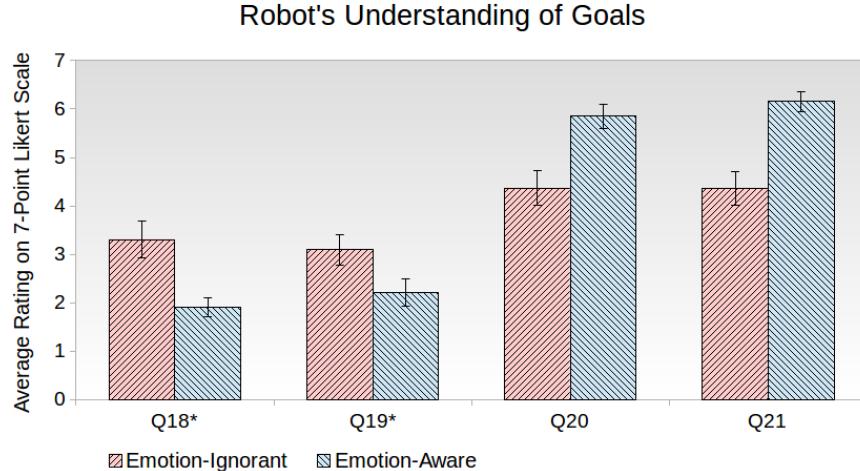


Figure 5.20: Results of the Likert scale survey for questions related to the robot’s understanding of goals. The p-value for the difference between the means for all questions is $\ll 0.001$, except Question 19, for which the p-value is 0.006.

commitment to the collaboration; for this measure, the average participant score assigned to the emotion-aware robot was 6.2 points out of a maximum of 7 points, indicating that the participants felt that the emotion-aware robot was strongly committed to the collaboration. The emotion-ignorant robot received an average rating of 4.4 points, indicating only moderate commitment. These results strongly support Hypothesis 5, which posits that humans will feel that the emotion-aware robot will better understand their goals than the emotion-ignorant robot.

Human’s Feeling about the Collaboration

Questions 22 through 28 were designed to gauge how the human participants felt about the partnership within the collaboration and the outcome of the collaboration. For each of the 7 questions, the participants ranked the emotion-aware robot as better than the emotion-ignorant robot, by a minimum, on average, of 0.8 points. Questions 24, 27 and 28 addressed whether the robot and the participant were working toward mutually agreed-upon goals and on the outcome of the collaboration; in the emotion-aware condition, participants rated the robot a minimum of 6.1 points,

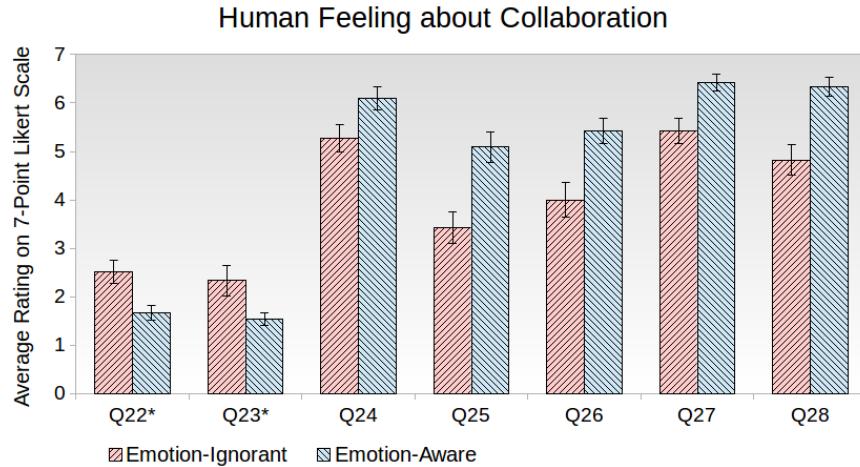


Figure 5.21: Results of the Likert scale survey for questions related to the human's feeling about the collaboration. The p-value for the difference between the means is $\ll 0.001$ for questions 22, 25, 26, and 28. The p-value for Questions 23, 24 and 27 are 0.02, 0.008 and 0.001, respectively.

on average, while rating the emotion-ignorant robot 1-1.6 points lower, indicating that the participants felt a very strong sense of collaboration with the emotion-aware robot, and only a moderate sense of collaboration with the emotion-ignorant robot. Questions 25 and 26 address whether the robot and the participant set the collaboration goals together; these two questions have lower scores than Questions 24, 27 and 28, for both the emotion-aware and the emotion-ignorant case. The lower overall scores are likely due to the fact that the robot decides the task order or action in the event of failure in both conditions; however, the higher score in the emotion-aware case may indicate that emotional awareness can increase a feeling of collaboration. These results support Hypothesis 6 that humans will feel a greater sense of mutual collaboration and understanding about the collaboration with the emotion-aware robot.

Human Perception of Mutual Satisfaction with Collaborative Partner

Questions 29, 30 and 31 were designed to measure the human's perception of the robot's satisfaction with the human, the human's satisfaction with the robot and

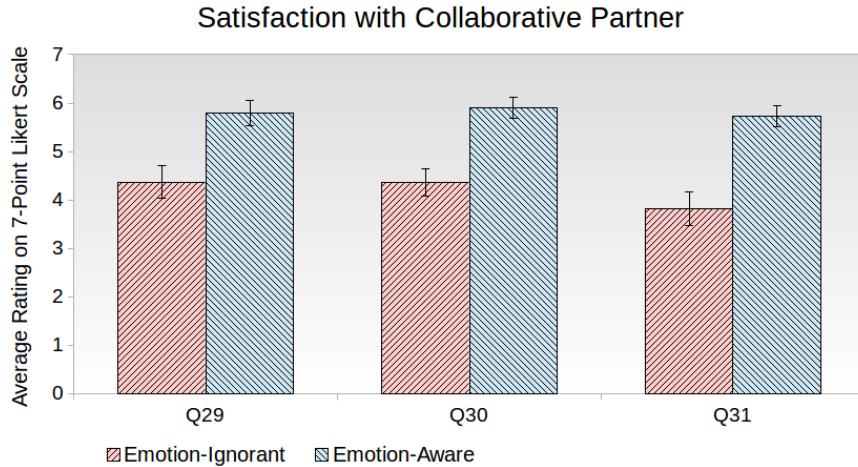


Figure 5.22: Results of the Likert scale survey for questions related to satisfaction with collaborative partner. The p-value for the difference between means is $\ll 0.001$ for all questions.

the mutual understanding between the human and the robot, respectively. The participants provided an average response in the emotion-aware condition of 5.8, 5.9 and 5.7 to Questions 29, 30 and 31, respectively, indicating a high level of mutual satisfaction; all three answers were about 1.4-1.9 points lower, on average, in the emotion-ignorant condition. These results indicate a higher level of satisfaction working with the robot in the emotion-aware condition, and strongly support Hypothesis 7, which posited that humans will feel a greater sense of mutual satisfaction with the emotion-aware robot than the emotion-ignorant robot.

Results from the Open-Ended Questionnaire

As described in Section 5.2.3, each participant answered an open-ended questionnaire at the end of the study. Figure 5.23 summarizes the questionnaire and which run users preferred for certain conditions (i.e. emotion-ignorant or emotion-aware). Note that some users chose not to state a preference regarding which run they preferred for certain conditions; because we were specifically interested in whether users preferred the emotion-aware case, we considered the ambiguous responses to be fail-

ures in the binomial analysis. The binomial analysis is based off of a population size of 33.

Question	Number of Participants Who Did Not Prefer One Run Over the Other *	Number of Participants Favoring Emotion-Aware Robot	p-value
Which of the two runs with the robot did you prefer?	0	33	0
In which of the two runs did the robot exhibit behavior that could be useful in a more complex task?	1	30	< 0.001
In which of two runs did the robot exhibit behavior that could prevent human error?	3	18	> 0.1
In which of the two runs did the robot exhibit behavior that could improve the efficiency of collaboration?	2	26	< 0.001
What was the most interesting behavior of the robot and in which run did it happen?	5	24	0.002

Figure 5.23: Open-ended questionnaire questions and results. (*Note: Because we are evaluating whether humans prefer an emotion-aware robot, these results are taken as negative test results when calculating the p-value using the binomial distribution. Only those participants who clearly indicated a preference for the emotion-aware robot are taken as positive test results.)

As shown in Figure 5.23, 100% of users unambiguously preferred the run with the emotion-aware robot. In general, this preference stemmed from a feeling of closeness and partnership, as seen in these responses: “the robot had emotions and responded to my emotions. Also, what it said about my failing was cute and aimed to make me feel better.” Another example is “I liked feeling needed and accounted for; I felt closer to the robot.” Finally, “I saw the changes in its feeling, which motivated me to care more about my act...I also liked that he asked me to correct its failure, although it could ask the supervisor.”

When asked in which of the two runs the robot exhibited behavior that could be useful in a more complex task, 90.9% chose the emotion-aware robot. In general, respondents thought that the emotion-aware robot was better at problem solving, more adaptable, and more capable of handling the social complexities that occur in collaboration, as shown in responses such as “The robot explained motives...which is important to keep a team communicating and on the same pace.” Also, “When we failed he initially switched to a new task and then came back to the originally failed task. It kept me from getting irritated and negative.” Finally, “The more

complex, the more necessary it is to understand how humans think and operate...an empathetic robot can adapt, encourage and help.” It is worth noting that one respondent preferred the emotion-ignorant case, saying “In a more complex task it might be better for the robot to take control and simply tell me what to do; trying to be understanding and collaborative wouldn’t be as important as doing the task correctly.”

The only question that did not provide statistically significant support in favor of the emotion-aware robot related to which case the robot exhibited behavior that could prevent human error. About 36.4% of respondents thought that the emotion-ignorant robot was more likely to prevent human error; however, all but one of these cited calling the supervisor as the main method of preventing human error, in spite of the fact that the instructions indicated that the robot’s need to call the supervisor counted against the collaboration. Of the 54.5% who thought that the emotion-aware robot was better at preventing human error, most cited the robot’s ability to console the human as the main behavior that could prevent human error. Respondents indicated that this enabled them to move on and feel better about the collaboration, as with this response: “The robot switched to a different task and we came back to an error later. This allowed my mind to move away from being frustrated. I was able to complete a different task which felt like a win - then come back and finish the error. Making my mind move away from frustration could definitely prevent more errors.”

When asked in which of the runs the robot exhibited behavior that could improve the efficiency of the collaboration, 78.8% responded with the emotion-aware case; of these, the vast majority stated that this was because of the robot’s ability to change the order of tasks in the event of a failure, and to ask the human for help.

Finally, when asked in which run the most interesting behavior occurred, 72.7% chose the emotion-aware condition. Of these respondents, 12 individuals stated that the robot’s attempt to console the human by saying “It was not your fault” in response to the human’s negative emotion that occurred as a consequence of the

human's failed task was the most interesting behavior, and a majority mentioned that it actually made them feel more positive. Six participants referred to the robot's ability to understand and express emotion. Several participants referred to the robot's ability to communicate, including the ability to ask questions. Of those who responded with the emotion-ignorant case, most found the ability to call the supervisor, and mechanical functions, such as gripping, to be most interesting.

Impact of Demographics

As mentioned in Section 5.2.3, we recorded certain demographic information from each participant, including age and gender. We also had each participant complete several personality questionnaires. Although it was not the primary purpose of the study, we investigated the Likert scale results to determine if there were any relevant trends based on the demographics and personalities of the participants. A close study of the results did not reveal any identifiable pattern based on gender or personality.

Age did reveal an interesting pattern. We divided the participants into two groups, below 30 years of age and 30 or above. While question-by-question comparisons revealed only a few statistically significant differences based on age, a general pattern emerged. For all but four of the 31 questions presented, the younger age group reported higher scores than the older age group (or lower, in the case of reverse-scored questions) for the emotion-aware robot. In the emotion-ignorant case, the younger group tends to score the robot nearer to the same value as the older age group for all but seven questions, leading to a pattern in which the score drop between the emotion-aware and the emotion-ignorant case was more for the younger group than for the older group; the seven questions that broke this pattern were 7, 9, 11, 12, 18, 19 and 22.

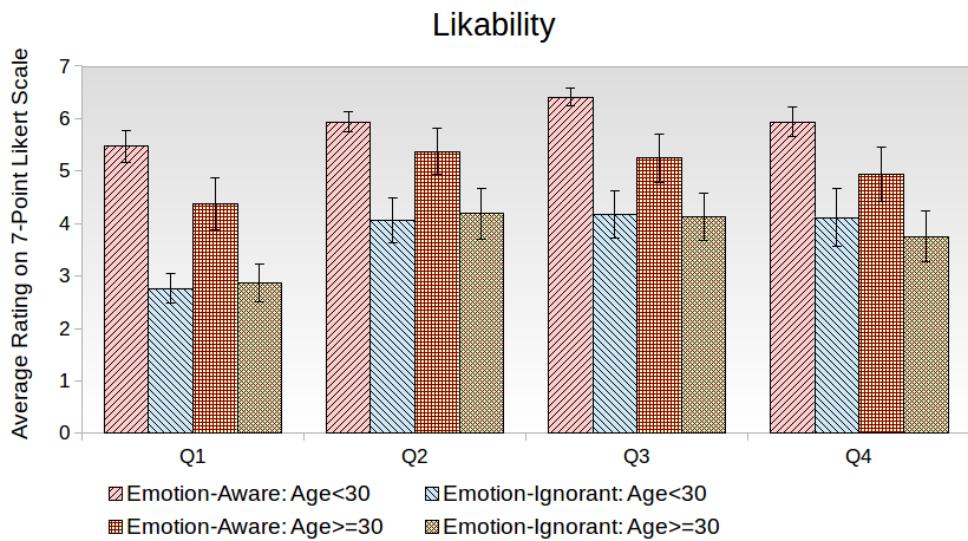


Figure 5.24: Impact of age on results of Likert scale questions related to likability.

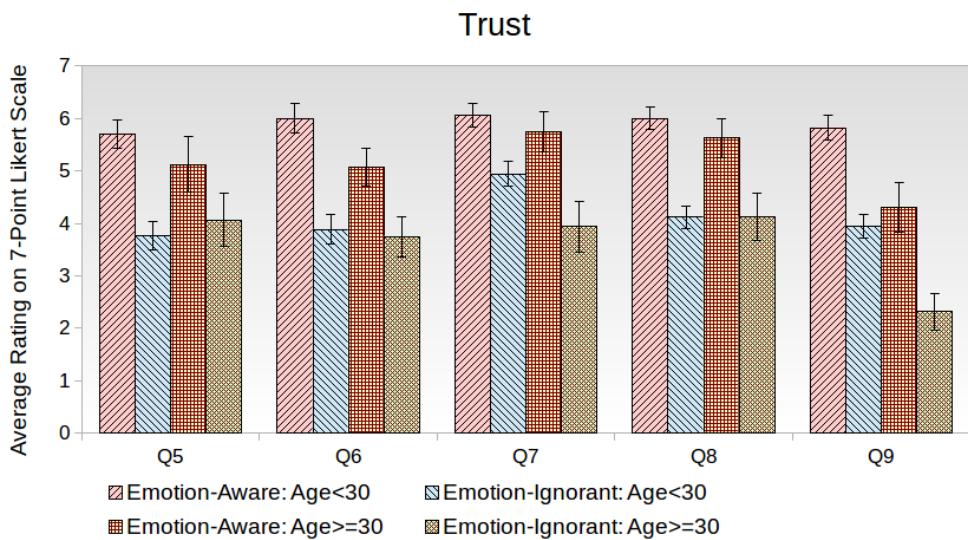


Figure 5.25: Impact of age on results of Likert scale questions related to trust.

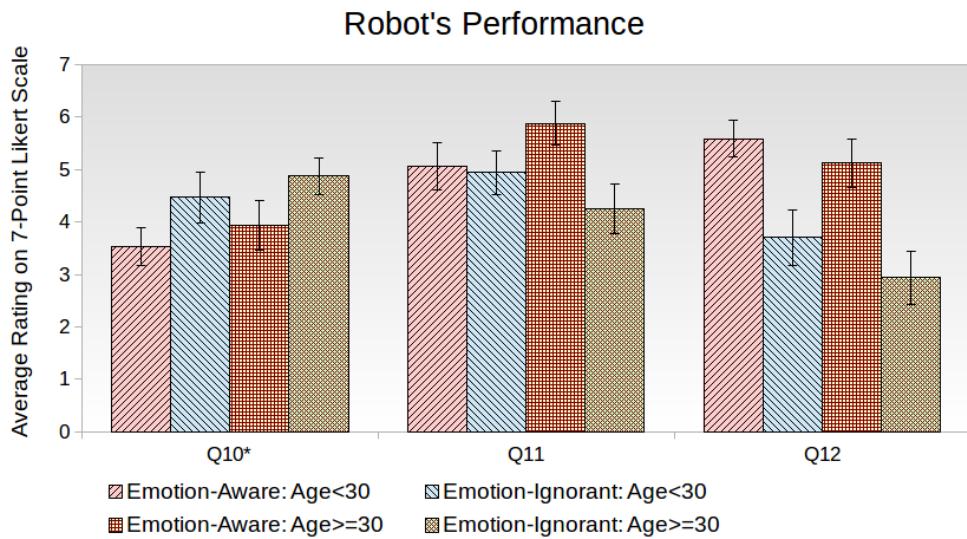


Figure 5.26: Impact of age on results of Likert scale questions related to performance.

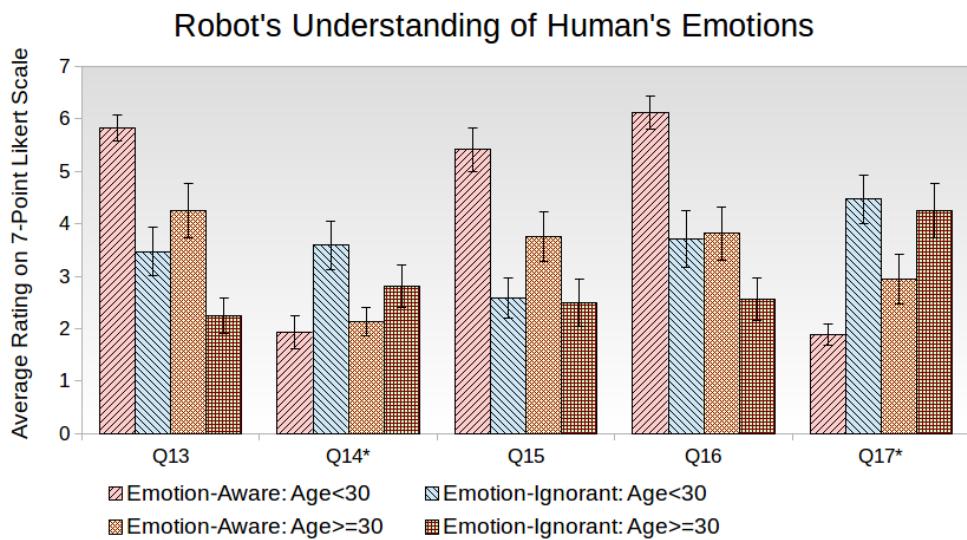


Figure 5.27: Impact of age on results of Likert scale questions related to robot's understanding of human's emotions.

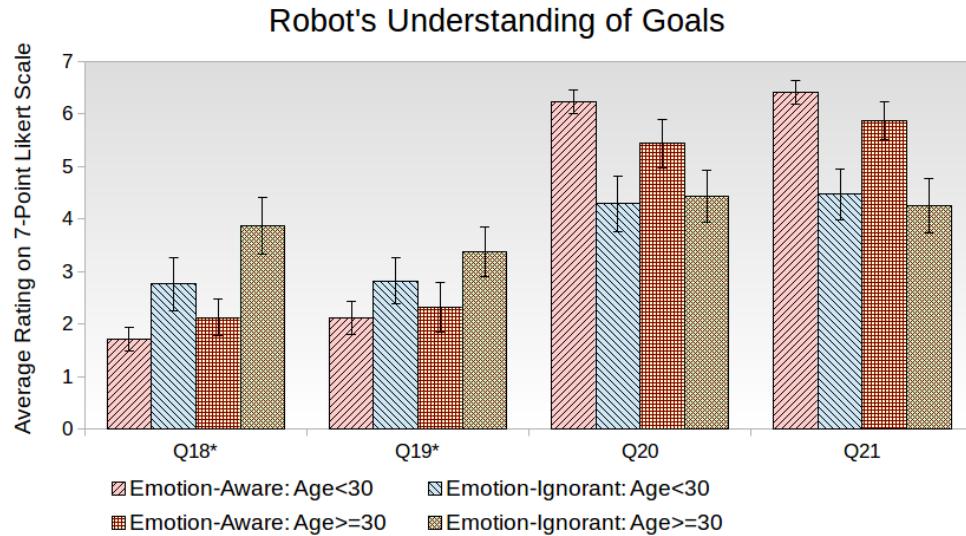


Figure 5.28: Impact of age on results of Likert scale questions related to robot's understanding of goals.

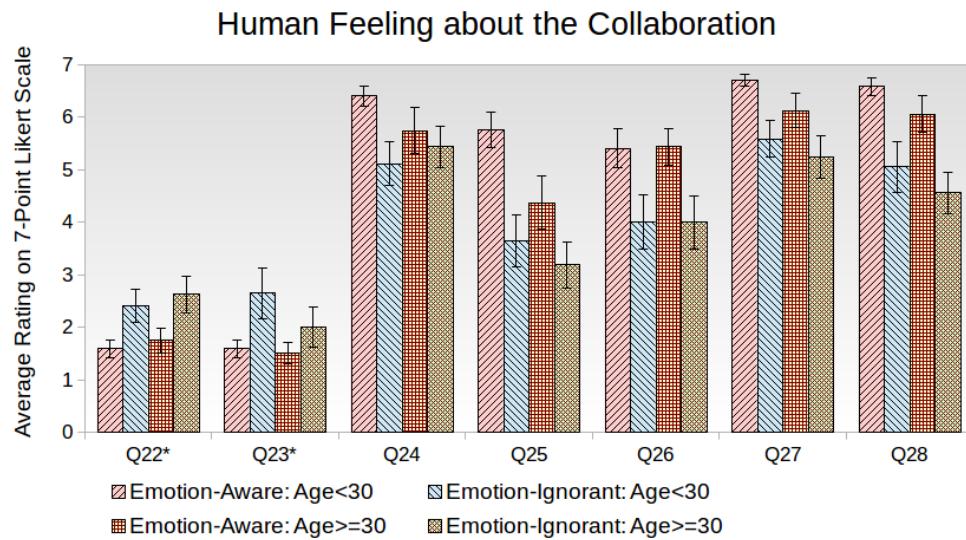


Figure 5.29: Impact of age on results of Likert scale questions related to human's feeling about collaboration.

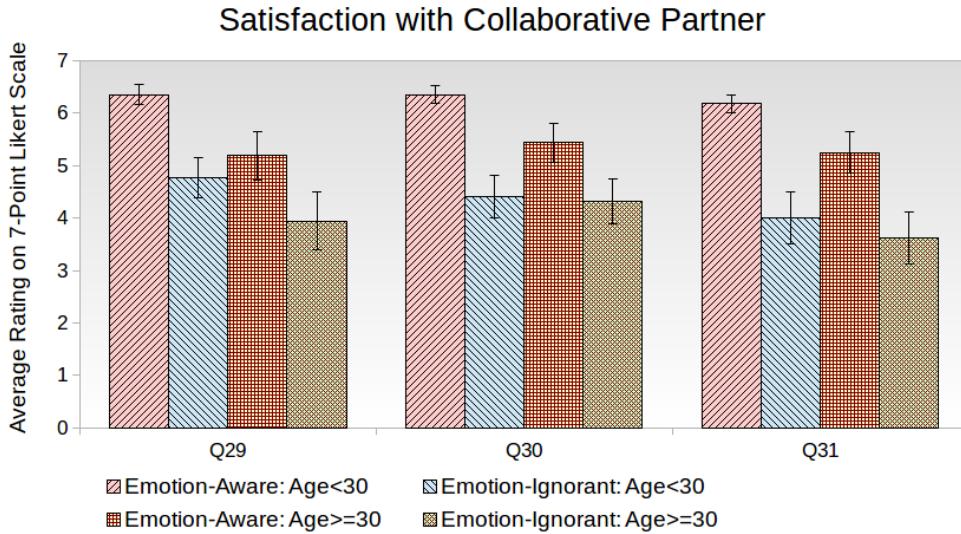


Figure 5.30: Impact of age on results of Likert scale questions related to satisfaction with collaborative partner.

5.2.5 Discussion

Based on the results, all participants prefer to work with the emotion-aware robot. Humans find the emotion-aware robot more likable and more trustworthy, as indicated in the Likert-scale responses and the open-ended questionnaire responses. Based on the responses, the emotional interaction with the robot can help create a sense of closeness and enjoyment that makes humans want to continue working with the robot.

The results also indicate that the emotion aware robot can better maintain a collaborative relationship. Both Likert-scale responses and Open-Ended Questionnaire responses indicate this. Humans felt a stronger sense of the robot's commitment to the collaboration, and greater understanding of their goals and emotions from the robot. Several open-ended responses also indicated that the robot was able to successfully motivate people and maintain their commitment to the collaboration, especially when tasks failed. Additionally, as shown in Section 5.2.4, humans rated the emotion-aware case much higher than the emotion-ignorant case when asked

which robot's decisions improved their performance, in essence acknowledging that their collaborator's (i.e. the robot's) decisions had a significant impact on their performance. As some of the open-ended responses indicated, successfully managing emotions within the collaboration can help keep the collaboration on track, and prevent distractions due to guilt and other negative emotions.

Finally, the emotion-aware robot developed a stronger sense of partnership through greater communication. The participants felt better understood by the emotion-aware robot, and felt that the goals were more mutually agreed-upon, refer to Section 5.2.4. As evidenced in the following response, the emotion-aware robot was successfully able to create a sense of partnership through its more open communication style: “Communication is very important. In the first run (i.e. emotion-aware) the robot states what tasks he is working on, it is clear and straight-forward. Also during the first run the robot cares about the human(me)'s feelings and cheers me up when I failed at the tasks, I think that could also improve efficiency of collaboration, because it would be more like a team or partnership.”

CHAPTER 6

CONCLUSION

6.1 Discussion

6.2 Future Work

BIBLIOGRAPHY

- [1] J. Andriessen, K. de Smedt, and M. Zock. Discourse planning: Empirical research and computer models. In A. Dijkstra and K. de Smedt, editors, *Computational psycholinguistics: AI and connectionist models of human language processing*, pages 247–278. Taylor & Francis, 1996.
- [2] G. E. M. Anscombe. *Intention*. NY: Cornell University Press, 1963.
- [3] T. Babaian, B. J. Grosz, and S. M. Shieber. A writer’s collaborative assistant. In *Proceedings of the International Conference on Intelligent User Interfaces (IUI2000)*, pages 7–14. ACM Press, 2002.
- [4] M. E. Bratman. *Intention, Plans, and Practical Reason*. Cambridge, Mass.: Harvard University Press, 1987.
- [5] M. E. Bratman. What is intention? In *Intentions in Communication*, Cognitive Technologies, pages 15–32. The MIT Press, Cambridge, MA, June 1990.
- [6] M. E. Bratman. Shared cooperative activity. *Philosophical Review*, 101(2):327–341, 1992.
- [7] M. E. Bratman, D. J. Israel, and M. E. Pollack. Plans and resource-bounded practical reasoning. *Computational Intelligence*, 4(3):349–355, 1988.
- [8] C. Breazeal, A. Brooks, J. Gray, G. Hoffman, C. Kidd, H. Lee, J. Lieberman, A. Lockerd, and D. Mulanda. Humanoid robots as cooperative partners for people. *Journal of Humanoid Robots*, 1(2):1–34, 2004.

- [9] C. Breazeal, A. Takanishi, and T. Kobayashi. *Social Robots that Interact with People*, pages 1349–1369. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [10] C. Burghart, R. Mikut, R. Stiefelhagen, T. Asfour, H. Holzapfel, P. Steinhaus, and R. Dillmann. A cognitive architecture for a humanoid robot: A first approach. In *5th IEEE-RAS International Conference on Humanoid Robots*, pages 357–362, 2005.
- [11] H.-N. Castañeda. *Thinking and Doing*. Dordrecht, Holland: D. Riedel, 1975.
- [12] C. Castelfranchi. Commitments: From individual intentions to groups and organizations. In *Proceedings of the first international conference on multiagent systems*, pages 41–48, 1995.
- [13] A. B. S. Clair and M. J. Matarić. Modeling action and intention for the production of coordinating communication in human-robot task collaborations. In *21st IEEE International Symposium on Robot and Human Interactive Communication: Workshop on Robot Feedback in HRI*, Paris, France, 2012.
- [14] W. J. Clancey. Roles for agent assistants in field science: Understanding personal projects and collaboration. *IEEE Transactions on Systems, Man and Cybernetics, special issue on Human-Robot Interaction*, 34(2):125–137, 2004.
- [15] P. Cohen, H. Levesque, and I. Smith. On team formation. In *Contemporary Action Theory. Synthese*, pages 87–114. Kluwer Academic Publishers, 1997.
- [16] P. Cohen and H. J. Levesque. *Teamwork*. SRI International, 1991.
- [17] P. R. Cohen and H. J. Levesque. Intention is choice with commitment. *Artificial Intelligence*, 42(2-3):213–261, 1990.
- [18] P. R. Cohen and H. J. Levesque. Persistence, intention, and commitment. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 33–69. MIT Press, Cambridge, MA, 1990.

- [19] P. R. Cohen, J. Morgan, and M. E. Pollack. *Intentions in Communication*. A Bradford Book, 1990.
- [20] C. L. Dancy. ACT-R ϕ : A cognitive architecture with physiology and affect. *Biologically Inspired Cognitive Architectures*, 6:40–45, 2013.
- [21] C. M. de Melo, J. Gratch, P. Carnevale, and S. J. Read. Reverse appraisal: The importance of appraisals for the effect of emotion displays on people’s decision-making in social dilemma. In *Proceedings of the 34th Annual Meeting of the Cognitive Science Society (CogSci)*, 2012.
- [22] N. Esau, L. Kleinjohann, and B. Kleinjohann. Integrating emotional competence into man-machine collaboration. In *Biologically-Inspired Collaborative Computing, September 8-9, Milano, Italy*, pages 187–198, 2008.
- [23] O. García, J. Favela, G. Licea, and R. Machorro. Extending a collaborative architecture to support emotional awareness. In *Emotion Based Agent Architectures (ebaa99*, pages 46–52, 1999.
- [24] E. Goffman. *The Presentation of Self in Everyday Life*. Anchor, 1959.
- [25] J. Gratch and S. C. Marsella. A domain-independent framework for modeling emotion. *Cognitive Systems Research*, 5(4):269–306, 2004.
- [26] S. A. Green, M. Billinghurst, X. Chen, and J. G. Chase. Human-robot collaboration: A literature review and augmented reality approach in design. *International Journal of Advanced Robotic Systems*, 5(1):1–18, 2008.
- [27] B. Grosz and S. Kraus. The evolution of shared plans. In *Foundations and Theories of Rational Agency*, pages 227–262, 1998.
- [28] B. J. Grosz. AAAI-94 presidential address: Collaborative systems. *AI Magazine*, 17(2):67–85, 1996.

- [29] B. J. Grosz. Beyond mice and menus. *Proceedings of the American Philosophical Society*, 149(4):523–543, 2005.
- [30] B. J. Grosz and L. Hunsberger. The dynamics of intention in collaborative activity. *Cognitive Systems Research*, 7(2-3):259–272, 2007.
- [31] B. J. Grosz, L. Hunsberger, and S. Kraus. Planning and acting together. *AI Magazine*, 20(4):23–34, 1999.
- [32] B. J. Grosz and S. Kraus. Collaborative plans for complex group action. *Artificial Intelligence*, 86(2):269–357, 1996.
- [33] B. J. Grosz and C. L. Sidner. Attention, intentions, and the structure of discourse. *Computational Linguistics*, 12(3):175–204, July 1986.
- [34] B. J. Grosz and C. L. Sidner. Plans for discourse. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 417–444. MIT Press, Cambridge, MA, 1990.
- [35] P. A. Heeman. *A Computational Model of Collaboration on Referring Expressions*. PhD thesis, University of Toronto, 1991.
- [36] R. W. Hill, Jr., J. Chen, J. Gratch, P. Rosenbloom, and M. Tambe. Intelligent agents for the synthetic battlefield: A company of rotary wing aircraft. In *Innovative Applications of Artificial Intelligence (IAAI-97)*, pages 227–262, 1997.
- [37] J. R. Hobbs, A. Sagae, and S. Wertheim. Toward a commonsense theory of microsociology: Interpersonal relationships. In *Formal Ontology in Information Systems - Proceedings of the Seventh International Conference*, pages 249–262, 2012.
- [38] L. Hunsberger and B. J. Grosz. A combinatorial auction for collaborative planning. In *In Proceedings of ICMAS*, 2000.

- [39] B. Jarvis, D. Jarvis, and L. Jain. Teams in multi-agent systems. In *Intelligent Information Processing III*, volume 228. Springer US, 2007.
- [40] N. R. Jennings. *Joint Intentions as a Model of Multi-Agent Cooperation in Complex Dynamic Environments*. PhD thesis, Department of Electronic Engineering, University of London, 1992.
- [41] N. R. Jennings. On being responsible. In E. Werner and Y. Demazeau, editors, *Proceedings of the Third European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, pages 93–102. North-Holland, 1992.
- [42] N. R. Jennings. Controlling cooperative problem solving in industrial multi-agent systems using joint intentions. *Artificial Intelligence*, 75(2):195–240, 1995.
- [43] N. R. Jennings and E. H. Mamdani. Using joint responsibility to coordinate collaborative problem solving in dynamic environments. In *10th National Conference on Artificial Intelligence (AAAI-92)*, pages 269–275, 1992.
- [44] C. L. John Robert Anderson. *The Atomic Components of Thought*. Lawrence Erlbaum Associates, 1998.
- [45] A. Kabil, C. D. Keukelaere, and P. Chavaillier. Coordination mechanisms in human-robot collaboration. In *Proceeding of the 7th International Conference on Advances in Computer-Human Interactions*, pages 389–394, 2014.
- [46] D. Kinny, M. Ljungberg, A. Rao, G. Tidhar, E. Werner, and E. Sonenberg. Planned team activity. In *Lecture notes in artificial intelligence*. Springer-Verlag, 1992.
- [47] H. Kitano, M. Asada, Y. Kuniyoshi, I. Noda, E. Osawai, and H. Matsubara. Robocup: A challenge problem for AI. *AI Magazine*, 18(1):73–85, 1997.
- [48] J. Laird. *The Soar Cognitive Architecture*. MIT Press, 2012.

- [49] N. Lesh, C. Rich, and C. L. Sidner. Collaborating with focused and unfocused users under imperfect communication. In M. Bauer, P. J. Gmytrasiewicz, and J. Vassileva, editors, *User Modeling 2001*, volume 2109, pages 64–73. Springer Berlin Heidelberg, 2001.
- [50] H. J. Levesque, P. R. Cohen, and J. H. T. Nunes. On acting together. In *AAAI*, pages 94–99. AAAI Press / The MIT Press, 1990.
- [51] D. J. Litman and J. F. Allen. Discourse processing and commonsense plans. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 365–388. MIT Press, Cambridge, MA, 1990.
- [52] K. E. Lochbaum. A collaborative planning model of intentional structure. *Computational Linguistics*, 24(4):525–572, 1998.
- [53] K. E. Lochbaum, B. J. Grosz, and C. L. Sidner. Models of plans to support communication: An initial report. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 485–490. AAAI Press, 1990.
- [54] C. Marinetti, P. Moore, P. Lucas, and B. Parkinson. Emotions in social interactions: Unfolding emotional experience. In *Emotion-Oriented Systems, Cognitive Technologies*, pages 31–46. Springer Berlin Heidelberg, 2011.
- [55] R. P. Marinier III, J. E. Laird, and R. L. Lewis. A computational unification of cognitive behavior and emotion. *Cognitive System Research*, 10(1):48–69, March 2009.
- [56] S. Marsella, J. Adibi, Y. Al-Onaizan, A. Erdem, R. Hill, G. A. Kaminka, Z. Qiu, and M. Tambe. Using an explicit teamwork model and learning in robocup: An extended abstract. In *RoboCup-98: Robot Soccer World Cup II*, volume 1604, pages 237–245. Springer Berlin Heidelberg, 1999.
- [57] S. C. Marsella and J. Gratch. EMA: A process model of appraisal dynamics. *Cognitive Systems Research*, 10(1):70–90, March 2009.

- [58] L. Matignon, A. B. Karami, and A.-I. Mouaddib. A model for verbal and non-verbal human-robot collaboration. In *AAAI Fall Symposium Series*, pages 62–67, 2010.
- [59] K. E. Merrick and K. Shafi. Achievement, affiliation, and power: Motive profiles for artificial agents. *Adaptive Behavior*, 19(1):40–62, 2011.
- [60] V. Montreuil, A. Clodic, M. Ransan, and R. Alami. Planning human centered robot activities. In *Proceedings of the IEEE International Conference on Systems, Man and Cybernetics*, pages 2618–2623, 2007.
- [61] B. Mutlu, A. Terrell, and C.-M. Huang. Coordination mechanisms in human-robot collaboration. In *Proceedings of the HRI 2013 Workshop on Collaborative Manipulation*, 2013.
- [62] S. Neale. Paul grice and the philosophy of language. *Linguistics and Philosophy*, 15(5):509–559, 1992.
- [63] S. Nikolaidis, P. A. Lasota, G. F. Rossano, C. Martinez, T. A. Fuhlbrigge, and J. A. Shah. Human-robot collaboration in manufacturing: Quantitative evaluation of predictable, convergent joint action. In *ISR*, pages 1–6, 2013.
- [64] S. Planalp. *Communicating Emotion: Social, Moral, and Cultural Processes*. Cambridge University Press, 1999.
- [65] M. E. Pollack. A model of plan inference that distinguishes between the beliefs of actors and observers. In *Proceedings of the 24th Annual Meeting on Association for Computational Linguistics*, pages 207–214. Association for Computational Linguistics, 1986.
- [66] M. E. Pollack. Plans as complex mental attitudes. In *Intentions in Communication*, pages 77–103. MIT Press, 1990.

- [67] D. V. Pynadath and M. Tambe. An automated teamwork infrastructure for heterogeneous software agents and humans. *Journal of Autonomous Agents and Multi-Agent Systems, Special Issue on Infrastructure and Requirements for Building Research Grade Multi-Agent Systems*, 7(1-2):71–100, 2003.
- [68] T. W. Rauenbusch and B. J. Grosz. A decision making procedure for collaborative planning. In *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems*, 2003.
- [69] I. Ravenscroft. *Folk Psychology as a Theory*. Stanford Encyclopedia of Philosophy, 2004.
- [70] C. Rich. Building task-based user interfaces with ANSI/CEA-2018. *IEEE Computer*, 42(8):20–27, July 2009.
- [71] C. Rich and C. L. Sidner. COLLAGEN: A collaboration manager for software interface agents. *User Modeling User-Adapted Interaction*, 8(3-4):315–350, 1998.
- [72] C. Rich, C. L. Sidner, and N. Lesh. COLLAGEN: Applying collaborative discourse theory to human-computer interaction. *AI Magazine*, 22(4):15–26, 2001.
- [73] J. Rickel, N. Lesh, C. Rich, C. L. Sidner, and A. Gertner. Collaborative discourse theory as a foundation for tutorial dialogue. In *Proceedings Sixth International Conference on Intelligent Tutoring Systems*, 2002.
- [74] P. Scerri, D. Pynadath, L. Johnson, P. Rosenbloom, M. Si, N. Schurr, and M. Tambe. A prototype infrastructure for distributed robot-agent-person teams. In *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems*, AAMAS '03, pages 433–440, New York, NY, USA, 2003. ACM.

- [75] P. Scerri, D. Pynadath, L. Johnson, P. Rosenbloom, M. Si, N. Schurr, and M. Tambe. A prototype infrastructure for distributed robot-agent-person teams. In *The Second International Joint Conference on Autonomous Agents and Multiagent Systems*, 2003.
- [76] K. R. Scherer, A. Schorr, and T. Johnstone. *Appraisal Processes in Emotion: Theory, Methods, Research*. Oxford University Press, 2001.
- [77] M. Scheutz and V. Andronache. Architectural mechanisms for dynamic changes of behavior selection strategies in behavior-based systems. *IEEE Transactions on Systems, Man, and Cybernetics, Part B*, 34(6):2377–2395, 2004.
- [78] J. R. Searle. Collective intentionality. In *Intentions in Communication*, pages 401–415. MIT Press, 1990.
- [79] C. Sidner. An artificial discourse language for collaborative negotiation. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 814–819. MIT Press, 1994.
- [80] C. L. Sidner and M. Dzikovska. A first experiment in engagement for human-robot interaction in hosting activities. In *Advances in Natural Multimodal Dialogue Systems*, volume 30 of *Cognitive Technologies*, pages 55–76. Springer Netherlands, 2005.
- [81] H. Skubch. *Modelling and Controlling of Behaviour for Autonomous Mobile Robots*. Springer Science Business Media, 2012.
- [82] D. Sofge, M. D. Bugajska, J. G. Trafton, D. Perzanowski, S. Thomas, M. Skubic, S. Blisard, N. Cassimatis, D. P. Brock, W. Adams, and A. C. Schultz. Collaborating with humanoid robots in space. *International Journal of Humanoid Robotics*, 2(2):181–201, 2005.

- [83] R. A. Subramanian, S. Kumar, and P. Cohen. Integrating joint intention theory, belief reasoning, and communicative action for generating team-oriented dialogue. In *AAAI*, pages 1501–1507. AAAI Press, 2006.
- [84] Suppressed for Anonymity. 2016.
- [85] M. Tambe. Agent architecture for flexible, practical teamwork. In *Proceedings of the National Conference on Artificial Intelligence*, pages 22–28, 1997.
- [86] M. Tambe. Towards flexible teamwork. *Journal of Artificial Intelligence Research*, 7:83–124, 1997.
- [87] S. van Hooft. Scheler on sharing emotions. *Philosophy Today*, 38(1):18–28, 1994.
- [88] B. Wilsker. A study of multi-agent collaboration theories. In *ISI Research Report*, pages 396–449, 1996.
- [89] J. Yen, J. Yin, T. R. Ioerger, M. S. Miller, D. Xu, and R. A. Volz. Cast: Collaborative agents for simulating teamwork. In *Proceedings of IJCAI2001*, pages 1135–1142, 2001.
- [90] J. Yin, M. S. Miller, T. R. Ioerger, J. Yen, and R. A. Volz. A knowledge-based approach for designing intelligent team training systems. In *Proceedings of the Fourth International Conference on Autonomous Agents*, pages 427–434. ACM, 2000.
- [91] Zamfirescu and Candea. On integrating agents into gdss. In *Preprints of the 9th IFAC / IFORS / IMACS / IFIP/ Symposium on Large Scale Systems: Theory and Applications*, 2001.

APPENDIX A