

Affective Motivational Collaboration Theory

by

Mohammad Shayganfar - mshayganfar@wpi.edu

A PhD Dissertation

Presented at

WORCESTER POLYTECHNIC INSTITUTE

in partial fulfillment of the requirements for the

DOCTOR OF PHILOSOPHY

in

Computer Science

January 2017

APPROVED

Professor Charles Rich, Thesis Advisor

Professor Candace L. Sidner, Thesis Co-Advisor

Professor John E. Laird, Thesis Committee Member

Professor Stacy Marsella, Thesis Committee Member

© Copyright by Mahni Shayganfar 2016
All Rights Reserved

ABSTRACT

Existing computational theories of collaboration explain some of the important concepts underlying collaboration, e.g., the collaborators' commitments and communication. However, the underlying processes required to dynamically maintain the elements of the collaboration structure are largely unexplained. Our main insight is that in many collaborative situations acknowledging or ignoring a collaborator's affective state can facilitate or impede the progress of the collaboration. This implies that collaborative agents need to employ affect-related processes that (1) use the collaboration structure to evaluate the status of the collaboration, and (2) influence the collaboration structure when required. This thesis develops a new affect-driven computational framework to achieve these objectives and thus empower agents to be better collaborators. Contributions of this thesis are: (1) Affective Motivational Collaboration (AMC) theory, which incorporates appraisal processes into Shared-Plans theory. (2) New computational appraisal algorithms based on collaboration structure. (3) Algorithms such as goal management, that use the output of appraisal to maintain collaboration structures. (4) Implementation of a computational system based on AMC theory. (5) Evaluation of AMC theory via two user studies to a) validate our appraisal algorithms, and b) investigate the overall functionality of our framework within an end-to-end system with a human and a robot.

ACKNOWLEDGMENTS

First and foremost I would like to thank all of my professors, and advisors throughout my higher-education. I thank Prof. Charles Rich for his guidance, teaching, and academic freedom as I explored my area of research and ultimately his support for this dissertation work. Prof. Rich continuously provided insightful discussions and suggestions about the research which ultimately shaped me as a person and a researcher. It has been an honor to be his Ph.D. student. I appreciate all his contributions of time, ideas, corrections and funding to make my Ph.D. experience productive and stimulating. In addition, I would like to thank Prof. Candace L. Sidner for not only serving as a committee member, but guiding me into the world of research. Furthermore, I thank Prof. Stacy Marsella (Northeastern University) and Prof. John E. Laird (University of Michigan) for their feedback as committee members whose advice and discussions are always much appreciated. My ability to complete this work would also not have been possible without the support of grants which help to fund me and my work including National Science Foundation award number IIS-1012083. Finally, I would like to thank my family. For my parents who raised me with love and supported me in all my pursuits. And most of all for my loving, supportive, encouraging, and patient wife Rudy whose faithful support during all stages of this Ph.D. has been greatly helpful and is so appreciated. Thank you.

Contents

Abstract	i
Acknowledgments	ii
List of Algorithms	xiii
1. Introduction	1
1.1 Motivation	1
1.2 Thesis Statement and Scope	3
1.3 Contributions	3
2. Background and Related Work	7
2.1 Introduction	7
2.2 Computational Theories of Collaboration	8
2.2.1 SharedPlans Theory	10
2.2.2 Joint Intentions Theory	16
2.2.3 STEAM – A Hybrid Approach	20
2.2.4 Other Approaches	22
2.2.5 Similarities and Differences	22
2.2.6 Applications of Collaboration Theories	24
2.3 Emotions and Affective Computing	29
2.3.1 Affect and Emotions	30
2.3.2 Emotion in Social Context	31
2.3.3 Communicating Emotions	34

2.3.4	Social Functions of Emotions	36
2.3.5	Artificial Emotions	38
2.3.6	Cognitive Architectures	40
2.4	Computational Models of Emotions	41
2.4.1	Appraisal Theory	41
2.4.2	Other (Non-Appraisal) Computational Models	48
2.4.3	Similarities and Differences	54
2.4.4	Applications in Autonomous Agents and Robots	60
2.5	Affect and Motives	64
2.5.1	Motives	65
2.5.2	Motivation Theory	65
2.6	Theory of Mind	67
2.7	Summary	68
3.	Affective Motivational Collaboration Theory	71
3.1	Overview of Mechanisms	73
3.2	Example Scenario	76
3.3	General Argument	79
3.4	Events	82
3.4.1	Collaboration Mechanism and Events	82
3.4.2	Appraising Events	83
3.4.3	Coping with Events	83
3.4.4	Motivation and Events	84
3.4.5	Theory of Mind and Events	84
3.5	Functions of Emotions	85
3.5.1	Action Selection	85
3.5.2	Adaptation	85
3.5.3	Social Regulation	86
3.5.4	Sensory Integration	86

3.5.5	Alarm	86
3.5.6	Motivation	87
3.5.7	Goal Management	87
3.5.8	Focus of Attention	87
3.5.9	Strategic Processing	88
3.5.10	Self Model	88
3.6	Components of the Architecture	88
3.6.1	Collaboration	89
3.6.2	Appraisal	90
3.6.3	Coping	91
3.6.4	Motivation	92
3.6.5	Theory of Mind	93
3.6.6	Perception	93
3.6.7	Action	94
3.6.8	Mental State and Emotion Instances	94
3.7	Attributes of Mental State Elements	97
3.7.1	Attributes of Beliefs	97
3.7.2	Attributes of Goals	99
3.7.3	Attributes of Motives	100
3.7.4	Attributes of Intentions	101
3.7.5	Emotion Instances	103
3.8	Summary	105
4.	Computational Framework	107
4.1	Introduction	107
4.2	Collaboration Mechanism	108
4.3	Appraisal Mechanism and Underlying Processes	110
4.3.1	Relevance	111
4.3.2	Desirability	116

4.3.3	Expectedness	117
4.3.4	Controllability	119
4.4	Goal Management	122
4.5	Coping Mechanism and Strategies	127
4.5.1	Planning	128
4.5.2	Active Coping	128
4.5.3	Seeking Social Support for Instrumental Reasons	130
4.5.4	Acceptance	131
4.5.5	Mental Disengagement	131
4.5.6	Shifting Responsibility	132
4.5.7	Activation of Coping Strategies	132
4.6	Motivation Mechanism	133
4.6.1	Satisfaction Motive	134
4.6.2	Achievement Motive	136
4.6.3	External Motive	138
4.7	Theory of Mind	138
4.8	Elicitation of Emotion Instances	140
5.	Evaluation	144
5.1	Implementation	144
5.2	Evaluating Appraisal Algorithms (Crowd Sourcing)	145
5.2.1	Experimental Scenario	145
5.2.2	Hypothesis and Methodology	146
5.2.3	Results	148
5.2.4	Discussion	154
5.3	End-to-End System Evaluation	154
5.3.1	Experimental Setup	155
5.3.2	Experimental Design	158
5.3.3	Hypotheses	163

5.3.4	Procedure	164
5.3.5	Results	166
5.3.6	Discussion	180
6.	Conclusion	182
6.1	Discussion	182
6.2	Future Work	184
Appendix A	216
Appendix B	235

List of Figures

1.1	A robotic arm collaborating with a human (in our end-to-end system evaluation) to achieve a shared goal using <i>Affective Motivational Collaboration</i> framework.	5
2.1	Plans for collaborative action [97].	13
2.2	Schematic view of the componential theory of emotion [118].	42
2.3	Comprehensive illustration of the CPM of emotion [220, 223].	44
2.4	A simple visualization of OCC model [178].	47
2.5	OCC taxonomy of emotion triggers and emotions [178].	49
2.6	Russell's suggested affective states based on core affect [209].	50
2.7	Three dimensional model of pleasure, arousal and dominance as tripartite view of experience [17].	51
2.8	Basic emotions and corresponding expressions.	52
2.9	Representing basic emotions within a dimensional framework [105].	56
2.10	A rough projection of emotion groups of OCC on the circumplex of affect [3].	59
3.1	Primary influence of mechanisms in <i>Affective Motivational Collaboration Theory</i>	74
4.1	Example of collaboration structure.	108
4.2	Using Collaboration structure in Appraisal (mechanisms in our framework).	111

4.3	Using Appraisals' outcome to influence Collaboration structure (mechanisms in our framework).	123
4.4	Cost values indicated by tuples with (second number) and without (first number) the influence of emotions.	124
4.5	Three functions of satisfaction motive for different values of valence. The x-axis indicates the satisfaction drive's delta value in [-1, +1], and the y-axis indicates the magnitude of satisfaction motive in [-1, +1].	135
4.6	Two functions of the achievement motive for different values of valence. The x-axis indicates the success probability value of achieving a goal which is in [0, +1], and the y-axis indicates the magnitude of achievement motive in [-1, +1].	137
4.7	Two functions of external motive for different values of valence. The x-axis indicates the success probability value of achieving a proposed goal which is in [0, +1], and the y-axis indicates the magnitude of the achievement motive in [-1, +1].	139
5.1	Collaboration task model for the evaluation.	146
5.2	Example expectedness question.	149
5.3	Example controllability question.	150
5.4	Example desirability question.	152
5.5	Example relevance question.	154
5.6	Experimental setup for end-to-end system evaluation.	155
5.7	Collaboration structure used as the task model.	156
5.8	The layout of the available spots for the human and the robot to place their pegs during the collaboration.	158
5.9	The Graphical User Interface (GUI) used during interaction.	159
5.10	Experimental setup.	161
5.11	Results of the Likert scale survey for Likability questions.	169

5.12 Results of the Likert scale survey for questions related to trust.	170
5.13 Results of the Likert scale survey for questions related to the robot's performance.	171
5.14 Results of the Likert scale survey for the questions related to the robot's understanding of human emotions.	172
5.15 Results of the Likert scale survey for questions related to the robot's understanding of goals.	173
5.16 Results of the Likert scale survey for questions related to the human's feeling about the collaboration.	174
5.17 Results of the Likert scale survey for questions related to satisfaction with collaborative partner.	175
5.18 Age and gender distribution of the user study participants.	178

List of Tables

4.1	Conditions for selecting candidate coping strategies	129
4.2	Appraisal values for relevance, desirability, expectedness and controllability.	141
4.3	Conditions for selecting emotion instances.	143
5.1	Number of participants	147
5.2	Expectedness results (the Equally Expected column indicates for which questions our algorithm provides option C as the response)	149
5.3	Controllability results (the Equally Controllable column indicates for which questions our algorithm provides option C as the response) . . .	150
5.4	Desirability results (the Equally Desirable column indicates for which questions our algorithm provides option C as the response)	151
5.5	Relevance results (the Equally Relevant column indicates for which questions our algorithm provides option C as the response)	153
5.6	The 31 Likert scale questions organized according to their categories (hypotheses).	168
5.7	Open-ended questionnaire questions and results. (*Note: Because we are evaluating whether humans prefer an affect-aware robot, these results are taken as negative test results when calculating the p-value using the binomial distribution. Only those participants who clearly indicated a preference for the emotion-aware robot are taken as positive test results.)	176
5.8	Multiple linear regression analysis results for the affect-aware condition.	179

5.9 Multiple linear regression analysis results for the affect-ignorant condition.	180
--	-----

LIST OF ALGORITHMS

1	Relevance Appraisal Process	113
2	Desirability Appraisal Process	117
3	Expectedness Appraisal Process	118
4	Controllability Appraisal Process	119

CHAPTER 1

INTRODUCTION

1.1 Motivation

The idea of robots and other intelligent agents sharing human environments has been a persistent aspiration in science fiction books, artificial intelligence and robotics laboratories. Collaborative robots are expected to become an integral part of human environments to accomplish industrial and household tasks. In many of these cases, humans will be involved in the robots' operations and decision-making processes. This involvement influences the efficiency of robots' interaction and performance, and makes the robots sensitive to human cognitive abilities and behaviors.

Current computational theories of collaboration are too task-driven. These theories explain many of the important concepts underlying collaboration, focusing on tasks, their constraints and their requirements, including the collaborators' commitments, and the necessity of communicating about mental state in order to maintain progress over the course of a collaboration. However, a key aspect of collaborative robots missing from those theories is being able to show behaviors that make them more likable, trustworthy, and understanding of human's feelings and goals. These aspects of collaborative behaviors can greatly influence the performance of a collaboration. Therefore, collaborative robots need to take into account human collaborator's affective state, and not only focus on executing different actions with respect to their plan to maintain the collaboration process.

According to [103] collaboration is a coordinated activity in which the participants

work jointly to satisfy a shared goal. We believe that in addition to the status of the shared plan, mutual beliefs and intentions, or other task-driven details of the collaboration, human's decisions are also influenced by affect-driven functions. Humans perceive, assess, and interpret their collaborator's activities in order to coordinate their own acts. Therefore, collaborators need mechanisms to a) perceive their counterpart's affective states (and perhaps their meanings), and b) communicate their own understanding of these perceptions. This aspect of reasoning is missing in the existing computational collaboration theories and their applications.

From a different point of view the most prominent computational collaboration theories, i.e., SharedPlans [101, 103] and Joint Intentions [54], explain only the structure of a collaboration. For example, in SharedPlans theory collaborators build a shared plan containing a collection of beliefs and intentions about the tasks in the shared plan. Collaborators communicate these beliefs and intentions. This communication leads to the incremental construction of a complete shared plan, and successful completion of the collaboration. Although these theories explain the important elements of a collaboration structure, the underlying processes required to dynamically create, use, and maintain the elements of this structure are largely unexplained. In particular, a general mechanism has yet to be developed that allows an agent to effectively integrate the influence of its collaborator's perceived or anticipated affective state into its own cognitive mechanisms to prevent shared task failures. Therefore, a process view of collaboration should inherently involve social interactions, since all collaborations occur between social agents, and it should contain a means of modifying the content of social interaction as the collaboration unfolds. The social functions of emotions explain some aspects of the underlying processes in collaboration. This thesis makes the case for affect-driven processes within collaboration and demonstrates how they further collaboration between humans and robots.

1.2 Thesis Statement and Scope

In this thesis, we develop and evaluate a framework called *Affective Motivational Collaboration* (AMC) which can improve the effectiveness of collaboration between agents/robots and humans. We address only two-participant collaboration; larger team collaboration is out of our scope. This thesis focuses on the reciprocal influence of the collaboration structure and the appraisal processes in a dyadic collaboration, specifically: a) the influence of affect-regulated processes on the collaboration structure, and b) the prediction and interpretation of the observable affective behaviors of the other during a collaborative interaction.

The AMC framework relies theoretically on the collaboration structure described by SharedPlans theory, and its implementation uses the collaboration manager, Disco. Disco is the open-source successor to COLLAGEN [203, 204] which incorporates algorithms based on SharedPlans theory for discourse generation and interpretation. AMC theory deals with the major affect-driven processes having an impact on the collaboration structure and ultimately a collaborative robot's behavior. This theory is informed by research in psychology and artificial intelligence, which is reviewed in Chapter 2. Our contribution, generally speaking, has been to synthesize prior work on appraisal¹, SharedPlans theory's description of collaboration, and motivation to provide a new computational theory of affect-regulated dyadic collaboration.

1.3 Contributions

Throughout this work we aim to show how a robot can leverage affect-driven processes, specifically appraisal algorithms, to improve collaboration with humans. As such, we first introduce our foundational theoretical concepts under the title Affective Motivational Collaboration theory, and then we introduce a novel computational

¹We have chosen appraisal-based modeling of emotions among several theories of emotions.

framework, based on this theory, which allows an agent to collaborate with a human incorporating underlying affect-driven processes and the affective expression of the human. The following summarize our contributions:

1. *Affective Motivational Collaboration (AMC) theory:*

(Chapter 3) The theoretical foundations of AMC framework are SharedPlans theory of collaboration [101, 103] and the cognitive appraisal theory of emotions [162] [223]. Applying cognitive appraisal theory in the collaboration context described by the SharedPlans theory is novel. AMC theory accounts for several key functions of affect in collaboration: *goal management*, *motivation*, *social regulation*, and *attentional focus*.

2. *New appraisal algorithms based on collaboration structure:*

(Chapter 4) We use SharedPlans description of the collaboration structure in four appraisal algorithms. i.e., *relevance*, *desirability*, *expectedness* and *controllability*, to compute the value of appraisal variables in a dyadic collaboration. These algorithms are inspired by [162], and as a novel approach these algorithms are designed based on different elements of mental state and human collaborator's affective state.

3. *New algorithms to influence collaboration structure:*

(Chapter 4) We use the evaluative nature of the appraisal to make reciprocal changes to the collaboration structure as required. We have developed new algorithms for different functions of emotions such as affect-driven goal management in the context of collaboration.

4. *Implementation of a computational system based on *Affective Motivational Collaboration* theory:*

We implemented a computational system which employs our models and algorithms in *Affective Motivational Collaboration* framework. Our computational

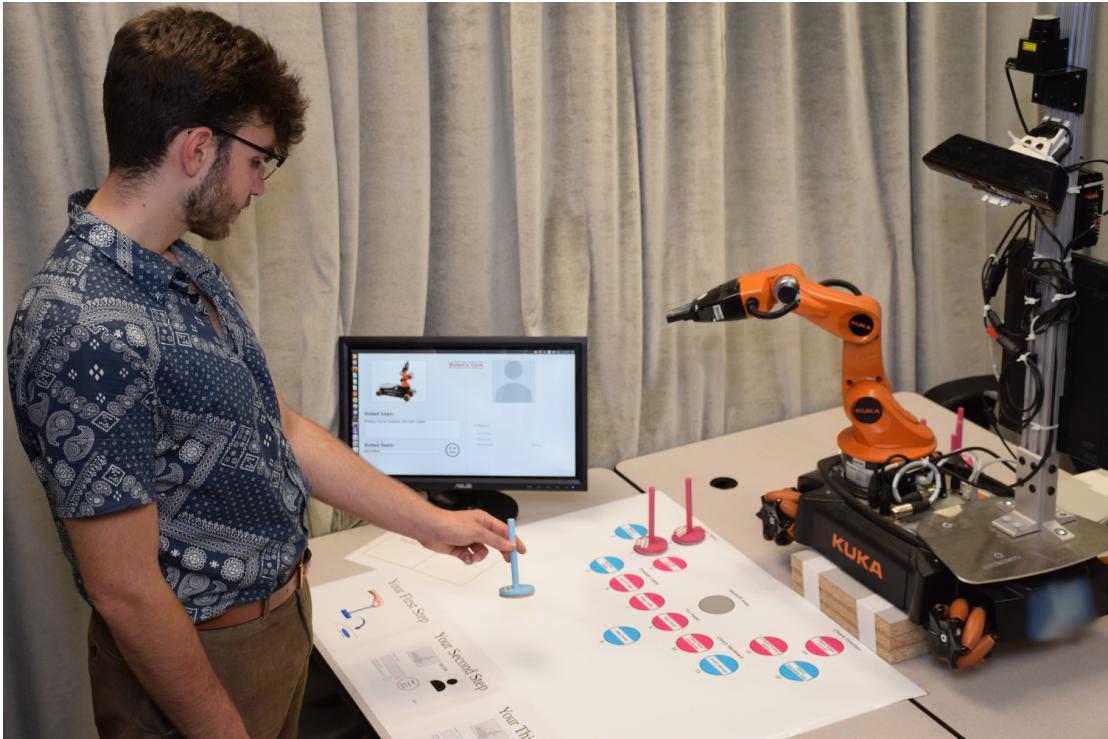


Figure 1.1: A robotic arm collaborating with a human (in our end-to-end system evaluation) to achieve a shared goal using *Affective Motivational Collaboration* framework.

system implements the key concepts related to *Affective Motivational Collaboration* theory as well as minimal implementation of other processes which are required for validation of the model but are not part of this thesis' contributions. We use Disco as our collaboration manager to receive and maintain the collaboration structure. The emphasis of the implementation is on the underlying cognitive processes of collaboration and appraisal, however the implementation also includes the Perception and the Action mechanisms.

5. Evaluation of *Affective Motivational Collaboration* theory:

(Chapter 5) We conducted two user studies to a) evaluate our appraisal algorithms before further development of our framework, and b) investigate the overall functionality of our framework within an end-to-end system evaluation with participants and a robot (see Figure 1.1). In the first user study,

we crowd-sourced questionnaires to test our hypothesis that our algorithms will resonate with humans' decisions by providing answers similar to humans' responses to questions related to different factors within our appraisal algorithms. In the second user study, we investigated the importance of affect-awareness in human-robot collaboration, and the overall functionality of the AMC framework with the participants in our study environment.

CHAPTER 2

BACKGROUND AND RELATED WORK

2.1 Introduction

In this chapter, we start by reviewing the background of prominent collaboration theories including SharedPlans theory [101, 103]. We discuss the similarities and differences between these collaboration theories (see Section 2.2.5) as well as related theoretical and practical work and applications. We continue by discussing the concept of affective computing and the social and communicative aspects of emotions from a psychological point of view. Understanding the social aspects of emotions is important in our work, since our work is focused on collaboration which is a social phenomenon in human environments. We also present the concept of artificial emotions and provide some examples of the existing computational models of emotions. Then, we provide the background of the cognitive appraisal theory of emotions as the second theoretical foundation in our work as well as other computational models of emotions and the related concepts such as some examples of cognitive architectures and the influence of affect in decision-making procedures. This chapter continues with a description of motives and related theories in psychology and Artificial Intelligence (AI). The role of motives as goal-driven affective components is crucial in our work, since the collaboration structure is based on the concept of a shared goal between collaborators. Finally, a brief description and the related work on theory of mind in psychology and artificial intelligence is provided as another concept used in a limited scale in our work.

2.2 Computational Theories of Collaboration

The construction of computer systems and robots that are intelligent, collaborative problem-solving partners is important in artificial intelligence and its applications. It has always been important to make computer systems better at helping humans do whatever these systems are designed for. To build collaborative systems, we need to identify the capabilities that must be added to individual agents so that they can work with humans or other agents. As Grosz says, collaboration must be designed into systems from the start; it cannot be patched on later [97].

Collaboration is a special type of coordinated activity in which the participants work jointly, together performing a task or carrying out the activities needed to satisfy a shared goal [101]. Collaboration involves several key properties at both the structural and functional levels: most collaborative situations involve participants who have different beliefs and capabilities; most of the time, collaborators only have partial knowledge of the process of accomplishing the collaborative activities; collaborative plans are more than the sum of individual plans; collaborators are required to maintain mutual beliefs about their shared goal throughout the collaboration; they need to be able to communicate with others effectively; they need to commit to the group activities and to their role in it; collaborators need to commit to the success of others; they need to reconcile between commitments to the existing collaboration and their other activities; and they need to interpret others' actions and utterances in the collaboration context [98]. These collaboration properties are captured by existing computational theories of collaboration.

As mentioned above, to be collaborative, partners, e.g., a robot and a human, need to meet the specifications stipulated by collaboration theories. These theories argue for an essential distinction between a collaboration and a simple interaction or even a coordination [96, 152]. This section briefly provides descriptions of major computational collaboration theories, their similarities and differences, and their application in AI and robotics. It primarily focuses on Joint Intention, SharedPlans

and hybrid theories of collaboration. In this section, we do not present the theories in formal language, but rather describe their features in general terms.

The prominent collaboration theories are mostly based on plans and joint intentions [54] [103] [150], and were strongly influenced by the BDI paradigm introduced by Bratman [26] which is fundamentally reliant on folk psychology [198]. The two theories, Joint Intentions [54] and SharedPlans [101, 103], have been used extensively to analyze and implement teamwork and collaboration.

The SharedPlans theory grew out of the theories of Bratman and Pollack [29, 189, 190], who outlined a mental-state view of plans in which having a plan is not just knowing how to do an action, but also having the intention to do the actions. Bratman's views of intention goes back to the philosophical views of Anscombe [6] and Castañeda [44] about intention. Also, as Grosz and Sidner mention in [103] the natural segmentation of discourse reflects intentional behaviors in each segment.

Cohen and Levesque also mention that in Joint Intentions theory their view of intention is primarily future-directed [55] which makes their view similar to Bratman's theory of intention [27], and contrary to Searle [230].

Commitment – One of the most important concepts in teamwork and collaboration is commitment. Collaboration theories are required to address the notion of commitment, otherwise the participants are just doing some coordinated activities. Since the prominent computational collaboration theories, reviewed in this paper, are based on Bratman's view of intention, we briefly provide his view of commitment here before describing these theories. Bratman defines certain prerequisites for an activity to be considered shared and cooperative [28]. He stresses the importance of:

- a) **Mutual commitment to joint activity** – which can be achieved by agreement on the joint activity, and prevention of abandoning the activity without involving teammates;

- b) **Mutual support** – which can be achieved by team members if they actively try to help teammate activity;
- c) **Mutual responsiveness** – which means team members should take over tasks from teammates if necessary.

In the following sections, we will see how each collaboration theory addresses the notion of commitment.

2.2.1 SharedPlans Theory

The SharedPlans theory of collaborative action, developed by Grosz and Sidner [100, 101, 103], aims to provide the theoretical foundations needed for building collaborative robots or agents [97]. SharedPlans is a general theory of collaborative planning that requires no notion of joint intentions (see Section 2.2.2), accommodates multi-level action decomposition hierarchies and allows the process of expanding and elaborating partial plans into full plans. SharedPlans theory explains how a group of agents can incrementally form and execute a shared plan which then guides and coordinates their activity towards the accomplishment of a shared goal¹. SharedPlans is rooted in the observation that collaborative plans are not simply a collection of individual plans, but rather a tight interleaving of mutual beliefs² and intentions of different team members. In [101] Grosz and Kraus use first-order logic to formalize SharedPlans.

Grosz and Sidner in [103] present a model of plans to account for how agents with partial knowledge collaborate in the construction of a domain plan. They are interested in the type of plans that underlie discourse in which the agents are collaborating in order to achieve a shared goal. They propose that agents are building a shared plan in which participants have a collection of beliefs and intentions about

¹In our framework we also have the notion of *private* goals (vs. *shared* goals) which are the goals that each collaborator might pursue privately; i.e., the other collaborator does not know about these goals.

²In our framework we also have the notion of *private* beliefs (vs. *shared* beliefs) which the other collaborator does not know about these beliefs.

the actions in the plan. Agents have a library of how to do their actions, i.e. recipes. These recipes may partially specify how an action is executed, or contributes to a goal. Then, each agent communicates its beliefs and intentions by making utterances about what actions they can contribute to the shared plan. This communication leads to the construction of a shared plan, and ultimately termination of the collaboration with each agent mutually believing that there exists one agent who is going to execute an action in the plan, and the fact that that agent has intention to perform the action, and that each action in the plan contributes to the goal [103] [153].

Later, we will see that to successfully complete a plan the collaborators must mutually believe that they have a common goal and have agreed on a sequence of actions for achieving that goal. They should believe that they are both capable of performing their own actions and intend to perform those actions while they are committed to the success of their plans.

Recipes

The SharedPlans theory differentiates between knowing how to accomplish a goal (a recipe) and having a plan, which includes intentions. The SharedPlans definition of mutual beliefs states that when agents have a shared plan for doing some action, they must hold mutual beliefs about the way in which they should perform that action [101, 103]. Following Pollack [190], the term recipe refers to what collaborators know when they know a way of doing an action. Recipes are specified at a particular level of detail. Although the agents need to have mutual beliefs about actions specified in the recipe, they do not need to have mutual beliefs about all levels of performing actions. Therefore, having mutual beliefs of the recipe means that the collaborators hold the same beliefs about the way in which an action should be accomplished. Consequently, the collaborators need to agree on how to execute an action. Recipes are aggregations of action-types and relations among them. Action-types, rather than actions, are the main elements in recipes. Grosz and Sidner in their earlier

work [103] have considered only simple recipes in which each recipe consisted of only a single action-type relation [153]. Recipes can be partial, meaning they can expand and be modified over time.

Grosz and Sidner propose that collaboration must have the following three elements:

1. the participants must have commitment to the shared activity;
2. there must be a process for reaching an agreement on a recipe for the group action;
3. there must be commitment to the constituent actions.

Shared plan is an essential concept in the collaboration context. The definition of the shared plan is derived from the definition of plans Pollack introduced in [189, 190] since it rests on a detailed treatment of the relations among actions and it distinguishes the intentions and beliefs of an agent about those actions. However, since Pollack's plan model is just a simple plan of a single agent, Grosz and Sidner extended that to plans of two or more collaborative agents. The concept of the shared plan provides a framework in which to further evaluate and explore the roles that particular beliefs and intentions play in collaborative activity [153]. However, Pollack's formulation of shared plans (a) could only deal with activities that directly decomposed into single-agent actions, (b) did not address the requirement for the commitment of the agents to their joint activities, and (c) did not adequately deal with agents having partial recipes [101]. Grosz and Kraus in [101], reformulate Pollack's definition of the individual plans [190], and also revise and expand the SharedPlans to address these shortcomings.

Figure 2.1 shows what we need to add to individual plans in order to have plans for group actions. The top of the figure lists the main components for individual plans. First, an individual agent needs to know the recipe for an action, whereas agents in a group need to have a mutual belief of a recipe for an action (bottom of

the figure). In the case of a group plan, having a mutual belief of a recipe, leads the agents to agree on how they are going to execute the action. Then, similar to individual agents that need to have the ability to perform the constituent actions in an individual plan and must have intentions to perform them, the participants in a group activity need to have individual or group plans for each of the constituent actions in the mutually agreed recipe [97, 103].

PLANS FOR COLLABORATIVE ACTION

- To have an individual plan for an act, need
 - knowledge of a recipe
 - ability to perform subacts in the recipe
 - intentions to do the subacts
- To have a group plan for an act, need
 - mutual belief of a recipe
 - individual or group plans for the subacts
 - intentions that group perform act
 - intentions that collaborators succeed

Figure 2.1: Plans for collaborative action [97].

As shown in Figure 2.1 (bottom), plans for group actions include two essential constituents that do not have correlates in the individual plan. First, the agents need to have a commitment to the group activity; All the agents need to intend that the group will do the action. For instance, when a robot and an astronaut are collaborating to install a solar panel, they need to have intentions to install the solar panel together. Among other things, these intentions will keep them both working on the panels until the panels are installed. Second, the participants need to have some commitment to the other agents to succeed in their own their actions. For instance, the robot must have an intention that the astronaut be able to measure the quality of installation successfully. This intention will prevent the robot from interrupting the astronaut's measurement action or prevent the robot from using the astronaut's measurement tool [97, 103].

Full Vs. Partial Shared Plan

The SharedPlans formalization distinguishes complete (full) plans and partial plans. A shared plan can be either a *Full Shared Plan (FSP)* or a *Partial Shared Plan (PSP)*. A *FSP* is a complete plan in which agents have fully determined how they will achieve a shared goal. A *PSP* definition provides a specification of the minimal mental state requirements for collaboration to exist and gives criteria governing the process of completing the plan.

A *FSP* to do α represents a situation where every aspect of a joint activity α is fully determined. This includes mutual belief and agreement in the complete recipe to do α . A recipe is a specification of a set of actions A_i , which constitutes the performance of α when executed under specified constraints. $FSP(\mathbf{P}, \Theta, \alpha, T_p, T_\alpha, \mathbf{R}_\alpha)$ denotes a group Θ 's plan \mathbf{P} at time T_p to do action α at time T_α using recipe \mathbf{R}_α .

Usually a team and its members do not possess a *FSP* to achieve their shared goal until it is done. SharedPlans uses the concept of *PSP* as a snapshot of the team's mental state in different situations, which further leads to communication and planning to fulfill the conditions of a *FSP*. The idea behind *PSP* is enabling the agents to modify the shared plan over the course of planning without impairing the achievement of the shared goals. Notice that for the same reason recipes also can be partial [101, 103].

Communicating Intentions

In SharedPlans, Grosz and Sidner are interested in the type of plans that underlie a discourse in which the agents collaborate to achieve a shared goal. Here we briefly present their view of discourse structure, since it is related to the intentions behind collaborators' actions. In [103], Grosz and Sidner argue that the SharedPlans theory recognises three interrelated levels of discourse structure, and the components of the discourse structure are a trichotomy of linguistic structure, intentions structure and the attention state. In their work, the linguistic structure of a discourse is a

sequence of utterances aggregating into discourse segments just as the words in a single sentence form constituent phrases. They also discuss the idea of the discourse purpose as the intention that underlies engagement in the particular discourse. They believe this intention is the reason behind performing a discourse rather than some other actions, and also the reason behind conveying a particular content of the discourse rather than some other contents. Finally, the third component in their theory, the attentional state, provides an abstraction of the agent’s focus of attention as the discourse unfolds. In short, the focusing structure is the central repository for the contextual content required for processing utterances during the discourse [103].

Intention-to and Intention-that

In Grosz and Sidner’s SharedPlans theory [103], two intentional attitudes are employed: *intending to* (do an action) and *intending that* (a proposition will hold). The notion of *intention to*, as an individual-oriented intention, models the intention of an agent to do any single-agent action while the agent not only believes that it is able to execute that action, but it also commits to doing so. In short, it is an intention to perform an action, similar to Bratman’s view of intention. In contrast with *intention to*, an *intention that*, as an intention directed toward group activity, does not directly imply an action. In fact, an individual agent’s *intention that* is directed towards its collaborators’ action or towards a group’s joint action. *Intention that* guides an agent to take actions (including communication), that enable or facilitate other collaborators to perform assigned tasks. This leads an agent to behave collaboratively. Therefore, agents will adopt intentions to communicate about the plan [101]. As another difference, *Intention to* commits an agent to means-end reasoning and acting [26] while *Intention that* does not necessarily entail this commitment. The key point about *Intention to* and *intention that* is that both commit an agent not to adopt conflicting intentions, and constrain replanning in case of failure. Further, an agent can *intention that* another agent achieve the specified proposition.

2.2.2 Joint Intentions Theory

Also starting with Bratman’s guidelines, Cohen and Levesque propose a different and more formal approach to building artificial collaborative agents. The Joint Intentions theory of Cohen and Levesque [54, 55, 56, 57, 146] represents one of the first attempts to establish a theory of collaboration expressed in formal logic, and due to its clarity, is a widely used teamwork theory.

A joint intention is a shared commitment to perform an action while in a group mental state [55]. Joint Intentions theory is based on individual and joint intentions (as well as commitments) to act as a team member. A joint intention is viewed not only as a persistent commitment of the team to a shared goal, but also implies a commitment on the part of all its members to a mutual belief about the state of the goal. In other words, Joint Intentions theory describes how a team of agents can jointly act together by sharing mental state about their actions where an intention is viewed as a commitment to perform an action.

In [54] Cohen and Levesque establish that a joint intention cannot be defined simply as individual intention with the team regarded as an individual. The reason is that after the initial formation of an intention, team members may diverge in their beliefs and their attitudes towards the intention. Instead, they first present a definition of individual persistent goal and individual intention. Then, they define team analogues of these concepts by presenting mutual belief in place of individual belief. The definition of joint persistent goal requires each team member to commit to informing other members, if it comes to believe that the shared goal is in its terminal status. As a result, in Cohen and Levesque’s theory, a team with a joint intention is a group that shares a common objective and a certain shared mental state [121].

In this theory, once an agent entered into a joint commitment with other agents, the agent should communicate its private beliefs with other team members if the agent believes that the joint goal is in its terminal status, i.e., either the joint goal

is achieved, or it is unachievable, or irrelevant [263]. Thus, as we mentioned above, team members are committed to inform other team members when they reach the conclusion that a goal is achievable, impossible or irrelevant. For instance, if a robot and an astronaut are collaborating to install a solar panel, and the robot reaches the conclusion that the welding tool has a deficiency, it is essential for the robot to have an intention to communicate with the astronaut and make this knowledge common. Therefore, according to this theory, in a collaboration, agents can count on the commitment of other members, first to the goal and then to the mutual belief of the status of the goal.

Individual Commitment

As we mentioned earlier, intentions and commitments are the basic ideas of Joint Intentions theory. Here, we provide the definition of “individual commitment” (also called *persistent goal*) by Cohen et. al. in [53]. According to their definition an agent has a persistent goal relative to q to achieve p only when:

1. agent believes that p is currently false;
2. agent wants p to be true;
3. it is true (and agent knows it) that (2) will continue to hold until the agent comes to believe either that p is true, or that it will never be true, or that q is false.

Note that the condition q is an “escape” clause, which can be omitted for brevity, or it can be used as a reason for the agent to drop a commitment, even though it could be quite vague.

Individual Intention

As we mentioned above, Joint Intention theory adopts Bratman’s view of future-directed properties of intention. In this theory, an intention is defined to be a

commitment to act in a certain mental state. In other words, an agent intends relative to some condition to do an action when it has a persistent goal or commitment (relative to that condition) of having done the action and, moreover, believing throughout that it is doing that action [54].

Intention inherits all the properties of commitment (e.g., consistency with mental state). Typically, an agent uses an intention as a decision within a goal hierarchy to do a particular action. For instance, initially, the agent commits to p becoming true without having any concern about who or how p is going to be accomplished. Then, the agent commits to x or y as a mean to accomplish p . Lastly, the agent selects one of the actions (e.g., x) and forms an intention to do it. This intention will be given up when for whatever reason p is accomplished.

An agent has a *Weak Achievement Goal* (WAG) relative to q and with respect to a team to bring about p if either of the following conditions holds:

- The agent has a normal achievement goal to bring about p ; that is, the agent does not yet believe that p is true and wants p to be true as a goal.
- The agent believes that p is true, will never be true, or is irrelevant, but has as a goal that the status of p be mutually believed by all the team members.

Joint Commitment and Joint Intention

A joint intention of a team is based on its joint commitment, which is defined as a *Joint Persistent Goal* (JPG)¹. A JPG to achieve a team action p , requires all team members to mutually believe that p is currently false and want p to eventually be true. A JPG guarantees that team members cannot decommit until p is mutually known to be *achieved*, *unachievable* or *irrelevant*. This commitment typically makes an agent communicate with its teammates [54].

¹Castelfranchi in [45] criticizes the necessary and sufficient conditions for the joint persistent goal. He argues that these conditions are not sufficient for the collaborators working as a team.

Therefore, an important consequence of achieving joint commitment in a team is that it predicts future communication which is critical within the course of a collaboration. Thus, this communication leads team members to attain mutual beliefs which is a fundamental concept in teamwork activities. Notice that the minimum mutual belief for team members to attain is the achievement or failure of the shared goal which terminates collaboration.

Joint intention is defined to be a joint commitment to the team members trying to do a joint action. Based on Cohen and Levesque's definition of joint intention, a team of agents jointly intends (relative to some escape condition) to do an action if and only if the members have a JPG (relative to that condition) of them having the action completed, and having it completed mutually believing throughout that they are doing it (knowingly) [54].

Teamwork & Communication

In summary, according to Joint Intentions theory, the notion of teamwork is characterized by joint commitment, also known as joint persistent goal (JPG). The definition of JPG states that the agents mutually believe they have the appropriate goal, and that they mutually believe a persistent weak achievement goal (which represents the one-way commitment of one agent directed towards another) to achieve it persists until the agents mutually believe that the goal has either been achieved, or become impossible or irrelevant.

Joint Intentions theory claims that an efficient collaboration requires communication. Sharing information through communication is critical given that collaborators have different capabilities, and each individual often has only partial knowledge relevant to solving the problem, and sometimes diverging beliefs about the state of the collaborative activity. Communication is important in coordinating team members' roles and actions to accomplish their goal. For instance, it can help team members to establish and maintain a set of mutual beliefs regarding the current state of the collaboration, and the respective roles and capabilities of each member.

2.2.3 STEAM – A Hybrid Approach

Tambe in [249] argues that teamwork in complex, dynamic, multi-agent domains requires the agents to obtain flexibility and reusability by using integrated capabilities. Tambe created STEAM (**S**hell **TEAM**work) based on this idea. STEAM’s operationalization in complex, real-world domains is the key in its success in addressing important teamwork issues, some of which are discussed in Section 2.2.6. STEAM is founded on the Joint Intentions theory and it uses joint intentions as the basic building block of teamwork but is also informed by key concepts from SharedPlans theory.

Building on the well developed theories of Joint Intentions and SharedPlans, the STEAM teamwork model was operationalized as a set of domain-independent rules that describe how teams should work together. According to Tambe, there are several advantages due to his use of Joint Intentions theory, including achieving a principled framework for reasoning about coordination and communication in a team. Another advantage is guidance for monitoring and maintenance of a team activity which the joint commitment concept in joint intention provides. And lastly, Tambe believes the joint intention in a team can facilitate reasoning about team activity and team members’ contribution to that activity.

However, he also believes that for a high level team goal, one single joint intention is not sufficient to achieve all these advantages. STEAM therefore borrows some of the concepts of SharedPlans theory. First, STEAM uses the concept of “intention that” (see Section 2.2.1) towards an activity as well as the fact that SharedPlans theory mandates team members’ mutual belief in a common recipe and shared plans for individual steps in the common recipe. Thus, in this case, SharedPlans helps STEAM to achieve coherency within the teamwork. In addition, STEAM uses joint intentions to ensure the teamwork coherency to build the mental attitudes of team members. In other words, as the recipe evolves, STEAM requires all team members to agree on the execution of a step and form joint intentions to execute it while

other joint intentions are formed, leading to a hierarchy. A second concept STEAM borrows from SharedPlans is the amount of information that a team member needs to know to perform an action. According to SharedPlans, team members require to know only that a recipe exists to enable them to perform actions (not the recipe details – see Section 2.2.1). Similarly in STEAM, team members only track the sub-team or individual team member responsible to perform a specific step; this tracking does not need detailed plan recognition. The third concept is parallel to what is called an unreconciled case in SharedPlans theory, which in STEAM is handled by replanning and communication between team members assigning the unassigned or unachieved task. The last concept is communication between team members which also borrows the concept of “intention that” from SharedPlans theory, to help the generalization of STEAM’s communication capabilities beyond what Joint Intentions theory offers.

In summary, STEAM builds on both Joint Intentions theory and SharedPlans theory and tries to overcome their shortcomings. Based on Joint Intentions, STEAM builds up hierarchical structures that parallel the SharedPlans theory. Hence, STEAM formalizes commitments by building and maintaining Joint Intentions, and uses SharedPlans to formulate the team’s attitudes in complex tasks.

In [249] Tambe argues that the novel aspects of STEAM relate to its teamwork capabilities. A key novelty in STEAM is team operators. In STEAM, when agents select a team operator for execution, they instantiate a team’s joint intentions. Team operators explicitly express a team’s joint activities, unlike the individual operators which express an agent’s own activities. Hence, STEAM agents maintain their own private (to apply individual operators) and team states, e.g., mutual belief about the world (to apply team operators).

Tambe added further practical concepts into STEAM’s architecture. For instance, STEAM has a team synchronization protocol to establish joint intention, and it has constructs for monitoring joint intentions which helps the agent to monitor team performance. STEAM facilitates this monitoring by exploiting its explicit

representation of team goals and plans. In particular, STEAM allows an explicit specification of monitoring conditions to determine achievement, unachievability or irrelevancy conditions of team operators. Finally, in STEAM, communication is driven by commitments embodied in the Joint Intentions theory, i.e., team members may communicate to obtain mutual belief while building and disbanding joint intentions. Thus, joint intentions provide STEAM with a principled framework for reasoning about communication. Also, STEAM addresses some practical issues not addressed in other teamwork theories. One of these issues is STEAM’s detailed attention to communication overheads and risks, which can be significant [248]. Furthermore, operationalization of STEAM is based on enhancements to the Soar architecture [137], plus a set of about 300 domain-independent Soar rules.

2.2.4 Other Approaches

There are other frameworks, approaches, and models focusing on teamwork and collaborative agents. For instance, Jennings developed the Joint Responsibility framework which is specified formally using modal temporal logic. Joint Responsibility stresses the role of joint intentions (based on Joint Intentions theory) specifying how both individuals and teams should behave whilst engaged in collaborative problem solving [122, 123, 124, 125]. Jennings has developed Generic Rules and Agent model Testbed Environment (GRATE) as a prototype system based on the Joint Responsibility framework. In [131] Kinny et. al. elaborate the concept of Planned Team Activity and introduce a language for representing joint plans for teams of agents and describe how agents can organize the formation of a skilled team to achieve a joint goal. They use joint intentions to capture the mental properties which characterize team activity.

2.2.5 Similarities and Differences

There are some similarities between SharedPlans and Joint Intentions theories:

1. Similar to SharedPlans theory, Joint Intentions theory specifies what it means for agents to execute actions as a team [246].
2. Both theories follow Bratman's basic ideas about the roles of intention in relational actions which prevent the collaborative agents from adopting conflicting intentions. Both theories also follow Bratman's BDI model.
3. Just as SharedPlans theory, Joint Intentions theory states that a joint action cannot be seen simply as a collection of individual actions, but rather as agents working together who need to share beliefs.
4. Both theories in their mature forms emphasize that agents are required to communicate to maintain collaboration. SharedPlans theory requires collaborators to communicate to establish and maintain the shared plan, which is crucial especially when collaborators only have a partial shared plan. Similarly in Joint Intentions theory, communication is an explicit requirement of collaborative agents until the shared goal is achieved, unachievable or irrelevant.
5. Both Joint Intentions and SharedPlans theories are concerned about commitment to the joint activity. However, these two theories use different concepts to fulfill the requirements of commitment during collaboration.

There are also differences between SharedPlans and Joint Intention theories:

1. The crucial components of the SharedPlans theory (see Section 2.2.1) lack the notion of a joint intention, which is the most significant notion within the Joint Intentions theory. For philosophocal reasons, Grosz and Sidner do not believe that such a phenomenon (joint intention) exists in a collaboration. They believe their notion of “intention that” and mutual beliefs about states of the collaboration can provide similar functionalities as described in Joint Intentions theory (see Section 2.2.2).

2. In SharedPlans theory teammates agree on the shared plan, whereas in Joint Intentions theory teammates agree on intentions.
3. In contrast to Joint Intentions, the SharedPlans theory employs hierarchical structures over intentions, thus overcoming the shortcoming of a single joint intention for complex team tasks.
4. The SharedPlans theory describes the way to achieve a common goal through the hierarchy of plans, whereas the Joint Intentions theory describes only this common goal [239].
5. Joint Intentions theory assumes that knowledge about the teammates is always available, whereas SharedPlans theory uses the concept of partial plan/recipe to make the process of dynamically achieving information possible throughout the collaboration.
6. Communication requirements are derived from “intention that” in SharedPlans theory, as opposed to being “hard wired” in Joint Intentions theory.

2.2.6 Applications of Collaboration Theories

There is significant practical research focusing on different aspects of collaboration based on different collaboration theories, i.e., SharedPlans, Joint Intentions, and hybrid theories of collaboration. In this section, we provide some examples of implemented homogeneous and heterogeneous agent/robot and human collaborations. Some of these works also integrate the concept of emotions, but neither of them focuses on using affect-driven process to maintain collaboration structure.

Some work focuses on the concepts of robot assistants [49], or teamwork and its challenges at cognitive and behavioral levels [177, 214]. Some researchers have taken an overall look at a collaboration concept at the architectural level. In [82] authors present a collaborative architecture, COCHI, to support the concept of emotional awareness. In [75] authors present the integration of emotional competence into

a cognitive architecture which runs on a robot, MEXI. In [244] authors discuss the challenges of integrating natural language, gesture understanding and spatial reasoning of a collaborative humanoid robot situated in space. The importance of communication during collaboration has also been considered by some researchers from human-computer interaction and human-robot collaboration [48, 166, 204] to theories describing collaborative negotiation, and discourse planning and structures [5, 102, 236]. There are other concepts such as joint actions and commitments [99], dynamics of intentions during collaboration [146], and task-based planning providing more depth in the context of collaboration [39, 202]. The concept of collaboration has also received attention in industries and academic robotic laboratories [93].

Applications of SharedPlans Theory – COLLAGEN [203, 204] is the first implemented general tool for collaboration based on the SharedPlans theory. It incorporates algorithms for discourse generation and interpretation, and is able to maintain a segmented interaction history, which facilitates the discourse between the human user and the intelligent agent. The model includes two main parts: (1) a representation of a discourse state and (2) a discourse interpretation algorithm for the utterances of the user and agent [205]. In [108] Heeman presents a computational model of how a conversational participant collaborates in order to make a referring action successful. This model is based on the view of language as goal-directed behaviour, and in his work, he refers to SharedPlans as part of the planning and conversation literature. In [153], Lochbaum and Sidner modify and expand the SharedPlan model of collaborative behavior [103]. They present an algorithm for updating an agent's beliefs about a partial shared plan and describe an initial implementation of this algorithm in the domain of network management. Lochbaum, in [152], provides a computational model (based on the collaborative planning framework of SharedPlans [101]) for recognizing intentional structure and utilizing it in discourse processing. CAST (Collaborative Agents for Simulating Teamwork) [266] [267] is a teamwork framework based on SharedPlans theory. CAST focuses on flexibility in dynamic environments and on proactive information exchange enabled

by anticipating what information team members will need. Petri Nets are used to represent both the team structure and the teamwork process, i.e., the plans to be executed. Researchers in [114] discuss developing an ontology of microsocial concepts for use in an instructional system for teaching cross-cultural communication. They believe being acquainted with one another is not a strong enough relationship from which to create a society. Hence, there is a need for commitment and shared plans (as the basis of social life) to achieve a shared goal. In this work, Grosz and Sidner's SharedPlans theory [103] is used to explain the concept of shared plans within the interpersonal relationships of societies in an industrial environment. In [119] Hunsberger and Grosz discuss the idea of how rational, utility-maximizing agents should determine commitment to a group activity when there is an opportunity to collaborate. They call this problem the "initial-commitment decision problem" (ICDP) and provide a mechanism that agents can use to solve the ICDP. They use the representation of action, act-types and recipes in the SharedPlans theory in this work. In [269] an integrated agent-based model for Group Decision Support Systems is proposed and discussed. The decisional model that authors outline in this paper is based on the SharedPlans theory. Rauenbusch and Grosz in [197] formally define a search problem with search operators that correspond to the team planning decisions. They provide an algorithm for making three types of interrelated decisions by recasting the problem as a search problem. Their model respects the constraints on mental state specified by the SharedPlans theory of collaboration. Babaian et. al. in [12] describe Writer's Aid, a system that deploys AI planning techniques to enable it to serve as an author's collaborative assistant. While an author writes a document, Writer's Aid helps in identifying and inserting citation keys and by autonomously finding and caching potentially relevant papers and their associated bibliographic information from various on-line sources. They believe the underlying concepts of SharedPlans are relevant since in collaborative interfaces like Writers Aid, the users establish shared goals with the system, and user and the system both take initiative in satisfying them. In [171] researchers address high-level robot planning issues for

an interactive cognitive robot that acts in the presence of or in collaboration with a human partner. They describe a Human Aware Task Planner (HATP) which is designed to provide socially acceptable plans to achieve collaborative tasks. They use notions of plans based on SharedPlans theory. In [237] Sidner and Dzikovska conducted an initial research on engagement in human-human interaction and applications to stationary robots performing hosting activities, such as tutoring and sales. They believe hosting activities are collaborative because neither party completely determines the goals to be undertaken nor the means of reaching the goal. To build a robot host, they rely on an agent built using COLLAGEN which is a tool based on the SharedPlans theory.

Applications of Joint Intentions Theory – In [131] authors introduce a language for representing joint plans for teams of agents. They describe how agents can organize the formation of a suitably skilled team to achieve a joint goal, and they explain how such a team can execute these plans to generate complex, synchronized team activity. In this work, authors adopt the underlying concepts of the Joint Intentions theory as the structure of their collaborative agents. Breazeal et. al. in [34] present an overview of their work towards building socially intelligent, cooperative humanoid robots, such as Leonardo, that can collaborate and learn in partnership with humans. They employ the Joint Intentions theory of collaboration to implement the collaborative behaviors while performing a task in collaboration with humans. In [246] the researchers' goal is to develop an architecture, based on the concepts of Joint Intentions theory, that can guide an agent during collaborative teamwork. They describe how a joint intention interpreter that is integrated with a reasoner over beliefs and communicative acts can form the core of a dialogue engine. Ultimately, the system engages in dialogue through the planning and execution of communicative acts necessary to attain the collaborative task at hand. Mutlu et. al. in [175] discuss key mechanisms for effective coordination toward informing the design of communication and coordination mechanisms for robots. They present two

illustrative studies that explore how robot behavior might be designed to employ these mechanisms (particularly joint attention and action observation) to improve measures of task performance in human-robot collaboration. Their work uses Joint Intentions theory to develop shared task representations and strategies for task decomposition. The GRATE* system by Jennings [124] is based on the Joint Intention theory. GRATE* provides a rule-based modelling approach to cooperation using the notion of Joint Responsibilities, which in turn is based on Joint Intentions. GRATE* is geared towards industrial settings in which both agents and the communication between them can be considered to be reliable.

Applications of Hybrid Theories – The domain independent teamwork system, STEAM, has been successfully applied to a variety of domains. From combat air missions [112] to robot soccer [134] to teams supporting human organizations [195] to rescue response [215], applying the same set of STEAM rules has resulted in successful coordination between heterogeneous agents. The successful use of the same teamwork model in a wide variety of diverse domains provides compelling evidence that it is the principles of team-work, rather than exploitation of specific domain phenomena, that underlies the success of teamwork based approaches. In [159] authors provide their RoboCup (robotics soccer testbed) in which their focus is on teamwork and learning challenges. Their research investigation in RobotCup is based on ISI Synthetic, a team of synthetic soccer-players. They also investigate the use of STEAM as their model of teamwork which is influenced by the Joint Intentions and SharedPlans theories. In [127] researchers propose a behavioral architecture C²BDI that allows the enhancement of the knowledge sharing using natural language communication between team members. They define collaborative conversation protocols that provide proactive behavior to agents for the coordination between team members. Their agent architecture provides deliberative and conversational behaviors for collaboration, and it is based on both the SharedPlans and Joint Intentions theories.

2.3 Emotions and Affective Computing

According to Picard [187], the term affective computing encompasses a new approach in artificial intelligence to building computers that show human affection. Studies show that the decision making of humans is not always logical [95], and in fact, not only is pure logic not enough to model human intelligence, but it also shows failures when applied in artificial intelligence systems [69].

If we want robots and virtual agents to be more believable and efficient partners for humans, we must consider the personal and social functionalities and characteristics of emotions; this will enable our robots to coexist with humans, who are emotional beings. To have a better understanding of applications of affective computing, we can categorize the existing literature of computational emotion modeling into four major categories: a) detecting and recognizing human emotions, b) interpreting and understanding human emotions, c) generating artificial emotions and applying the underlying processes to exploit emotion functions, and d) expressing human-perceivable emotions during interaction.

The major approaches to model emotions are *appraisal*, *dimensional* and *discrete (basic)*, some of which have corresponding computational models, e.g., EMA [162] and WASABI [22, 23]. These models have been used in different domains including AI and robotics. Applying these models can help robots and virtual agents to achieve communicative, evaluative, interpretive, and regulatory aspects of emotions in some or all of the four categories mentioned above.

This section provides descriptions of the major computational emotion theories, their comparison, and their applications in AI and robotics. It includes the existing influential computational emotion theories as well as the underlying psychological theories; we mainly focus on appraisal theory since it emphasizes and explains the connection between affective state and cognition, although we also discuss dimensional theory, and we briefly cover other approaches, e.g. discrete (basic) emotions.

2.3.1 Affect and Emotions

Affect is the representation of an agent’s emotional state in some form, which can be the current instance of the agent’s affective state. Emotion influences not only what people do, but also the way they do it [60]. Aristotle in *The Nicomachean Ethics* reveals his idea about emotions. He says “Anyone can become angry—that is easy. But to be angry with the right person, to the right degree, at the right time, for the right purpose, and in the right way—this is not easy [7].”

Intelligence is a set of mental abilities that enables a human to comprehend, reason and adapt in the environment, and as a result, act effectively and purposefully in that environment. Emotions play a crucial role in scientists’ explanation of humans’ intelligent behaviors. Emotions significantly impact the processes of action generation, execution, control, and interpretation [272] in different environments. Emotions are conceptualized as ongoing processes rooted in dynamic social contexts, which can shape both implicit and explicit emotional responses [156]. An emotion is a dynamic episode that not only involves changes in cognitive states, but also produces a sequence of response patterns on body movements, posture, voice and face [222]. Emotions typically occur in response to an event, usually a social event, real, remembered, anticipated, or imagined. Emotions are associated with distinctive relational meanings [185]. These relations can be with the individual’s past experience, the individual’s surrounding objects and environment, or the other individuals with or without mutual beliefs in a dyadic or a group setting. Emotions are evaluative and responsive patterns that serve the function of providing appraisal about whether the ongoing event is harmful, threatening or beneficial for the well-being of an individual [272]. Consequently, intelligence and emotional processes have an integral and a supportive relationship, rather than an antagonistic or a conflicting one.

A better question than what emotions are, is the question of what they can do, and how they impact human life. Emotions impact fundamental parts of cognition

including perception, memory, attention and reasoning [50]. This impact is caused by the information emotions carry about the environment and event values. The influence of emotions depends on an individual’s focus of attention. For instance, a positive affect can cause a positive attitude towards an object if the individual’s focus is on the object, whereas the same positive affect can be interpreted as a positive feedback towards one’s partner during the course of a collaboration. As another example, a positive feedback can promote certain cognitive processes, or it can inhibit other cognitive processes according to the conditions in the environment [51]. In both cases, emotions play a regulatory role for cognitive processes [94]. Some of these effects flow from underlying shifts in the way people perceive and think under the influence of emotion.

2.3.2 Emotion in Social Context

In this section, we discuss the importance of studying emotions within a social context. This perspective is important in our research because our work is focused on collaboration as a particular social setting between individuals. Understanding the dynamics of collaboration requires one to understand influential underlying components.

Emotions are involved in developing social contexts. Humans are social and most of the causations and constitutions of their emotions are social. Brian Parkinson in [183] argues that many of the causes of emotions are interpersonal and communicative rather than internal and reactive phenomena. There are different social aspects of emotions influenced by various factors such as social context and social relationship type. For instance, a dominant-submissive social relationship can cause and contain different emotions with different intensities compared to a reciprocal or a friendship social relationship. As another example, an emotion can be interpreted in a certain way when an individual is situated in an environment with other people who are expressing a particular emotion.

As mentioned earlier, the social context is an important factor influencing one’s

emotions. A dyadic interaction is one type of social setting [52]. Dyadic interaction tasks make it possible to examine how individuals experience and express emotions during social interactions and how emotions shape and are shaped by the reciprocal interactions between individuals. In addition, eliciting and monitoring emotional processes yields useful information about the role emotion plays in interpersonal relationships. Compared with other emotion-eliciting events, events in a dyadic interaction can better help us study an ongoing emotional relationship between two individuals in addition to their internal emotional and cognitive processes. Dyadic interaction tasks are ideal for studying a range of emotional responses because of the fairly unstructured conversations between the individuals. Thus, dyadic interaction tasks will generate a wide range of emotions in comparison with the controlled emotion-eliciting events.

There are numerous ways that emotions can be social [254]. There is a consensus on the fact that social events and entities surrounding the individual play an essential role in the generation of emotion. There are several ways in which other people elicit emotional responses in us. One is that we feel the emotions of those around us. Also, we have emotions about actions of those people around us. Another is we have emotions about the things that happen to other people. Yet another is our concern about our relationship with others that elicits emotion in us. The groups to which we belong can also elicit our emotions. Moreover, we can feel emotion about the success and failure of our own group or of other groups. In addition, groups or individuals may make salient cultural concerns or societal expectations that can elicit our emotions.

Beside the fact that social context can elicit emotions in individuals, social context provides information about what emotion should be expressed, by whom, and in what situations. For instance, people are well aware of the inappropriateness of expressing too much emotion to acquaintances [254]. However, the social knowledge of emotion expression is only partially delivered in an explicit fashion. There are studies on the regulatory role of society and social relationships on emotions,

showing that people's emotions become socialized in implicit and unconscious ways. From this perspective, social context can control and direct our attention toward certain types of events and away from others.

Humans are emotional and social beings. Their emotions and the social context in which they are involved have mutual impacts on each other. Humans can share their emotions with others just as they share their thoughts, resources and their environment. Sharing an emotion with others may alter the experience of an event. For instance, according to the nature of the relationship between the individuals, the expression of emotions can either restrain them from further interactions or improve their relationship. Furthermore, individuals sharing emotions might possess a shared understanding of their environment. Socially shared and regulated emotions also provide social meanings to the events happening in the environment [264]. For instance, people are likely to make social inferences based on the presence or absence of particular emotions in their social environment. Moreover, emotions can provide a basis for judgment depending on the individual's relationships with others. In other words, emotions can associate or disassociate an individual, therefore, they can change or maintain the individual's social relationships [254].

Emotions can also play the role of a motivator in a social context. There is a subset of social emotions delineated as role-taking emotions. In [233] Shott provides two categories of *reflexive* (e.g., shame or pride) and *empathic* (e.g., empathy or pity) role-taking emotions. Reflexive emotions can motivate the individual's self-control which depends on the anticipated reactions of others to the individual's behaviors. For instance, guilt might lead the individual to behave altruistically to restore a positive social stance for that individual. Empathic (or vicarious) emotions are based on an individual mentally placing himself in other's situation to understand how the other feels in that situation. These emotions motivate prosocial behaviors to maintain an individual's internal well-being [252].

2.3.3 Communicating Emotions

Humans need to communicate their emotions within a social context for different reasons. In [85] Goffman argues that human behaviors around others are performative; i.e., they are often intended to convey information to others. When a human's actions are visible in a social context, they behave differently [268]. The social life of an individual is comprised of the individual's internal cognitive competencies and his interactions in the society. Lazarus says, if society is a fabric, then emotion is its color [140]. Although emotions undeniably have personal aspects, they are usually experienced in a social context and acquire their significance in relation to this context [156].

There are several events that can elicit emotions in social contexts. For instance, during the interaction the cause of an emotion can be verbal (an utterance during conversation), nonverbal (someone's gesture), personal thoughts (interpretation of an event), or even emotions themselves (e.g., happiness for a partner's sense of pride). An utterance can include content and relational meaning. The content carries the information about the topic or the subject of the interaction, and the relational meaning reveals the meaning between the speaker and the hearer. An emotion might seem to be elicited by the content of the utterance, but in fact it is an individual's response to the relational meaning [188].

The interpretation of these relational meanings are handled by the appraisal of the events. Appraisal processes (see Section 4.3) also give us a way of viewing emotion as social [258]. Meaning is created by an individual's social relationships and experiences in the social world. Individuals communicate these meanings through utterances. Utterances in emotionally charged conversations, by their very nature, are supposed to inform the others about something novel. Novelty is an essential component of an event for appraisal. Conversations also possess the concept of consistency, because utterances with consistent meaning constitute the individual's underlying beliefs. Relevancy is another component of an event that can be assessed

by appraisal. The degree to which the individual's personal and mutual beliefs are strong and related controls emotionally rich social contexts. In other words, the more divergent the individual's beliefs, the more effort is required to converge (to be understood) which leads to more emotional responses in individuals. Human speech carries emotional information in the semantics and in the speech prosody. The semantics or the content of what an individual says includes obvious expression of emotion. The prosody holds more detailed emotional information by combining non-semantic cues in spoken language (e.g., rhythm and intonation) [154].

Interpretation of the events in a social context requires a baseline for the individual's assessment process. Goals, as the pillar of collaborative interactions, can provide this baseline for an individual. Goals are crucial in relational meanings of the events in a social context. The facilitation, interference and inhibition of goals are each correlated with certain type of emotions. In most conversations during collaboration, goals can be categorized into three different groups: goals related to accomplishing a task, goals to reveal one's personal beliefs, and goals to regulate one's social relationships [188]. For instance, for task-related goals, utterances related to accomplished tasks reveal joyful relational meaning; utterances related to impeded tasks reveal disappointing relational meanings which can lead to anger, and utterances related to tasks with no or little progress reveal the frustration of the individuals. Lastly, all these emotional responses in a social context will not only regulate or maintain individual's actions to reveal or hinder an intention, but also can control the way that action should be taken.

A successful and effective emotional communication necessitates ongoing reciprocal adjustments between interacting individuals that can happen by interpreting each other's behaviors [156]. It not only requires proper interpretation of the other's expressions, but also correct assessment of the extent to which others can read an individual's expressions. In emotional communication, individuals are constantly exchanging messages about their mental state, and modifying each other's emotional responses as they occur. Individuals perceive others' emotional states by processing

verbal and nonverbal messages during the interaction. Communication dynamics represent the temporal relationship between these communicative messages. The verbal and nonverbal messages from one participant are interpreted inside the context including the history and the ongoing messages from the other individuals. Interpersonal dynamics (also known as micro-dynamics in sociology) represent this influence of relationships between individuals [172].

2.3.4 Social Functions of Emotions

Humans are able to communicate their emotions in a social context. The social functions of emotions are the reason behind why humans try to communicate their emotions. Ekman in [70] asserts that the primary function of emotions is to mobilize the organism to deal with important interpersonal encounters. Darwin in [62] argues the significance of social communicative functions of emotions. Emotions describe interpersonal dynamics in a way that they can constitute individuals' relationships [183, 254]. One aspect of expressing and communicating emotion in a social context is to express one's social motives and intentions [110]. Another aspect of communicating emotions is to reveal the underlying mental state of an individual [184]. In other words, emotions constitute two different functionalities of expressing communicative signals associated with one's social motives and intentions as well as expressing one's internal states and how one feels about something. In [135] Van Kleef has discussed the idea of inferential processes with which individuals can infer information about others' feelings, relational orientations and behavioral intentions based on their emotional expressions. He also argues that emotional expressions can impact social interactions by eliciting others' affective responses.

Functional accounts of emotions vary according to the kind of system being analyzed. Thus, functional approaches to the emotions vary by level of analysis. Social functions of emotions can be analyzed in *individual*, *dyadic*, *group* and *cultural* levels. The focus of this research is on social functions in dyadic interaction (more specifically collaboration); these functions are also considered at the individ-

ual's level especially when interpreting the other collaborator's behaviors. Studies in all these levels share a few assumptions: a) individuals are social by nature and pursue solutions to survival problems in social relationships, b) individuals apply their emotions to coordinate their social interactions and relationships to address these survival problems, c) emotions are processes mediating the individuals' relations to their dynamic environment [129]. In dyadic interactions, studies focus on how emotions impact the interactions of individuals in meaningful relationships. In [129] Keltner and Haidt discuss that in a dyadic setting, researchers mostly focus on communication of emotion (e.g. Scherer [217], DePaulo [66]), properties (e.g. emotion contingency, emotion synchrony) of dyadic emotions (e.g. Levenson & Gottman [144]), discourse (e.g. Bretherton [36]), and attachments (e.g. Hazan & Shaver [107]).

Examples of Social Emotions:

There are many different types of emotions, only some of which are considered social, since they appear and provide meaning in social context. Here, we provide four examples of such emotions as well as their social functions to show how social functions of emotions impact individuals and the groups they belong to, and what causes them to be expressed by an individual.

Guilt – The function of guilt is to positively direct our behavior toward our group. We feel guilt when we hurt someone in our group, or when we fail to reciprocate care or kindness. Guilt motivates us to not hurt people in our group and to give back to others who have given to us, and in this way we strengthen the survival prospects of both the group and ourselves.

Shame – The function of shame is twofold. On the one hand, it keeps us within the rules and norms of society by informing us when we have done something dishonorable, disgraceful, or in some way condemned by our group. On the other hand, it informs the other members of our group that we know that we have dishonored

ourselves. The main difference between guilt and shame is that guilt is focused on a behavior, whereas shame is focused on ourselves.

Embarrassment – Embarrassment is related to shame, but includes some important differences. Embarrassment can only happen in public, whereas shame can happen when we are alone. We can feel embarrassment about very minor issues that have no moral implications, such as body odor, whereas shame typically concerns more grave issues with moral implications.

Pride – The function of pride is to reinforce when we or another person has done or represented something the group finds excellent. In this way, group values are reinforced and incentivized, which again helps the group to function better and motivates us to do things the group values. There is a negative form of pride in which our internal appraisal of our worth is inflated compared to the opinions of others, which is more correctly called hubris.

2.3.5 Artificial Emotions

Emotions, as an integral part of rational behavior, provide adaptive values for an artificial creature. They can control an agent's *attention* to focus on the most salient and relevant stimulus to solve the immediate problem. They can also help an agent to *monitor its own performance* so that the agent can make alterations on goals and plans. Emotions can act as a *memory filter* allowing a better recall of the events that are congruent with current cognitive and emotional states [30]. *Assisting the reasoning process* is another role of emotions; they assist the reasoning process by directing the cognitive information processes to the perceptual cues. Emotions impact the transformation of the agent's *decision-making behavior* [84] leading to a particular type of actions in a certain type of environment [272]. Emotions can *govern behavior tendencies* by providing immediate emotional responses, e.g., avoidance of elaborate reasoning because of lack of time or an unconcerned situation. Furthermore, emotions *provide support for social interactions* by helping the agent to understand others' behaviors as well as making expressions of the agent's internal

states more perceivable during the interaction [81].

The importance of these values of emotions for designing social agents having artificial emotions is apparent. However, the question is what problems are we facing in designing an effective social agent? In [63] authors discuss some of these problems and provide references speculating on the nature, function and mechanisms of emotions. Also, the importance of emotions and the incorporation of emotions in intelligent systems as well as implementation of emotions in several multi-agent systems are presented in [164]. Scheutz discusses the role of emotions in artificial intelligence and how we can determine the utility of emotions for the design of an artificial agent [225]. In [25] the authors present a definition and theory of artificial emotions viewed as a sequential process comprising the appraisal of the agent's global state; they also show how emotions are generated, represented and used in the Salt and Pepper architecture for autonomous agents. From the behavior perspective, appropriately timed and clearly expressed emotions are a central requirement for believable agents [19].

There are several architectures modelling emotions for the purpose of enhancing the believability and effectiveness of the agents and robots. But the question is how do we model emotions? Hudlika in [117] deconstructs the concept of emotion modelling into: (a) fundamental categories of processes for emotion generation and emotion effects, and (b) identification of some of the fundamental computational tasks required to implement these processes. These building blocks can be helpful as a guideline for the systematic development of new computational models, or for the assessment of existing computational models of emotions as discussed in [148] and [163]. There are also logical formalizations of emotions and emotional attitudes (including speech acts) and corresponding mental state to provide a systematic analysis of computational models of emotions [1, 88, 104].

From another perspective, the necessity of employing emotions in robotics and more specifically social robots has been argued in [186] and [242]. Social robotics and cognitive robotics have many overlapping concepts, especially when they focus

on interaction between a robot and a human. The relationship between cognition and emotion receives more attention due to the mutual influences they have on each other [173, 228]. For instance, in [81] the authors employ emotions in the learning procedure of a robot, and in [38] and [224] authors discuss the importance of emotions in the action selection procedure of an agent or a robot, impacting the behavior arbitration and self-adaptation mechanisms. Ultimately, employing artificial emotions will impact the context of human-robot/computer interaction [115] and how humans and robots understand each other’s emotions in a social environment [132, 176]. In [207] the authors selected twelve autonomous agents that incorporate an emotion mechanism into the action selection procedure to compare. They introduced a framework based on correlations between emotion roles performed and aspects of emotion mechanisms used to perform those roles. Gratch and Marsella also present a method to evaluate a computational model of emotion in [90] which compares behavior of the model against human behavior.

2.3.6 Cognitive Architectures

There are several integrated cognitive architectures that try to produce all aspects of behavior as a single system in different domains [126, 137]. Surveys on the comparison of underlying philosophy and functional description of the most prominent cognitive architectures introduce several criteria to evaluate such architectures [47, 139, 253]. The necessity of integrating these cognitive architectures into robots has been discussed from the perspective of developmental psychology [10, 67, 128]. There are also many examples emphasizing the importance of cognitive robotics from this perspective, while some of them also incorporate the concept of affect in their design [40]. Some of these cognitive architectures are biologically inspired, e.g., *eBICA* [211], or [20] and [199], while some others are inspired by psychological theories, e.g., *ACT – RΦ* [61], or [170] and [68].

2.4 Computational Models of Emotions

There are different types of computational theories of emotion such as appraisal and dimensional theories. These theories differ in the type of relationships between their components and whether a particular component plays the crucial role in an individual emotion. For instance, the basic component of an emotion can be the behavioral tendencies, the cognitive elements, or the somatic processes. Emotion theories can also differ based on their representational distinctions.

2.4.1 Appraisal Theory

Appraisal theories of emotion were first formulated by Arnold [9] and Lazarus [140] and then were actively developed in the early 1980s by Ellsworth and Scherer and their students [206] [212] [216] [221] [223]. In general, the emotional experience is the experience of a particular situation [79]. Appraisal theory describes the cognitive process by which an individual evaluates the situation in the environment with respect to the individual’s well-being and triggers emotions to control internal changes and external actions. In this section, we are going to review sequential and structural approaches incorporating the appraisal concept.

Componential Approach

This approach emphasizes the distinct components of emotions, and is often called the *componential* approach [145]. The “components” referred to in this approach are the components of the cognitive appraisal process. These are referred to as *appraisal variables*, and include *novelty*, *valence*, *goal relevance*, *goal congruence*, and *coping abilities* (later in this section some of the appraisal variables used in computational models are introduced) [216, 223]. A stimulus, whether real or imagined, is analyzed in terms of its meaning and consequences for the agent, to determine the affective reaction. The analysis involves assigning specific values to the appraisal variables. Once the appraisal variable values are determined by the organism’s eval-



Figure 2.2: Schematic view of the componential theory of emotion [118].

uative processes, the resulting vector is mapped onto a particular emotion, within the n-dimensional space defined by the n appraisal variables. The semantic primitives for representing emotions within this model are thus these individual appraisal variables. Figure 2.2 shows the relationship of the individual appraisal dimensions to the broader categories of evaluations taking place during appraisal (Relevance, Implications, etc.).

Component Process Model

The Component Process Model (CPM) is Scherer's influential appraisal theory of emotions [218, 223]. This theory focuses on the dynamic unfolding of emotions. The CPM suggests that an event and its consequences are appraised with a set of criteria on multiple levels of processing (the appraisal component). The result of the appraisal will generally have a motivational effect, often changing or modifying the motivational state (see Section 2.5) before the occurrence of the event. Based on the appraisal results and the motivational changes, some effects will occur in the

autonomic and somatic nervous system. The CPM considers emotions as the synchronisation of many different cognitive and physiological components. Emotions are identified with the overall process whereby low level cognitive appraisals, in particular the processing of relevance, trigger bodily reactions, behaviours and subjective feelings. The model suggests that there are four major appraisal objectives required to adaptively react to a salient event [220]:

- a) **Relevance:** How relevant is this event for the agent? Does it directly affect the agent or its social reference group?
- b) **Implications:** What are the implications or consequences of this event and how do they affect the agent's well-being and its immediate or long-term goals?
- c) **Coping Potential:** How well can the agent cope with or adjust to these consequences?
- d) **Normative Significance:** What is the significance of this event for the agent's self-concept and for social norms and values?

To attain these objectives, the agent evaluates the event and its consequences on a number of criteria or *Stimulus Evaluation Checks* (SECs), with the results reflecting the agents subjective assessment of consequences and implications on a background of personal needs, goals, and values [223]. Figure 2.3 shows the postulated sequence, the cognitive and motivational inputs and the effects on response systems. Also, the bidirectional effects between appraisal and other cognitive functions are illustrated by the arrows in the upper part of Figure 2.3.

Appraisal Process

According to appraisal theory, appraisals are separable antecedents of emotion, that is, the individual first evaluates the environment and then feels an appropriate emotion [223]. The appraisal procedure begins with the evaluation of the environment



Figure 2.3: Comprehensive illustration of the CPM of emotion [220, 223].

according to the internalized goals and is based on systematic assessment of several elements [218]. The outcome of this process triggers the appropriate emotions. In many versions of appraisal theory, appraisals also trigger cognitive responses often called *coping strategies*. In fact, the coping mechanism manages the individual's action with respect to the individual's emotional state and the existing internal and/or external demands [77]. The majority of computational models of emotions take this approach. An individual can also use knowledge about the emotional reactions of others to make inferences about the others. According to the appraisal patterns, different emotions can be experienced and expressed. Since expression of emotions reflects one's intentions through the appraisal process, the *reverse appraisal* mechanism helps one to infer others' mental state based on their expressions. [64, 106].

The appraisal process is typically viewed as the cause of emotion and the cognitive and behavioral changes associated with emotion. For instance, a particular

pattern of the appraisal variables (i.e., individual judgements) will elicit a certain emotion or emotional expressions. Some of the (computational) appraisal variables include [162]:

- **Relevance:** A relevant event has non-zero utility for an agent. This relevancy can either be based on a negative influence of an event on the agent or a positive one.
- **Perspective:** The point of view in which an event will be judged, e.g. self or other.
- **Desirability:** A desirable event advances the utility for an agent whose perspective is being taken, or if it is an undesirable event, inhibits that.
- **Likelihood:** A measure of how likely is the outcome.
- **Expectedness:** The extent to which the truth value of a state could have been predicted from causal interpretation.
- **Causal Attribution:** The agent who deserves the credit/blame.
- **Controllability:** Whether the outcome can be altered by the agent whose perspective is taken (this variable is related to the coping process).
- **Changeability:** Whether the outcome can be altered by some other causal agent (this variable is related to coping process).

Coping Process

Another key process involved in appraisal is coping. This process determines whether and how the agent should respond with respect to the outcome of appraising the events. There are several coping strategies that computational models such as EMA [92] use as control signals. These control signals enable or suppress the cognitive processes that operate on the causal interpretation of the appraisal patterns. The coping

process controls the congruency of the actions according to these patterns. As it is shown below, coping strategies are organized into two categories: *problem-focused* and *emotion-focused*. Problem-focused coping strategies can be applied when the agent must do something with respect to the problem, whereas emotion-focused coping works by changing one's interpretation of circumstances. The following is a short list of a broad range of coping strategies [92]:

Problem-focused coping

- **Active coping:** Taking steps to remove or circumvent the stressor,
- **Planning:** Coming up with action strategies,
- **Seeking social support for instrumental reasons:** Seeking advice, assistance, or information.

Emotion-focused coping

- **Seeking social support for instrumental reasons:** Getting sympathy, moral support or understanding,
- **Acceptance:** Accepting the stressor and learning to live with it,
- **Restraint coping:** Waiting till the appropriate opportunity (holding back).

OCC, a Structural Appraisal Model of Emotion

OCC (Ortony, Clore and Collins) model, similar to Lazarus' [141] and Scherer's [216] cognitive views, considers emotions to arise from affective or valenced reactions subsequent to the appraisal of a stimulus as being beneficial or harmful to one's concern [178]. OCC model categorizes emotions based on their underlying appraisal patterns. These patterns are fundamental criteria a person employs for evaluating a situation. They involve the person's focus of attention, her concern, and her



Figure 2.4: A simple visualization of OCC model [178].

appraisal preceding an affective reaction. Figure 2.4 shows the main building blocks of OCC model.

As shown in Figure 2.4, a person could alternatively have three types of focus. These types of focus are the consequence of events, actions of agents, and aspects of objects. A person evaluates the significance of the causes behind these three types of focus based on her personal concerns. As a result, an affective reaction will be elicited, resulting in an emotion. Various combinations of the elements depicted in Figure 2.4 create specific patterns resulting in six main groups of emotions in which all emotion types in a group share the same cognitive pattern (see Figure 2.5). Emotion groups are *fortune-of-others*, *prospect-based*, *well-being*, *attribution*, *well-being/attribution- compound*, and *attraction*. The OCC model introduces 22 emotion types. Each of these emotions is introduced as a representative of a family of similar emotions with various intensities (since relying on a list of discrete emotions that is understood by everyone equally is impossible due to people's language barriers and various interpretations of the actual words)¹. For instance, while they all share the same eliciting conditions, happiness can be referred to by many other emotion terms such as joy, cheerfulness, gladness and delighted. Thus the emotion types used in

¹In stark contrast to basic emotion theories discussed in Section 2.4.2.

the model (e.g., relief, love, pride, and shame) are meant to represent an emotional experience rather than a lexical taxonomy.

For instance, as shown in Figure 2.4, the appraisal criterion for consequences of events is their *desirability* for achieving one's goals. This generates the affective reaction of being *pleased* in positive cases, or *displeased* in negative ones. Figure 2.5 shows the resulting emotion groups in OCC model such as *fortune-of-others* (e.g., gloating, pity), *prospect-based* (e.g., satisfaction, relief), and *well-being* (e.g., joy, distress) [178]. The appraisal of the praiseworthiness of the actions of an agent against one's personal standards, as well as the appealing aspects of objects happens in the same way as shown in Figure 2.4.

Finally, the OCC model introduces some global variables of an emotion's intensity to distinguish all types of emotions that a person could experience when encountering events, agents or objects. These variables are as follows

1. Sense of reality (representing the degree to which the event, agent or object in focus appear real to the person),
2. Proximity variable (representing the psychological closeness of an event, agent or object),
3. Unexpectedness (representing how surprising an event is for one, either positive or negative),
4. Arousal (representing how stimulating an event, agent or object is).

2.4.2 Other (Non-Appraisal) Computational Models

Constructivist (Dimensional) Emotion Theories

The components and dimensions of emotions have been the subject of much speculation since the 19th century. Dimensional models of emotion attempt to conceptualize human emotions by defining where they lie in two or three dimensions. Dimensional

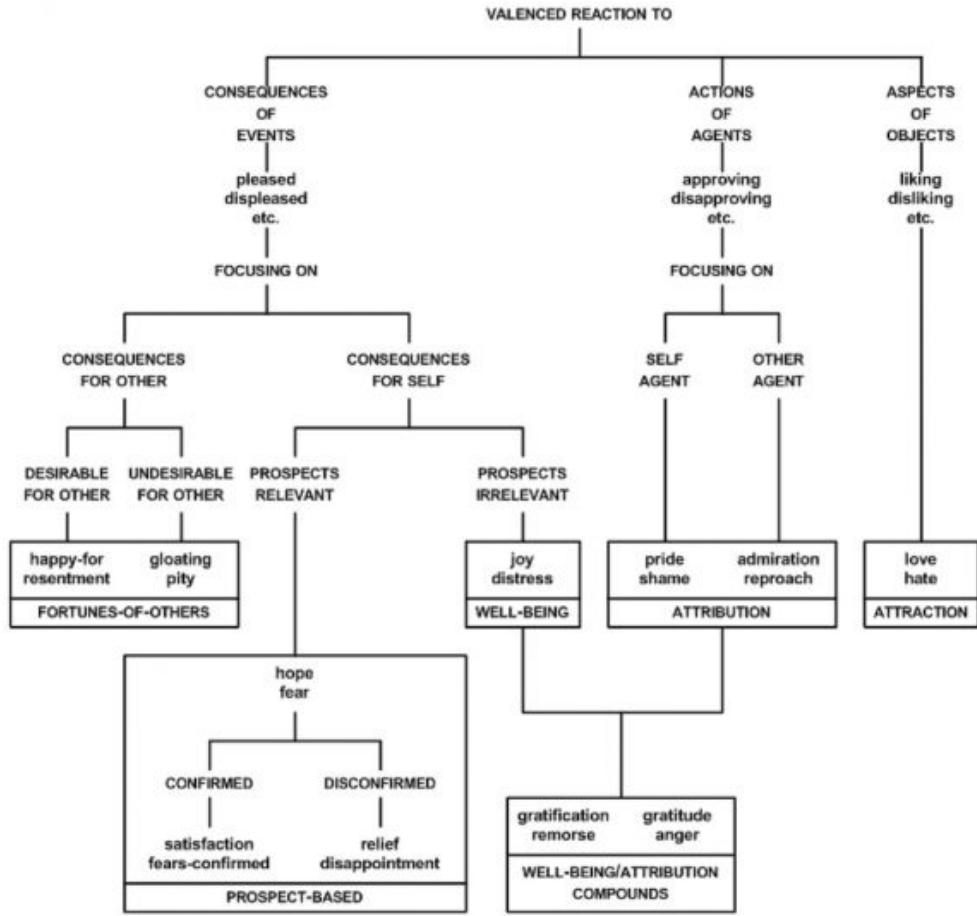


Figure 2.5: OCC taxonomy of emotion triggers and emotions [178].

theories of emotion argue that emotions should be conceptualized as points in a continuous dimensional space, rather than looking at them as discrete entities [43] [168] [209] [261].

Two dimensions that are commonly proposed to describe emotions are valence and physiological arousal [9] [141] [208]. Models based on dimensional theories introduce the emotion concept as a non-relational construct summarizing a unique overall state of the individual. The models based on dimensional theories contrast theories of basic emotion (see Section 2.4.2), which propose that different emotions

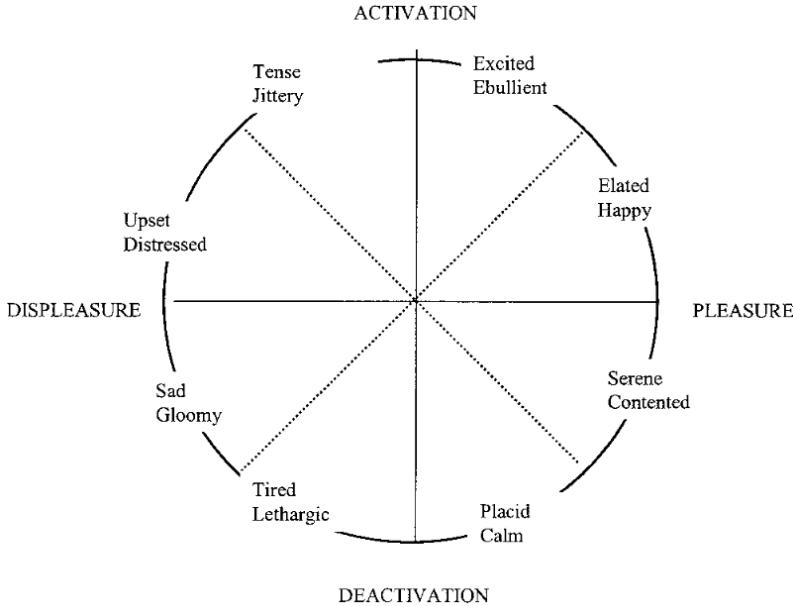


Figure 2.6: Russell’s suggested affective states based on core affect [209].

arise from separate neural systems [192]. Also, models based on dimensional theories contrast appraisal theories, which propose that appraisals are the relational constructs characterizing the relationship between one’s mental state and a specific stimulus. Many dimensional theories argue that discrete emotion categories (e.g., sadness, fear and anger) have no “reality” in that there are no specific brain regions or functions that correspond to specific emotions [18]. Therefore, dimensional theories do not emphasize the term emotion.

One of the most prominent two-dimensional models is Russell’s circumplex model [208]. Russell suggested that affective states are all related to each other systematically through what is called core affect [208, 209] (see Figure 2.6) and emotions are best described as a change in core affect which, in turn, is describable as a point in a space of two dimensions. One dimension is *valence* or how good or bad objects and events are for a being (ranging from pleasant to unpleasant). The other dimension is *arousal*, ranging from calm to excited. Russell put a number of affective states around a circular space between those two dimensions (see Figure 2.6) which is also known as *circumplex*, representing the variety of core affects [208, 209]. Since some-

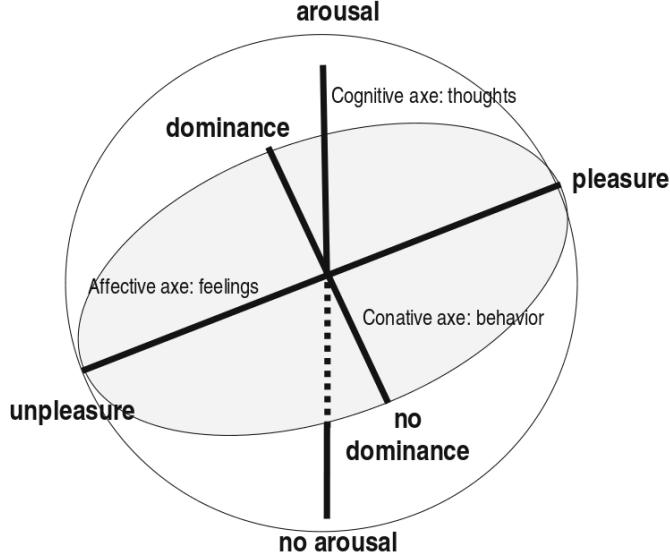


Figure 2.7: Three dimensional model of pleasure, arousal and dominance as tripartite view of experience [17].

times two-dimensional space cannot easily differentiate among emotions that share the same values of arousal and valence, e.g., anger and fear (both characterized by high arousal and negative valence), some of the dimensional models incorporate valence and arousal as well as *intensity*, or *dominance* or *stance* dimensions. Many computational dimensional models build on the three dimensional PAD model of Mehrabian and Russell [168] where these dimensions correspond to pleasure (a measure of valence), arousal (indicating the level of affective activation) and dominance (a measure of power or control). Figure 2.7 shows these three dimensions.

Basic (Discrete) Emotion Theories

Basic emotion theories are inspired by Tomkins' [255] rediscovery of Darwin's work [62, 110] which were later developed further by Ekman [70] and Izard [120]. These theories emphasize a small set of discrete and fundamental emotions. The underlying assumption of this approach is that these emotions are mediated by associated neural circuitry, with a hardwired component [70]. Different emotions are then characterized by stable patterns of triggers, behavioral expression, and associ-

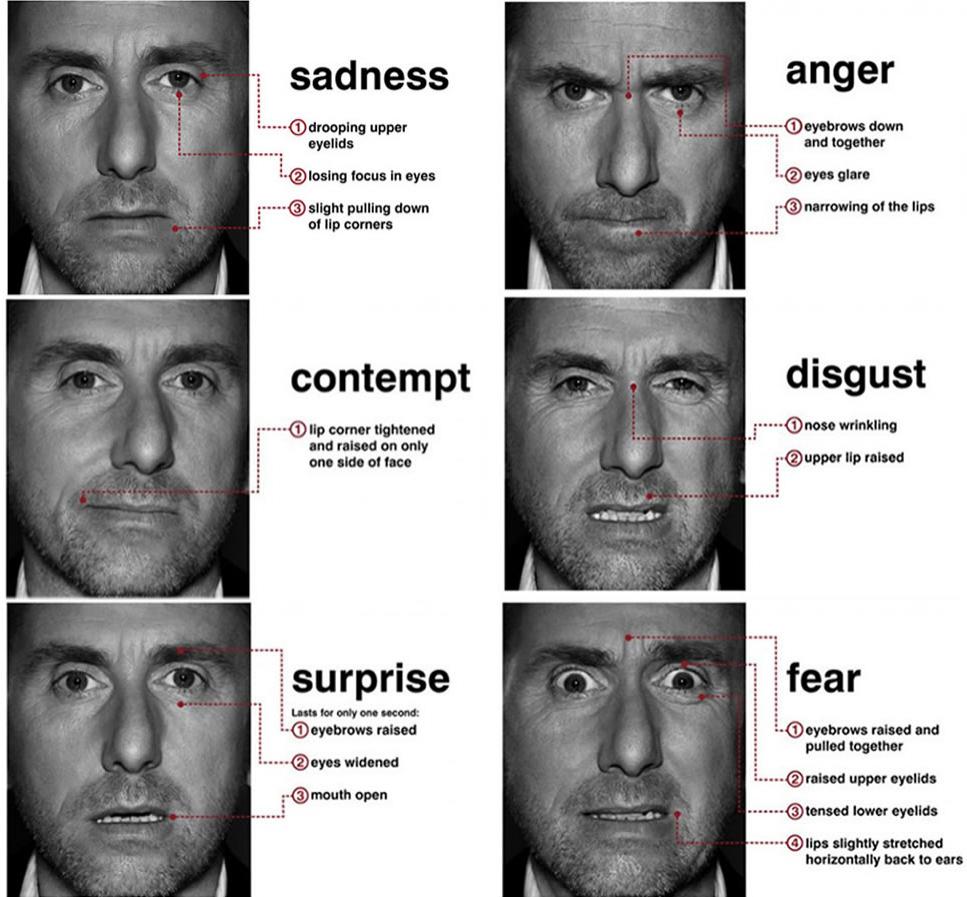


Figure 2.8: Basic emotions and corresponding expressions.

ated distinct subjective experiences. The emotions addressed by these theories are typically called the *basic* emotions. Emotions including happiness, sadness, fear, anger, surprise, and disgust are often considered to comprise the most prototypical basic emotions [70]. The theory of basic emotions holds that there is a set of emotions shared by all humans that evolved to deal with ancestral life challenges [70]. For instance, disgust evolved to address the challenge of avoiding noxious stimuli, and fear evolved to address the challenge of avoiding dangers. Because of the emphasis on discrete categories of states, this approach is also termed the *categorical* approach [182]. Much of the supporting evidence offered for the theory comes from experiments that show how certain facial expressions are universally associated with specific basic emotions, regardless of the observer's cultural background.

This universality has a production side and a recognition side. On the production side, a particular emotional state is said to elicit a facial expression comprised of a specific set of facial muscles. On the recognition side, observers are able to infer the emotional state of the person who expresses an emotion, due to the direct correspondence between emotional states and the facial expressions they cause. Computational models inspired by the basic emotions or discrete approach often focus on low-level perceptual-motor tasks and encode a two-process view of emotion that argues for a fast, automatic, undifferentiated emotional response and a slower, more differentiated response that relies on higher level reasoning processes (e.g., [8]).

There are other approaches that different researchers take based on their emphasis on the applicability of emotions in their systems.

Rational Approaches

Rational approaches start from the question of what adaptive functions emotions serve and then attempt to incorporate these functions into a model of intelligence. Emotion, within this approach, is simply another set of processes and constraints that have adaptive value. Models of this sort are most naturally directed towards the goal of improving theories of artificial intelligence [4] [227] [238].

Communicative Approaches

Communicative theories of emotion argue that emotion processes function as a communicative system. They can function first, as a mechanism for informing other individuals of one's mental state (thereby facilitating social coordination), and second, as a mechanism for requesting/demanding changes in the behavior of others. Communicative theories emphasize the social-communicative function of expressions [89]. Computational models inspired by communicative theories focus on machinery that decides when an emotional expression can have a desired effect on a human

counterpart.

2.4.3 Similarities and Differences

Different theoretical perspectives should not be viewed as competing for a single truth. Instead, they should be seen as perspectives arising from particular research areas (e.g., biological vs. social psychology), focusing on different sets of affective phenomena, considering different levels of resolution and fundamental components (e.g., emotions vs. appraisal variables). These different perspectives provide different degrees of support for the various processes of emotion, e.g., the componential theories provide extensive details about cognitive appraisals [118]. Therefore, this section provides a pairwise comparison between these fundamental theories. Note that a separate pairwise comparison will not be provided for appraisal vs. discrete (basic) emotion theories as the important points are adequately covered in the comparisons presented below.

Dimensional Vs. Discrete (Basic) Emotion Theories

The fundamental assumption of the basic emotion theory is that a specific type of event triggers a specific affect program corresponding to one of the basic emotions and producing characteristic expression patterns and physiological response configurations [221]. Dimensional theories' main criticism of basic emotions theory is based on the observation that affective phenomena appear to be both qualitatively and quantitatively diverse.

Russell in [209] argues the labels such as “fear”, “anger”, “happiness” do not capture this diversity. For instance, one might say: a) a person being chased by an assailant brandishing a knife, b) a person who retreats from an insect moving across the floor, and c) a person who is concerned they will never find a fulfilling career, are all in a state of fear. On the basic emotions account, an emotional episode involves fixed patterns of neurophysiological and facial expression changes in response to

an eliciting stimulus that are distinct between emotions, but are the same within the same emotional category [70]. If this were the case, one would expect that the three individuals described above would respond to their eliciting stimuli in the same way, yet a similarity of behavioral responses between these three cases seems unlikely. Dimensional theorists, in contrast, would argue that the individuals in the above three cases are applying the concept of fear to experience, despite the fact that each individual has a unique core affect. While basic emotion theorists would hold that since all three individuals are experiencing fear, they would perform the same behavioral responses to the stimuli, dimensional theorists would argue this is not the case, as each individual bears a core affective state that is distinguished from the other two. For instance, the individual's arousal in response to an armed assailant should be higher than the individual in response to an insect, as the former case poses a threat to their life. As a result, the individual in the first case would likely make every effort to escape from the assailant, including trying to negotiate and plead with the assailant, while the individual in the second case would be relatively less dedicated to escaping the insect.

In sum, a dimensional theory is compatible with the differences in the behavioral responses to eliciting stimuli, while basic emotions theory only allows for a single fixed behavior of responses to a given emotion. Furthermore, dimensional theories can represent instances of basic emotions (see Figure 2.9), for example, fear elicited by a snake (green rectangle), in terms of variation along affective dimensions, i.e., arousal and valence.

Also, basic emotion theory fails to account for affect that lacks object-directedness [209]. In the basic emotions approach, an emotion is supposed to have an intentional object it is directed towards (e.g., being angry at someone, or being sad for someone). The dimensional theory argues that emotion may not necessarily be aimed at a particular object. For instance, an individual can experience a certain type of emotion (e.g., anger) without knowing of anything in particular that has offended her. Dimensional models of emotion are therefore capable of accounting for a wider

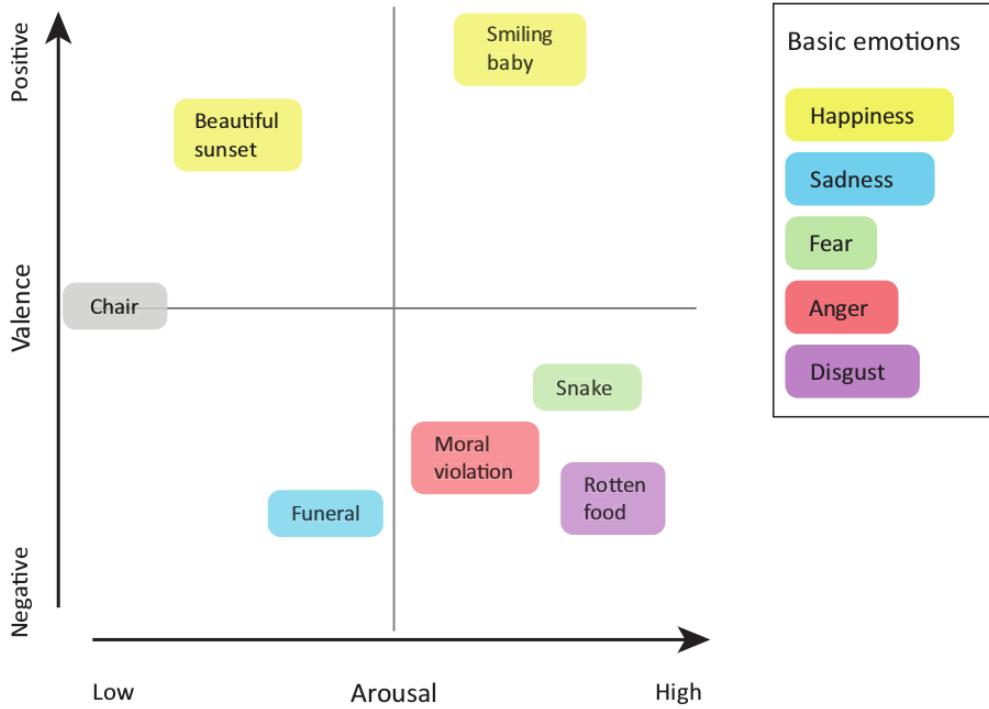


Figure 2.9: Representing basic emotions within a dimensional framework [105].

range of affective phenomena than basic emotions theory.

Another difference between dimensional and basic emotion theories is that the basic emotion categorization of emotions captures facets of the experience of an emotion not conveyed by the dimensional description, such as elicitation of a facial expression of the emotion. In fact, this attribute of the basic emotions theory is one of the major differences with all other emotion theories. It is argued in basic emotion theory that basic emotions are hard-wired to their corresponding facial expressions. Ekman, who elaborated the concept of basic emotions, developed the *Facial Action Coding System* (FACS) which encodes movements of individual facial muscles and it is a common standard to systematically categorize the physical expression of emotions [71].

Appraisal Vs. Dimensional Emotions Theories

Dimensional theories struggle to adequately distinguish emotions because of the existence of limited dimensions.

To compare the appraisal and dimensional theories of emotion, we argue that there is a relationship between the dimensions in the dimensional theories of emotion and the appraisal dimensions. For instance, the pleasure dimension roughly maps onto appraisal dimensions that characterize the valence of an appraisal-eliciting event, (e.g., intrinsic pleasantness –desirability–, or goal congruence), dominance roughly maps onto the appraisal dimension of coping potential, and arousal can be considered as a measure of intensity. However, these dimensions and corresponding appraisal variables have quite different meanings. Appraisal (as mentioned earlier) is a relational construct characterizing the relationship between some specific object/event in the environment and the individual's mental constructs including beliefs, motives and intentions. Also, several appraisals may be simultaneously active, whereas emotions in dimensional emotion theory are non-relational constructs, each summarizing a unique overall state of the individual.

Furthermore, dimensional emotion theories emphasize different components of emotion than appraisal theories and link these components quite differently. In contrast to appraisal theories, dimensional emotion theories do not address affect's antecedents in detail. Dimensional theorists question the tight causal linkage between appraisal and emotion that is central to appraisal accounts. As mentioned earlier, dimensional theorists believe that the emotion is not necessarily about some object (as in “I am angry at him”). In such theories, many factors may contribute to a change in emotion including intentional judgments (e.g., appraisal). However, in dimensional emotion theories, the link between any preceding intentional meaning and emotion is broken and most of the time cannot be recovered correctly. For example, Russell argues for the following sequence: some event occurs (e.g., a bear walks out of the forest), it is perceived in terms of its affective quality; this per-

ception results in a crucial change in core affect; this change is attributed to some “object” (e.g., the bear); and only then is the object cognitively appraised in terms of its goal relevance, causal antecedents and future prospects [161].

We can also compare dimensional emotion theories to the OCC model as a cognitive appraisal model. The major similarity between these two models is that they both consider emotions to descend from valenced reactions to the stimuli. Furthermore, they acknowledge the role of arousal in determining emotional reactions. As mentioned in Section 2.4.2, Russell considered arousal as one of the two key dimensions of emotions which could be used to partially discriminate emotional states [208]. In a different manner, the OCC model recognizes arousal as a necessary condition for eliciting emotions, and regards arousal as a major determinant of the elicited emotion’s intensity which distinguishes among various emotions of a particular type (e.g., fearful versus scared). In [219] Scherer speculates that the arousal dimension in dimensional models gives little information about the underlying appraisal of the elicited emotion and proposes to replace it with coping potential, which is an appraisal dimension referring to the individual’s perceived control in a given situation.

Models based on dimensional emotion theories pursue the idea of eliciting an emotion according to the joint features in circumplex space (2D or 3D – see Section 2.4.2), while OCC or other models of appraisal theory are based on patterns of antecedents of emotions. This is the fundamental difference between OCC, or appraisal theories in general, and the circumplex approach of Russell [208] or Mehrabian’s PAD model [17, 168]. Also, models based on appraisal theory of emotions employ causation, attribution and eliciting conditions in order to distinguish emotions, while the eliciting conditions are not directly accessible from a dimensional approach. A dimensional model may fall short in establishing why certain emotions are elicited. However, when the objective is to identify the generated emotions and their level of pleasantness and intensity, a circumplex model presents an excellent opportunity [3].

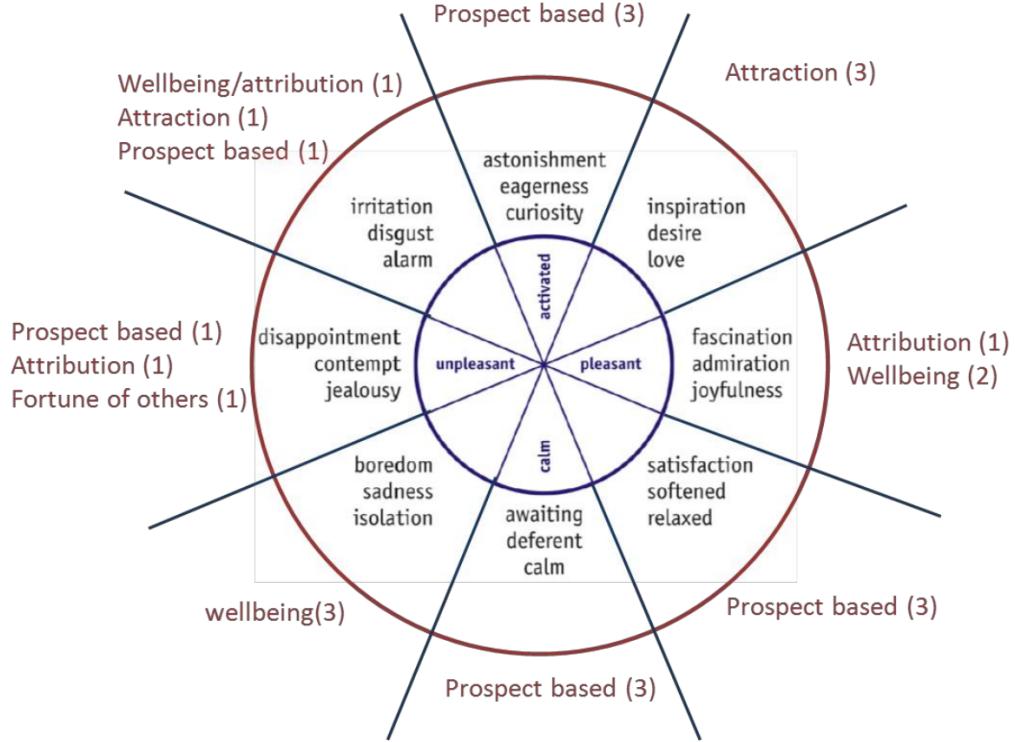


Figure 2.10: A rough projection of emotion groups of OCC on the circumplex of affect [3].

Finally, we consider how a model based on dimensional emotions (i.e., Russell's 2D circumplex) relates to a cognitive model based on appraisal theory (i.e., OCC). Figure 2.10 shows the relationship between Russell's circumplex and the OCC model in terms of the categorization of actual emotions. The number of emotions in a section of Russell's circumplex that fall into an emotion group of OCC is shown in parentheses. For instance, all three emotions in the top section (highly excited, neutrally valenced emotions) fall into prospect-based emotion group. Or, as another example, emotions in the left section (neutral arousal value, negative valenced emotions) make a one to one relationship between disappointment and the prospect based emotion group, contempt and attribution emotion group, and jealousy and fortune of others emotion group (hence number (1) is indicated in front of each).

2.4.4 Applications in Autonomous Agents and Robots

There are many research areas, including robotics and autonomous agents, that employ the structure and/or functions of emotions in their work with a variety of reasons for modeling emotions [262]. Some of this work is inspired by specific psychological theories, while some are freely using the concept of emotion without a theoretical grounding in social sciences; some are using a combination of concepts from different psychological theories. For instance, in PECS [257], which is designed for modeling human behaviors, the agent architecture is not based on any specific social or psychological emotion theory. In fact, it is intentionally designed and described in a way which enables the integration of a variety of theories. The PECS' design enables an integrative modeling of physical, emotional, cognitive and social influences within a component-oriented agent architecture. Also, in [165] the computational architecture, which is designed to provide information about the possible overall behavior of a work team, is not based any specific theory. As mentioned earlier, some researchers apply combinations of emotion theories in their work [133]. For instance, in [41] Cañamero shows how an agent can use emotions for activity selection while taking into account both dimensional and discrete approaches in an action selection mechanism.

We can also see the application of emotion theories in designing companion robots capable of expressing emotions and social behaviors, as well as robots which can convey certain types of emotion products, e.g., empathy [33] [142] [181] [232]. Robots also use emotion theories for automatic affect recognition using different modalities [109] [270]. Moreover, in some work, researchers have explored the user's affective state as a mechanism to adapt the robot's behaviors during the interaction [32] [151].

Applications of Appraisal Theory – The emphasis of models derived from appraisal theories of emotion is on making appraisal the central process. Compu-

tational appraisal models often exploit elaborate mechanisms for deriving appraisal variables such as decision-theoretic plans [92] [162], reactive plans [196] [201] [245], Markov-decision processes [72] [234], or detailed cognitive models [158]. However, emotion itself is sometimes treated less elaborately, simply as a label to which behavior can be attached [74]. Appraisal is usually modeled as the cause of emotion being derived via simple rules applied to a set of appraisal variables.

Computational appraisal models have been applied to a variety of uses including contributions to psychology, robotics, AI, and HCI. For instance, Marsella and Gratch have used EMA [162] to generate specific predictions about how human subjects will appraise and cope with emotional situations and argue that empirical tests of these predictions have implications for psychological appraisal theory [91] [160]. There are other examples in artificial intelligence and robotics of applying appraisal theory [2] [130] [162]. In robotics, appraisal theory has been used to establish and maintain a better interaction between a robot and a human. For instance in [130] researchers use a computational model of emotion generation based on appraisal theory to have a positive human-robot interaction experience. In [212] the authors describe a system approach to appraisal processes based on Scherer's work on appraisal and the CPM [216]. They show how the temporal unfolding of emotions can be experimentally tested. They also lay out a general domain-independent computational model of appraisal and coping. In [260] researchers consider their robot, INDIGO's, emotion, speech and facial expressions as key features to establish effective communication between the robot and a human during their interaction. They apply concepts of appraisal theory in INDIGO's emotion modeling. MAGGIE, a sociable robot, also applies an appraisal theory of emotions to consider fear in its decision making system [86]. Velasquez developed Cathexis, which is a distributed computational model for generation of emotions and their influence in the behavior of the autonomous agents [259]. The emotion model in this work is based on Roseman's work on appraisal theory. Marinier and Laird in [157] focus on the functional benefits of emotion in a cognitive system. In this work, they integrate their emotion

theory (which is based on appraisal theory) with the Soar cognitive architecture, and use emotional feedback to drive reinforcement learning. In [116] Hudlicka provides a generic mechanism mediating the affective influences on cognition based on cognitive appraisal. This model is implemented within a domain-independent cognitive-affective architecture (MAMAID).

In the virtual agents community, empathy is a research topic that has received much attention in the last decade [30] [167] [179] [193] [250]. In [191] researchers developed an agent with the capability of affective decision-making based on appraisal theory to establish a relationship with its users. Then, they compared the performance of their agent with a human (based on a WoZ study) in a speed-dating experiment. In HCI, appraisal theory has been primarily used for the creation of interactive characters that exhibit emotions in order to make characters more believable [200], more realistic [155] [256], more capable of understanding human motivational states [58] or more able to induce desirable social effects in human users [180].

Applications of Dimensional Theory – The emphasis of models influenced by dimensional theories is on processes associated with core affect which is usually represented as a continuous time-varying process, and can be represented at a given time by a point in a 2D or 3D-space as a response to the eliciting events. Generally, there are detailed mechanisms in computational dimensional models which determine how this point changes over time, e.g., decay to some resting state, and incorporating the impact of dispositional tendencies such as personality or temperament [83] [161]. Models based on dimensional theories have also been used in robotics. For instance, researchers in [147] apply PAD’s three-dimensional space to rate the pleasure, arousal and dominance of their Multimodal Emotional Intelligence robot (MEI) in each interaction with human subjects by understanding and expressing emotions in voice, gesture and gait. In [271] researchers want to understand the effect of different interface features for a service robot. They use valence and arousal dimensions in their questionnaires to assess the perceived anthropomorphism of their

service robot by subjects. In [136] researchers introduce the implementation of a dynamic personality for a robot based on a dimensional emotion model. They use WASABI’s architecture [22, 23] as their emotional model. In [149] Lisetti describes an affective knowledge representation scheme to be used in the design of a socially intelligent artificial agent. Lisetti uses the valence-arousal two dimensional model of emotion in this work. This model has been applied in an emotion-based architecture of Lisetti’s autonomous robots as well as a multimodal affective user interface agent. ROMAN, an expressive robotic head, uses a behavior-based emotional control architecture. The approach to the emotional component of the architecture is based on the dimensional emotion theory [113].

Comparison of Applications of Emotion Theories – Researchers often use computational dimensional models for behavior generation of animated characters. The reason might be because it is easier to translate emotions to a limited number of dimensions that can be readily mapped to continuous features of behavior such as the spatial extent of a facial expression. For example, PAD models describe all behavior in terms of only three dimensions of pleasure, arousal and dominance, whereas researchers using appraisal models would need to either associate each behavior with a large number of appraisal variables [222] [243], or try to map appraisal variables into a limited and small number of discrete expressions [74]. For a similar reason, dimensional models are also frequently used as a representational framework for systems that attempt to recognize human emotional behavior. There is some evidence that they may better discriminate user affective states than approaches that rely on discrete labels [18].

There is also a relationship between dimensional and appraisal theories. Some of the computational models of emotion that incorporate dimensional theories have viewed appraisal as the mechanism that initiates changes to core affect. For instance, ALMA [83] includes OCC inspired appraisal rules [178], and WASABI [22] includes appraisal processes inspired by Scherer’s sequential-checking theory into a PAD-based emotion model. Moreover, some computational models explore how core affect

in dimensional models can influence cognitive processes. For example, HOTCO 2 [251] allows explanations to be biased by dimensional affect [161].

2.5 Affect and Motives

Motives are essential mental components in decision-making procedures and applying them in an affect-driven collaborative agent is part of this thesis' contribution. In this section, we review related work on computational models of motivation and discuss the nature of motives. We also explain three of the important social motives which will be used in our work. Finally, we discuss how humans' beliefs, emotions and motives are related and influence each other.

Motivation principles and mechanisms, as the reasons behind one's intentions and actions, and the influences of motives on cognition have been discussed in philosophy, neuroscience, psychology and artificial intelligence [15, 24, 37, 238, 240]. There are several examples in AI of computational models for different psychological theories of motivation. Bach's MicroPsi agent architecture describes the interaction of emotion, motivation and cognition of agents based on Dietrich Dörner's Psi theory [13, 14, 15, 16]. Merrick and Shafi provide a computational model for motivation based on Henry Murray's theory [174] describing the three important social motivations of *achievement*, *affiliation* and *power*. They focus on the role of motivation in a goal-selection mechanism [169]. There are other examples focusing on the impact of motives on different cognitive processes in robots and artificial agents [31, 46, 65, 231, 259, 265]. The motivation mechanism in our work is also inspired by Murray's theory and Bach's approach on Dörner's theory. It is focused on the role of motives in cognitive processes, e.g., intention formation in coping, during collaboration, which will be discussed in Chapters 3 and 4. Motives play an important role in maintaining the congruency of intentions due to the fact that they are regulated by emotions (which are the outcome of one's internal states, e.g., beliefs), and oriented by goals (which are the desired outcome one wants to achieve).

2.5.1 Motives

A motive consists of an urge (that is, a need indicating a demand) and a goal that is related to this urge [14]. Motives shape cognition and behavior [229]. To be motivated means to be moved to do something [210]. Motives direct behaviors towards particular goals, which makes the agent more persistent in actions it takes. They also affect cognitive processes by increasing the level of attention. Motive, as the outcome of the motivation process, initiates, directs and maintains goal-oriented behaviors.

Motives are goal-driven and they move the agent towards the attainment of corresponding sets of intentions. In other words, motives as an essential part of affect can lead the agent to empower an intention. They are essentially mechanisms that in light of beliefs tend to produce, modify or select between actions and their reciprocal intentions. Some of motives are transient, like helping the Astronaut to hold the panel, while some are long term, like reaching to the shared goal of the collaboration in our ongoing example, installing solar panels and satisfying the Astronaut's needs in the field constitutes the shared goal (see Section 3.2).

2.5.2 Motivation Theory

There are several prominent motivation theories in psychology [21, 87, 138], some of which have received attention as the basis for computational models. In [174], Murray described and studied 20 different human motives, of which three have received attention in psychology and artificial intelligence as social motives [169, 273]. Our work on motives has been inspired by some of these works in the literature. However, we have developed our own motive types and their computational models (see Section 4.6). The following is a brief description of these three social motives, *achievement*, *affiliation* and *power* [11, 273] which will be used in this thesis:

- **Achievement motivation:** Achievement motivation drives humans to strive for excellence by improving on personal and societal standards of performance.

It involves a concern for excellence, for doing one's best. In artificial agents, achievement motivation has potential roles in focusing agent behavior and driving the acquisition of competence.

- **Affiliation motivation:** Affiliation refers to a class of social interactions that seek contact with formerly unknown or little known individuals and maintain contact with those individuals in a manner that both parties experience as satisfying, stimulating, and enriching. It involves a concern with developing friendly connections with others through the two contrasting emotional components of hope of affiliation and fear of rejection. These two components become more crucial in the collaboration domain due to the importance of social emotions and their impact on beliefs and intentions.
- **Power motivation:** Power can be described as a domain specific relationship between two individuals, characterized by the asymmetric distribution of social competence, access to resources, or social status. It involves concern with having an impact on other people or on the world at large. There are different aspects of fear or avoidance of power which channel and moderate the expression of power into socially acceptable behavior, working as inhibitions to unseemly tendencies. Power motivation can be considered with respect to the probability of success which makes it relevant to the cognitive appraisal of emotions during collaboration.

In [273] it is shown that failure of a power goal is associated with anger, confusion and disgust; success at an affiliation goal is associated with interest, happiness and feeling loved; and success at an achievement goal is associated with interest, surprise, happiness, excitement and a sense of focus. In other words, succeeding at a particular motive is associated with experiencing particular emotions.

2.6 Theory of Mind

Theory of Mind (ToM), as a crucial component in human's social interaction, plays an important role in our computational model. It concerns one's beliefs about others as intentional agents. Beside the immediate effect, an individual's action also depends on her beliefs about other's perception of that action as well as the reaction they take. In this thesis, we use the ToM concept whenever the agent needs to anticipate a human's mental state. We will also use the term *user model* as a standard collection of properties to describe others.

The concept of theory of mind has received much attention in social psychology and artificial intelligence. Eligio et al. explore what collaborators understand about each other's emotions and conclude that being aware of each other's emotions helps collaborators to improve their performance [73]. Fussell and Kraus discuss the importance of perspective taking in a successful communication in a social setting [80]. Scassellati discusses the importance of attribution of beliefs, goals and desires to others. He presents two psychological theories on the development of theory of mind in humans and their potential application in building robots with similar capabilities [213]. Hiatt and Trafton present a cognitive model which borrows mechanisms from three different postulates of theory of mind and show that their model produces behaviors in accordance with various theories of experiences [111]. Si, Marsella and Pynadath discuss PsychSim, an implemented multi-agent-based simulation tool for modeling social interaction, which has its own beliefs about its environment and a recursive model of other agents [194]. They also investigate the computational modeling of appraisal in a multi-agent decision-theoretic framework using POMDP based agents [235, 234]. Since applying the concept of theory of mind is crucial in social interaction and collaboration, this thesis includes a simple ToM mechanism inspired by this previous work.

2.7 Summary

In this chapter, we started by defining the concept of collaboration based on Grosz and Sidner’s work [103], and listed a number of collaboration properties. Then, we provided the background of two prominent computational theories of collaboration which helped develop a better understanding of the theories and how they relate to each other. Next, we presented the SharedPlans theory and its main properties, e.g., partial shared plan, recipe, and two notions of intention. Afterwards, we discussed the key concepts of the Joint Intentions theory including joint commitment and joint intention. Then, we continued with the hybrid approach of modeling collaboration and discussed one of the most well-established models, STEAM. We also briefly mentioned some other approaches. Later, we presented two different lists to compare similarities and differences between SharedPlans and Joint Intentions collaboration theories. We ended this document with different categories of applications of these theories in agent/robot and human collaboration areas.

We believe the SharedPlans and Joint Intentions collaboration theories are the most well-defined and well-established theories in computer science. We found SharedPlans theory more convincing than the other major and subordinate approaches, with respect to its inclusive explanation of the collaboration structure and its association to discourse analysis which directly improves the communicative aspects of a collaboration theory. We also understand the value of Joint Intentions theory due to its clarity and closeness to the foundations of collaboration concepts. These specifications of the Joint Intentions theory can make it applicable in multi-agent system designs and human-robot collaboration. We also consider hybrid approaches valuable, such as STEAM, because they clearly understand drawbacks with existing theories and successfully achieve better collaborative agents by infusing different concepts from different theories. Although all these theories are well-defined and properly introduce collaboration concepts, they mostly explain the structure of a collaboration and they lack the underlying domain-independent pro-

cesses with which collaborative procedures could be defined more systematically and effectively in different applications.

We have also discussed the description of affective computing and the importance of the concept of emotion in general and in social context. We reviewed the importance of communicating emotions as well as emotions' social functions. Then, we provided some examples of agents and robots using artificial emotions in their decision making process. We also briefly provided a few examples of cognitive architectures producing different aspects of behaviors in robotics.

There are major theories of emotions explaining the concept of emotion. We discussed these major theories in detail separately, providing their psychological background and underlying concepts. Following the explanation of these theories, we were able to discuss the similarities and differences between these major theories. Finally, we provided applications of these theories in robotics and AI.

In general, we believe the key components provided by prominent computational collaboration theories could help us in better understanding of collaboration concepts and ultimately developing our computational models to be compatible with the collaboration structure. However, we believe these theories need to be expanded by a better understanding of a) the underlying processes leading to these structures, and b) the involvement of cognitive constructs that could be involved in these processes. Furthermore, we believe theories explaining the affect-driven processes could successfully model nebulous concepts in psychology, to some degree. We have developed our work based on computational models of emotions, because it is good to follow well-established (in comparison with others) theoretical foundations. These theories can be a guideline for our computational models, and they can explain more details of the structure or the processes involved in affective situations. However, we do not necessarily think that the computational models must exactly follow only one theory and its descriptions. Meaning, different aspects of models can represent different theories. For instance, appraisal theory is a good representation of the interpretive aspect of emotions and basic emotion theories provide detailed systematic

methods for expressive application. More importantly, we believe the interpersonal functions of emotions should be our first concern and we should try to relate them to the structure of our domain, i.e., collaboration. In conclusion, we can see the importance of interpretive, communicative and regulatory aspects of emotion functions in this proposed work.

CHAPTER 3

AFFECTIVE MOTIVATIONAL COLLABORATION THEORY

Current computational collaboration theories specify the structure of collaborative activities, but are weak on the underlying processes that generate and maintain these structures. In addition, current computational models of emotion (specifically appraisal models) provide antecedents (e.g., beliefs and goals) of appraisal processes and clearly distinguish between different appraisals, but do not describe the differences of appraisal in a collaboration context. We argue that affective states are crucial to these underlying processes and we have developed a new computational theory, called *Affective Motivational Collaboration* theory, that combines affect-based processes, such as appraisal and coping, with collaboration processes, such as planning, in a single unified framework. In this chapter, we provide a general argument about our AMC theory, major functions of emotions that can be applied in a collaboration context, as well as the components in our architecture and how each component (mechanism) deals with the events in a collaborative environment. We also provide the definition of elements of mental state and their attributes. Later in Chapter 4, we use all of this information in the new algorithms we have developed, e.g., appraisal processes, as part of a new overall computational model.

This work is implemented to build robots capable of generating and recognizing emotions in order to be better collaborators. We have investigated the mutual influences of affective and collaborative processes in a cognitive theory to support interaction between humans and robots or virtual agents. We build primarily on the

cognitive appraisal theory of emotions and the *SharedPlans* theory of collaboration to investigate the structure, fundamental processes and functions of emotions in a collaboration.

Although existing collaboration theories specify the important elements of a collaboration structure, the underlying processes required to dynamically create, use, and maintain the elements of this structure are largely unexplained. For instance, a general mechanism has yet to be developed that allows an agent to effectively integrate the influence of its collaborator's perceived or anticipated emotions into its own cognitive mechanisms to prevent shared task failures while maintaining collaborative behavior. Therefore, a process view of collaboration must include certain key elements. It should inherently involve social interactions since all collaborations occur between social agents, and it should contain a means of modifying the content of the social interaction as the collaboration unfolds. The underlying processes of emotions possess these two properties, and social functions of emotions explain some aspects of the underlying processes in collaboration.

There is also a communicative aspect of emotions. For instance, emotions are often intended to convey information to others [85]. Emotions are also involved in verbal behaviors. For instance, an utterance can include both content and relational meaning. An affect might appear to be elicited by the content of the utterance, but in fact be an individual's response to the relational meaning [188]. The interpretation of these relational meanings are handled by the appraisal of events. Appraisal processes give us a way to view affect as social [258]. Meaning is created by an individual's social experiences in the social world, and individuals communicate these meanings through utterances. Consequently, the meaning of these utterances and the emotional communication change the dynamic of social interactions. A successful and effective emotional communication necessitates ongoing reciprocal adjustments between interactants that can happen based on interpretation of each other's behaviors [156]. This adjustment procedure requires a baseline and an assessment procedure. While the components of the collaboration structure, e.g., shared plan,

provide the baseline, affect-related processes (e.g., appraisal) provide the assessment procedure.

Affective Motivational Collaboration theory is about the interpretation and prediction of the observable behaviors in a dyadic collaborative interaction. The theory focuses on the processes regulated by emotional states. These observable behaviors represent the outcome of processes related to the interpretation of the self's relationship to the collaborative environment. The processes are triggered by the events occurring in the collaborative environment. Thus, *Affective Motivational Collaboration* theory explains how emotions regulate the underlying processes when the events occur during collaboration.

Emotion-regulated processes operate based on the self's mental state including, the anticipated Mental State of the other. These elements of mental state include beliefs, intentions, goals, motives and emotion instances. Each of these mental state possesses multiple attributes impacting the underlying processes of collaboration and ultimately the relation between cognition and behavior of the agent. The nature of these attributes will be discussed in Section 3.7.

There are several theories discussed in Chapter 2, which describe the underlying structure of a collaboration based on mental state of the collaborators. The collaboration structure of *Affective Motivational Collaboration Theory* is based on the SharedPlans theory [96]. *Affective Motivational Collaboration* theory focuses on the processes that generate, maintain and update this structure based on mental state. The collaboration structure is important because social agents/robots ultimately need to co-exist with humans, and therefore need to consider humans' mental state as well as their own internal states and operational goals.

3.1 Overview of Mechanisms

In this section, we introduce all mechanisms and the connections between them in our architecture. The *Perception* and the *Action* mechanisms are only the source (to

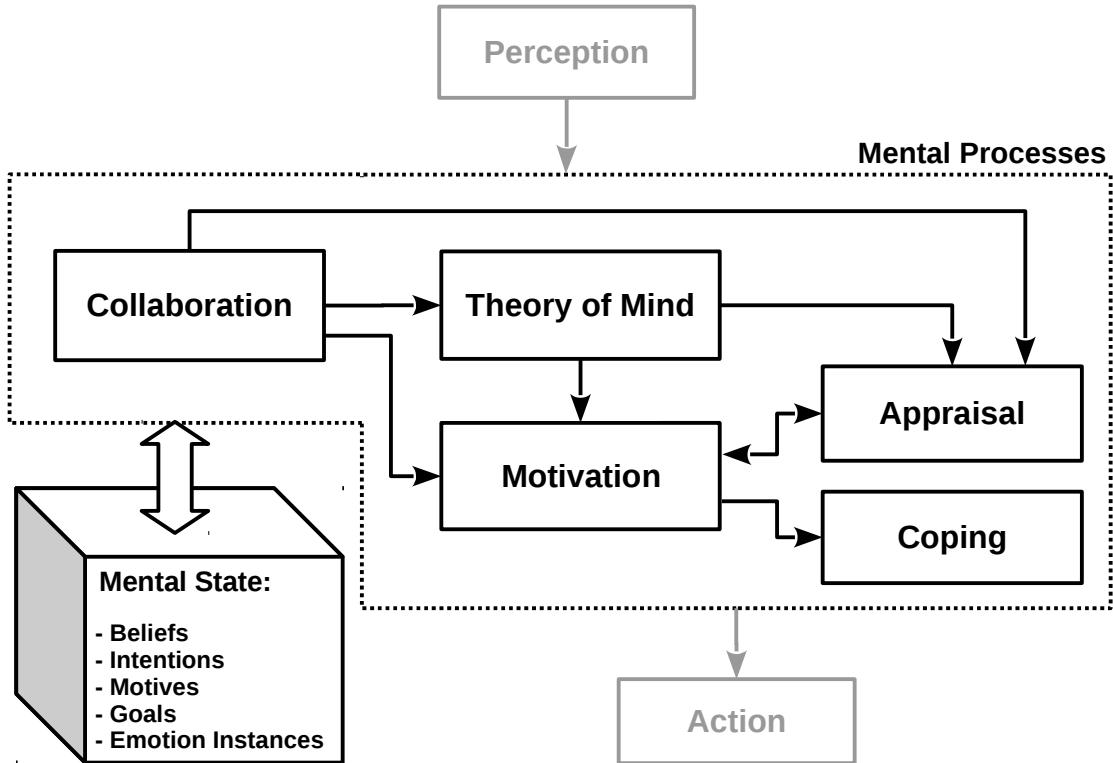


Figure 3.1: Primary influence of mechanisms in *Affective Motivational Collaboration Theory*.

produce sensory information) and the sink (to show proper behavior) of data in our architecture. These two mechanisms function based on the same events introduced in Section 3.4.

The *Appraisal* and *Coping* mechanisms are two major mechanisms in our computational framework that are tightly coupled. The *Appraisal* mechanism is responsible for evaluating changes in the self's Mental State, the anticipated Mental State of the other, and the state of the collaboration environment. Consequently, the *Appraisal* mechanism is connected to a) the *Theory of Mind* mechanism, to serve as an evaluator whenever the self applies the *Appraisal* mechanism to the Mental State attributed to the human collaborator, b) the *Collaboration* mechanism, to interpret the progress and changes in the collaboration plan and associated Mental State, and to make changes to the shared plan if required, c) the *Motivation* mechanism, to generate and assess the self's new goal-driven motives whenever a new motive

or intention is required, e.g., following the failure of a task, and d) the Perception mechanism, to interpret the events from the collaboration environment. In general, our appraisal algorithms are developed based on the data available in a collaboration structure. This feature makes our algorithms unique and different than other computational appraisal models. The details about our appraisal algorithms are provided in Section 4.3. The Coping mechanism provides the self with different coping strategies associated with changes in the self’s Mental State with respect to the state of the collaboration. In other words, the Coping mechanism produces cognitive responses by forming new intentions based on the appraisal patterns. The Coping mechanism is also inspired by other work; however, it is developed to enable an agent to cope with the events in a collaborative environment which also distinguishes our work from other computational models of coping. Details about our coping mechanism can be found in Section 4.5.

The *Motivation* mechanism provides motives to influence the coping process in terms of the collaborative agent’s or the human collaborator’s needs. The Motivation mechanism uses the Appraisal mechanism to compute attributes (see Section 3.7) of competing motives. Also, the Motivation mechanism can serve the Theory of Mind mechanism by helping the self to infer the motive behind the other’s current action. Moreover, the Motivation mechanism applies the beliefs associated with the Appraisal mechanism to generate and compare a new set of motives related to the status of the collaboration. The outcome of the Motivation mechanism is involved in forming a new intention to cope with the current event. As a result, the self can take an action based on the new intention to sustain the collaboration progress.

The *Theory of Mind* mechanism is responsible for inferring a model of the other’s anticipated Mental State. The self will progressively update this model during the collaboration. The refinement of this model helps the self to anticipate the other’s mental state more accurately, which ultimately impacts the quality of the collaboration and the achievement of the shared goal. Furthermore, the self can make inferences about the motive (or intention) behind the other’s actions using the Mo-

tivation mechanism. This inference helps the self to update its own beliefs about the other’s mental state. In the reverse appraisal process (see Sections 2.4.1 and 4.7), the self also applies the Appraisal mechanism together with updated beliefs about the other’s Mental State to make inferences about the other’s current mental state based on the other’s emotional expression. Finally, the Collaboration mechanism provides the collaboration structure, including status of the shared plan with respect to the shared goal and the mutual beliefs to the Theory of Mind mechanism. Consequently, any change to the self’s model of the other will update the self’s mental state.

The *Collaboration* mechanism maintains constraints on actions. These constraints include constraints on task states and on the ordering of tasks. The Collaboration mechanism also provides processes to update and monitor the shared plan. These processes depend on the Appraisal mechanism to evaluate the current Mental State with respect to the current status of the collaboration. The self also shifts its focus of attention according to the outcome of the Appraisal mechanism. Moreover, the Collaboration mechanism can help the self to identify the failure of a task. The Appraisal and Motivation mechanisms provide interpretation of task failure and the formation of new elements of Mental State (e.g. intentions) respectively. Ultimately, the Coping mechanism allows the self to perform behavior appropriate to the current state of the collaboration.

3.2 Example Scenario

We now provide the following scenario in a robotic domain to illustrate a collaborative interaction. In the scenario, there is an astronaut, who has had a high success rate in accomplishing space missions. She is capable of operating the manipulator system and supporting equipment. She works as a commander in the field during the operation. She is trained to collaborate with general-purpose field operation robots.

There is also a robot which is assigned to the mission to provide services to the astronaut. It has been tested in extreme environmental conditions and has a low failure rate. It is capable of communicating with the astronaut and understanding the astronaut's nonverbal behavior. It has the ability to identify and assess its own emotions and those of the astronaut.

The robot and the astronaut will collaborate with each other to achieve their shared goal, which is to install two solar panels. They will face various difficulties, ranging from the task being unpleasant and challenging to conflicts of their private and/or shared goals occurring because of a blocked or a protracted sub-task. The robot and the astronaut will go through a series of assessment processes to figure out a) how did the current blocking happen? b) why is the current task blocked? and c) what is the next action they are going to take? The robot uses its cognitive and affective abilities and its communication skills to overcome these problems and to motivate the astronaut to propose alternative tasks. The following is part of an interaction between the astronaut and the robot during their collaboration on installing solar panels.

1. **Astronaut:** Please hold the panel on this structure.

[Robot holds the panel and Astronaut begins to work on the panel.]

[Both the Robot and the Astronaut continue their collaboration to achieve their shared goal.]

2. **Astronaut:** At this point you should be careful how you hold the panel. Turn the right side 45 degrees towards me.

3. **Robot:** Is this what you want?

4. **Astronaut:** Yes, do not move it.

[Astronaut finishes attaching the panel onto the structure and checks the connectors to make sure they are working.]

5. **Astronaut:** The connectors on this panel have problems and we might not be able to finish this task.
6. **Robot:** Don't worry! I can replace the connectors in about 4 minutes. We definitely can finish this task after that.
7. **Astronaut:** Okay, go ahead and fix the connectors.

[Robot fixes the issue with the connectors and passes them to the Astronaut. Astronaut connects the wires to the connectors.]
8. **Astronaut:** I need you to begin welding this panel and also prepare the measurement tool for me.
9. **Robot:** Do you want me to prepare the measurement tool first? Then, I can begin welding afterwards.
10. **Astronaut:** Yes, that's fine!

[Astronaut waits for the Robot to weld the panel and prepare the measurement tool for him. Robot finishes the welding task after a long time, then prepares and passes the measurement tool to the Astronaut. But, the measurement tool has an accuracy problem.]

11. **Astronaut:** Oh no! Finishing the quality check of our installation with this measurement problem is so frustrating. I think we should stop now!
12. **Robot:** I see. But, I can help you with the measurement and we can finish the task as originally planned.
13. **Astronaut:** That would be great!

[Robot helps the Astronaut to finish the measurement task with its own measurement tool.]

[Then, the Robot goes back to its own goal, which is to fetch the second panel to finish the overall task.]

3.3 General Argument

Affective Motivational Collaboration theory focuses on affect-regulated processes involved in collaboration and builds on two well-established theories in this context. The first is Grosz and Sidner’s SharedPlans collaboration theory, which is based on the concepts of mutual belief and shared plans [96, 103]. Secondly, we build on the computational model of the appraisal theory of emotions by Marsella and Gratch [92, 90, 162, 163] which explains how emotions arise from an individual’s interpretation of its relationship with the environment, and specifies the dimensions of appraisal and the appraisal patterns characteristic of different emotions [223]; although our algorithms are inspired by their work, our focus is to employ other concepts that are involved in appraisal of an event due to the collaboration context. For instance, as discussed in Section 4.3.1, we believe not only is the relevance of an event influenced by the utility of the event, but it is also influenced by the perceived affective state of the other collaborator. The details of these differences are described in Section 4.3.

Existing collaboration theories (including SharedPlans) consider the nature of a collaboration to be more than a set of individual acts. These theories argue for an essential distinction between a collaboration and a simple interaction or even a coordination in terms of commitments [96, 152]. We believe there is also a need for a computational theory to specify and characterize the underlying cognitive processes of collaborative activities. The study of these cognitive processes helps explain why and how humans collaborate with each other. For instance, SharedPlans theory can describe our scenario in Section 3.2 in terms of fundamental Mental State, such as mutual beliefs, intentions, and shared plans. However, it does not explain the underlying processes leading to these elements of Mental State. *Affective Motivational Collaboration* theory extends the SharedPlans theory by describing these processes. As another example, all of the prominent computational models of appraisal are able to clearly define the antecedents of appraisals, and how each

appraisal process can evaluate a particular aspect of an event. However, no appraisal model provides the required factors for each appraisal process in the context of collaboration. Furthermore, emotions, due to their evaluative and regulatory nature, provide fundamental functions (see Section 3.5) each of which plays an essential role in maintaining a collaboration’s structure and status. In other words, these functions explain the dynamics of a collaboration structure.

Affective Motivational Collaboration theory specifies the processes involved in the progress of a collaboration and how they impact the collaboration’s underlying structure. For example in the exchange below, the Robot needs to respond appropriately to the Astronaut’s new request, to maintain progress during collaboration. The emotion function, i.e., goal management, is involved in such situations. Our computational model starts with high-level semantic representation of events (including utterances), i.e., natural language processing is out of the scope of this work:

8. **Astronaut:** I need you to begin welding this panel and also prepare the measurement tool for me.

9. **Robot:** Do you want me to prepare the measurement tool first? Then, I can begin welding afterwards.

What is the nature of the processes involved in a collaboration? For example, in the exchange below, the Robot changes its focus of attention to something important to the Astronaut because of its perception of the Astronaut’s negative affective state:

5. **Astronaut:** The connectors on this panel have problems and we might not be able to finish this task.

6. **Robot:** Don’t worry! I can replace the connectors in 4 minutes. We definitely can finish this task after that.

And, how do these processes impact the social characteristics of a collaboration?

For instance, in the exchange below, emotions and the Appraisal mechanism can influence the self's awareness during collaboration:

11. **Astronaut**: Oh no! Finishing the quality check of our installation with this measurement problem is so frustrating. I think we should stop now!

12. **Robot**: I see. But, I can help you with the measurement and we can finish the task as originally planned.

Finally, *Affective Motivational Collaboration* theory incorporates motivation as an affect-regulated and goal-driven mechanism, by which the self can form a new intention based on its own beliefs about self and the other, as well as the result of an Appraisal mechanism. In general, a new motive can be involved in formation of a new intention and the self can take a new action based on the new intention. The Motivation mechanism also connects the outcome of the Appraisal mechanism and the Collaboration mechanism by applying the self's belief structure and appraisal patterns. The result of this process generates a set of competing motives, each of which can influence the formation of self's intention. The self can store its own motives as well as the other's motives along with their corresponding attributes which can impact the Appraisal mechanism. In the following example extracted from the scenario, the Astronaut informs the Robot of a new problem, and the Robot forms a new intention to solve the problem:

5. **Astronaut**: The connectors on this panel have problems and we might not be able to finish this task.

6. **Robot**: Don't worry! I can replace the connectors in 4 minutes. We definitely can finish this task after that.

In the same example that we saw earlier, the Astronaut faces a problem in his own task and informs the Robot of his decision. The Robot forms a new intention

to help the Astronaut to overcome his problem and ultimately, make progress in their collaboration:

11. **Astronaut**: Oh no! Finishing the quality check of our installation with this measurement problem is so frustrating. I think we should stop now!
12. **Robot**: I see. But, I can help you with the measurement and we can finish the task as originally planned.

3.4 Events

The *events* occurring in a collaborative environment include a) *utterances* spoken by the collaborators, b) *primitive actions* executed, deferred, or aborted, and c) observable *affective states*. These events are the events that our affective collaborative agent perceives. We will discuss below the operation of individual processes in our theory based on these events. Each of the following five sub-sections describe how an individual mechanism in Figure 3.1 handles these events.

3.4.1 Collaboration Mechanism and Events

The Collaboration mechanism is responsible for maintaining the internal structure of a collaboration, including the focus of attention, constraints on actions, updating the shared plan and, in general, monitoring the collaboration. All of these structures require updating each time the self perceives an event. For instance, an utterance by the other can impact the self's focus of attention during the collaboration, or the effect of a primitive action can influence the self's view of an impasse on a task. As another example, the perception of the other's affective state can cause significant changes in the self's collaboration monitoring.

3.4.2 Appraising Events

The other's *utterances*, the effect(s) of the collaborators' *primitive actions*, and the other's *affective states* (expressed verbally or nonverbally) are the three types of events perceived by the self during collaboration. The Appraisal mechanism receives the output of the Perception and Collaboration mechanisms as well as the requisite Mental States related to the current event. It appraises the event, in terms of appraisal variables using the collaboration structure and the history of the self's related Mental State. The collaboration structure contains information about the collaboration's shared plan and the collaborators' shared goal, the temporal and the hierarchical constraints of the tasks, and the current focus of attention. Moreover, the self progressively generates and updates various elements of Mental State (discussed in Section 3.7) during collaboration. The occurrence of a new event causes a change in the self's Mental State. The construct of the new mental state, e.g., beliefs, are semantically connected to the older ones. The Appraisal mechanism uses the history of the Mental State to consistently evaluate a new event.

3.4.3 Coping with Events

Events do not directly cause the self's Coping mechanism to operate. Instead, it is the formation of Mental State that cause the Coping mechanism to choose an appropriate cognitive response to these events. The cognitive responses (also known as "coping strategies") are considered to act upon the self's relationship to the world and its own Mental State. Events also trigger other processes, which impact the self's Mental State. The changes in Mental State cause the Coping mechanism to provide consistent and appropriate cognitive responses to the world. For instance, suppose the self perceives an utterance and evaluates it in terms of the appraisal variables. The values of these variables and the corresponding affective states will cause new beliefs and intentions to be formed, which then cause the Coping mechanism to appropriately choose the self's behavior.

3.4.4 Motivation and Events

The Motivation mechanism acts to regulate the self's Mental State and goal-directed behaviors for internal and social purposes. The Appraisal mechanism evaluates the state of self, the environment, or the anticipated mental state of the other. In each of these cases, the outcome of the Appraisal mechanism might indicate the need for internal or behavioral regulation. In such cases, the Motivation mechanism uses the elements of Mental State associated with the state of self, the environment or the other's anticipated Mental State as well as the pattern provided by the Appraisal mechanism to generate motives aligned with private or shared goals. Thus, via the Appraisal mechanism the Motivation mechanism implicitly responds to the events. The attributes of the generated motives (see Section 3.7) will be updated every time a new event occurs. For instance, the Appraisal mechanism may evaluate the outcome of the current task as *unexpected*, *undesirable*, *uncontrollable* and *urgent* which is indicative of the failure of a task. Then, the Motivation mechanism provides goal-directed motives, each of which can influence the formation of an intention.

3.4.5 Theory of Mind and Events

Theory of Mind operates when an event occurs and the self wants to infer and interpret the other's mental state. Thus, Theory of Mind helps the self to choose the behavior best matched to the other's anticipated Mental State. The Theory of Mind mechanism infers the mental state of the other, which helps the self to update the user model of the other. The Motivation and the Appraisal mechanisms are also involved in this procedure. For instance, the self can infer the other's mental state through a reverse appraisal procedure (see Sections 2.4.1 and 4.7). The Motivation mechanism includes another inverse procedure to infer the other's active motives, which can lead to inferring the other's goal, beliefs, motives and intentions.

3.5 Functions of Emotions

We have talked about the crucial role of emotions in communicating Mental States, motivating actions, and evaluating and interpreting internal states and the environment. Emotions, generally speaking, provide a set of intra- and interpersonal functions which regulate internal processes and the self’s relationship to the other during the collaboration. Emotions have meanings in a social context which can be interpreted by an observer. The self uses the meanings of affective states to trigger appropriate emotion functions with respect to the current social context. Ultimately, the elicited affective state’s functions impact the self’s Mental State and consequently behaviors. In the rest of this section, we briefly describe how ten different emotion functions are related to the collaboration context. There are other emotion functions, such as learning and memory control, which are outside of the scope of this thesis. We have implemented some of the emotion functions in this section such as social regulation, motivation, focus of attention, and goal management in our computational framework (see Chapter 4). Other functions are described here with their relationship to the collaboration context, but they are beyond the implementation of this thesis.

3.5.1 Action Selection

Action selection is the function in which affective states influence choosing the most appropriate action out of a repertoire of possible actions at a point in time. This function influences the Coping mechanism and results in consistency of the self’s actions based on anticipated emotional responses of the other and the satisfaction of the shared goal.

3.5.2 Adaptation

Adaptation is the raison d’être of emotions. It helps the self to properly respond to changing challenges in a dyadic social context by adjusting its behavior. Adaptation

is a specialized problem-solving technique implicating the necessity of the self's emotional states for short and long term behavior changes during collaboration.

3.5.3 Social Regulation

Social regulation by emotions is the process which enables the self to communicate internal Mental State through the expression of emotions in a social context. It can assist the self to regulate various social interactions required in the course of a collaboration, such as conflict resolution and negotiation. Emotional expressions influence the other's behavior by triggering the other's inferential processes and emotional reactions [135].

3.5.4 Sensory Integration

Sensory integration can guide the self through the course of a collaboration by sustaining rich-sensory tasks to demonstrate more effective collaborative behaviors. It benefits the self by anticipating a certain type of inferential process to the other's mental and emotional states. For instance, perceiving fear in the other can lead to an increased focus of attention on the ongoing task, or discerning anger can raise the probability of avoiding current events (generated by the self) by the other.

3.5.5 Alarm

The alarm function is a purely reactive and pattern-driven process [241]. It accounts for persuading the self that an undesired or unsatisfactory condition happened in the past, and since then, has persisted in the self's Mental State. The alarm function also provides the self with a rapid reaction to the events. The self will be able to interrupt deliberative processes and show quick behavioral reactions. For instance, the self can consider corrective actions when a high probability of anticipated failure occurs during the collaboration.

3.5.6 Motivation

Motivation is a goal-driven emotion function associated with the self's behaviors. There is a motive behind every intentional action created by the Motivation mechanism. This motive is computed based on underlying beliefs relying on the evaluative role of emotions. Therefore, the motive behind any behavior carries an anticipated value of the future consequence for that behavior. It also reveals the belief foundation of a behavior. Consequently, the self can apply this function of emotions to a) cope with certain types of problems, and b) infer the other's mental state based on each action.

3.5.7 Goal Management

The goal management function identifies the existence or the need for a high priority goal for the self. These goals include both private goals and shared goals. Emotions provide an evaluation mechanism for the self to choose or reprioritize goals at each point in time. This function of emotions can impact the self's behavior with respect to the dynamics of interaction during the course of a collaboration.

3.5.8 Focus of Attention

Emotion instances and the patterns generated by the Appraisal mechanism are directly linked to the focus of attention of the self. Both positive and negative results of a cognitive evaluation of events can change, maintain, or intensify the self's focus of attention. For instance, negative emotions, e.g., fear or anger, can influence the self's focus of attention by orienting the self towards the events [76]. Positive emotions, e.g., happiness, can broaden or expand the self's focus of attention from details of the events to their general features [78].

3.5.9 Strategic Processing

The occurrence of new events can lead the self to rapid and/or strategic responses. The Coping mechanism contains various strategies associated with different elements of the Mental State, e.g., belief or intention-related strategies. The content of the self's Mental State changes as time passes, which causes the Coping mechanism to choose an appropriate action. The Appraisal mechanism allows the self to demonstrate a rapid response or strategically prioritize the current events generated based on the changes in the Mental State. For instance, is a mild, reactive facial expression an adequate response to the other's current utterance or does the self need to show a stronger behavior? Is it the new belief about the current state of the collaboration, or is it the new intention pursuing the self's private goal that the self needs to cope with? Thus, appraisal patterns and affective states impact the self's strategic processing.

3.5.10 Self Model

Emotions can be a representation of how the self interprets the collaboration environment. The self can generate or update beliefs about its self-model when faced with unambiguous events and apply the same self-model when confronted with events possessing more ambiguity and uncertainty. Creating a self-model can also help the self to demonstrate more consistent and coherent behaviors when similar situations occur during the collaboration. This reliability in the self's behavior can help the other to predict the self's responses during collaboration.

3.6 Components of the Architecture

Affective Motivational Collaboration theory consists of seven mechanisms (see Figure 3.1) most of which directly store and fetch the data in the Mental State. The Mental State will keep all the required data about the self (agent), other (human) and the

environment (including events). In this section we explain each of the mechanisms and the elements of Mental State in more detail.

3.6.1 Collaboration

- **Input:** The input to the *Collaboration* mechanism includes all the data that affects the execution of individual tasks in the collaboration plan. This data will be provided via the different elements of Mental State including beliefs, intentions and goals. These elements of Mental State will establish the agent's initial plan and will be updated throughout the collaboration process by the Perception mechanism and other processes.
- **Output:** The output of *Collaboration* includes all the data that is modified or created during execution of a plan in the form of elements of Mental State. These elements of Mental State will be used by the internal processes in the Theory of Mind mechanism. Additionally, the Appraisal mechanism will use these elements of Mental State to evaluate the events during collaboration. These elements of Mental State also will be used by other processes, e.g. goal management, for the purpose of maintaining the collaboration.
- **Function:** The *Collaboration* mechanism will construct a hierarchy of tasks and also manage and maintain the constraints and other required details of the collaboration specified by the plan. These details include the inputs and outputs of individual tasks, the *preconditions* specifying whether it is appropriate to perform a task, and the *postconditions* specifying whether a just-completed task was successful (which can be used as an indication of an impasse or failure). *Collaboration* also keeps track of the focus of attention, which determines the salient objects, properties and relations at each point of the collaboration. Moreover, *Collaboration* has the ability to shift the focus of attention during the collaboration. All the other mechanisms in the overall *Affective Motivational Collaboration* theory are influenced by changes in the collaboration

plan. The *Collaboration* mechanism in general performs various logical deductions required by other processes in our computational model. It is designed to ameliorate the shortcomings of the existing Collaboration theories by providing required inferences such as dynamic planning based on the recent changes in the collaboration environment and the internal changes in the agent's Mental State. For instance, in our scenario (see Section 3.2), when the Astronaut interrupts the Robot asking for a new and urgent task, the Robot needs to alter the collaboration plan to continue. *Collaboration* also supports essential monitoring processes during the collaboration such as event monitoring.

3.6.2 Appraisal

- **Input:** The most significant part of *Appraisal*'s input data is based on the activity of the Collaboration mechanism. This data includes all the required elements of Mental State associated with the Collaboration mechanism. For instance, the beliefs and their strengths will be used by algorithms inside of *Appraisal* to compute the value of the appraisal variables. *Appraisal* also receives data from the Theory of Mind mechanism. This data helps the agent use *Appraisal* for inferring the human's intentions and motives based on a reverse appraisal procedure. The input data from the Perception mechanism is generally needed to support the evaluation of the events. *Appraisal* also uses the information about the motives in the underlying processes.
- **Output:** The output of *Appraisal* can directly and indirectly impact other mechanisms. The Motivation mechanism uses this data to generate and maintain motives based on the current appraisal of the environment.
- **Function:** *Appraisal* is a subjective evaluation process based on individual processes each of which computes the value of the appraisal variables used in our computational model. The Collaboration mechanism needs the evaluative assistance of *Appraisal* for various reasons. The course of a collaboration is

based on a full or a partial plan which needs to be updated as time passes and collaborators achieve, fail at or abandon a goal assigned to them. The failure to achieve a goal should not destroy the entire collaboration. Appraising the environment and the current events helps the agent to update the collaboration plan and avoid further critical failures during collaboration. *Appraisal* also helps the agent to have a better understanding of the human's behavior by making inferences based on appraisal variables. Furthermore, in order to collaborate successfully, a collaborator cannot simply use the plan and reach to the shared goal; there should be an adaptation process not only for updating the plan but also the underlying Mental State. For instance, there are beliefs about the appraisal of the self and the other that augment the model of what collaborators have done, and what and how they are planning to achieve the current shared goal based on their emotional states. This process will be discussed in more detail under the Motivation mechanism (see Section 3.6.4). Additionally, the beliefs formed based on the appraisals can impact other mechanisms such as the Theory of Mind, Motivation and Coping, essentially including the whole computational model.

3.6.3 Coping

- **Input:** The *Coping* mechanism operates based on the data stored in different aspects of the Mental State. This data includes changes in the agent's beliefs as well as the agent's intentions (whether they are created or altered during the collaboration), and the private or shared goals.
- **Output:** The output of the *Coping* mechanism is the data specifying the intention for a behavior which the agent needs to perform based on the current state of the collaboration.
- **Function:** The *Coping* mechanism is responsible for interpreting ongoing changes in the Mental State and adopting the appropriate behavior with re-

spect to these changes. This component includes rules categorized into four coping strategies which are *Belief-related*, *Intention-related*, *Attention-related* and *Desires-related* strategies [162]. These rules will apply to the attributes and structures of the elements of Mental State to cope with the internal changes as well as the demands of the environment. For example, the *Coping* mechanism will utilize certain beliefs about the self to regulate the agent's internal states, while using mutual beliefs to maintain progress in the existing collaboration. As another example, motives' attributes can guide the *Coping* mechanism by voting for a particular behavior.

3.6.4 Motivation

- **Input:** The most essential part of the input to *Motivation* is the Mental State, and more specifically the private and shared goals associated with the collaboration. *Motivation* also uses data from two other mechanisms, namely, Theory of Mind and Appraisal. Input from Theory of Mind is used by *Motivation* whenever new motives need to be generated or compared according to the shared goal. Input from Appraisal is used whenever the motive attributes are involved in the internal processes of the *Motivation*.
- **Output:** The output of *Motivation* includes the data required to form new intentions to reach the private or the shared goals. The motives which are the output of the *Motivation* mechanism are also used by the Coping mechanism to choose appropriate behavior according to the goals of the collaboration plan.
- **Function:** The *Motivation* mechanism works closely with the Appraisal mechanism. The purpose of this component is to generate new motives which will be added to the Mental State. These motives are generated based on what the agent believes about the environment including self and the other collaborator and the corresponding appraisals. The agent uses these motives to achieve a

private or shared goal according to new conditions, to interact better with a human who needs social interactions, or to evaluate the success of task performances. The *Motivation* mechanism consists of several processes. These processes generate several motives with respect to the agent's current Mental State. Then, these motives will be used to make a decision to form an intention in the Coping mechanism.

3.6.5 Theory of Mind

- **Input:** *Theory of Mind* receives its input from the Mental State as well as the Collaboration and Perception mechanisms. This mechanism uses the current Mental State to infer the other's Mental State (which is simpler than the Mental State associated with self). The Collaboration mechanism provides the structure of the collaboration plan, including the constraints which can be used in the internal inference processes of *Theory of Mind*, such as reverse appraisal. The Perception mechanism also helps *Theory of Mind* with the input data from the sensory system.
- **Output:** The output of *Theory of Mind* will be stored in the Mental State. The Motivation mechanism can use this output to generate new motives according to the current state of the collaboration.
- **Function:** The agent uses the *Theory of Mind* mechanism to infer and attribute beliefs, intentions, motives and goals to its collaborator. The agent can also infer the Mental State of the other based on the reverse appraisal of the other's behavior. Another internal process of the *Theory of Mind* is to infer the other's motives on the basis of his behavior.

3.6.6 Perception

We consider the *Perception* mechanism only as a source of data to our computational model (see Figure 3.1). Thus, our computational model starts with high-level

semantic representation of events (including utterances), i.e., natural language processing is out of the scope of this work.

- **Output:** Predefined utterances will be used for verbal communication with the agent. These utterances will be a part of the output data of the *Perception* mechanism. The output of the *Perception* mechanism will be given to the Collaboration, Theory of Mind and Appraisal mechanisms. We will provide a unified perception representation across all of these mechanisms.
- **Function:** The *Perception* mechanism is responsible for producing the sensory information used by other processes in our model.

3.6.7 Action

We consider the *Action* mechanism only as a sink of data in our computational model (see Figure 3.1).

- **Input:** The input to the *Action* mechanism is provided by the Coping mechanism. This data will cause the *Action* mechanism to execute an appropriate behavior of the agent. This data has the same level of abstraction as the output of the Perception mechanism, i.e., it includes agent's utterances, primitive actions and emotional expressions.
- **Function:** The *Action* mechanism functions whenever the agent needs to show a proper behavior according to the result of the internal processes of the collaboration procedure.

3.6.8 Mental State and Emotion Instances

The Mental State shown in Figure 3.1 comprises the knowledge base required for all the mechanisms in the overall model. This Mental State includes: beliefs, intentions, motives, goals, and emotion instances for both self and the other.

Beliefs are a crucial part of the Mental State. We have two different perspectives on categorization of beliefs. In one perspective, we categorize beliefs based on whether they are shared between the collaborators. The SharedPlans [103] theory is the foundation of this categorization, in which, for any given proposition the agent may have: a) private beliefs (the agent believes the human does not know these), b) the inferred beliefs of the human (the agent believes the human collaborator has these beliefs), and c) mutual beliefs (the agent believes both the self and the human have these same beliefs and both of them believe that). From another perspective, we categorize beliefs based on who or what they are about. In this categorization, beliefs can be about the self, the other, or they can be about the environment. Beliefs about the environment can be about the outcomes of a new appraisal, or even the human's offer, question or request, and general beliefs about the environment in which the agent is situated. Beliefs can be created and updated by different processes. They also influence how these processes function as time passes.

Intentions are mental constructs directed at future actions. They play an essential role in: a) taking actions according to the collaboration plan, b) coordination of actions with the human collaborator, c) formation of beliefs about self and anticipated beliefs about the other, and d) behavior selection in the Coping mechanism. First, taking actions means that the agent will intend to take an action for primitive tasks that have gained the focus of attention, possess active motives, have satisfied preconditions, and for which required temporal predecessors have been successfully achieved. Second, intentions are involved in action coordinations in which the human's behavior guides the agent to infer an anticipated behavior of the human. Third, intentions play a role in belief formation mainly as a result of the permanence and commitment inherent to intentions in subsequent processes, e.g., appraisal of the human's reaction to the current action and self regulation. And lastly, intentions are involved in selecting intention-related strategies, e.g., planning, seeking instrumental support and procrastination; these strategies are an essential category of the strategies in the Coping mechanism. Intentions possess a set of attributes,

e.g. *Involvement*, *Certainty*, *Ambivalence* (see Section 3.7.4), which moderate the consistency between intention and behavior. The issue of consistency between the intentions (in collaboration) and the behaviors (as a result of the Coping mechanism in the appraisal cycle) is important because neither of these two mechanisms alone provides a solution for consistency.

Motives are mental constructs which can initiate, direct and maintain goal-directed behaviors. They are created by the affect-regulated Motivation mechanism. Motives can cause the formation of a new intention for the agent according to: a) its own emotional states (how the agent feels about something), b) its own private goal (how an action helps the agent to make progress), c) the collaboration goal (how an action helps to achieve the shared goal), and d) the other's anticipated beliefs (how an action helps the other). Motives also possess a set of attributes, e.g., *Insistence* or *Failure Disruptiveness* (see Section 3.7.3). These attributes are involved in the comparison of newly generated motives based on the current state of the collaboration. Ultimately, the agent forms or updates a belief about the winning motive in the Mental State.

Goals help the agent to create and update its collaboration plan according to the current private and shared goal content and structure, i.e., the *Specificity*, *Proximity* and *Difficulty* of the goal. Goals direct the formation of intentions to take appropriate corresponding actions during collaboration. Goals also drive the Motivation mechanism to generate required motive(s) in uncertain or ambiguous situations, e.g., to minimize the risk of impasse or to reprioritize goals.

Emotions instances that are elicited by the Appraisal mechanism (see Section 3.7.5 for list of emotion types used in this model). These emotion instances include the agent's own emotions as well as the anticipated emotions of the other which are created with the help of the processes in the Theory of Mind mechanism.

3.7 Attributes of Mental State Elements

Mental states are conscious states of the mind providing the content for cognitive processes. As we discussed *Affective Motivational Collaboration* theory operates with the following Mental State: beliefs, intentions, motives, goals and emotion instances. This Mental State possesses attributes, each of which provides a unique interpretation of the related cognitive entities. These Mental State's attributes are used in different cognitive processes such as the Appraisal mechanism and the Motivation mechanism. We provide more details about these attributes in this section. Attributes that are implemented in our computational framework in Chapter 4 are marked with an asterisk throughout this section. All of the attribute values have a scale, e.g., [0, 1], in our computational framework.

3.7.1 Attributes of Beliefs

The attributes of a belief are involved different processes in *Affective Motivational Collaboration* theory. They impact the evaluation of an event by the Appraisal mechanism, generation of new motives, updates on the collaboration plan, activation of coping strategies and ultimately the self's behavior. The following six attributes of beliefs are related to *Affective Motivational Collaboration* theory.

- **Strength***: Belief strength is about how rigorously the self holds salient beliefs about an object, an entity, or an anticipated behavior. It can be measured through scales, for instance, how probable or likely that belief is, or just whether it is true or false. The strength of a belief can impact the self's appraisal processes, e.g. relevance (see Relevance algorithm in Chapter 4). A belief may be strong, but not necessarily accurate, and vice versa.
- **Saliency***: The saliency of a belief is a cognitive attribute that pertains to how easily the self becomes aware of a belief. This property of a belief has a prominent influence on the self's attention during collaboration. It directs the

self's focus of attention to the most pertinent spatio-temporal salient event (see Relevance algorithm in Chapter 4).

- **Persistence***: It is argued that beliefs form and change due to cognitive and social considerations [42]. Persistent beliefs are very resistant to these changes. However, even persistent beliefs can change. Persistence of goal-related belief(s) influences the appraisal of the relevancy of an event (see Relevance algorithm in Chapter 4).
- **Recency**: The recency of a belief refers to how temporally close a particular belief is to the current state of collaboration. The recency attribute of the self's belief can bias (recency effect) the evaluation processes of the cognitive mechanism during collaboration. It can create a tendency to weight recent events more than earlier ones whenever it is required according to self's Mental State (see satisfaction drive in Chapter 4). The recency of a belief can ultimately impact adopting an appropriate Coping mechanism.
- **Accuracy**: Accuracy of a belief is the relation between that belief and the fact which that belief is about. The accuracy of a belief can be measured by looking at how closely that belief can relate to the truth. The accuracy of a belief as a gradational property can be used in evaluative processes of the self, i.e., Appraisal. It can also impact the self's other goal-driven processes and triggering of an emotion function.
- **Frequency**: The frequency of a belief is related to how regularly it appears as the result of the occurrence of an event. The frequency of beliefs can impact attributes of the self's other Mental State. For instance, beliefs forming or maintaining intentions with direct experiences (see Section 3.7.4) are more likely to occur frequently.

3.7.2 Attributes of Goals

The attributes of a goal impact the processes in *Affective Motivational Collaboration Theory*, especially the processes involved in Motivation and Appraisal mechanisms. The attributes of a goal are important because the Motivation and the Appraisal mechanisms in this theory are goal-driven and attribution of the goals according to the self's standards provides coherency of the processes and their outcomes. We discuss three relevant goal attributes used in our computational framework in this section.

- **Proximity*:** Goals can be distinguished by how far they project into the future during the collaboration. Proximal (short-term) goals result in more related motives and subsequently better self and social-regulation than temporally distant goals. Proximal goals can impact the self's behaviors by influencing the goal management process (see Section 4.4 in Chapter 4). As a result, the self can determine and maintain the collaboration progress towards the shared goal more accurately while operating based on proximal goals.
- **Specificity*:** Goals incorporating specific performance standards are more likely to enhance the self's self-evaluation than general goals. Specific goals raise the self-evaluation performance, because they provide a more accurate baseline for the mechanisms, e.g., Appraisal or Collaboration (see Section 4.4 in Chapter 4), or any arbitration process that the self needs for self-evaluation during collaboration. Consequently, by increasing the self-evaluation performance, the self can improve the level of satisfaction within the collaboration. As an example, holding an object A in a particular position with respect to an object B for a certain amount of time and welding them with a material C is a more specific goal than a general goal of installing an object on another one.
- **Difficulty*:** Goals that are moderately difficult have the most impact on the self and social regulation processes of the self. Conversely, overly easy

or impossible goals usually do not motivate an individual to achieve the goal. Difficult goals increase the probability of a motive's *failure disruptiveness*, and overly easy goals decrease the *importance* of the related motive; in both cases the goals have less chance to be pursued. The existence of a partial shared plan, dependency on the other to perform a task, the failure of the same or similar task in the past all increase the difficulty level of a goal. See Section 4.4 in Chapter 4 for the influence of difficulty of a goal on the goal management process.

3.7.3 Attributes of Motives

According to Sloman, motives can be compared on various dimensions [240]. This comparison is based on motive attributes. In *Affective Motivational Collaboration* theory motives are formed based on the self's existing Mental State under the influence of the Appraisal mechanism. Different elements of Mental State, and the results of self appraisal as well as the reverse appraisal of the other can cause a variety of motives to be formed. The Motivation mechanism needs a set of attributes to compare newly generated motives and choose the one which is most related to the current state of the collaboration. We have chosen the following five motive attributes as most related to the collaboration context.

- ***Importance****: The importance of a motive is determined by the corresponding beliefs about the effects of achieving or not achieving the associated goal (see Section 4.3.1 in Chapter 4). It is a function of belief attributes (e.g., saliency) and the current goal. For instance, if a motive is supported by a belief about the current goal with relatively high attribute values, that motive will become important for the self.
- ***Urgency****: The urgency of a motive defines how much time the self has to acknowledge and address that motive before it is too late. The urgency of a motive is a function of beliefs about the other's mental and emotional state

(see Section 4.3.1 in Chapter 4). For instance, the self responds to an urgent motive due to the existence of an important anticipated outcome for the other, and limited time to accomplish the corresponding tasks, even if those tasks are not important for the self.

- **Insistence:** The insistence of a motive defines the “interrupt priority level” of the motive, and how much that motive can attract the self’s focus of attention. This dimension of motive is associated with what the Appraisal mechanism considers as *relevance* and *desirability* when evaluating an event. Beliefs about successive subgoals and the other’s anticipated Mental State influence the insistence attribute of a motive. Insistent motives have higher priority and are able to interrupt the self’s ongoing tasks.
- **Intensity:** The intensity of a motive determines how actively and vigorously that motive can help the self to pursue the goal if adopted. Motives with higher intensity will motivate the self to apply certain types of coping processes for an obstructed goal to avoid termination of the collaboration. Motives with higher intensity cause the self to find alternative solutions for the problem rather than abandoning the goal and ultimately the collaboration.
- **Failure Disruptiveness:** The failure disruptiveness attribute of a motive determines how significant an interruption failure is to achieving the corresponding goal. In other words, it gives the self a measure of the pleasantness of achieving a related goal. This attribute directs the self’s behavior toward positive and negative outcomes during collaboration.

3.7.4 Attributes of Intentions

The attributes of an intention influence several processes in *Affective Motivational Collaboration Theory*. They can be involved in mechanisms such as Appraisal and Coping. One of the most important uses of intention attributes is to moderate

the intention-behavior relations [59]. Ultimately, the self can show more consistent behavior with respect to its own preceding behaviors and current state of the collaboration. We decided to include the following five intention attributes extracted from the psychology literature in *Affective Motivational Collaboration Theory*.

- **Temporal Status:** The temporal status of an intention can be defined as the extent to which an intention remains consistent over time. The self needs to maintain the stability of its intentions as time passes until the task is performed. Temporally stable intentions helps the other to accurately predict the self's behavior. The anticipated cognitive load of perceiving the self's task by the other impacts the temporal stability of the self's intention. In other words, the temporal stability of an intention moderates the intention-behavior relation of the self during collaboration.
- **Direct Experience:** The direct experience of an intention refers to whether the self previously has performed a task based on a similar intention. The self can refer to the corresponding intention directly experienced in the past before taking a new action. The Mental State associated with the prior experience of an intention can influence the appraisal of a new event requiring the self to perform the same task. For instance, the existence of a direct experience of an intention can impact the degree of the *expectedness* and *controllability* of an event during the collaboration which ultimately guides the Coping mechanism to produce an appropriate behavior.
- **Certainty:** The certainty of an intention is determined by the quality of the underlying motive and the beliefs associated with that motive. The more *strong*, *accurate*, *frequent*, *recent*, *salient* and *persistent* a set of pertinent beliefs of the self are, the more chance the related motive has to be selected. Since the certainty of an intention depends on the associated motive, the nature of the pursued goal also implicitly impacts the certainty of that intention. A goal with a higher *specificity* (see Section 3.7.2) value influences the certainty of the

affiliated intention. The certainty of an intention is an important moderator of the self's intention-behavior consistency.

- **Ambivalence:** The Mental State of the self might contain contradictory intentions towards the pursuit of the same goal, which makes those intentions ambivalent. For instance, the self might already have an intention to perform a task according to the shared plan, while the Appraisal and the Motivation mechanisms dynamically cause formation of a new opposing intention. Furthermore, ambivalent intentions can occur because of the contrast between the self's private goal and the shared goal during the collaboration. The ambivalence attribute of an intention is inversely related to the intention-behavior consistency of the self.
- **Affective-Deliberative Consistency:** The self's intentions possess an affective and a deliberative component. The affective component refers to the affective state and in general the affective evaluation of the self's intention towards its own behavior. However, the deliberate component refers to the self's actual intention which is formed either based on the existing shared plan or under the influence of a new motive generated by the Motivation mechanism. For instance, as an example of affective-deliberative inconsistency, the self can appraise an event as an *urgent* and *uncontrollable* one (which leads the self's emotion towards anger), despite the fact that pursuing the goal related to this intention is required for the satisfaction of the shared plan. In general, mutually consistent affective and deliberate components of an intention positively impacts the consistency of the self's intention and behavior.

3.7.5 Emotion Instances

Each emotion has its own functionality at either the intrapersonal or interpersonal level. Emotions not only regulate the self's internal processes, but also assist the self to anticipate the other's Mental State. In this section, we provide the description

of the emotions that can be elicited during collaboration, and are involved in our scenario (see Section 3.2). Some of these emotions are modeled in Chapter 4 (see Table 4.3). In this theory, to avoid the controversial issue of whether virtual agents or robots can feel emotions, we are going to use the convention of having emotions by the agent or the robot. The agent can also possess beliefs about an emotion instance which is similar to having beliefs about any other proposition.

- ***Joy:*** Joy is the state of an individual’s well-being and is associated with the sense of successful achievement of a goal. Joy reveals one’s sense of pleasure which implies an impending gain for the individual.
- ***Anger:*** Anger can be elicited by an unfair obstacle, hindering the individual’s goal attainment and it is usually triggered by some event (e.g., threat) which provokes a behavioral reaction. Anger functions to set boundaries or escape from dangerous situations, and implies an urgent desire for justice.
- ***Hope:*** Hope is the result of an optimistic evaluation of an event by an individual having expectations of positive and desirable future outcomes related to that event. It is usually a poignant assimilation of the present discontent and the future content implying an imagined or anticipated successful future goal state.
- ***Guilt:*** Guilt is based on self-condemnation in response to a negative outcome of one’s self performance evaluation. It is caused by the violation of others’ beliefs about the self, and others’ standards and bearing significant responsibility for that violation. The occurrence of guilt usually implies the desire to atone in social context.
- ***Pride:*** Pride is a product of the satisfied sense of one’s own actions or decision outcomes. It implies the self-approval of the evaluation oucomes of one’s own actions. Pride is associated with the achievement motivation (see Section

2.5.2) wherein succeeding at a particular goal motivates the corresponding action.

- **Shame:** Shame is produced when one evaluates one's own actions or behaviors and attributes failure to oneself. The individual focuses on specific features of the self which led to failure. Shame implies the existence of remorse.
- **Worry:** Worry is one's emotional attempt to avoid anticipated potential threats or unidentified undesirable events. The individual's concern can be about a real or an imagined issue. Worry implies a fear of a future failure about which one should make a decision or take an action at present.

3.8 Summary

The purpose of *Affective Motivational Collaboration* theory is to explain how to integrate affect-driven processes into collaboration context. *Affective Motivational Collaboration* theory reveals the reciprocal influence of collaboration structure and the appraisal processes. In other words, the appraisal processes can evaluate a collaboration environment based on the information in collaboration structure. Reciprocally, the outcome of appraisals can be used by other processes to maintain progress during collaboration. We believe functions of emotions can be used as these processes to maintain the collaboration structure. Therefore, the required computational framework should include at least Collaboration and Appraisal mechanisms to implement this concept. However, a collaborative agent also needs to be able anticipate and interpret the human collaborator's internal states, e.g., beliefs, emotions. The Theory of Mind mechanism's role in our architecture is to fulfill this concept. Furthermore, a collaborative agent should not only act based on a) emotional state, and b) capability of doing an action at a particular time. We believe the Motivation mechanism is required to reflect the collaborators' need into the formation of the agent's intentions. Motives are crucial mental constructs since they are

affect-regulated (hence, they are influenced by emotions' antecedents, e.g., beliefs), and they are goal-oriented (hence, they can impact the congruency of intentions and actions). Finally, the Coping mechanism works conjointly with Motivation and Appraisal mechanisms to form appropriate intentions with respect to the collaboration status. All of these mechanisms operate based on a knowledge base which includes different elements of mental state as we discussed in this chapter. However, all the processes involved in the different mechanisms can process or update certain aspects of an element of the mental state. Thus, *Affective Motivational Collaboration* theory includes definition of the attributes for all of these elements. In this chapter, we also discussed how different mechanisms can handle events during the collaboration. We believe, verbal and nonverbal behaviors are loaded with information that can impact collaborators' decisions when they occur. Thus, we consider both types of behaviors to be interpreted by our architecture as events during the collaboration. In Chapter 4, we use these theoretical concepts to discuss the implementation of our computational framework.

CHAPTER 4

COMPUTATIONAL FRAMEWORK

4.1 Introduction

Our computational framework includes all of the mechanisms discussed in Chapter 3. The emphasis of our implementation is on the Appraisal, Coping, Collaboration, and Motivation mechanisms, and in general the reciprocal influence of the Appraisal and Collaboration mechanisms (see Section 4.4). In this chapter, we provide concrete algorithms for the theoretical concepts discussed in Chapter 3. These algorithms have been implemented as part of the AMC framework. We also evaluated these algorithms and the overall system in an end-to-end system evaluation user study (see Chapter 5).

There are several appraisal models (e.g., EMA [162]) contributing in different applications such as social sciences, virtual agents, and robotics. However, none of these models have focused on the appraisal processes during collaboration. We believe appraisal plays a key role in collaboration due to its regulatory and evaluative nature. Also, collaboration induces some changes to appraisal processes due to its unique nature. For instance, although the appraisal models mostly use utility to compute the relevance of an event, we have found new cognitive components involved in determining utility because of the influence of the collaboration. These components, such as the recurrence of a belief by the human collaborator or the influence of the human collaborator’s perceived affective state on the robot’s decisions emphasize the fact that collaboration requires additional procedures in appraisal

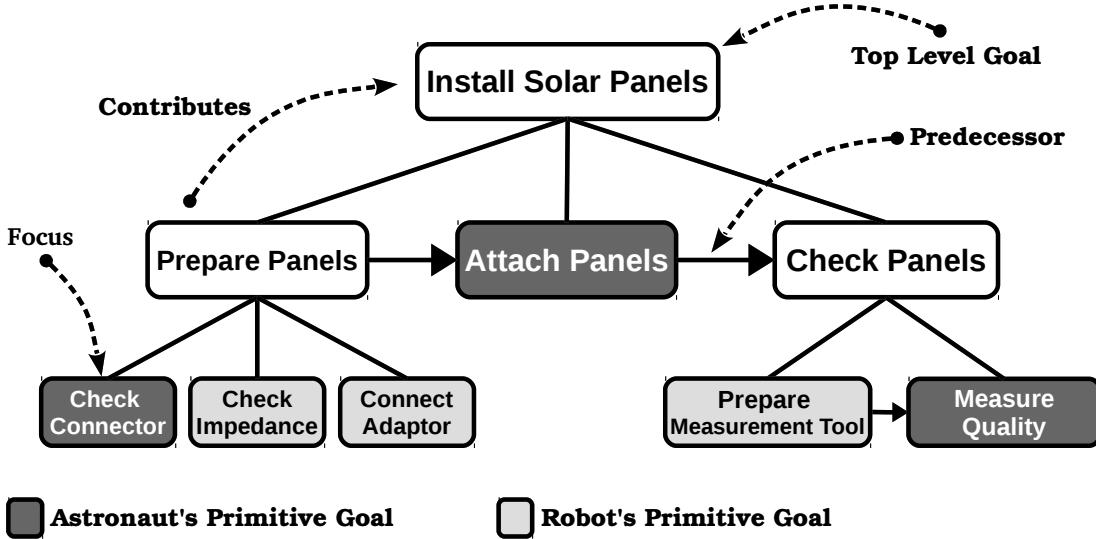


Figure 4.1: Example of collaboration structure.

processes. One of our contributions is to ground general appraisal concepts in the specific context and structure of collaboration.

Furthermore, we believe collaboration and appraisal have reciprocal influences on each other (see Figure 4.3). In this chapter, we also talk about the influence of appraisal on collaboration through the goal management process. Also, we discuss our coping mechanism and strategies within the collaboration context. Then, we provide our computational model of three different motives used in our framework. Finally, we briefly discuss other mechanisms in our framework.

4.2 Collaboration Mechanism

The Collaboration and Appraisal mechanisms (see Figure 3.1) have reciprocal influences on each other. In this section, we focus on information about the collaboration structure which will be incorporated in appraisal processes in Section 4.3. We describe some of the methods in our Collaboration mechanism which are used to retrieve information about the collaboration structure.

The Collaboration mechanism constructs a hierarchy of goals associated with tasks in the form of a hierarchical task network (see Figure 4.1), and also manages the

constraints and other required details of the collaboration including the inputs and outputs of individual tasks, the *preconditions* (specifying whether it is appropriate to perform a task), and the *postconditions* (specifying whether a just-completed task was successful). Collaboration also keeps track of the focus of attention, which determines the salient objects, properties and relations at each point, and shifts the focus of attention during the interaction. For example in Figure 4.1, “Check Connector” is the current (focused) goal¹.

Here, we describe the methods which retrieve information about the collaboration structure, and are used in our algorithms to compute the values of appraisal variables. Some of these methods use the *focus stack* which includes a stack of goals and the top goal on the focus stack represents the current pursuing goal. In these methods, ε_t is the event corresponding to time t , and g_t is a given goal at time t .

- $getPrimitiveGoal(\varepsilon_t)$ returns the unique primitive goal to which the given event (action, utterance, or emotional expression) *directly* contributes; it is only one goal since the robot can only do one primitive action at a time in our collaboration model, i.e, in the goal tree, a given primitive action can only directly contribute to one parent goal. The method returns AMBIGUOUS if it does not find a goal in the plan².
- $getGoalStatus(g_t)$ returns whether g_t ’s status is ACHIEVED, FAILED, BLOCKED, INAPPLICABLE, PENDING, or INPROGRESS.
- $getTopLevelGoal(g_t)$ returns g_t ’s top level goal.
- $precondStatus(g_t)$ returns the status of the precondition for the given goal; whether it is SATISFIED, UNSATISFIED or UNKNOWN. For instance, the precondition for attaching a panel is whether the panel is appropriately located on its frame.

¹The focused goal is the goal that the robot currently pursues.

²Ambiguity introduces some extra complexities which are beyond scope of this thesis.

- $isLive(g_t)$ returns *true* iff all the predecessors of g_t are ACHIEVED and all the preconditions are SATISFIED $\stackrel{\text{def}}{=} \text{PENDING} \vee \text{INPROGRESS}$
- $isFocusShift(g_t)$ returns *true* iff the given goal was not the previous focus (at time t-1).
- $isNecessaryFocusShift(g_t)$ returns *true* iff the status of the previous focus was ACHIEVED [143].
- $isPath(g_1, g_2)$ returns *true* iff there is a path between g_1 and g_2 in a plan tree structure.
- $getContributingGoals(g_t)$ returns g_t 's children in plan tree.
- $getPredecessors(g_t)$ returns g_t 's predecessors in plan tree.
- $getInputs(g_t)$ returns all required inputs for g_t . For example, the goal “Attach Panels” requires the inputs *welding tool* and *panel*.
- $isInputAvailable(g_t)$ returns whether the given input is available. For instance, whether the *welding tool* is available for the goal “Attach Panels”.
- $isFocused(g_t)$ returns whether g_t is the current focus.
- $getResponsible(g_t)$ returns responsible agent(s) for g_t . In a dyadic collaboration, both of the agents jointly can be responsible for a nonprimitive goal, while only one agent (self or other) is responsible for each primitive goal. For instance, both the Robot and the Astronaut are responsible for the nonprimitive goal of “Install Solar Panels”, whereas it is only the Robot who is responsible for the primitive goal of “Prepare Measurement Tool”.

4.3 Appraisal Mechanism and Underlying Processes

In this section, we focus on the specific problem of appraising the *Relevance* (since other appraisals are only computed for relevant events), *Desirability* (since it dis-

crimines facilitating and inhibitory events towards the collaboration progress), *Expectedness* (since it underlies a collaborative robot's attention), and *Controllability* (since it is associated with the agent's coping ability) of events within a collaborative interaction. There are other appraisal variables introduced in psychological [223] and computational literature [92]. We believe most of these variables can be straightforwardly added to our appraisal mechanism whenever they are required. All of the algorithms in this section use the mental state of the robot (discussed in Section 3.6.8) which is formed based on the collaboration structure (see Figure 4.2). These algorithms use the corresponding primitive goal to the most recent event at each turn.

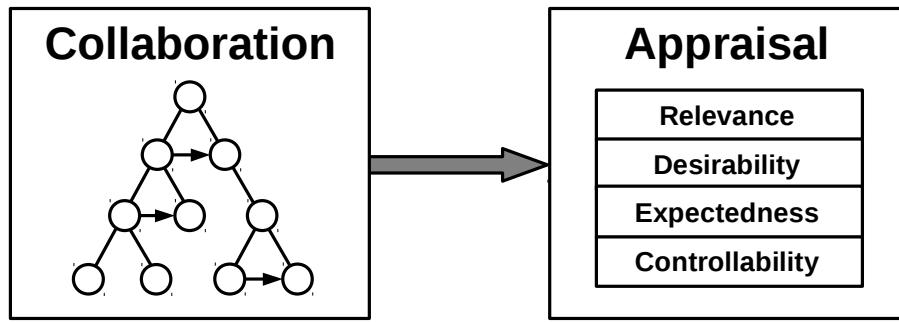


Figure 4.2: Using Collaboration structure in Appraisal (mechanisms in our framework).

4.3.1 Relevance

Relevance is a key appraisal variable since the other appraisal variables are meaningful only for relevant events. Relevance as an appraisal variable measures the significance of an event for the self. An event can be evaluated to be relevant if it has a non-zero utility [162]. However, the utility of an event is also influenced by the other collaborator's emotional expressions as the reflection of the other collaborator's mental state with respect to the status of the collaborative environment. Other appraisal models only consider the utility of an event based on the self's goal and plan.

Algorithm 1 determines the relevance of the given event with respect to the current mental state. The relevance of the event depends on the significance of the event with respect to the collaboration status, which is determined based on the utility of the event as presented in [92, 162]. Our algorithm for computing the relevance of an event during collaboration involves other factors that other appraisal models do not consider. For instance, the human’s perceived affective state, recurrence of a belief, or occurrence of a belief about an unrelated goal by the human play important roles by influencing the utility of an event during collaboration. As a result, evaluating the relevance of events can cause a collaborative robot to respond effectively which can positively impact the status of the shared goal, without dedicating all its resources to every event.

After perceiving an event, the belief about that event represents the event in the robot’s mental state. *getPrimitiveGoal* returns the goal to which the current event contributes, unless it is AMBIGUOUS; g_t represents the shared goal at time (turn) t within the shared plan.

Utility of an Event

We compute the utility ($-1 \leq \mathcal{U} \leq 1$) of the event using the values of the attributes associated with the existing beliefs, and the attributes of the motive associated with the recognized goal (see details below). We use three belief attributes (i.e., *Strength*, *Saliency*, and *Persistence* – see Section 3.6.8) to compute the belief-related part of the utility:

We provide the utility function (\mathcal{U}) in Equation 4.1. This function uses: saliency (S) and persistence (P) of the belief related to the recognized goal, the recognized goal’s status (v), and the aggregation of belief and motive attributes (Ψ) according to Equation 4.1.

$$\mathcal{U}(g_t) = \begin{cases} vP \cdot S^\Psi & \Psi > 0 \\ 0 & \Psi = 0 \end{cases} \quad (4.1)$$

Algorithm 1 Relevance Appraisal Process

```
1: function ISEVENTRELEVANT(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{getPrimitiveGoal}(\varepsilon_t)$ 
3:    $\mathcal{U} \leftarrow \text{GETEVENTUTILITY}(g_t)$ 
4:    $\tau_t \leftarrow \text{GETEMOTIONALTHRESHOLD}(g_t)$ 
5:   if ( $\tau_t \leq |\mathcal{U}|$ ) then
6:     return RELEVANT
7:   else
8:     return IRRELEVANT
```

Intuitively, we use v to generate positive and negative utility values. The v 's value becomes +1 if the status of the corresponding goal is ACHIEVED, PENDING, or IN PROGRESS, and v 's value becomes -1 if the status of the corresponding goal is FAILED, BLOCKED, or INAPPLICABLE. The P influences the value of utility only as a coefficient since recurrent beliefs are not formed frequently during collaboration. The Ψ value indicates the magnitude of the influence of beliefs and motives using their attributes. Hence, the Ψ value impacts the saliency value of beliefs exponentially, helping to differentiate between beliefs.

- *Strength:* The extent to which the preconditions (α), postconditions (β), predecessors (λ), and contributing goals (μ) of a goal are known (SATISFIED or UNSATISFIED) makes beliefs about the goal stronger. An UNKNOWN pre and postcondition status of a goal and its predecessors and contributing goals forms weaker beliefs. For instance, if one knows all predecessors of a pursued goal (e.g., “Check Panels”) are SATISFIED (i.e., “Fix Panels” and “Prepare Panels”), failure of the pursued goal will elicit one’s negative affective state (due to the strong beliefs related to the goal); whereas not knowing the status of the goal-related factors (e.g., whether the Astronaut could find the tool to fix a panel) causes one to form weaker beliefs about the goal.

In equation 4.2, the subscript k refers to the *known* goal-related factors (SATISFIED

or UNSATISFIED); whereas the subscript *all* includes both *known* and *unknown* goal-related factors. In this equation, both urgency (γ) and importance (η) attributes of motives can impact the outcome of the goal-related belief attributes' ratio, and ultimately the Ψ value.

$$\begin{aligned} \Psi &= \frac{\alpha_k + \beta_k + \lambda_k + \mu_k}{\alpha_{all} + \beta_{all} + \lambda_{all} + \mu_{all}} + \eta + \gamma & (4.2) \\ \eta, \gamma &\in \mathbb{N}, & \eta, \gamma \geq 0 \\ \alpha_k, \beta_k, \lambda_k, \mu_k &\in \mathbb{N}, & \alpha_k, \beta_k, \lambda_k, \mu_k \geq 0 \\ \alpha_{all}, \lambda_{all}, \mu_{all} &\in \mathbb{N}, & \alpha_{all}, \lambda_{all}, \mu_{all} \geq 0 \\ \beta_{all} &\in \mathbb{N}, & \beta_{all} \geq 1 \end{aligned}$$

- *Saliency (S)*: Beliefs related to the focused goal are more salient than beliefs related to any other goal in the plan; according to Figure 4.1, if one of the collaborators is preparing a solar panel, beliefs related to all of the other *live* (PENDING or IN PROGRESS) goals (e.g. “Connect Adaptor”) will be less salient than beliefs related to the focused goal, i.e., “Check Connector”. Beliefs’ saliency decreases according to their corresponding *live* goal’s distance from the focused goal in the shared plan. *Non-live* goals will not be salient.
- *Persistence (P)*: The recurrence of a belief over time (turns) increases the persistence of the belief. Beliefs occurring only once have the lowest value of persistence. For instance, if the Astronaut repeatedly says that she can not find the measurement tool to check the connector, the Robot could pursue a new goal outside of the shared plan to acknowledge Astronaut’s concern.

We also use two motive attributes discussed in Section 3.6.8 to compute the motive related part of the utility (\mathcal{U}):

- *Urgency (γ)*: There are two factors impacting the urgency of a motive: a) whether the goal directing the given motive is the predecessor of another goal

for which the other collaborator is responsible, and b) whether achieving the goal directing the given motive can mitigate the other collaborator’s negative valenced affective state. For instance, if the Robot has a private goal to fetch another panel while the Astronaut is waiting for the Robot to connect the adaptor, connecting the adaptor will be more urgent than Robot’s private goal.

- *Importance (η):* A motive is important if failure of the directing goal causes an impasse in the shared plan (i.e., no further goal is available to achieve), or achievement of the directing goal removes an existing impasse. For example, if the Robot cannot find the adaptor (an impasse to connect the adaptor), and the Astronaut provides another adaptor (external motive), the new motive becomes important to remove the impasse in the shared plan.

The significance of an event in a collaborative environment is based on the utility of the event and the human’s perceived affective state. The human’s perceived affective state influences the relevance of the event in the form of a threshold value τ_t in Algorithm 1. In Equation 4.3, we use the valence of the perceived affective state (\mathcal{V}_{e_h}) to compute τ_t .

$$\tau_t = -0.5\mathcal{V}_{e_h} + 0.5 \quad (4.3)$$

$$\mathcal{V}_{e_h} \in \mathbb{R}, \quad -1 \leq \mathcal{V}_{e_h} \leq 1$$

Hence, perceiving human’s positive affective state (e.g., happiness) reduces the threshold value which makes the robot find an event RELEVANT with even a slightly positive utility. Similarly, an event can be considered IRRELEVANT even though the utility has a relatively positive value, because of perceiving the human’s negative affective state.

4.3.2 Desirability

Desirability characterizes the value of an event to the robot in terms of whether the event facilitates or thwarts the collaboration goal. Desirability captures the valence of an event with respect to the robot’s preferences [92]. In a collaborative robot, preferences are biased towards those events facilitating progress in the collaboration. Desirability plays an important role in the overall architecture; it makes the processes involved in the other mechanisms (e.g., Motivation and Theory of Mind) and consequently the robot’s mental state, congruent with the collaboration status which is a collaborative robot’s desire. Therefore, it causes the robot to dismiss events causing inconsistencies in the robot’s collaborative behavior. Moreover, desirability is also crucial from the collaboration’s point of view.

Algorithm 2 defines a process in which the desirability of an event is computed with regard to the status of the shared goal; i.e., it operates based on whether and how the event changes the status of the current shared goal. It distinguishes between the top level goal and the current goal because the top level goal’s change of status attains a higher positive or negative value of desirability. For instance, failure of the top level goal (e.g., installing solar panel) is more undesirable than failure of a primitive goal (e.g., measuring the quality of the installed panel).

A top level goal’s status must be ACHIEVED (i.e., SATISFIED postcondition) to consider the event MOST-DESIRABLE. When the goal’s status is FAILED (i.e., UNSATISFIED postcondition) or BLOCKED, the associated event has the MOST-UNDESIRABLE or UNDESIRABLE values respectively. A goal is BLOCKED if any of the required goals or goals recursively through the parent goal are not ACHIEVED. An INAPPLICABLE goal is also considered as UNDESIRABLE. A goal is INAPPLICABLE if any of its predecessors are not ACHIEVED, and/or its preconditions are not SATISFIED. For PENDING and INPROGRESS top level goals, the status of the current goal associated with the top level goal determines the status of the event ε_t . Only a non-primitive goal can have INPROGRESS status, if it has been started but is not yet completed.

Algorithm 2 Desirability Appraisal Process

```
1: function IsEVENTDESIRABLE(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{getPrimitiveGoal}(\varepsilon_t)$ 
3:    $g_{top} \leftarrow \text{getTopLevelGoal}(g_t)$ 
4:   if ( $\text{getGoalStatus}(g_{top}) = \text{ACHIEVED}$ ) then
5:     return MOST-DESIRABLE
6:   else if ( $\text{getGoalStatus}(g_{top}) = \text{FAILED}$ ) then
7:     return MOST-UNDESIRABLE
8:   else if ( $\text{getGoalStatus}(g_{top}) = \text{BLOCKED}$ ) or
9:     ( $\text{getGoalStatus}(g_{top}) = \text{INAPPLICABLE}$ ) then
10:    return UNDESIRABLE
11:   else if ( $\text{getGoalStatus}(g_{top}) = \text{PENDING}$ ) or
12:     ( $\text{getGoalStatus}(g_{top}) = \text{INPROGRESS}$ ) then
13:       if ( $\text{getGoalStatus}(g_t) = \text{ACHIEVED}$ ) then
14:         return DESIRABLE
15:       else if ( $\text{getGoalStatus}(g_t) = \text{FAILED}$ ) then
16:         return MOST-UNDESIRABLE
17:       else if ( $\text{getGoalStatus}(g_t) = \text{BLOCKED}$ ) or
18:         ( $\text{getGoalStatus}(g_t) = \text{INAPPLICABLE}$ ) then
19:           return UNDESIRABLE
20:         else if ( $\text{getGoalStatus}(g_t) = \text{PENDING}$ ) or
21:           ( $\text{getGoalStatus}(g_t) = \text{INPROGRESS}$ ) then
22:             return NEUTRAL
```

A goal can be PENDING if it is live, or if it is a non-primitive goal that has not been started yet. ACHIEVED current goals mark an event (ε_t) as DESIRABLE, while FAILED or BLOCKED current goals render the event associated with them as MOST-UNDESIRABLE and UNDESIRABLE respectively. PENDING or INPROGRESS current goals mark their associated events as NEUTRAL.

4.3.3 Expectedness

Expectedness is the extent to which the truth value of a state could have been predicted from a causal interpretation of an event. In the collaboration context, the expectedness of an event evaluates the congruency of the event with respect

Algorithm 3 Expectedness Appraisal Process

```
1: function IsEVENTEXPECTED(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{getPrimitiveGoal}(\varepsilon_t)$ 
3:    $g_{top} \leftarrow \text{getTopLevelGoal}(g_t)$ 
4:   if ( $\text{isLive}(g_t)$ ) then
5:     if ( $\neg\text{isFocusShift}(g_t)$  or  $\text{isNecessaryFocusShift}(g_t)$ ) then
6:       return MOST-EXPECTED
7:     else
8:       return EXPECTED
9:   else
10:    if ( $\text{isPath}(g_t, g_{top})$ ) then
11:      return UNEXPECTED
12:    else
13:      return MOST-UNEXPECTED
```

to the existing knowledge about the shared goal. Thus, expectedness underlies a collaborative robot’s attention. The collaboration mechanism uses expectedness to maintain the robot’s attention and subsequently its mental state with respect to the shared goal. Reciprocally, the appraisal mechanism uses the underlying information of the collaboration structure to evaluate the expectedness of an event [247].

In Algorithm 3 we define the process of computing the expectedness based on the shared plan and status of the shared goal. The key point in this algorithm is the status of the current shared primitive goal (g_t), which is associated with the event ε_t and its relationship with the top level goal (g_{top}).

The intuition captured here is that one expects the current goal to be finished before undertaking another activity, but the goals that can be the next focus of attention are also to be expected. Therefore, if the goal is live, the algorithm checks whether the goal has not changed, or whether the interpretation of the last event results in a necessary focus shift. Shifting the focus to a new goal is necessary when the former goal is achieved and a new goal is required. Consequently the new event is the MOST-EXPECTED one. However, even if the focus shift is not necessary, the new event can be considered as EXPECTED, since the corresponding goal is already

Algorithm 4 Controllability Appraisal Process

```
1: function IsEventControllable(Event  $\varepsilon_t$ )
2:    $g_t \leftarrow \text{getPrimitiveGoal}(\varepsilon_t)$ 
3:    $\mathcal{M} \leftarrow \text{GetAgencyRatio}(g_t)$ 
4:    $\mathcal{R} \leftarrow \text{GetAutonomyRatio}(g_t)$ 
5:    $\mathcal{P} \leftarrow \text{GetSuccPredecessorsRatio}(g_t)$ 
6:    $\mathcal{I} \leftarrow \text{GetAvailableInputs}(g_t)$ 
7:    $\mathcal{V}_{e_h} \leftarrow \text{GetEmotionValence}(g_t)$ 
8:    $\omega \leftarrow \text{GetWeights}(g_t)$ 
9:    $\mathcal{X} \leftarrow \frac{\omega_0 \cdot \mathcal{M} + \omega_1 \cdot \mathcal{R} + \omega_2 \cdot \mathcal{P} + \omega_3 \cdot \mathcal{I}}{\omega_0 + \omega_1 + \omega_2 + \omega_3} + \mathcal{V}_{e_h}$ 
10:  if ( $\mathcal{X} > 0$ ) then
11:    return CONTROLLABLE
12:  else
13:    return UNCONTROLLABLE
```

live. For goals that have not yet been started (that is, are not live), the algorithm must determine how unexpected it would be to pursue one now; if the goal is at least in the plan, i.e., on the path to the top level goal, it is just UNEXPECTED while any others are MOST-UNEXPECTED.

4.3.4 Controllability

Controllability is the extent to which an event can be influenced; it is associated with a robot's ability to cope with an event [92]. Thus, a robot can determine whether an event's outcome can be altered by actions under either of the collaborators' control. In other words, controllability is a measure of a robot's ability to maintain or change a particular state as a consequence of an event.

Controllability is important for the overall architecture. For instance, the robot can choose to ask or negotiate about a collaborative task which is not controllable, or form a new motive to establish an alternative goal for the current uncontrollable

event. In general, other mechanisms in the architecture use the controllability output in their decision-making processes; while controllability uses information from the collaboration structure, e.g., predecessors of a goal.

An important determinant of one’s emotional response is the sense of control over occurring events. This sense of subjective control is based on one’s reasoning about the self’s power. For instance, the robustness of one’s plan for executing actions can increase one’s sense of power and subsequently the sense of control. In the collaboration context, we have translated the sense of control into a combination of four different factors including a) *agency* and b) *autonomy* of the robot, as well as the ratios of c) *successful predecessors*, and d) the *available inputs* of a given goal (i.e., g_t) in the shared plan.

In Algorithm 4, we partially compute the controllability of an event based on the above four factors. We use weighted averaging of these factors to determine their impact on the controllability of an event (line 9). The value of all these weights are set to *1.0* for the purpose of simplicity (**getWeights**). These weights can be adjusted after further investigating the influence of these factors, and implementing other mechanisms in the overall architecture. We believe that the human’s perceived affective state also impacts the controllability of an event (**getEmotionValence**). The ($-1.0 \leq V_{e_h} \leq 1.0$) is the valence value of the human’s perceived affective state. Positive emotions, e.g., happiness, possess positive values, and negative emotions, e.g., anger, have negative values. The magnitude of this value can change with respect to the intensity of the perceived affective state. Thus, a positive controllability value indicates that an event is **CONTROLLABLE**; otherwise **UNCONTROLLABLE**.

GetAgencyRatio: *Agency* is the capacity of an individual to act independently in a given environment. In a collaborative environment collaborators are sometimes required to act independently of each other. Hence, they need to have some internal motives that are formed based on their own Mental State rather than motives that are reinforced by the other. These internal motives will lead the collaborators to acquire new intentions when required. If the robot’s mental state possesses only

an internal motive supporting the recognized goal, we consider a maximum agency value denoted as \mathcal{M} in Algorithm 4 (i.e., $\mathcal{M} = 1.0$); otherwise we consider the minimum agency value (i.e., $\mathcal{M} = 0.0$).

GetAutonomyRatio: *Autonomy* is the ability to make decisions without the influence of others, and implies acting on one's own and being responsible for that. In a collaborative environment, tasks are delegated to the collaborators based on their capabilities. Therefore, each collaborator is responsible for the delegated task and the corresponding goal. In Algorithm 4, \mathcal{R} denotes the value of autonomy with regard to the goal g_t . This value ($0.0 \leq \mathcal{R} \leq 1.0$) is the ratio of the number of goals contributing to g_t for which the robot is responsible over the total number of contributing goals, if the goal associated with the current event is a nonprimitive goal. However, if the associated goal of the current event corresponds to a primitive goal the value of \mathcal{R} would be 0.0 (if the human is responsible) or 1.0 (if the robot is responsible). In general, higher autonomy leads to a more positive value of controllability.

GetSuccPredecessorsRatio: The structure of a shared plan contains the order of the required *predecessors* of a goal. Predecessors of a goal, g_t , are goals that the collaborators should achieve before trying to achieve goal g_t . We use the ratio of successfully achieved predecessors of the associated primitive goal over the total number of predecessors of the same goal. If all of the predecessors of the given goal are achieved, then $\mathcal{P} = 1.0$ which is the maximum value for \mathcal{P} . On the contrary, failure of all of the predecessors will lead to $\mathcal{P} = 0.0$. Therefore, a higher \mathcal{P} value positively impacts the value of controllability for the current event.

GetAvailableInputs: Finally, *inputs* of a task are the required elements that the collaborators use to achieve the specified goal of the task. These inputs are also part of the structure of a shared plan. We compute the ratio of the available required inputs over the total required inputs of the goal associated with the current event. This value (denoted as \mathcal{I} in Algorithm 4) will be set between 0.0 and 1.0. Similar to the other factors in the controllability process, the closer the value of \mathcal{I}

gets to 1.0, the more positive impact it has on the overall controllability value of the event.

In summary, the output of these four appraisal processes serves as critical input for the other mechanisms of the Affective Motivational Collaboration framework, shown in Chapter 3. By providing adequate interpretation of events in the collaborative environment, the appraisal mechanism enables the robot to carry out proper collaborative behaviors.

4.4 Goal Management

A collaborative robot needs to be able to regulate and manage shared goals during collaboration. Emotion has a crucial influence on this goal management process. In this section, we provide a cost function that we use to choose the goal in the shared plan with the lowest cost value out of a set of alternative goals. This cost function is a) based on the goal attributes, b) with respect to the reverse appraisal of the perceived affective state, and c) the appraisal of the collaborative environment. Adding goal management process to Collaboration mechanism is one of our contributions.

Goals represent a key part of the context during collaboration. However, not all goals are appropriate to pursue at the moment, depending on conditions. In fact, it can be destructive for a collaboration to pursue a plausible goal in a poor context. Therefore, a collaborative robot must be able to manage shared goals during collaboration. The goal management process has a critical influence on a collaborative robot's behavior by maintaining or shifting the focus of attention to an appropriate goal based on the collaboration status.

Changes in a collaboration environment alter the relative importance of alternative goals. These changes can reflect the collaborators' internal changes and the influence of their actions. In a collaboration environment, emotions represent the outcome of underlying mental processes of the collaborators. Emotions have many

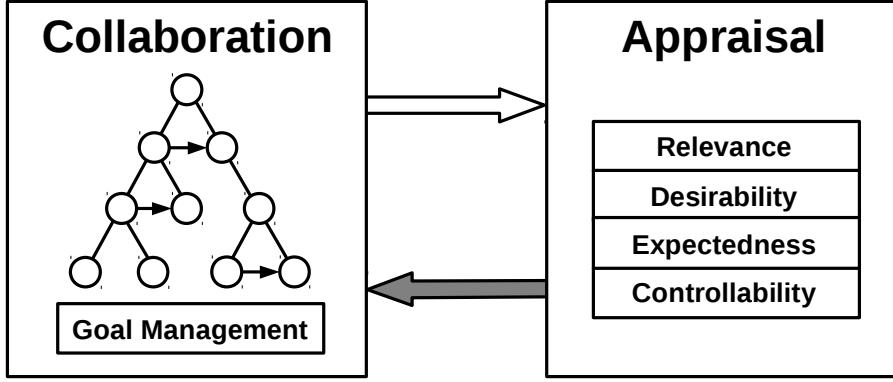


Figure 4.3: Using Appraisals' outcome to influence Collaboration structure (mechanisms in our framework).

different functions [226] including goal management. Goal-oriented emotions such as anger, frustration and worry regulate the mental processes influenced by one's internal goals. In our ongoing example, a robot and an astronaut are collaborating to install solar panels. When one of the astronaut's goals is blocked, the robot must manage the shared goals in order to prevent failure of the collaboration. By using reverse appraisal [64] of the astronaut's affective state and its own appraisal of individual goals, the robot is able to successfully shift the focus of attention from the blocked goal (eliciting worry in the astronaut) to an appropriate one to maintain the collaboration. A similar example is provided in Chapter 5.

Here, we describe the goal management process in our framework using an astronaut-robot collaboration example. We introduce the goal management process based on a cost function including the influence of affective appraisal and reverse appraisal processes. Goal management is a crucial part of our investigation of the reciprocal influence of appraisal on a collaboration structure (see Figure 4.3).

As we mentioned earlier, we use four appraisal variables including: relevance, desirability, expectedness and controllability. The outcome of each appraisal process is a specific value for the corresponding appraisal variable. The vector containing these appraisal variables can be mapped to a particular affective state at each point in time when required (see Algorithms in Section 4.3). Moreover, the functions

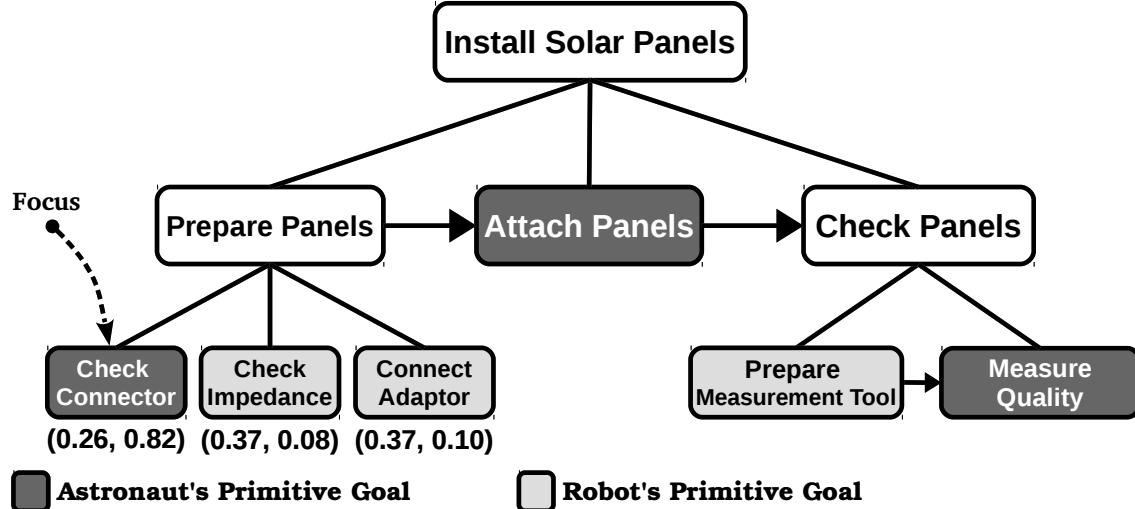


Figure 4.4: Cost values indicated by tuples with (second number) and without (first number) the influence of emotions.

of emotions, such as goal management, in a social setting and the meaning of the collaborator’s perceived affective state in collaboration context are also important.

A collaboration structure provides a hierarchy and constraints of the shared goals in the form of a shared plan which contains both the robot and the human collaborator’s goals. The robot pursues the goals for which the robot is responsible in the shared plan. However, there can be several live goals available for the robot to pursue at each point in time during collaboration. A goal is live *iff* all of its predecessors are achieved and all of its preconditions are satisfied. Therefore, a collaborative robot requires a mechanism to choose between a set of live goals. We believe appraisal processes are crucial to choose between the available live goals, since the appraisals are the immediate outcome of the robot’s assessment of the collaboration environment.

For instance, Figure 4.4 shows a non-primitive “Prepare Panels” goal decomposed into three unordered primitive goals. Therefore, if “Prepare Panels” is live, its primitive goals can be pursued by the responsible agent. In our example, the astronaut is responsible for the “Check Connector” goal; the robot is responsible for the remaining two primitive goals. According to the collaboration mechanism in

our overall framework, “Check Connector” is in focus, with the astronaut pursuing this goal. Suddenly, however the astronaut tells the robot that she can not find the connector and she is *worried* about failure of this goal. The robot’s response to this situation will be explored below as we discuss details of our cost function.

Equation 4.4 shows the function to calculate the cost of each live goal. Goal management algorithm chooses the minimum cost goal. The base in the equation calculates the cost of pursuing any given goal. The three functions used to calculate the cost are: *proximity* $P(g)$, *difficulty* $D(g)$, and *specificity* $S(g)$ (see equations 4.6 to 4.8).

$$Cost(g) = (\omega_0.P(g) + \omega_1.D(g) + \omega_2.\frac{1}{S(g)+1})^\Gamma \quad (4.4)$$

For simplicity in this example, we assume equal values for the weights: $\omega_i=1$.

$$\Gamma = -C[(R_r + 1)D_r + \alpha(R_h + 1)D_h] \quad (4.5)$$

The exponent part of our cost function (Equation 4.5) captures a) the influence of the human’s perceived emotional instance, and b) the influence of self appraisal of the given goal. $R_h \in [0, 1]$ and $D_h \in [-1, 1]$ are the relevance and desirability values respectively, which are based on the *reverse* appraisal of the human’s perceived affective state. For instance, if the astronaut is *worried*, D_h is negative, e.g., -0.8 (depending on how undesirable the event is according to reverse appraisal), and R_h will be 1 for the active goal and its value descends to 0 for other live goals depending on their distance to the active goal in the shared plan (e.g., 0.1).

$R_r \in [0, 1]$ and $D_r \in [-1, 1]$ are relevance and desirability values, provided by the *self* appraisal functions for all of the live goals. For instance, for the active goal for which the astronaut was *worried*, D_r can be positive, e.g., 0.8 (depending on the self’s desirability appraisal function); R_r can be 1, since the active goal is relevant for the robot. These values will change for the other live goals depending on how relevant they are with respect to the collaboration status (e.g., 0.9 and 0.8).

Finally, $C \in [1, \infty)$ is a constant (e.g., 2) used to control the influence of affect on cost value. It is negative since undesirability (negative values) should increase the cost. $\alpha \in [1, \infty)$ is another constant (e.g., 3) used to control the importance of reverse appraisal relative to self appraisal.

The *proximity* of a goal indicates how far the goal is from the current active goal in the shared plan. It is calculated by the distance function (Equation 4.6) which returns the number of edges between the current active goal g_{act} , and the given goal g in the shared plan. In our example, $P(g)$ is 2 for both “Check Impedance” and “Connect Adaptor” goals.

$$P(g) = \max\{1, \text{distance}(g_{act}, g)\} \quad (4.6)$$

The *difficulty* of a goal is a function of three parameters (Equation 4.7) which consider the difficulty based on a) topology of the shared plan tree (domain independent), and b) the amount of effort required to pursue a given goal (domain dependent). The $\sum pred_e(g)$ is the sum of efforts that all the *predecessors* of a given goal g require. The $\sum desc_e(g)$ is the sum of efforts that all the *descendants* of a given goal g require. The effort values represent the amount of effort for the goals with respect to the domain. In our example, we assume the values of all the goal efforts are 1 for simplicity. The $H(g)$ is the height of the given goal g . The heights of all primitives under “Prepare Panel” goal are 0 in our example.

$$D(g) = (H(g) + 1) \times \left[\sum_{m=0}^M pred_e(g) + \sum_{n=0}^N desc_e(g) \right] \quad (4.7)$$

The *specificity* of a goal is the function of *depth* (distance from the root) and *degree* (number of children in the graph) of a given goal g . The first non-primitive goal (root) is the least specific goal, and the primitives (leaves) are the most specific goals. As calculated based on Figure 4.4, the values of $S(g)$ for the three primitives under the “Prepare Panels” are 2.

$$S(g) = \frac{\text{depth}(g)}{\text{degree}(g) + 1} \quad (4.8)$$

The tuples below the three leftmost primitive goals in Figure 4.4 indicate cost values of each goal: the first number in each tuple is the normalized cost value without the influence of the affective part of the cost function, i.e., the exponent is equal to 1 in Equation 4.4; the second number of each tuple indicates the normalized value of the cost including the influence of affective appraisal and the astronaut’s perceived affective state.

Based on our cost function, the cost of completing the primitive goal “Check Connector” is 0.82 (see Figure 4.4). As shown, when affect is not considered the cost is 0.26; the negative affective state of the astronaut (worry) significantly increases the cost of the current goal, and also impacts the other two primitive live goals under the same parent. Therefore, instead of insisting on pursuing the same blocked goal which has caused the astronaut’s negative affective state, the robot can mitigate the astronaut’s emotions by adapting to her worry. The robot shifts the focus of attention to “Check Impedance” to maintain progress and prevent failure of the collaboration. We use our proposed cost function in our goal management algorithm to integrate affective appraisal into the collaboration mechanism in our framework.

4.5 Coping Mechanism and Strategies

We have developed an algorithm for the Coping mechanism to determine how the agent would respond to events using our framework. Coping operates based on the antecedents of the appraisals and includes strategies that the agent chooses to make changes, directly or indirectly, that would have a desired impact on the appraisal. In our computational framework, the output of Coping is one or multiple intentions. Our Coping mechanism includes a set of coping strategies that can be triggered based on different conditions (see Table 4.1). All of these coping strategies are known in the literature, however, none of these strategies are applied in col-

laboration context. Some of our coping strategies, i.e., *planning*, *active coping* and *seeking social support for instrumental reasons*, are categorized as problem-focused and others, i.e., *acceptance*, *mental disengagement*, and *shifting responsibility*, are categorized as emotion-focused strategies as described in [92]. We implemented these six coping strategies because they let our agent demonstrate distinct behaviors with respect to the output of the appraisal mechanism and the agent’s mental state in our framework.

4.5.1 Planning

The *planning* coping strategy works based on the shared plan and the task structure introduced as an input to our framework. The task structure includes the hierarchy and ordering of the tasks, the required inputs of each task as well as the preconditions and postconditions of individual tasks. We use this task structure to create our shared plan which includes the primitive and non-primitive goals that our agent and its collaborator want to achieve throughout their collaboration. Therefore, our agent executes actions related to its own goals based on this shared plan, and uses the same shared plan to associate goals and their status with the human collaborator. To achieve a goal the agent is required to execute an action, and to execute an action the agent needs to have the right intention. In our framework, whenever this coping strategy is activated the Coping mechanism provides the selected intention to the Action mechanism. The Action mechanism executes an action based on the given intention to achieve the corresponding goal in the shared plan.

4.5.2 Active Coping

The *active* coping strategy can provide one or all of the following three different intentions with respect to whether this coping strategy is activated and the required conditions are provided. Firstly, this coping strategy can provide an intention to *acknowledge* the human’s emotions. For instance, if the human expresses an affective

Table 4.1: Conditions for selecting candidate coping strategies

Coping Strategy	Emotions (AND)		Need [a AND (b OR c)]			Ability Controllability
	Other	Self	Satisfaction Motive (a)	Achievement Motive (b)	External Motive (c)	
Planning	Neutral Positive	Any	-/+	high +	high +	High
Active Coping	Any	Neutral Negative	-/+	med +	med +	High
Seeking Social Support for Instrumental Reasons	Neutral Positive	Any	-/+	low +	low +	Low
Acceptance	Negative	Negative	high -	high -	high -	No
Mental Disengagement	Neutral Negative	Neutral Negative	low/med -	low/med -	low/med -	No
Shifting Responsibility	Neutral Positive	Negative	high -	-/+	-/+	No Low

state with negative valence, the agent can acknowledge human's negative affective state accordingly. Secondly, the active coping strategy can provide an intention to *respond* to the human if the human asks a question. Currently, in our framework, the agent can respond to the human if the human asks the agent: a) what input is required to achieve a goal, b) how to do a task to achieve a goal, c) to achieve a goal, d) who is responsible to achieve a given goal. For instance, if the human asks the agent to achieve a goal, the active coping strategy forms an intention to either accept the human's proposal (if achieving the given goal is controllable for the agent), or reject the human's proposal (if it is not controllable for the agent). Thirdly, the active coping strategy can form an intention to *delegate* a task to the human collaborator. The intention for task delegation can be formed if the agent fails to achieve its own goal, and the human's perceived affective state is not negative. As mentioned earlier, any or all of these intentions can be formed if active coping is selected. The agent acts accordingly by passing these intentions to the Action mechanism. For instance, if the human is frustrated about a failure that occurred when using a tool to perform its own task and asks the agent whether the agent can provide its own tool, the active coping strategy forms a new intention to acknowledge the human's frustration and responds to the human by providing the right tool (input) to use and fulfill the task. In this example, there will be no new intention to delegate a new goal to the human since the agent perceives the human's negative affective state.

4.5.3 Seeking Social Support for Instrumental Reasons

The *seeking social support for instrumental reasons* strategy forms new intentions for the agent whenever the agent needs the human's help and needs to ask questions from the human collaborator to make progress in collaboration. The questions that our agent can ask are the reciprocal of those questions that the human can ask and the human can respond as we mentioned in Section 4.5.2. Therefore, our agent can ask a) what input is required to achieve a goal, b) how to do a task to achieve a

goal, c) the human to achieve a goal, d) who is responsible to achieve a given goal. Reciprocally, again, the agent expects the human collaborator to accept or reject the agent’s proposals. In our framework, whenever this strategy is activated the agent considers human’s perceived affective state. For instance, if the human is worried about the outcome of a task failure, the agent does not form an intention to ask questions about any of the above cases and consequently prevents asking for more help.

4.5.4 Acceptance

The *acceptance* coping strategy forms an intention to drop the intention of pursuing a goal. In our framework, if this strategy becomes activated, the intention to pursue the current goal will be dropped; see Table 4.1. For instance, if the human has failed to achieve a goal due to the lack of a required input, and the agent is not able to pursue another goal and the agent is not able to provide the required input, this strategy becomes activated. The acceptance strategy also forms an intention to inform the human collaborator about the agent’s decision on not pursuing the current goal.

4.5.5 Mental Disengagement

The *mental disengagement* coping strategy forms new intention to lower the negative emotional intensity associated with a goal in the event of a failure or an impasse. We use our goal management algorithm (see section 4.4) as the result of selecting this strategy to dissociate from the current goal in the collaboration process and subsequently disengage the collaborator from a negative event (e.g., failure to achieve a goal). This disengagement helps the agent to lower the utility of an unsuccessful goal achievement attempt and focus on other achievable goals with respect to their costs to facilitate progress of collaboration. In our framework, this coping strategy forms an intention to run the goal management process. As the result of mental

disengagement activation, the Coping mechanism also forms another intention to inform the human about the outcome of the goal management process, i.e., whether the agent proposes switching to pursue another goal with lower cost, or if there is not much the agent can do since there is no other goal with a lower cost to pursue. The process and example of choosing another goal with a lower cost are shown in Section 4.4.

4.5.6 Shifting Responsibility

The *shifting responsibility* strategy forms a new intention to shift the blame from the agent to another entity. In our framework, we use this strategy to mitigate the influence of negative events causing negative emotions in the agent or the human collaborator. For instance, if this strategy becomes activated as a result of a failure, a new intention will be formed to blame the other collaborator, or the third person who provided the input (if the task needed a tool as an input). It can also form an intention to give the credit to the human collaborator to mitigate human's negative emotions.

4.5.7 Activation of Coping Strategies

In our Coping mechanism, there are three activation criteria for each coping strategy. The first criterion is the conjunction of emotion valences of the self and the other collaborator (see Emotion Valence column in Table 4.1). For instance, if the valence of the human collaborator's affective state is *negative* **and** the valence of the agent's affective state is also *negative*, the active coping (2nd row), the acceptance (4th row), and the mental disengagement (5th row) coping strategies are the coping strategy candidates that have potential to become activated if the other activation criteria also exist for any of them. For example, if the human collaborator is frustrated and the agent's elicited emotion is guilt, the three above mentioned coping strategies become potential candidates to be selected as the agent's active coping strategy.

The second criterion is the need for the agent to cope with an event. The values of our three different motives (i.e., *satisfaction*, *achievement*, and *external*) are involved in the decision of whether there is a need for a particular coping strategy to become activated. We use conjunction of satisfaction motive's value with the disjunction of achievement and external motives. For instance, if we have highly negative values for all three motives for the potential candidates of coping strategies based on the example we mentioned above, the acceptance coping strategy will be selected as the strategy with the highest need for the agent. For example, this kind of condition can occur when the agent fails doing its own task and pursuing the current goal (negative satisfaction motive), and can not find another goal to overcome the impasse (negative achievement motive). The details about how the motive values are computed is presented in Section 4.6. Finally, the ability to cope with an event is the third criterion that impacts the decision of whether the selected coping strategy can be activated. The controllability of an event represents whether the agent is able to control the situation occurring with the given event. In our example, if the agent finds the event uncontrollable, the acceptance coping strategy becomes activated (see Table 4.1).

4.6 Motivation Mechanism

As we discussed in Chapters 2 and 3, motives are goal-driven affect-regulated constructs indicating an urge related to their goal. There are several motives in psychological and computational literatures as we reviewed in Chapter 2. However, none of these computational models have particularly focused on the application of motives in the collaboration context. We believe motives have a key role to fill the gap between the Appraisal and Coping mechanisms in a collaborative environment. In fact, motives can improve the intention formation process with respect to the urge of pursuing a goal by considering the emotional states of the collaborators. As shown in Table 4.1, it is not enough to choose a particular coping strategy only by

knowing how controllable it is to pursue the given goal. For instance, motive values can help the agent to choose between *Acceptance* and *Mental Disengagement* when pursuing a goal is not controllable.

As mentioned in Chapter 2, we provided three types of prominent motives in the literature; i.e., *achievement*, *affiliation* and *power*. However, due to the fact that not all of these motives fit in to the dyadic collaboration context, we developed our own computational models of motives in our framework, including: *satisfaction*, *achievement*, and *external* motives. Our approach in general is inspired by Merrick and Shafi's work in [169] modeling motives using curves generated by sigmoid functions. In our work, curves are influenced by the valence of the human collaborator's perceived affective state. This section provides more details about how different curves model different motives in our computational framework. We use the values of these three motives in other mechanisms including the Coping mechanism as we described in Section 4.5 and show in Table 4.1.

4.6.1 Satisfaction Motive

The satisfaction motive indicates the satisfaction level with the collaboration for the agent and its human collaborator. The satisfaction motive process maintains the value of *satisfaction drive* throughout the collaboration. The satisfaction drive is the quantitative weighted accumulation of desirability values between -1 and +1 over time. For instance, if the desirability values of the agent's appraisal over three consecutive turns are $\{0.75, 0, -0.25\}$, and their corresponding weights are $\{0.25, 0.5, 1.0\}$, the satisfaction drive value will be $(0.25)(0.75) + (0.5)(0) + (1.0)(-0.25)$ which is -0.0625. Notice that the latest desirability values get higher weights. Intuitively, it is because older desirable events have less influence on overall desirability and consequently the satisfaction level of the collaboration. The same process computes the satisfaction drive values for the agent and the human collaborator. Only the sources of desirability values are different, i.e., appraisal for the agent and reverse appraisal for the human collaborator. Then, the satisfaction motive process com-

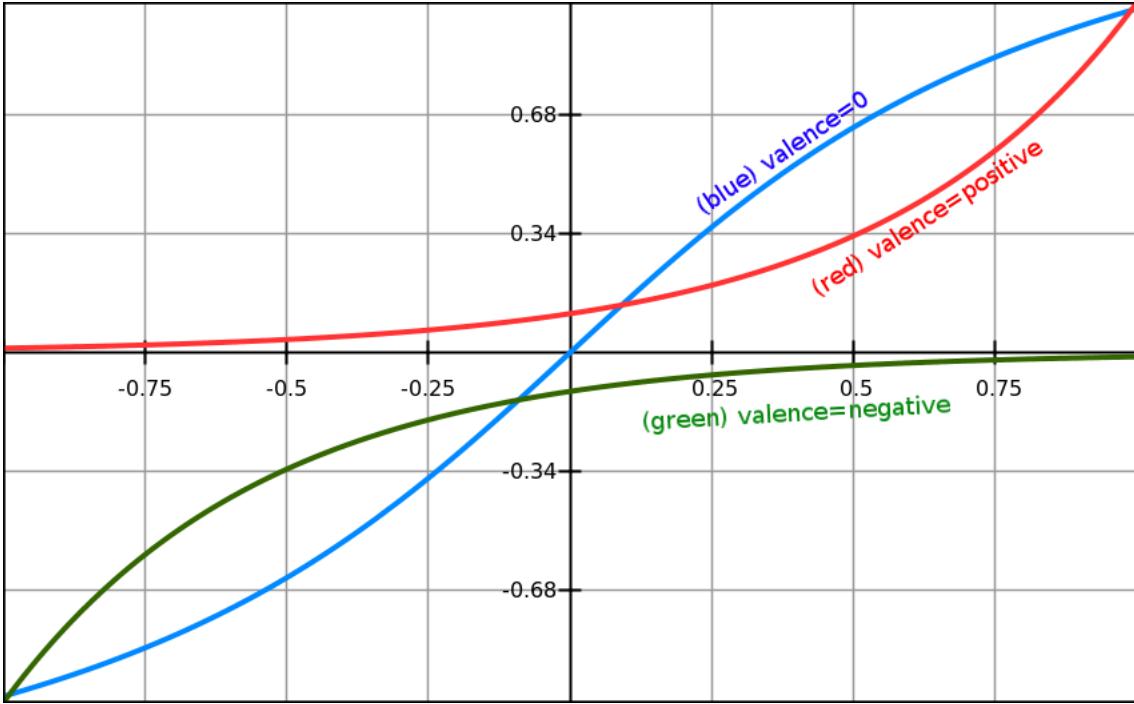


Figure 4.5: Three functions of satisfaction motive for different values of valence. The x-axis indicates the satisfaction drive's delta value in $[-1, +1]$, and the y-axis indicates the magnitude of satisfaction motive in $[-1, +1]$.

putes the difference between the current and the previous ($t-1$) satisfaction drives, called the delta of satisfaction drive value, δ_{sat} . As shown in equation 4.9, we use the δ_{sat} value in all three functions to compute the overall satisfaction motive's value \mathcal{M}_{sat} . We also use three different functions with respect to the valence value of the the human collaborator's perceived affective state. Our satisfaction motive's model has three domain dependent parameters $\mathcal{S}_{sat} \in [0, 1.5]$, i.e. strength of motive, $\mathcal{B}^{\mathcal{L}}$ where \mathcal{B} is the base parameter of the function in $(1, \infty)$ and \mathcal{L} is the exponential parameter of the same function in $(0, \infty)$; together \mathcal{B} and \mathcal{L} define *unsatisfiability* value. Currently we set \mathcal{S}_{sat} value to 1.5, \mathcal{B} to 3.0, and \mathcal{L} to 2.0.

$$\mathcal{M}_{sat}(\varepsilon_t) = \begin{cases} \arctan(\mathcal{S}_{sat} \times \delta_{sat}) & \text{valence} = 0 \\ \mathcal{B}^{\mathcal{L} \times (\delta_{sat}-1)} & \text{valence} > 0 \\ -\mathcal{B}^{-\mathcal{L} \times (\delta_{sat}+1)} & \text{valence} < 0 \end{cases} \quad (4.9)$$

The curves, shown in Figure 4.5, suitably represent the change in magnitude of satisfaction motive based on different valence values of the human collaborators affective state. Intuitively, if the human collaborator does not express any affective state, the satisfaction motive’s value can vary between -1 and +1 (blue curve in Figure 4.5). However, if the agent perceives positive affective state, there will be no negative satisfaction value since the other collaborator is in positive state of mind (red curve in Figure 4.5), and in contrast, if the agent perceives negative affective state, the satisfaction motive value only changes between -1 and 0 (green curve in Figure 4.5) with respect to how satisfied the agent is according to the status of its own goals during collaboration.

4.6.2 Achievement Motive

The achievement motive drives the agent’s need to achieve a goal during the collaboration. According to the literature, e.g. [169], the achievement motive is based on the estimation of success probability and the difficulty of achieving a goal. In our framework, we compute the probability of success as the product of the *controllability* and *expectedness* appraisal values. Intuitively, the more controllable and expected the events are, the probability of successful achievement of their related goal is higher.

In our framework we use two sigmoid-based functions to compute the achievement motive’s value. These functions values change based on the probability of success and valence of the human collaborator’s affective state. We use Equation 4.10 when the perceived affective state of the human has positive or zero valence value, and we use Equation 4.11 when the perceived affective state of the human has a negative valence value. As shown in Figure 4.6, when the value of the valence changes between 0 and +1, the output of \mathcal{M}_{ach}^+ function changes between the red and the blue lines respectively. Conversely, when the value of the valence changes between -1 and a small negative number (close to zero), the output of \mathcal{M}_{ach}^- function changes between the green and the orange lines.

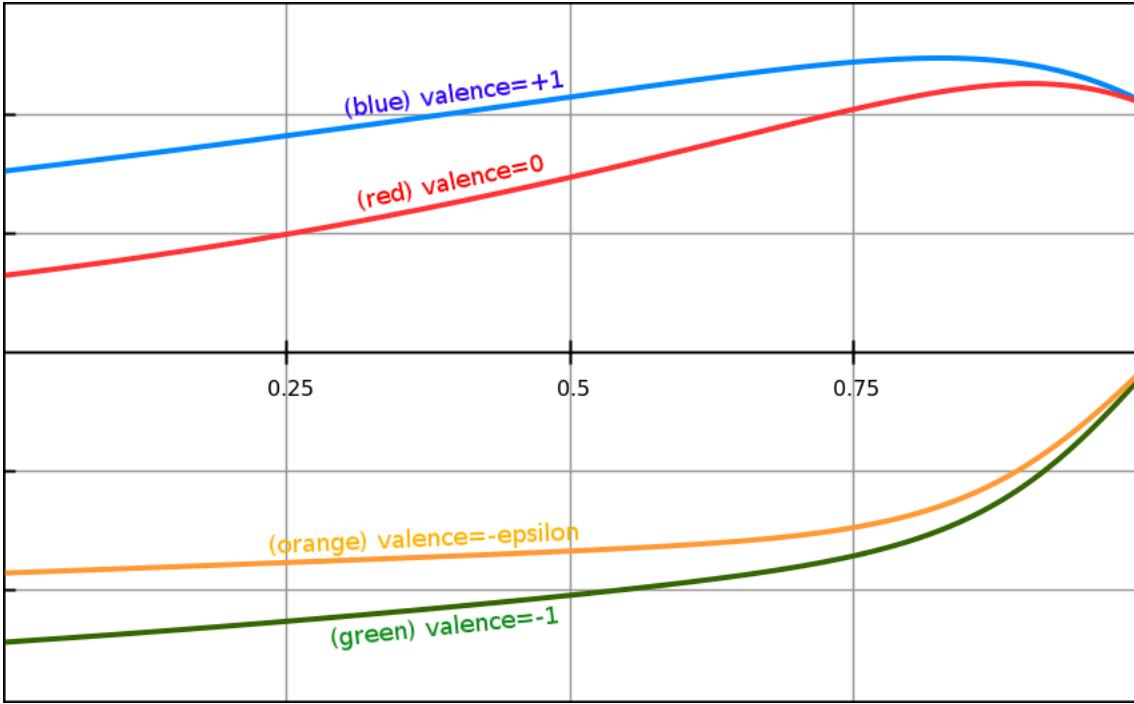


Figure 4.6: Two functions of the achievement motive for different values of valence. The x-axis indicates the success probability value of achieving a goal which is in $[0, +1]$, and the y-axis indicates the magnitude of achievement motive in $[-1, +1]$.

$$\mathcal{M}_{ach}^+(\varepsilon_t) = \frac{2.0}{1 + e^{(2.0 - valence) \times (1.05 - p(success))}} - \frac{1.0}{1 + e^{(12.0 - valence) \times (1.2 - p(success))}} \quad (4.10)$$

$$\mathcal{M}_{ach}^-(\varepsilon_t) = \frac{1.0}{1 + e^{(0.5 + valence) \times (1.05 - p(success))}} - \frac{1.0}{1 + e^{(12.0 + valence) \times (p(success) - 1.02)}} \quad (4.11)$$

By intuition, as the probability of success increases the agent is more motivated to achieve a goal and this motive gets higher when the human's affective state is positive or at least neutral. The human's negative emotions cause lower values of achievement motive since taking care of and acknowledging the human's negative affective state should have higher priority for a collaborative agent than achieving

a goal.

4.6.3 External Motive

The external motive drives the agent’s need to achieve a proposed goal by the human collaborator during the collaboration. In our framework, the external motive is also based on the estimation of success probability and the difficulty of achieving a goal, but this goal is proposed by the human collaborator. The probability of success for the external motive is computed the same way as the achievement motive’s probability of success, i.e. the product of *controllability* and *expectedness* appraisal values.

The only difference from the achievement motive is that we use Equations 4.10 and 4.11 in reverse order for the external motive; i.e., we use Equation 4.11 when the valence of human’s perceived affective state is positive, and Equation 4.10 when the valence of the human’s perceived affective state is negative or zero.

Intuitively, when the human proposes a new goal while expressing a negative affective state the agent should be more motivated to acknowledge human’s proposal and pursue the proposed goal to mitigate human’s negative affective state and maintain the collaboration. For example, when the human collaborator is worried about the failure of attaching a solar panel due to a malfunctioning tool, and proposes that the robot attach the panel, the high value of the external motive causes the robot to accept the human’s proposal.

4.7 Theory of Mind

The Theory of Mind mechanism uses the collaboration structure and functions described in Section 4.2 as well as appraisal processes to form anticipated beliefs about the human’s mental and emotional states. In other words, since our agent knows about the human’s goals (as part of the shared plan), it can apply the human goals to the same algorithms during the human’s turn of the collaboration. The agent

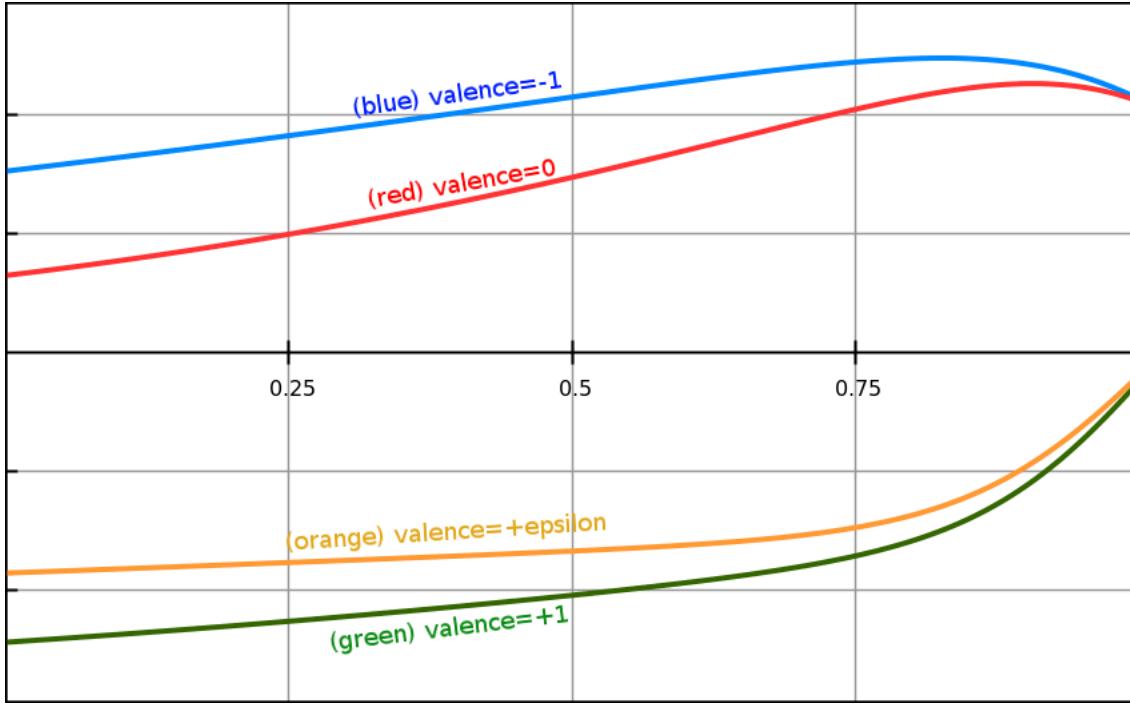


Figure 4.7: Two functions of external motive for different values of valence. The x-axis indicates the success probability value of achieving a proposed goal which is in $[0, +1]$, and the y-axis indicates the magnitude of the achievement motive in $[-1, +1]$.

uses the collaboration structure during the human's turn to compute appraisal values with respect to the human's current emotional state and the current goal in the shared goal structure. The outcome of the reverse appraisal forms beliefs about the anticipated mental and emotional state of the human collaborator.

Reverse Appraisal

We use the same *relevance*, *expectedness* and *controllability* algorithms for reverse appraisal as those algorithms we described in Section 4.3. In these three algorithms the Theory of Mind mechanism substitutes the agent's required goal and its corresponding constraints and information with the human's goal and its corresponding information which is provided to the agent within the shared plan structure. However, only for the reverse appraisal of *desirability* we chose to simply use the valence

value of the human’s perceived affective state and interpret negative, neutral and positive valence values as undesirable, neutral and desirable values respectively. In this way, our agent could directly infer whether the occurrence of the current event and its corresponding goal is desirable for the human. The outcome of all of these processes is a vector of reverse appraisal values that could be used by other mechanisms in our framework.

4.8 Elicitation of Emotion Instances

We have modeled 10 different emotion instances that can be elicited by the agent or anticipated from the human during collaboration in our framework (see Table 4.3). We chose these 10 emotions because we believe they are good examples of social emotions that can occur during a collaboration. These emotion instances have meanings in social context and more specifically in collaboration. There are two components involved in selecting a particular emotion: appraisal variables and collaboration context.

We use the outcome of the four appraisal processes discussed in section 4.3 to determine the potential emotion instance to be elicited (if the agent wants to express an emotion), or to anticipate a potential emotion from the human collaborator (if the human response is anticipated). The outcome of appraisal processes can be one of the values presented in Table 4.2 with respect to the corresponding process. For instance, relevance can only obtain either of two values, i.e., RELEVANT or IRRELEVANT., or the controllability can obtain one of the three values in Table 4.2; i.e., HIGH_CONTROLLABLE, LOW_CONTROLLABLE or UNCONTROLLABLE.

We also use the collaboration context as our second determinant to select a particular emotion. We define the collaboration context based on: *goal achievement* (HUMAN_ACHIEVED and AGENT_ACHIEVED), *goal failure* (HUMAN_FAILED and AGENT_FAILED), *proposal of a goal* (HUMAN_PROPOSED and AGENT_PROPOSED), *acceptance of the proposed goal* (HUMAN_ACCEPTED and AGENT_ACCEPTED), and

Table 4.2: Appraisal values for relevance, desirability, expectedness and controllability.

Appraisal Variable	Relevance	Desirability	Expectedness	Controllability
Values	RELEVANT	HIGH_DESIRABLE	MOST_EXPECTED	HIGH_CONTROLLABLE
		DESIRABLE	EXPECTED	LOW_CONTROLLABLE
	IRRELEVANT	NEUTRAL	UNEXPECTED	
		UNDESIRABLE	MOST_UNEXPECTED	UNCONTROLLABLE
		HIGH_UNDESIRABLE		

rejection of the proposed goal (HUMAN_REJECTED and AGENT_REJECTED). All of these situations can occur by either of the collaborators, i.e., agent or human (see Table 4.3). There is only one exception and it is when the desirability value is neutral the associated emotion to the event is always neutral without considering the collaboration context and the values of other appraisal variables (see first row in Table 4.3)¹. In summary, the outcome of four appraisal processes and the inferred context of a collaboration can lead the agent to elicit its own emotion or anticipate the human collaborator’s emotion.

The following interaction is based on our example scenario in Section 3.2:

5. **Astronaut**: The connectors on this panel have problems and we might not be able to finish this task.

6. **Robot**: Don’t worry! I can replace the connectors in 4 minutes. We definitely can finish this task after that.

The agent finds the Astronaut’s goal UNCONTROLLABLE, UNEXPECTED, UNDESIRABLE and RELEVANT (see all possible values of appraisal variables in Table 4.2). Also, the agent finds the current context of collaboration as HUMAN_PROPOSED; therefore, the agent infers that Astronaut’s perceived negative emotion instance can be *worry*. Thus, since the agent has access to working connectors (required inputs to the Astronaut’s task), first, the agent acknowledges the Astronaut’s negative emo-

¹Empty cell in Table 4.3 indicate that the value of the cell does not influence the selection of the emotion in the corresponding row.

tion, then informs the Astronaut with a proper solution to mitigate the Astronaut's negative emotion.

Table 4.3: Conditions for selecting emotion instances.

#	Emotion Instance	Context	Relevance	Desirability	Expectedness	Controllability
1	Neutral	human <i>HUMAN_ACHIEVED</i>	RELEVANT	DESIRABLE <i>HIGH_DESIRABLE</i>	EXPECTED <i>MOST_EXPECTED</i>	
2	Joy	agent <i>AGENT_ACHIEVED</i>	RELEVANT	UNDESIRABLE <i>HIGH_UNDESIRABLE</i>	EXPECTED <i>MOST_EXPECTED</i>	UNCONTROLLABLE
3	Sadness	human <i>AGENT_FAILED</i>	RELEVANT	DESIRABLE <i>HIGH_DESIRABLE</i>	EXPECTED <i>MOST_EXPECTED</i>	UNCONTROLLABLE
4	Gratitude	human <i>AGENT_ACCEPTED</i> <i>AGENT_ACHIEVED</i>	RELEVANT	DESIRABLE <i>HIGH_DESIRABLE</i>	EXPECTED <i>MOST_EXPECTED</i>	
5	Positive Surprise	human <i>AGENT_PROPOSED</i> <i>AGENT_ACCEPTED</i> <i>AGENT_ACHIEVED</i>	RELEVANT	DESIRABLE <i>HIGH_DESIRABLE</i>	MOST_UNEXPECTED	
6	Negative Surprise	agent <i>HUMAN_PROPOSED</i> <i>HUMAN_ACCEPTED</i> <i>HUMAN_ACHIEVED</i>	RELEVANT	UNDESIRABLE <i>HIGH_UNDESIRABLE</i>	MOST_UNEXPECTED	
7	Anger	human <i>AGENT_REJECTED</i> <i>AGENT_FAILED</i>	RELEVANT	HIGH_UNDESIRABLE	EXPECTED <i>MOST_EXPECTED</i>	UNCONTROLLABLE
8	Worry	agent <i>HUMAN_REJECTED</i> <i>HUMAN_FAILED</i>	RELEVANT	UNDESIRABLE <i>HIGH_UNDESIRABLE</i>	UNEXPECTED	UNCONTROLLABLE
9	Frustration	human <i>AGENT_PROPOSED</i> <i>AGENT_FAILED</i>	RELEVANT	UNDESIRABLE	EXPECTED <i>MOST_EXPECTED</i>	UNCONTROLLABLE
10	Guilt	agent <i>HUMAN_FAILED</i>	RELEVANT	UNDESIRABLE <i>HIGH_UNDESIRABLE</i>	EXPECTED <i>MOST_EXPECTED</i>	LOW_CONTROLLABLE HIGH_CONTROLLABLE

CHAPTER 5

EVALUATION

In this chapter, we provide the explanation and results of two different user studies. The first user study (see Section 5.2) was conducted online to evaluate our appraisal algorithms. Specifically, the goal of this study was to validate the effectiveness of the factors involved in our appraisal algorithms. We prepared online questionnaires and asked participants to tell us what their decision would be in the simple situations provided. The participants' answers to our questionnaires were compared with the results of our algorithms for the given situations. The results are provided in Section 5.2. The second user study (see Section 5.3) was conducted in the laboratory. The goal of this user study was to provide an end-to-end system evaluation using our overall framework. We provided pre- and post-study questionnaires as well as an open-ended questionnaire to study the humans' evaluation of a robot collaborating using our framework. The results are provided in Section 5.3.

5.1 Implementation

As described in Chapter 3, the Perception and Action mechanisms are not part of our theoretical work. Therefore, we only implemented these mechanisms to the extent to which they could help us to run and test our framework. The Perception mechanism only redirects the input values from the system's users to the framework. For instance, in our user study described in Section 5.3, the Perception mechanism only receives the valence of human's affective state from the input and provides it to

the framework. On the other hand, the Action mechanism executes some functions based on the intentions formed and provided by the Coping mechanism described in this section. We group all of these functions into three categories in our framework. The first group of functions includes all of the functions capable of executing some actions with respect to the domain. The second category includes all of the functions involved in revealing the agent’s utterances by writing on the screen or conveying through the agent’s voice and text to speech systems. The last category includes all of the functions to express the agent’s affective state. The emotions can be expressed through colors, emoticons, voice and text. For example, in the user study described in Chapter 5, we expressed the agent’s emotions by using emoticons and utterances through the text on the screen as well as the agent’s voice.

5.2 Evaluating Appraisal Algorithms (Crowd Sourcing)

In this section, we present a crowd-sourced user study and the results, which we conducted to validate the components of our appraisal processes.

5.2.1 Experimental Scenario

We developed an experimental scenario in which participants were asked to envision a sequence of hypothetical collaborative tasks between themselves and an imaginary friend, Mary, in order to accomplish their shared goal. To minimize the background knowledge necessary for our test subjects, we used a simple domestic example of preparing a peanut butter and jelly sandwich, and a hard boiled egg sandwich for a hiking trip. The tasks did not require the participants to do any deep problem solving; rather, the tasks were part of simple daily activities that should be familiar to all participants.

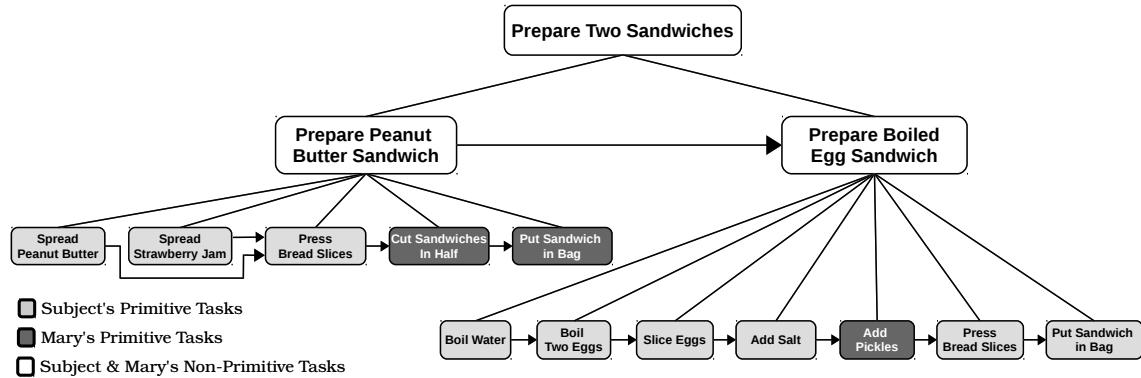


Figure 5.1: Collaboration task model for the evaluation.

5.2.2 Hypothesis and Methodology

Hypothesis

We conducted this user study to test our hypothesis that humans and our algorithms will provide similar answers to questions related to different factors used to compute four appraisal variables: relevance, desirability, expectedness, and controllability.

Procedure

We conducted a between-subject user study using an online crowdsourcing website – CrowdFlower¹. We had a questionnaire for each appraisal variable. There were 12 questions (including 2 test questions) in the controllability and expectedness questionnaires, 14 questions (including 2 test questions) in the desirability questionnaire, and 22 questions (including 3 test questions) in the relevance questionnaire.

We provided textual and graphical instructions for all questionnaires; Figure 5.1 shows the corresponding task model². The instructions, provided in the Appendix A, presented a sequence of hypothetical collaborative tasks to be carried out by the test subject and an imaginary friend, Mary, in order to accomplish their goal of preparing two sandwiches. We also provided a simple definition and an example of each appraisal variable. The collaboration structure and the instructions were

¹<http://www.crowdflower.com>

²Figure 5.1 was not given to the participants.

Table 5.1: Number of participants

appraisal variables	# of participants
Relevance	29
Desirability	35
Expectedness	33
Controllability	33

the same for all questionnaires. The questions introduced specific situations related to the shared plan, which included blocked tasks and failure or achievement of a shared goal. Each question provided three answers which were counterbalanced in the questionnaire. We provided an option like C in all questions (see Figure 5.2), because we did not want to force participants to choose between two options when they did not have a good reason. We derived two questions for different factors involved in each algorithm (see Section 4.3). For instance, we prepared two questions about the influence of the strength of a belief as a key factor involved in relevance algorithm. The questions were randomly placed in the questionnaire. Figure 5.2 shows an example question from the relevance questionnaire which was designed to test whether participants perceive saliency as a factor in relevance. The input for our algorithms was the task model depicted in Figure 5.1.

Participants

Each participant group originally had 40 participants. We limited the participant pools to those with the highest confidence level on the crowdsourcing website in the United States, Britain, and Australia. Test questions were included to check the sanity of the answers. We eliminated participants providing wrong answers to our sanity questions, and participants with answering times less than 2 minutes. The final number of accepted participants in each group is provided in Table 5.1.

5.2.3 Results

Each question in our questionnaires was designed based on different factors that we use in our algorithms (see Section 4.3). For each of the four questionnaires we provide an example question, and describe how each question relates to a specific factor within the corresponding algorithm. The input for our algorithms was the task model depicted in Figure 5.1. The complete list of questions is provided in the Appendix A. Additionally, we provide the p-value for each question, using a binomial distribution, with a probability of success of 0.33, which is the probability of selecting the right answer if the participant is simply guessing.

Expectedness

Figure 5.2 shows an example question from the expectedness questionnaire. In this example, with respect to Algorithm 3 (line 6), option A is more expected because the task related to this option provides the next available task in the focus stack (see the task model in Figure 5.1). Although the task in option B is part of the existing task model, it is considered as UNEXPECTED by our algorithm, since it is not live in the plan. We provided option C to determine whether the participants will differentiate between these two options. This question was presented to the participants to determine whether their decision for the expectedness of this event is similar to the output of the expectedness algorithm. For this question, the human decision was 97% similar to the algorithm’s output.

Results for the expectedness questionnaire are presented in Table 5.2. As shown in this table, the results are statistically significant; in fact, for questions 1-6 and 9-10, human participants showed between 67 and 100 % agreement with our algorithms, with p-values of $\ll 0.001$. Questions 7 and 8 were the only two questions that did not show a statistically significant p-value. It should be noted that these questions are comparing equally expected or equally unexpected situations, none of which our algorithms would consider most-expected or most-unexpected.

Table 5.2: Expectedness results (the Equally Expected column indicates for which questions our algorithm provides option C as the response)

Question	Factor	Equally Expected	Percentage of Matching Answers	p-Value
1	Live goal vs. Necessary focus shift	No	94%	« 0.001
2	Live goal vs. Not part of shared plan	No	97%	« 0.001
3	Live goal vs. Not part of current branch	No	82%	« 0.001
4	Necessary focus shift vs. Not part of shared plan	No	100%	« 0.001
5	Necessary focus shift vs. Not part of current branch	No	97%	« 0.001
6	Not part of shared plan vs. Not part of current branch	No	73%	« 0.001
7	Live goal	Yes	42%	0.093
8	Not part of current branch	Yes	42%	0.093
9	Necessary focus shift	Yes	67%	« 0.001
10	Not part of shared plan	Yes	88%	« 0.001

Controllability

Figure 5.3 shows an example question from the controllability questionnaire. The algorithm's output is option B, and is determined by Algorithm 4 (line 3), similarly to the expectedness example above. In this example, option B is more controllable than option A, because the self over total ratio of the responsibility of the predecessors of the given task (see *Autonomy* in Section 4.3.4) is higher than the ratio in

Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Which of the following two actions is **more expected**?

A. Mary puts the given sandwich into a zip lock bag after cutting it in half.
B. Mary puts some pickles on another slice of bread.
C. Equally expected.

Figure 5.2: Example expectedness question.

Table 5.3: Controllability results (the Equally Controllable column indicates for which questions our algorithm provides option C as the response)

Question	Factor	Equally Controllable	Percentage of Matching Answers	p-Value
1	Agency	No	85%	« 0.001
2	Autonomy (contributors)	No	52%	0.009
3	Autonomy (predecessors)	No	91%	« 0.001
4	Succeeded predecessors ratio	No	58%	0.001
5	Available inputs	No	91%	« 0.001
6	Agency	Yes	91%	« 0.001
7	Autonomy (contributors)	Yes	73%	« 0.001
8	Autonomy (predecessors)	Yes	55%	0.003
9	Succeeded predecessors ratio	Yes	70%	« 0.001
10	Available inputs	Yes	76%	« 0.001

option A, i.e., self is responsible to spread peanut butter on one slice of bread and strawberry jam on another slice of bread. In this question, the humans decision was 90% in agreement with the algorithm's output.

Results for the controllability questionnaire are presented in Table 5.3. As shown in the table, the p-value is <0.01 for each of the ten questions. The two questions with the lowest human agreement with the algorithms both relate to autonomy (Questions #2 and #8) of the participants with 52% and 55%.

Imagine you want to make a peanut butter sandwich. Which of the following two actions is **more controllable**?

A. You can spread the peanut butter on one slice of bread and you need Mary to spread strawberry jam on the second slice of bread.

B. You can spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread.

C. Equally controllable.

Figure 5.3: Example controllability question.

Table 5.4: Desirability results (the Equally Desirable column indicates for which questions our algorithm provides option C as the response)

Question	Factor	Equally Desirable	Percentage of Matching Answers	p-Value
1	Top level goal is failed	No	100%	« 0.001
2	Top level goal is achieved	No	83%	« 0.001
3	Predecessors or preconditions of the top level goal	No	100%	« 0.001
4	Focus is achieved	No	98%	« 0.001
5	Focus is failed	No	100%	« 0.001
6	Predecessors or preconditions of the focus	No	100%	« 0.001
7	Pending or in-progress focus	Yes	46%	0.040
8	Top level goal is failed	Yes	66%	« 0.001
9	Predecessors or preconditions of the top level goal	Yes	54%	0.003
10	Focus is achieved	Yes	57%	0.001
11	Focus is failed	Yes	60%	« 0.001
12	Predecessors or preconditions of the focus	Yes	77%	« 0.001

Desirability

Figure 5.4 shows an example question from the desirability questionnaire. The output based on the Algorithm 2 (line 14) is option C, since in both option A and option B, the focus goal has been achieved successfully. Therefore, in this example, both options A and B are desirable. The humans' decision was 77% in agreement with the algorithm's output in this question.

The results of the desirability questionnaire are presented in Table 5.4. As shown in the results table, the p-value is less than 0.05 for all of the desirability questions. However, an interesting trend is that human participants had a level of agreement of 83%-100% when the algorithm's output selected one alternate as more desirable than another alternate. When the algorithm's output chose option C (i.e. rating

Which of the following two actions is **more desirable**?

- A. Imagine you pressed two slices of bread together with peanut butter and strawberry jam on them, and passed them to Mary. Mary cuts the peanut butter sandwich in half and puts them in the zip lock bag.
- B. Imagine you want to make the egg sandwich. You have sliced the eggs, put them on one slice of bread, salted them, and waiting for Mary to put some pickles on your eggs. Mary puts some pickles on your eggs.
- C. Equally desirable.

Figure 5.4: Example desirability question.

two situations as equally desirable), the human participants only showed 46%-77% agreement. This may indicate that a higher level of granularity is required in the algorithm when evaluating options with similar levels of desirability.

Relevance

In the example shown in Figure 5.5, with respect to Algorithm 1, option A is relevant because of Mary's perceived negative affective state (see Equation 4.1). Although option B is relevant (since it achieves the next goal in the shared plan), 83% of participants consider it as less relevant than option A; we believe this is due to the effect of Mary's perceived negative affective state which also generates a higher utility value in our relevance algorithm. Another question also tested belief saliency. However, the options provided only related to the shared plan (i.e., no human emotions in the options). In this case 87% of participants chose the option that accomplished the next goal in the shared plan. Interestingly, when confronted with a negative affective state from their collaborator, human participants deviated from the shared plan and found their collaborator's affective state more relevant than the original plan. It is noteworthy that in both the absence and the presence of emotions the participants chose the more salient option with respect to our definition of saliency, which was not referenced or provided in the questionnaire.

The complete summary of results for the relevance questionnaire is provided in

Table 5.5: Relevance results (the Equally Relevant column indicates for which questions our algorithm provides option C as the response)

Question	Factor	Equally Relevant	Percentage of Matching Answers	p-Value
1	Belief Saliency	No	86%	« 0.001
2	Belief Strength	No	45%	0.063
3	Belief Recency	No	97%	« 0.001
4	Motive Insistence	No	86%	« 0.001
5	Motive Urgency	No	66%	« 0.001
6	Motive Intensity	No	72%	« 0.001
7	Goal Proximity	No	69%	« 0.001
8	Goal Specificity	No	79%	« 0.001
9	Belief Saliency	Yes	90%	« 0.001
10	Belief Strength	Yes	76%	« 0.001
11	Belief Recency	Yes	72%	« 0.001
12	Motive Insistence	No	90%	« 0.001
13	Motive Urgency	Yes	100%	« 0.001
14	Motive Intensity	Yes	100%	« 0.001
15	Goal Proximity	Yes	83%	« 0.001
16	Goal Specificity	Yes	90%	« 0.001
17	Belief Saliency	No	59%	« 0.001
18	Motive Insistence	No	10%	0.995
19	Goal Proximity	No	14%	0.982

Table 5.5. As shown in the table, all questions show 59%-100% agreement with our algorithms and statistically significant p-values except for questions 2, 18 and 19. Question 2 addresses belief strength. This question presents a situation in which participants must choose whether a self related goal, or collaborator's goal is more relevant. Questions 18 and 19 address motive insistence and goal proximity, respectively; both of these questions present situations in which participants must choose whether an intense emotional circumstance, or adherence to the collaboration plan is more relevant (the questionnaire is provided in the Appendix A). Our algorithms

Imagine you have made the peanut butter sandwich and passed it to Mary to cut it in half. Which of the following two actions is **more relevant**?

- A. Mary starts crying since she cut her finger with a knife.
- B. You begin to boil the water to boil the eggs for your second sandwich.
- C. Equally relevant.

Figure 5.5: Example relevance question.

choose that the strong emotional circumstance will be more relevant; however, human participants generally selected adherence to the collaboration plan to be more relevant.

5.2.4 Discussion

As shown in the preceding results tables, the human participants agreed 100% on some questions, while on some other questions there was a much lower level of agreement. Our results indicate that people largely performed as our hypothesis predicted. The very small p -values indicate that the data set is not random; in fact, the high percentage of similarity confirms our hypothesis and shows that the algorithms can help us to model appraisal in a collaboration. The very low level of agreement on a handful of questions may indicate algorithm components that require further refinement before implementation; therefore, we made limited changes to our algorithms in light of this study.

5.3 End-to-End System Evaluation

As mentioned earlier, collaborative robots need to take into account humans' internal states while making decisions during collaboration. Humans express emotions to reveal their internal states in social contexts including collaboration [35]. Due to the existence of such expressions, robots' emotional-awareness can improve the quality of collaboration in terms of humans' perception of performance and pref-

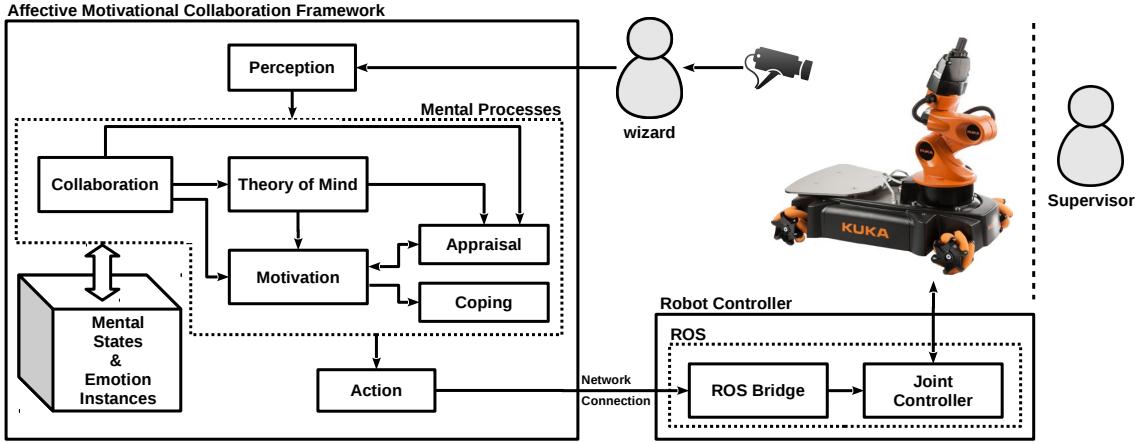


Figure 5.6: Experimental setup for end-to-end system evaluation.

erences. Hence, collaborative robots need to include affect-driven mechanisms in their decision-making processes to be able to interpret and generate appropriate responses and behaviors. Our aim in this experiment was to study the importance of emotional awareness and the underlying affect-driven processes in human-robot collaboration. We examined how emotional-awareness impacts different aspects of humans' preferences by comparing the results from our participants collaborating with an affect-aware versus an affect-ignorant robot.

5.3.1 Experimental Setup

The setup of this user study included three separate parts. The first part incorporated the Affective Motivational Collaboration framework consisting of all Mental Processes (see left-side of Figure 5.6) as we described in Chapters 3 and 4. The second part was implemented to receive action commands from the framework and forward them to the robot to control joints and actuators (see right-side of Figure 5.6). A wizard was the third part of this setting. The wizard did nothing but inform the robot/framework whether the current task performed by either the robot or the participant was achieved successfully. The wizard was completely invisible to the participants, and the wizard had no impact on the robot's decision other than providing input regarding tasks' failure or success.

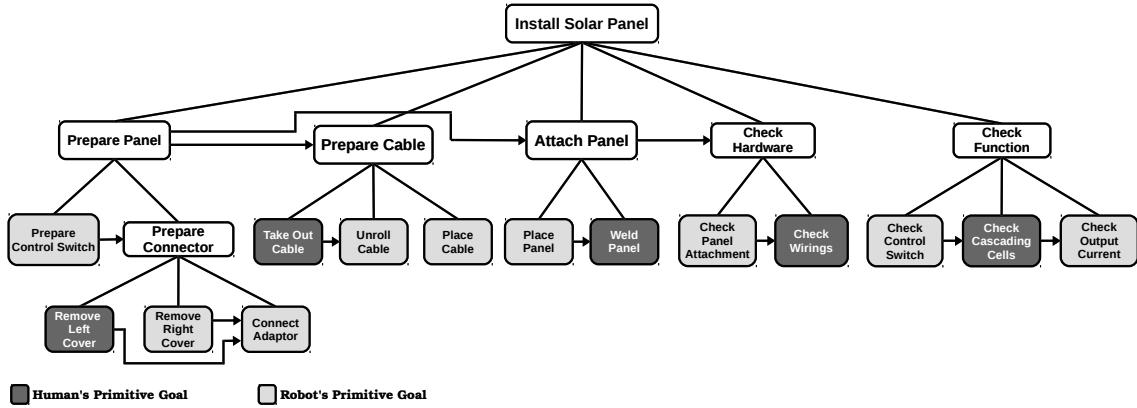


Figure 5.7: Collaboration structure used as the task model.

Environment and Tasks

The environment was set up in a laboratory and included the robot, the collaboration board on top of a desk, and the participant standing in front of the robot on the other side of the board (see Figure 5.10). One of the experimenters monitored the interactions using a live stream of a camera in a different room. The experimenter provided only the required perception, i.e., decision on success or failure of the tasks for the robot, through the entire time of the collaboration (see Section 5.3.2). We also had a supervisor, whose purpose was to provide help to the collaborators when requested after a task failure.

The tasks were defined based on the collaboration structure shown in Figure 5.7 and were executed in a turn-taking fashion by either of the collaborators¹. For each task either the robot or the participant was responsible for picking up one of the corresponding pegs from their own inventory and placing it on the right spot which was colored and tagged the same as the associated peg. Some pegs and corresponding spots on the board had hidden magnets which prevented the pegs from standing upright. Any peg that fell over was considered a failed task.

¹Figure 5.7 was not given to the participants.

Framework

The framework includes all of the mechanisms depicted as mental processes in Figure 5.6 along with the Mental State. The mental states shown in Figure 5.6 comprise the knowledge base required for all of the mechanisms in the overall model. The details about these mental processes and Mental State are described in Chapters 3 and 4. We use Disco as the basis of the Collaboration mechanism. Disco is the open-source successor to COLLAGEN [203, 204] which incorporates algorithms based on SharedPlans theory for discourse generation and interpretation. Disco is able to maintain a segmented interaction history, which facilitates the collaborative discourse between a human and a robot. In this user-study, the Collaboration mechanism uses a hierarchy of goals associated with tasks in the hierarchical task network structure depicted in Figure 5.7.

The Robot

We conducted our experiment with a KUKA Youbot (see Figure 5.10). The robot was stationary on top of a desk and was able to pick up and place available pegs corresponding to the robot's task. The robot was operated based on Robot Operating System (ROS – indigo) and was receiving commands through the ROS-bridge from our Affective Motivational Collaboration framework (see Figure 5.6). We provided a simple GUI using a touch-screen monitor (see Figure 5.9) to a) express the robot's positive, negative or neutral affective state through an emoticon, b) display the robot's utterances, c) control turn-taking process of the collaboration, and d) let the participants express (report) their positive, negative or neutral affective state for each turn. The GUI was identical in both affect-aware and affect-ignorant conditions (see Section 5.3.2). The robot used the MaryTTS¹ text-to-speech platform to provide corresponding speech for its utterances in English.

¹<http://mary.dfki.de/>

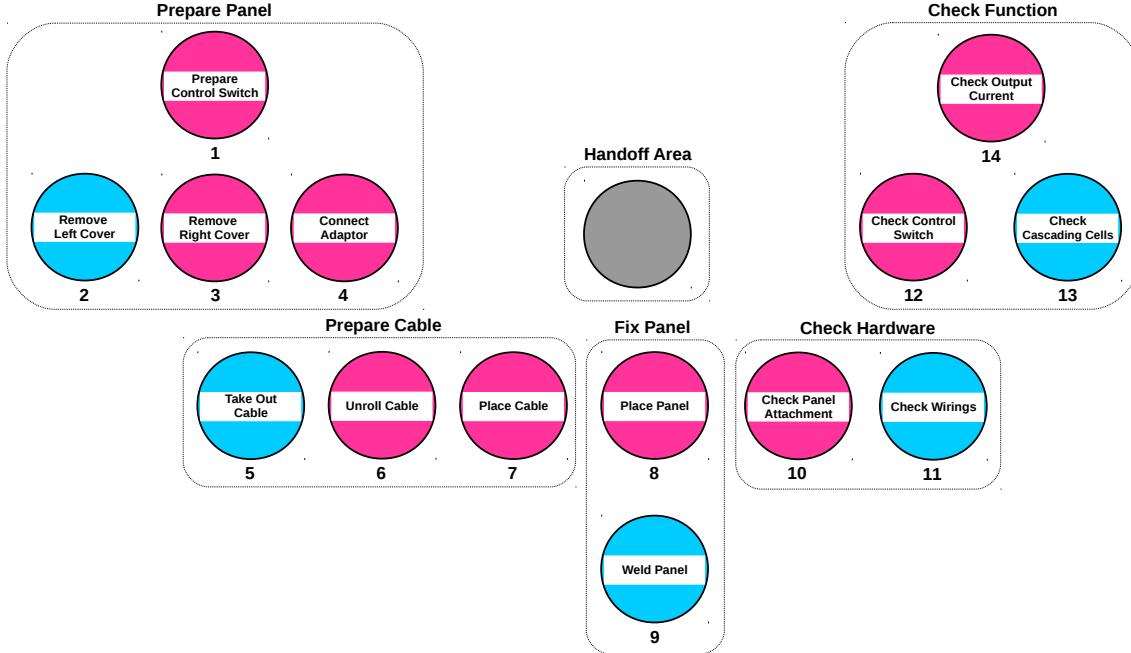


Figure 5.8: The layout of the available spots for the human and the robot to place their pegs during the collaboration.

Robot Controller

The robot controller is comprised of two major components: 1) ROS-bridge and 2) joint controller (see Figure 5.6). ROS-bridge¹ provides an API to ROS functionality for non-ROS programs which enables us to send action commands from our framework (implemented in JAVA) to the robot's joint controller. The joint controller receives action commands and translates them into actual joint and actuator commands and sends them to the robot (see Figure 5.10).

5.3.2 Experimental Design

Our scenario was based on a table top turn-taking game that we designed to simulate the installation of a solar panel. Participants collaborated one-on-one with our robot to complete all the given tasks required to install the solar panel. Each participant worked with the robot in two conditions, in a within subject study. Each primitive

¹http://wiki.ros.org/rosbridge_suite

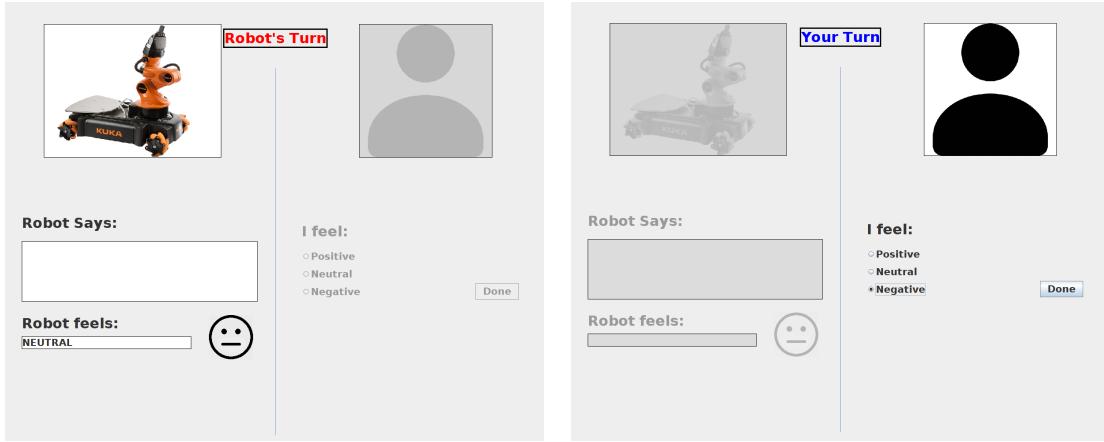


Figure 5.9: The Graphical User Interface (GUI) used during interaction.

task consisted of picking up and placing pegs on predefined spots on the board (see Figure 5.8). Each pick-and-place was associated with the robot’s or the participant’s task. The robot and the participants had their own unique primitive tasks that they had to accomplish in their own turns. The final goal of installing a solar panel required the robot and the participants to accomplish their own individual tasks. Failure of any task could create an impasse during the collaboration.

Interaction Paradigms

At the beginning of each collaboration the robot asked each participant to achieve the overall shared goal, i.e., “installing the solar panel”. Then, before working towards a new goal, the robot informed the participant about the higher level non-primitive goal (e.g. Prepare Panel – see Figure 5.7) of which the primitives were going to be working towards. The same procedure was used by the robot if there was a decision to switch to another nonprimitive due to the failure of a task in achieving the current goal. After achieving a new primitive goal, the robot either informed the human that it would pursue the next goal, or it informed and passed the turn to the human to execute the next task with respect to the human’s goal. In case of the human’s turn, the robot waited for the human to do a task, then the wizard let the robot know whether the human’s goal was achieved or not. Af-

terwards the robot made a decision about which goal to pursue and informed the human accordingly. There were two conditions of the robot: 1) affect-aware and 2) affect-ignorant. The same procedure was applied to both conditions. In the next section (i.e., Algorithmic Trace) we provide the algorithmic trace to show how goal management algorithm works in both affect-aware and affect-ignorant conditions.

The robot interacted via a) speech, b) the corresponding utterance on the screen, c) negative, positive and neutral expression of affective state through an emoticon on the screen. The robot used only neutral expression in the case of affect-ignorance. The interaction was controlled autonomously by the framework we discussed in Section 5.3.1 in both the affect-ignorant and the affect-aware cases. The reasoning about which task should be done and controlling the robot was entirely autonomous. Only the perception of the task failure or achievement by the robot or by the participant was done by a wizard monitoring the collaboration outside of the test area. The interaction was structured based on the same collaboration structure (see Figure 5.7) for both conditions. The robot used the same utterances in both conditions. In the affect-aware condition the robot used a different behavior in comparison with the affect-ignorant condition only if the participant was expressing a negative affective state in the event of a failure; i.e., the robot's utterances were identical in affect-ignorant and affect-aware cases if in the latter the participant reported (expressed) a positive or a neutral affective state.

Three different behaviors could be generated only in the affect-aware condition as a result of the Coping mechanism (planning, active coping, seeking social support for instrumental reasons, and mental disengagement coping strategies were involved to generate these three behaviors). These three behaviors were 1) mitigating the human's negative affective state and postponing its own task to help the human, 2) goal-management to switch to another goal which had lower cost with respect to the human's negative affective state, and 3) task delegation to the human to overcome the impasse. In each run, the human had two predetermined task failures, and the robot had one. If the human expressed negative affective state after the

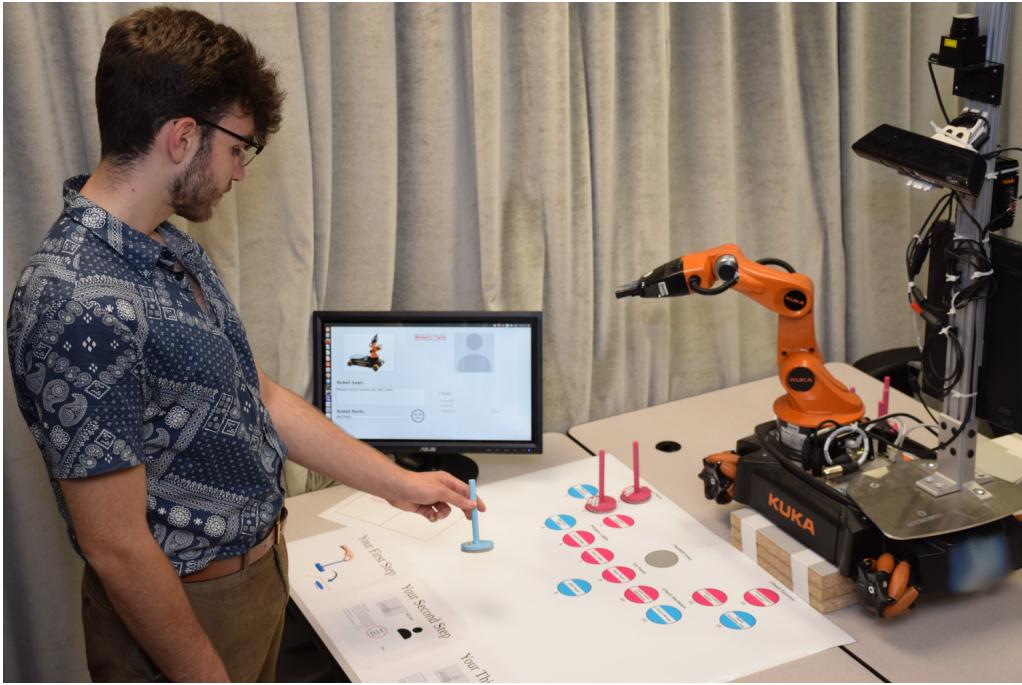


Figure 5.10: Experimental setup.

first human task failure, the robot responded by mitigating the human’s negative affective state by saying “It was not your fault. I can help you with this task” and helping the human by providing a peg to fulfill the human’s task. If the human expressed negative affective state after the second human task failure, the robot informed the human that they could proceed with another task to save time while simultaneously requesting a new peg (i.e., help) from the supervisor. If the human expressed negative affective state as a result of the robot’s task failure, the robot requested help from the human (who had the correct peg). In the event that the human expressed positive or neutral affective state during these three failures, the robot behaved identically in the affect-ignorant and the affect-aware cases, by asking the supervisor for help.

Algorithmic Trace

In this section we provide an algorithmic trace for one of the three behavioral changes in our user study when the robot, besides acknowledging human’s perceived negative

affective state, activates a coping strategy called *mental disengagement*. As shown in Table 4.1, the mental disengagement strategy can be activated when the robot and the human feel neutral or negative, all three values of motives (see Section 4.6) have obtained low or medium negative values, and the robot evaluates the controllability of the corresponding goal of the event as uncontrollable one. As the result of activation of the mental disengagement strategy, our framework applies the goal management algorithm to lower the effect of the stressor which is the failure of a goal in the plan.

The following algorithmic trace shows how different mechanisms (see Chapter 4) in our framework are involved to generate proper behavior in response to the perceived goal-failure of the human in the affect-aware condition:

1. As a result of perceiving failure to achieve a goal by the human collaborator (i.e., “Take Out Cable” – See Figure 5.7), AMC framework uses the **Appraisal mechanism** to appraise this event. To be able to appraise the event, the Appraisal mechanism needs the current information of the collaboration structure.
2. Thus, the **Collaboration mechanism** provides the information required to appraise the current event. This information includes the status of the preconditions, postconditions, hierarchical and temporal constraints, inputs, outputs, status of the predecessors and contributing goals in the shared plan. All of this information appears as an element of Mental State, i.e., belief.
3. Then, the **Theory of Mind mechanism** provides the anticipated beliefs about the human collaborator’s mental state based on reverse appraisal of the human’s perceived affective state, i.e., negative (valence).
4. As a result, the **Appraisal mechanism**’s outcome will be *relevant, undesirable, unexpected and controllable*.
5. The **Motivation mechanism** uses the collaboration structure (Collaboration mechanism), anticipated beliefs of the human collaborator (Theory of Mind mechanism) and the outcome of the appraisal (Appraisal mechanism) to

- compute the intensity of three motives (see Section 4.6), i.e., low or medium negative values for satisfaction, achievement, and external motives.
6. The **Coping mechanism** receives the values of the three motives and activates two coping strategies, i.e., Active Coping and Mental Disengagement, with respect to the conditions shown in Table 4.1. As a result of the activation of these strategies, first, active coping strategy forms a new intention to acknowledge human's negative affective state, and then, mental disengagement strategy uses the goal management algorithm to lower the effect of current stressor (goal failure) and overcome an impasse. The goal management algorithm uses the current collaboration structure and the human's perceived affective state to form a new intention to switch to another goal which has the lowest cost (see Section 4.4), i.e., "Place Panel".
 7. At the end, the **Action mechanism** receives two new intentions of acknowledging human's negative affective state and pursuing achievement of a new goal (i.e., Place Panel). As a result, the robot will say: "Don't worry! To manage the time let's switch to another task. We can come back and finish this later."

In affect-ignorant condition, AMC framework bypasses the **Appraisal, Theory of Mind, Motivation, and Coping mechanisms**. AMC framework forms a new intention based on the goal that the SharedPlans' collaborator manager, i.e., Disco, provides. In this condition the agent's decision is not influenced by the human's perceived negative affective state, even if the human reports a negative affective state. As a result, the robot asks the supervisor to come and provide help to overcome an impasse in response to the human's goal failure.

5.3.3 Hypotheses

The non/social functions of emotions impact a collaboration process. Human collaborators prefer to collaborate with others whose behaviors are influenced by these

functions of emotions depending on the context. We developed seven specific hypotheses regarding the positive influence of affect-awareness and the usefulness of emotion function during collaboration:

Hypothesis 1. Participants will feel closer to the affect-aware robot rather than the affect-ignorant robot.

Hypothesis 2. Participants will find the affect-aware robot to be more trustworthy than the affect-ignorant robot.

Hypothesis 3. Participants will find the affect-aware robot to have better performance in collaboration than the affect-ignorant robot.

Hypothesis 4. Participants will find the affect-aware robot to be more understanding of their feelings than the affect-ignorant robot.

Hypothesis 5. Participants will find the affect-aware robot to be more understanding of their goals than the affect-ignorant robot.

Hypothesis 6. Participants will feel more satisfied about the collaboration when working with the affect-aware robot rather than affect-ignorant robot.

Hypothesis 7. Participants will perceive higher level of mutual satisfaction with the affect-aware robot than affect-ignorant robot.

5.3.4 Procedure

Participants were first given a brief description of the purpose of the experiment. After the short introduction, they were asked to review and sign a consent form. Participants were then provided with a written instruction of their task and the rules for collaborating with the robot, provided in Appendix B. Then, one of the experimenters lead them into the experiment room and went through all the details of the instructions with the participants standing in front of the collaboration board and the robot. The experimenter confirmed participants' correct understanding of the tasks and informed them of the types of task failures (i.e., fallen peg) that might occur during the collaboration. Participants were told that researchers were developing a collaborative robot and would like their help in evaluating their de-

sign. Participants were provided with identical instructions and randomly assigned to complete either the affect-aware or the affect ignorant condition first. They were told that, after their collaboration with the robot, they would be asked to answer a questionnaire on their experience. After completing the first round of collaboration, participants answered a post-experiment questionnaire that measured their perceptions of the robot, the task, and the collaboration procedure; the questions are provided in Table 5.6. After answering the first post-experiment questionnaire, participants were told that they were going to collaborate with the robot one more time and the robot might not necessarily have the same collaborative behavior. After completing the second round of collaboration, participants were asked to answer the second post-experiment questionnaire which consisted of the same questions as the first post-experiment questionnaire. Finally, participants were asked to answer an open-ended questionnaire which measured their perception of difference between two runs, their preference of collaborative robot between two runs, and their reasons of preference; the questions are provided in 5.7.

Measurements

In our study two basic conditions of the robot were tested: a) the affect-ignorant condition, b) the affect-aware condition. We measured participants' recall of the collaborative behaviors presented by the robot using an open-ended post-experiment questionnaire. We also specifically asked the participants what behavior of the robot they liked during their collaboration. We also evaluated participants' levels of satisfaction, trust, goal achievement, mutual understanding of goals, mutual understanding of feelings, mutual agreement, and also participants' beliefs about the efficiency of collaboration and their feeling of robot's collaborative behaviors. Seven-point Likert scales were used in these questionnaire items.

Participants

A total of 37 participants participated in the experiment in 74 trials. Participants were recruited from Worcester Polytechnic Institute's students and staff as well as other people recruited from outside of the campus. The ages of the participants varied between 19 and 74 with an average of 34.2 years before our screening of 4 participants based on our sanity check questions. After this screening, the ages of the participants varied between 19 and 54 with an average of 30.8 years old. Of the 33 participants, 21 were female and 12 were male. Each participant participated in 2 trials. In one trial the robot was aware of human's affective state and in the second trial the robot was ignoring human's affective state. The order of these two trials were randomly assigned to each participant. Overall, we used affect-ignorant robot first in 16 experiments, and affect-aware robot first in 17 experiments.

5.3.5 Results

As discussed in Section 5.3.4, results of the user study were gathered through a 31-question Likert-scale survey that was given to each participant after each run with the robot, and through a 5-question open-ended summary questionnaire at the end of the experiment.

7-Point Likert Scale Survey Results

As mentioned previously, the 7-point Likert scale survey was administered at the end of the affect-ignorant run and at the end of the affect-aware run for each participant. The 31 questions are generally categorized in accordance with the seven hypotheses listed in Section 5.3.3 to evaluate the humans' perceptions of the following seven categories, with 3-7 questions per group: (1) the likability of the robot (2) the level of trust the human feels in the robot (3) the human's perception of the robot's performance (4) the human's perception of the robot's understanding of human's emotions (5) the human's perception of the robot's understanding of human's and

collaboration's goals and objectives (6) the human's feeling about the collaboration and (7) the human's perception of the human's and robot's mutual satisfaction with each other as collaborative partners.

Table 5.6: The 31 Likert scale questions organized according to their categories (hypotheses).

Question Category (Hypothesis)	Question	Question Number	p-value
Likability	I felt close to the robot.	Q1	« 0.001
	I would like to continue working with the robot.	Q2	« 0.001
	I like the robot.	Q3	« 0.001
	The robot was interesting.	Q4	« 0.001
Trust	I trust the robot.	Q5	« 0.001
	It was easy to express myself to the robot.	Q6	« 0.001
	I trust the robot to perform appropriately in our collaboration.	Q7	« 0.001
	I am confident in the robot's ability to help me.	Q8	« 0.001
	I trust the robot to assess my feelings appropriately in our collaboration.	Q9	« 0.001
Robot's Performance	The robot was repetitive.	Q10	0.002
	The robot made efficient decisions.	Q11	0.040
	The robot's decisions improved my performance during the collaboration.	Q12	« 0.001
Robot's Understanding of Human's Emotions	The robot understood my emotions.	Q13	« 0.001
	The robot is sometimes confused about what I feel about our activities.	Q14	0.005
	I feel that the robot, in its own unique ways, is genuinely concerned about me.	Q15	« 0.001
	The robot understands some of my feelings and takes them into account in our collaboration.	Q16	« 0.001
	The robot does not understand how I feel during our collaboration.	Q17	0.001
Robot's Understanding of Goals	The robot does not understand what we are trying to accomplish.	Q18	0.002
	The robot does not understand what I am trying to accomplish.	Q19	0.007
	The robot perceives accurately what my objectives are.	Q20	« 0.001
	The robot was committed to the collaboration.	Q21	« 0.001
Human Feeling about Collaboration	I find what the robot and I are doing is unrelated to my goals.	Q22	0.001
	I find what I am doing with the robot confusing.	Q23	0.026
	The robot and I are working towards mutually agreed-upon goals.	Q24	0.008
	The robot and I collaborate on setting goals for us to work on.	Q25	« 0.001
	The robot and I agree on what is important for us to work on.	Q26	« 0.001
	I believe that the robot and I achieved the goals we set.	Q27	0.001
	I am satisfied with the outcome of our collaboration.	Q28	« 0.001
Satisfaction of Collaborative Partner	The robot was satisfied with my collaborative behavior.	Q29	« 0.001
	I was satisfied with the robot.	Q30	« 0.001
	I understand the robot, and I think it understands me, at least in the best way it can.	Q31	« 0.001

The results distributions were often skewed, due to the natural limits imposed by the Likert scale. Because the study was a within-subject study, and due to the skewness of the data, the results were analyzed using the Wilcoxon Signed Rank Test, with the normal approximation for large samples. For all questions, the Wilcoxon Signed Rank Test indicated that participants' rating of the affect-aware condition was statistically significantly higher than the affect-ignorant condition (except in reverse-scored questions, when it was statistically significantly lower). The questions and their p-values are provided in Table 5.6. Analysis of the results revealed no statistically significant difference or consistent pattern based on which condition the participant completed first. In the following presentation of results, bar charts with sample means are presented for the purpose of illustration.

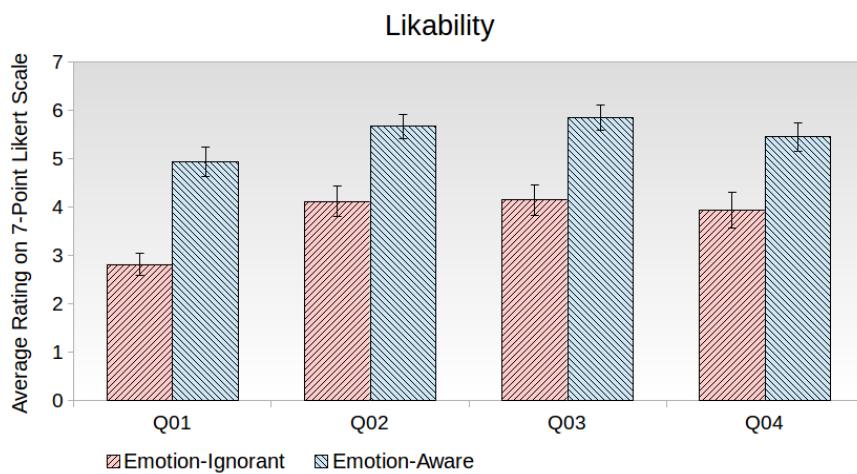


Figure 5.11: Results of the Likert scale survey for Likability questions.

Likability of the Robot

Questions 1 through 4 addressed the likability of the robot. As shown in Figure 5.11, participants rated the affect-aware robot 1.5-2.1 points higher than the affect-ignorant robot. These results indicate that participants felt closer with and preferred working with the affect-aware robot; these results support Hypothesis 1, which stated that humans would prefer to work with the affect-aware robot over the affect-ignorant robot.

Human Trust in the Robot

Questions 5-9 were designed to measure the degree of trust that the human participants felt in the robot. As shown in Figure 5.12, participants trusted the affect-aware robot, on average, a minimum of 1.4 points more than the affect-ignorant robot, both in general and in terms of collaboration performance. In Question 5, participants rated a general statement of trust 1.5 points higher in the affect-aware case. Additionally, in Question 7, participants rated their trust in the affect-aware robot to perform appropriately during collaboration an average of 5.9 on a 7-point Likert scale, where 7.0 would indicate maximum trust; this indicates an acceptable level of trust in the robot’s collaborative abilities. These results support Hypothesis 2, that posits that human participants would find the affect-aware robot to be more trustworthy than the affect-ignorant robot.

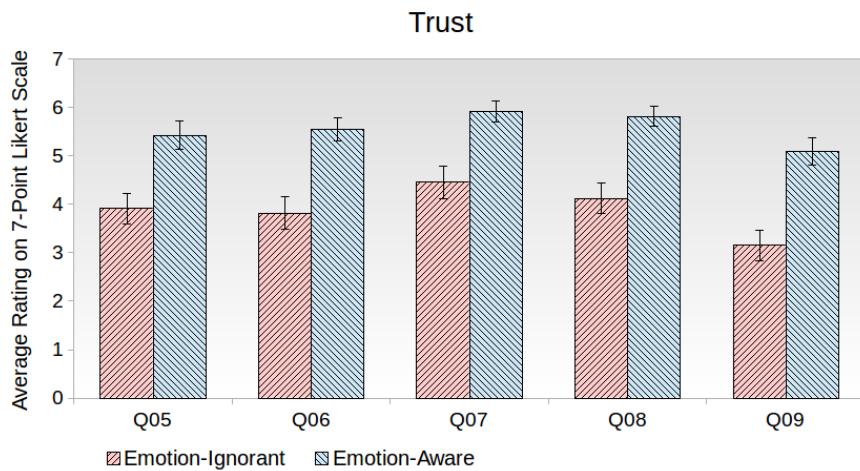


Figure 5.12: Results of the Likert scale survey for questions related to trust.

Perception of the Robot’s Performance

Question 10 (which is reverse-scored) measures the participant’s perception of repetitiveness in the robot during the collaboration. In both conditions, participants rated the robot as moderately repetitive, with the affect-ignorant robot’s average response




Figure 5.13: Results of the Likert scale survey for questions related to the robot’s performance.

being about 1.1 points higher than the affect-aware. This result correlates with several of the open-ended responses which described the affect-aware robot’s behaviors as “cute” and “interesting”, refer to Section 5.3.5. Question 11, which asks about the efficiency of the robot’s decisions is the question with the highest p-value of the 31 questions. This correlates with the result of the open-ended question asking which condition of the robot exhibited behaviors that could prevent human error (refer to 5.3.5); in response to this question, several respondents stated that it may be quicker or simpler to call the supervisor in the event of a task failure, rather than changing the order of the tasks. According to the results from Question 12, the participants felt that the affect-aware robot’s decisions during collaboration improved their own performance, with an average rating of 5.4, while the affect-ignorant robot only received an average rating of 3.3, indicating that participants felt it was not able to interact in such a way as to increase the human’s performance; refer to results from Question 6. These results support Hypothesis 3, which posited that humans will perceive the affect-aware robot as being more capable than the affect-ignorant robot.

Robot’s Understanding of Human Emotions

In Questions 13 through 17, participants evaluate the robot’s understanding of hu-

mans' emotions. In questions 13, 15, and 16, participants rated the affect-aware robot, on average, a minimum of 1.8 points higher than the affect-ignorant robot. In response to questions 14 and 17, which are reverse-scored, participants ranked the affect-ignorant robot 1.2 and 2.0 points higher, respectively, than then affect-aware robot. The results of all five questions in this category support Hypothesis 4.




Figure 5.14: Results of the Likert scale survey for the questions related to the robot's understanding of human emotions.

Robot's Understanding of Human and Collaboration Goals

Questions 18 and 19 were reverse-scored questions intended to determine whether the humans felt that the robot understood the shared collaboration goal and the human's personal goal, respectively. For both conditions of the robot, the average scores were lower than 3.5, indicating that the human's perceived the robot as having some understanding of the goals. For both questions, the affect-ignorant robot's average score was significantly higher than the affect-aware robot's score. Similarly, Question 20 was a measure of whether the human perceived that the robot correctly perceived the human's goal. On average, participants provided an average rating for the affect-aware robot that was 1.5 points higher than that for the affect-ignorant robot. Question 21 measured the human perception of the robot's commitment to the collaboration; for this measure, the average participant score

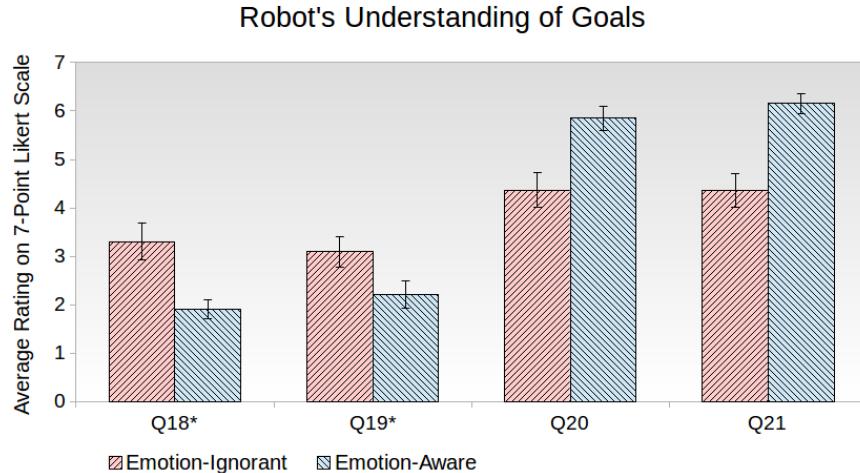


Figure 5.15: Results of the Likert scale survey for questions related to the robot’s understanding of goals.

assigned to the affect-aware robot was 6.2 points out of a maximum of 7 points, indicating that the participants felt that the affect-aware robot was strongly committed to the collaboration. The affect-ignorant robot received an average rating of 4.4 points, indicating only moderate commitment. These results strongly support Hypothesis 5, which posits that humans will feel that the affect-aware robot will better understand their goals than the affect-ignorant robot.

Human’s Feeling about the Collaboration

Questions 22 through 28 were designed to gauge how the human participants felt about the partnership within the collaboration and the outcome of the collaboration. For each of the 7 questions, the participants ranked the affect-aware robot as better than the affect-ignorant robot, by a minimum, on average, of 0.8 points. Questions 24, 27 and 28 addressed whether the robot and the participant were working toward mutually agreed-upon goals and on the outcome of the collaboration; in the affect-aware condition, participants rated the robot a minimum of 6.1 points, on average, while rating the affect-ignorant robot 1-1.6 points lower, indicating that the participants felt a very strong sense of collaboration with the affect-aware robot, and




Figure 5.16: Results of the Likert scale survey for questions related to the human's feeling about the collaboration.

only a moderate sense of collaboration with the affect-ignorant robot. Questions 25 and 26 address whether the robot and the participant set the collaboration goals together; these two questions have lower scores than Questions 24, 27 and 28, for both the affect-aware and the affect-ignorant case. The lower overall scores are likely due to the fact that the robot decides the task order or action in the event of failure in both conditions; however, the higher score in the affect-aware case may indicate that emotional awareness can increase a feeling of collaboration. These results support Hypothesis 6 that humans will feel a greater sense of mutual collaboration and understanding about the collaboration with the affect-aware robot.

Human Perception of Mutual Satisfaction with Collaborative Partner

Questions 29, 30 and 31 were designed to measure the human's perception of the robot's satisfaction with the human, the human's satisfaction with the robot and the mutual understanding between the human and the robot, respectively. The participants provided an average response in the affect-aware condition of 5.8, 5.9 and 5.7 to Questions 29, 30 and 31, respectively, indicating a high level of mutual satisfaction; all three answers were about 1.4-1.9 points lower, on average, in the

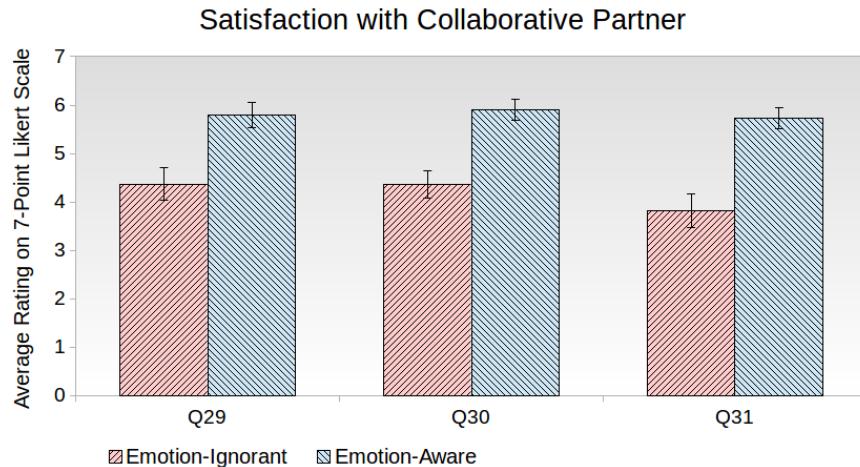


Figure 5.17: Results of the Likert scale survey for questions related to satisfaction with collaborative partner.

affect-ignorant condition. These results indicate a higher level of satisfaction working with the robot in the affect-aware condition, and strongly support Hypothesis 7, which posited that humans will feel a greater sense of mutual satisfaction with the affect-aware robot than the affect-ignorant robot.

Results from the Open-Ended Questionnaire

As described in Section 5.3.4, each participant answered an open-ended questionnaire at the end of the study. Table 5.7 summarizes the questionnaire and which condition users preferred for certain conditions (i.e. affect-ignorant or affect-aware). Note that some users chose not to state a preference regarding which condition they preferred for certain conditions; because we were specifically interested in whether users preferred the affect-aware case, we considered the ambiguous responses to be failures in the binomial analysis. The binomial analysis is based on a population size of 33.

As shown in Table 5.7, 100% of users unambiguously preferred the run with the affect-aware robot. In general, this preference stemmed from a feeling of closeness and partnership, as seen in these responses: “the robot had emotions and responded

Table 5.7: Open-ended questionnaire questions and results. (*Note: Because we are evaluating whether humans prefer an affect-aware robot, these results are taken as negative test results when calculating the p-value using the binomial distribution. Only those participants who clearly indicated a preference for the emotion-aware robot are taken as positive test results.)

Question	Number of Participants Who Did Not Prefer One Condition Over the Other *	Number of Participants Favoring Emotion-Aware Robot	p-value
Which of the two runs with the robot did you prefer?	0	33	0
In which of the two runs did the robot exhibit behavior that could be useful in a more complex task?	1	30	≤ 0.001
In which of two runs did the robot exhibit behavior that could prevent human error?	3	18	> 0.1
In which of the two runs did the robot exhibit behavior that could improve the efficiency of collaboration?	2	26	≤ 0.001
What was the most interesting behavior of the robot and in which run did it happen?	5	24	0.002

to my emotions. Also, “what it said about my failing was cute and aimed to make me feel better.” Another example is “I liked feeling needed and accounted for; I felt closer to the robot.” Finally, “I saw the changes in its feeling, which motivated me to care more about my act...I also liked that he asked me to correct its failure, although it could ask the supervisor.”

When asked in which of the two runs the robot exhibited behavior that could be useful in a more complex task, 90.9% chose the affect-aware robot. In general, respondents thought that the affect-aware robot was better at problem solving, more adaptable, and more capable of handling the social complexities that occur in collaboration, as shown in responses such as “The robot explained motives...which is important to keep a team communicating and on the same pace.” Also, “When we failed he initially switched to a new task and then came back to the originally failed task. It kept me from getting irritated and negative.” Finally, “The more complex, the more necessary it is to understand how humans think and operate...an empathetic robot can adapt, encourage and help.” It is worth noting that one respondent preferred the affect-ignorant case, saying “In a more complex task it might be better for the robot to take control and simply tell me what to do; trying to be understanding and collaborative wouldn’t be as important as doing the task

correctly.”

The only question that did not provide statistically significant support in favor of the affect-aware robot related to which case the robot exhibited behavior that could prevent human error. About 36.4% of respondents thought that the affect-ignorant robot was more likely to prevent human error; however, all but one of these cited calling the supervisor as the main method of preventing human error, in spite of the fact that the instructions indicated that the robot’s need to call the supervisor counted against the collaboration. Of the 54.5% who thought that the affect-aware robot was better at preventing human error, most cited the robot’s ability to console the human as the main behavior that could prevent human error. Respondents indicated that this enabled them to move on and feel better about the collaboration, as with this response: “The robot switched to a different task and we came back to an error later. This allowed my mind to move away from being frustrated. I was able to complete a different task which felt like a win - then come back and finish the error. Making my mind move away from frustration could definitely prevent more errors.”

When asked in which of the runs the robot exhibited behavior that could improve the efficiency of the collaboration, 78.8% responded with the affect-aware case; of these, the vast majority stated that this was because of the robot’s ability to change the order of tasks in the event of a failure, and to ask the human for help.

Finally, when asked in which run the most interesting behavior occurred, 72.7% chose the affect-aware condition. Of these respondents, 12 individuals stated that the robot’s attempt to console the human by saying “It was not your fault” in response to the human’s negative affective state that occurred as a consequence of the human’s failed task was the most interesting behavior, and a majority mentioned that it actually made them feel more positive. Six participants referred to the robot’s ability to understand and express affective state. Several participants referred to the robot’s ability to communicate, including the ability to ask questions. Of those who responded with the affect-ignorant case, most found the ability to call the supervisor,

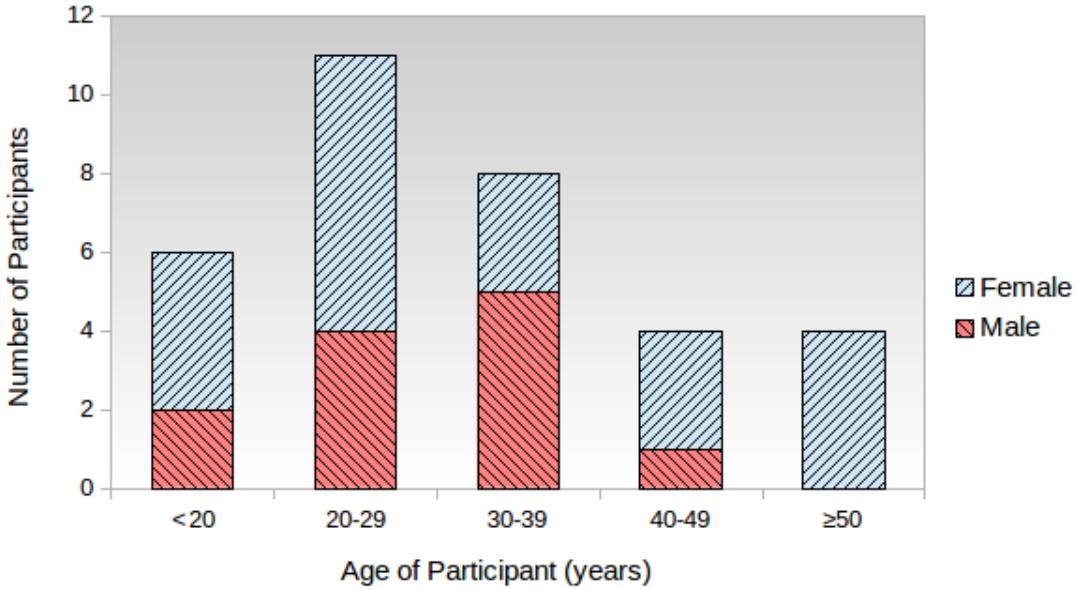


Figure 5.18: Age and gender distribution of the user study participants.

and mechanical functions, such as gripping, to be most interesting.

Impact of Demographics

As mentioned in Section 5.3.4, we recorded the age and gender of each participant; the distribution is provided in Figure 5.18. The distribution presented in this figure are the results presented in this thesis; however, four participants were removed from the results, due to the participant breaking study rules. Two of the removed participants were women over 40 years old and two were men over 50 years old. Although it was not the primary purpose of the study, we investigated the Likert scale results to determine if there were any relevant trends based on the demographics of the participants.

To investigate the impact of demographics, we carried out Multiple Linear Regression analysis for each question, under each condition of the robot (i.e. affect-aware and affect-ignorant). For each question, the dependent variable was the participant rating; each question was analyzed with four combinations of dependent

Table 5.8: Multiple linear regression analysis results for the affect-aware condition.

Question Number	Model Significance ($F(df_{\text{effect}}, df_{\text{error}}) = F\text{-value}, p\text{-value}$)	Adjusted R ² Value	Age (coeff., p-value)	Gender (coeff., p-value)	Age*Gender (coeff., p-value)
Q2	$F(3,29)=4.37, p=0.012$	0.240	-0.152, $p=0.003$	-3.600, $p=0.024$	0.152, $p=0.005$
Q3	$F(3,29)=7.42, p<0.001$	0.376	-0.202, $p<0.001$	-4.725, $p=0.003$	0.191, $p<0.001$
Q7	$F(3,29)=4.09, p=0.015$	0.224	-0.124, $p=0.007$	-2.747, $p=0.058$	0.126, $p=0.011$
Q8	$F(3,29)=5.31, p=0.005$	0.288	-0.140, $p=0.001$	-3.513, $p=0.010$	0.147, $p=0.002$
Q9	$F(1,31)=6.87, p=0.013$	0.155	-0.060, $p=0.013$	-	-
	$F(3,29)=3.88, p=0.019$	0.213	-0.164, $p=0.006$	-3.185, $p=0.093$	0.122, $p=0.056$
Q12	$F(3,29)=3.79, p=0.021$	0.207	-0.160, $p=0.008$	-3.629, $p=0.059$	0.164, $p=0.013$
Q13	$F(1,31)=6.55, p=0.016$	0.148	-0.064, $p=0.016$	-	-
Q16	$F(1,31)=7.40, p=0.011$	0.167	-0.076, $p=0.011$	-	-
	$F(3,29)=4.65, p=0.009$	0.255	-0.202, $p=0.005$	-3.123, $p=0.166$	0.141, $p=0.065$
Q17*	$F(1,31)=6.48, p=0.016$	0.146	0.054, $p=0.016$	-	-
Q18*	$F(1,31)=5.65, p=0.024$	0.127	-	-0.929, $p=0.024$	-
Q20	$F(3,29)=3.05, p=0.044$	0.161	-0.129, $p=0.018$	-2.935, $p=0.091$	0.133, $p=0.024$
Q23*	$F(3,29)=3.08, p=0.043$	0.163	0.064, $p=0.021$	1.962, $p=0.031$	-0.079, $p=0.010$
Q28	$F(1,31)=6.04, p=0.020$	0.136	-	0.917, $p=0.020$	-
	$F(3,29)=5.42, p=0.004$	0.293	-0.106, $p=0.006$	-2.165, $p=0.074$	0.110, $p=0.009$
Q30	$F(2,30)=4.46, p=0.020$	0.178	-0.040, $p=0.029$	0.958, $p=0.027$	-
	$F(3,29)=6.68, p=0.001$	0.348	-0.143, $p<0.001$	-2.589, $p=0.048$	0.124, $p=0.006$
Q31	$F(1,31)=5.42, p=0.027$	0.121	-	1.102, $p=0.027$	-
	$F(2,30)=5.60, p=0.009$	0.223	-0.039, $p=0.032$	1.194, $p=0.008$	-
	$F(3,29)=6.15, p=0.002$	0.326	-0.126, $p=0.004$	-1.782, $p=0.188$	0.104, $p=0.025$

variables: (1) age only, (2) gender only, (3) age and gender, and (4) age, gender and age*gender (effect modification). The results for the affect-aware case are presented in Table 5.8 and the results for the affect-ignorant case are presented in Table 5.9. The tables contain all models that were statistically significant, except those that revealed a decrease in the R^2 when the number of parameters increased; parameters not included in a specific regression are left blank in the tables. Questions that are excluded from the table did not have any statistically significant models based on age and gender.

As shown in Table 5.8 and Table 5.9, for many questions, age and gender only provided statistically significant regression models when effect modification was included; while this may actually indicate effect modification, we suspect that this may also be related to the demographic distribution. As mentioned previously, the

Table 5.9: Multiple linear regression analysis results for the affect-ignorant condition.

Question Number	Model Significance ($F(df_{\text{effect}}, df_{\text{error}}) = F\text{-value}, p\text{-value}$)	Adjusted R ² Value	Age (coeff., p-value)	Gender (coeff., p-value)	Age*Gender (coeff., p-value)
Q9	$F(2,30)=4.97, p=0.014$	0.199	-0.076, $p=0.006$	0.985, $p=0.118$	-
Q13	$F(2,30)=8.42, p=0.001$	0.317	-0.077, $p=0.002$	1.476, $p=0.009$	-
Q16	$F(2,30)=3.42, p=0.046$	0.131	-0.065, $p=0.031$	1.198, $p=0.095$	-
Q18*	$F(3,29)=2.946, p=0.049$	0.154	-0.031, $p=0.691$	-4.725, $p=0.071$	0.120, $p=0.165$
Q24	$F(1,31)=4.76, p=0.037$	0.105	-	1.214, $p=0.037$	-
Q31	$F(1,31)=9.08, p=0.005$	0.202	-	1.940, $p=0.005$	-

only two male participants over the age of 50 were removed from the data pool due to breaking the rules with the robot. The resulting demographic distribution contains eight participants over 40 years old, only one of whom is a male.

When effect modification is either not included or does not improve the explanatory power of the model(s), two consistent patterns emerge across the questions that are impacted by age and/or gender: (1) as age increases, participants' ratings of the robot tends to decrease, in both the affect-aware and the affect-ignorant condition and (2) female participants' ratings tend to be higher than male participants' ratings in both the affect-aware and the affect-ignorant condition.

Finally, while it is clear that age and gender impact the participants' ratings of the robot, the explanatory power of the regression models remains low; the highest R² value was 0.376, which occurred for Question 3, when effect modification was included. Age and gender of the participant impact the participants' ratings of the robot, but on their own, do not explain most of the variation observed in the Likert scale surveys.

5.3.6 Discussion

Based on the results, all participants prefer to work with the affect-aware robot. Humans find the affect-aware robot more likable and more trustworthy, as indicated in the Likert-scale responses and the open-ended questionnaire responses. Based on the responses, the emotional interaction with the robot can help create a sense

of closeness and enjoyment that makes humans want to continue working with the robot.

The results also indicate that the affect-aware robot can better maintain a collaborative relationship. Both Likert-scale responses and Open-Ended Questionnaire responses indicate this. Humans felt a stronger sense of the robot’s commitment to the collaboration, and greater understanding of their goals and emotions from the robot. Several open-ended responses also indicated that the robot was able to successfully motivate people and maintain their commitment to the collaboration, especially when tasks failed. Additionally, as shown in Section 5.3.5, humans rated the affect-aware case much higher than the affect-ignorant case when asked which robot’s decisions improved their performance, in essence acknowledging that their collaborator’s (i.e., the robot’s) decisions had a significant impact on their performance. As some of the open-ended responses indicated, successfully managing emotions within the collaboration can help keep the collaboration on track, and prevent distractions due to guilt and other negative emotions.

Finally, the affect-aware robot developed a stronger sense of partnership through greater communication. The participants felt better understood by the affect-aware robot, and felt that the goals were more mutually agreed-upon, refer to Section 5.3.5. As evidenced in the following response, the affect-aware robot was successfully able to create a sense of partnership through its more open communication style: “Communication is very important. In the first run (i.e. affect-aware) the robot states what tasks he is working on, it is clear and straight-forward. Also during the first run the robot cares about the human(me)’s feelings and cheers me up when I failed at the tasks, I think that could also improve efficiency of collaboration, because it would be more like a team or partnership.”

CHAPTER 6

CONCLUSION

6.1 Discussion

This thesis presents the Affective Motivational Collaboration (AMC) theory and our computational framework. The AMC is built on the SharedPlans theory of collaboration [103] and the cognitive appraisal theory of emotions [162] [223]. Our motivation to develop AMC was the lack of a theory describing the processes involved in a dyadic collaboration as well as their relationship and influences on each other. In particular, in this thesis we emphasized the reciprocal influence of the collaboration structure and the appraisal processes in collaboration. We provided algorithms to compute appraisal variables and their influence on collaboration processes, e.g., goal management. In general, our contribution in this thesis was to provide a theory which describes affect-regulated goal-driven behaviors within a dyadic collaboration. A further contribution of this thesis is to account for the influence of motives on the coping processes in collaboration. We validated our individual appraisal algorithms as well as our overall computational framework by conducting an online crowd-sourcing user study and a laboratory end-to-end system user study, respectively. The first study investigated whether humans and our appraisal algorithms provide similar answers to questions with respect to factors involved in our appraisal algorithms. The second study investigated a) the importance of emotional-awareness in collaboration, and b) the overall functionality of the AMC framework to autonomously control interactions of a collaborative robot.

After our introduction in Chapter 1, we presented the theoretical background on the two major foundations of our theory in Chapter 2. First, we reviewed the prominent computational collaboration theories including SharedPlans and Joint Intentions. We focused on the main concepts characterizing requirements of a collaboration introduced by these theories. We also analyzed the similarities and differences between these theories in terms of their essential concepts as well as their theoretical and practical applications. These applications involved the fields of robotic and artificial agents. Then, we continued discussing what emotions are and more importantly how they can influence one's cognition and social life. We also discussed the role of emotions in communicating one's internal states to others as the basic rationale behind different social emotions. We confined our discussion about emotion to artificial emotions in social robots or agents. Third, we reviewed existing computational models of emotions, including appraisal theory, analyzing their similarities, differences and applications in robotics and artificial agents. Finally, we reviewed the concept of motives, work in related fields and described three social motives based on the psychological theories.

Next, we introduced Affective Motivational Collaboration theory in Chapter 3. We discussed all the mechanisms involved in our theory. These mechanisms include various processes, each of which provides particular information required in overall operation of the system. Among these mechanisms our focus was on the Collaboration, Appraisal and Coping mechanisms. However, other mechanisms such as Motivation also play important roles in influencing the overall behavior of the agent using our framework. We also discussed the events that we consider in a collaborative environment including utterances, primitive actions and observable behaviors. We described how each mechanism handles these events. We believe it is important to focus on functions of emotions and their influence on collaboration processes. Therefore, we briefly described a set of emotion functions and how they are related to the collaboration context. Then, we continued by explaining the input, output and function of each mechanism involved in our architecture as well as the

Mental State, knowledge-base required by all the mechanisms, containing beliefs, intentions, motives, goals and emotion instances. We presented different attributes of each element of mental state.

In Chapter 4, we introduced our computational framework based on AMC theory in more detail. As a major part of our contribution, we explained our algorithms to compute the values associated with the appraisal variables. We use these algorithms to compute the values of relevance, desirability, expectedness and controllability of an event occurring during the collaboration. All of these algorithms process data provided by the collaboration structure. Reciprocally, we provided the details of how we use the outcome of the appraisals to influence the collaboration structure, specifically by providing inputs to our algorithm for goal management. Then, we explained details of the coping strategies, e.g. Active Coping, involved in our Coping Mechanism and the underlying processes associated with these coping strategies. We also included details about the Motivation mechanism, the types of motives we considered and how we compute their values. Finally, we describe the elicitation of different emotion instances in our framework and how they are interpreted according the different contexts during collaboration.

We carried out two user studies which validated our framework. The first study was designed to test whether humans and our appraisal algorithms perceive certain factors in our algorithms similarly. The results, which validated our algorithms, are presented in Chapter 5. Our second study, which was designed to test the overall functionality of AMC framework, was also presented in Chapter 5.

6.2 Future Work

This work paves the way for a number of potential extensions. In particular, we believe that extensions to the system could be made by exploring how to employ the emotion functions we discussed in Chapter 4 with respect to the human's emotion. The AMC framework currently employs emotion functions including *social regulation*,

motivation, goal management and focus of attention. By acquiring these or other emotion functions, the agent improves the quality of collaboration. However, each emotion can have a different impact on different emotion functions in a given collaboration context. For instance, the agent can interpret the meaning of the perceived emotion, e.g., anger, while the collaborators are negotiating (context) pursuing or abandoning a given goal. This interpretation can be different for goal management and motivation as emotion functions; i.e., human's perceived anger can cause the agent to choose a relatively easier goal to pursue, and to postpone motivating the human to pursue more difficult goals for the moment. Since the meaning of many social emotions are defined in the field of psychology, the agent can use these meanings to improve likability or other important factors in the collaboration.

Secondly, while our main contribution focused on enabling an agent to improve different aspects of collaboration (e.g., collaborator's satisfaction, likability, trust, etc.), another possible area of future work would be to explore ways in which the agent can improve a chosen aspect of the collaboration such as trust or performance at any given time. In particular, an interesting extension would be to enable the agent to perceive which aspect of the collaboration is suffering (e.g., lack of trust) and needs to be improved. For example, if the human collaborator is losing her trust in the agent to achieve a given goal, the agent can try to improve the human collaborator's sense of trust by showing an appropriate behavior, e.g., improving the precision or helping the human to achieve a goal.

Finally, in addition to expanding the adaptability of the agent to the human collaborator's internal state, future extensions are also possible in the other mechanisms as well. For example, an interesting area to explore is the ability to employ more elaborate computational models of motivation and theory of mind. These extensions could provide more information about the human collaborator and help the agent to act on a more accurate model of the human's internal state.

BIBLIOGRAPHY

- [1] C. Adam, A. Herzog, and D. Longin. A logical formalization of the OCC theory of emotions. *Synthese*, 168(2):201–248, 2009.
- [2] C. Adam and E. Lorini. A BDI emotional reasoning engine for an artificial companion. In *Workshop on Agents and multi-agent Systems for AAL and e-HEALTH (PAAMS)*, volume 430, pages 66–78. Springer, 2014.
- [3] N. Ahmadpour. Occ model: application and comparison to the dimensional model of emotion. In *In Proceedings of International Conference on KANSEI Engineering and Emotion Reseach*, pages 607–617, 2014.
- [4] J. R. Anderson and C. Lebiere. The newell test for a theory of cognition. *Behavioral and Brain Sciences*, 26(5):587–640, 2003.
- [5] J. Andriessen, K. de Smedt, and M. Zock. Discourse planning: Empirical research and computer models. In A. Dijkstra and K. de Smedt, editors, *Computational psycholinguistics: AI and connectionist models of human language processing*, pages 247–278. Taylor & Francis, 1996.
- [6] G. E. M. Anscombe. *Intention*. NY: Cornell University Press, 1963.
- [7] Aristotle. *The Nicomachean Ethics*. George Bell and Sons, 2009.
- [8] Armony, Servan-Schreiber, Cohen, and Ledoux. Computational modeling of emotion: Explorations through the anatomy and physiology of fear conditioning. *Trends in Cognitive Sciences*, 1:28–34, 1997.

- [9] M. B. Arnold. *Emotion and personality*. Cassell Co., 1960.
- [10] M. Asada, K. F. Macdorman, H. Ishiguro, and Y. Kuniyoshi. Cognitive developmental robotics as a new paradigm for the design of humanoid robots. *Robotics and Autonomous Systems*, 37:185–193, 2001.
- [11] J. W. Atkinson. *Motives in fantasy, action, and society: a method of assessment and study*. Van Nostrand, 1958.
- [12] T. Babaian, B. J. Grosz, and S. M. Shieber. A writer’s collaborative assistant. In *Proceedings of the International Conference on Intelligent User Interfaces (IUI2000)*, pages 7–14. ACM Press, 2002.
- [13] J. Bach. The MicroPsi Agent Architecture. In *Proceeding of ICCM-5*, pages 15–20, 2003.
- [14] J. Bach. *Principles of Synthetic Intelligence PSI: An Architecture of Motivated Cognition*. Oxford University Press, Inc., 2009.
- [15] J. Bach. A motivational system for cognitive ai. In J. Schmidhuber, K. R. Thórisson, and M. Looks, editors, *Artificial General Intelligence*, volume 6830 of *Lecture Notes in Computer Science*, pages 232–242. Springer Berlin Heidelberg, 2011.
- [16] J. Bach. Micropsi 2: The next generation of the micropsi framework. In *Proceedings of the 5th International Conference on Artificial General Intelligence, AGI’12*, pages 11–20. Springer-Verlag, 2012.
- [17] I. Bakker, T. van der Voordt, P. Vink, and J. de Boon. Pleasure, arousal, dominance: Mehrabian and russell revisited. *Current Psychology*, 33(3):405–421, 2014.
- [18] L. F. Barrett. Are emotions natural kinds? *Perspectives on Psychological Science*, 1(1):28–58, 2006.

- [19] J. Bates. The role of emotion in believable agents. *Commun. ACM*, 37(7):122–125, 1994.
- [20] P. Baxter, J. de Greeff, and T. Belpaeme. Cognitive Architecture for Human-Robot Interaction: Towards Behavioural Alignment. *Journal of Biologically Inspired Cognitive Architectures*, 6:30–39, 2013.
- [21] R. Beck. *Motivation: Theories and Principles*. Pearson, 2003.
- [22] C. Becker-asano. *WASABI: Affect Simulation for Agents with Believable Interactivity*. PhD thesis, University of Bielefeld, 2008.
- [23] C. Becker-Asano. Wasabi for affect simulation in human-computer interaction: Architecture description and example applications. In *Proceedings of ERM4HCI workshop in conjunction with ICMI 2013*. Springer, 2013.
- [24] K. C. Berridge. Motivation concepts in behavioral neuroscience. *Physiology & Behavior*, 81(2):179–209, 2004.
- [25] L. M. Botelho and H. Coelho. Machinery for artificial emotions. *Cybernetics and Systems*, 32(5):465–506, 2001.
- [26] M. E. Bratman. *Intention, Plans, and Practical Reason*. Cambridge, Mass.: Harvard University Press, 1987.
- [27] M. E. Bratman. What is intention? In *Intentions in Communication*, Cognitive Technologies, pages 15–32. The MIT Press, Cambridge, MA, June 1990.
- [28] M. E. Bratman. Shared cooperative activity. *Philosophical Review*, 101(2):327–341, 1992.
- [29] M. E. Bratman, D. J. Israel, and M. E. Pollack. Plans and resource-bounded practical reasoning. *Computational Intelligence*, 4(3):349–355, 1988.

- [30] S. Brave and C. Nass. Emotion in human-computer interaction. In J. A. Jacko and A. Sears, editors, *The Human-computer Interaction Handbook*, pages 81–96. L. Erlbaum Associates Inc., 2003.
- [31] C. Breazeal. A motivational system for regulating human-robot interaction. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence and Tenth Innovative Applications of Artificial Intelligence Conference, AAAI 98, IAAI 98, July 26-30, 1998, Madison, Wisconsin, USA.*, pages 54–61, 1998.
- [32] C. Breazeal. *Designing Sociable Robots*. MIT Press, 2002.
- [33] C. Breazeal. Role of expressive behaviour for robots that learn from people. *Philosophical transactions of the Royal Society of London. Series B, Biological sciences*, 364(1535):3527–38, 2009.
- [34] C. Breazeal, A. Brooks, J. Gray, G. Hoffman, C. Kidd, H. Lee, J. Lieberman, A. Lockerd, and D. Mulanda. Humanoid robots as cooperative partners for people. *Journal of Humanoid Robots*, 1(2):1–34, 2004.
- [35] C. Breazeal, A. Takanishi, and T. Kobayashi. *Social Robots that Interact with People*, pages 1349–1369. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [36] I. Bretherton, J. Fritz, C. Zahn-Waxler, and D. Ridgeway. Learning to talk about emotions: A functionalist perspective. *Child Development*, 57(3):529–548, 1986.
- [37] N. Brody. *Human motivation: commentary on goal-directed action*. Academic Press New York, 1983.
- [38] J. Broekens, W. A. Kosters, and F. J. Verbeek. On affect and self-adaptation: Potential benefits of valence-controlled action-selection. In: *J. Mira and J.R. lvarez (Eds.), IWINAC 2007, Part I, LNCS*, pages 357–366, 2007.

- [39] C. Burghart, R. Mikut, R. Stiefelhagen, T. Asfour, H. Holzapfel, P. Steinhaus, and R. Dillmann. A cognitive architecture for a humanoid robot: A first approach. In *5th IEEE-RAS International Conference on Humanoid Robots*, pages 357–362, 2005.
- [40] Z. Cai, B. Goertzel, C. Zhou, Y. Zhang, M. Jiang, and G. Yu. Dynamics of a computational affective model inspired by dörners PSI theory. *Cognitive Systems Research*, 17–18:63–80, 2012.
- [41] L. D. Canamero. Designing emotions for activity selection in autonomous agents. In R. Trappi, P. Petta, and S. Payr, editors, *Emotions in Humans and Artifacts*, pages 115–148. MIT Press, 2003.
- [42] K. Carley. On the persistence of beliefs. working paper. *Department of Social and Decision Sciences, Carnegie Mellon University*, 1990.
- [43] C. S. Carver. Affect and the functional bases of behavior: On the dimensional structure of affective experience. *Personality and Social Psychology Review*, 5(4):345–356, 2001.
- [44] H.-N. Castañeda. *Thinking and Doing*. Dordrecht, Holland: D. Riedel, 1975.
- [45] C. Castelfranchi. Commitments: From individual intentions to groups and organizations. In *Proceedings of the first international conference on multiagent systems*, pages 41–48, 1995.
- [46] D. Cañamero and W. V. de Velde. Socially emotional: Using emotions to ground social interaction. In *In Working Notes of the AAAI'97 Fall Symposium on Socially Intelligent Agents*, 1997.
- [47] H. Q. Chong, A. H. Tan, and G. W. Ng. Integrated cognitive architectures: A survey. *Artificial Intelligence Review*, 28(2):103–130, 2007.

- [48] A. B. S. Clair and M. J. Matarić. Modeling action and intention for the production of coordinating communication in human-robot task collaborations. In *21st IEEE International Symposium on Robot and Human Interactive Communication: Workshop on Robot Feedback in HRI*, Paris, France, 2012.
- [49] W. J. Clancey. Roles for agent assistants in field science: Understanding personal projects and collaboration. *IEEE Transactions on Systems, Man and Cybernetics, special issue on Human-Robot Interaction*, 34(2):125–137, 2004.
- [50] G. L. Clore and J. R. Huntsinger. How emotions inform judgment and regulate thought. *Trends in Cognitive Sciences*, 11(9):393–399, 2007.
- [51] G. L. Clore and J. Palmer. Affective guidance of intelligent agents: How emotion controls cognition. *Cognitive System Research*, 10(1):21–30, 2009.
- [52] J. A. Coan and J. J. Allen. *Handbook of Emotion Elicitation and Assessment*. Oxford University Press, USA, 2007.
- [53] P. Cohen, H. Levesque, and I. Smith. On team formation. In *Contemporary Action Theory. Synthese*, pages 87–114. Kluwer Academic Publishers, 1997.
- [54] P. Cohen and H. J. Levesque. *Teamwork*. SRI International, 1991.
- [55] P. R. Cohen and H. J. Levesque. Intention is choice with commitment. *Artificial Intelligence*, 42(2-3):213–261, 1990.
- [56] P. R. Cohen and H. J. Levesque. Persistence, intention, and commitment. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 33–69. MIT Press, Cambridge, MA, 1990.
- [57] P. R. Cohen, J. Morgan, and M. E. Pollack. *Intentions in Communication*. A Bradford Book, 1990.

- [58] C. Conati and H. Maclare. Evaluating a probabilistic model of student affect. In *7th International Conference on Intelligent Tutoring Systems*, 2004.
- [59] R. Cooke and P. Sheeran. Moderation of cognition-intention and cognition-behaviour relations: A meta-analysis of properties of variables from the theory of planned behaviour. *British Journal of Social Psychology*, 43:159–186, 2004.
- [60] R. Cowie, N. Sussman, and A. Ben-Ze’ev. Emotion: Concepts and definitions. In *Emotion-Oriented Systems*, Cognitive Technologies, pages 9–30. Springer-Verlag Berlin Heidelberg, London, New York, March 2011.
- [61] C. L. Dancy. ACT-R ϕ : A cognitive architecture with physiology and affect. *Biologically Inspired Cognitive Architectures*, 6:40–45, 2013.
- [62] C. Darwin. *The expression of the emotions in man and animals*. New York; D. Appleton and Co., 1916.
- [63] J. S. de Freitas and J. Queiroz. Artificial emotions: Are we ready for them? In F. A. e Costa, L. M. Rocha, E. Costa, I. Harvey, and A. Coutinho, editors, *ECAL*, volume 4648, pages 223–232. Springer, 2007.
- [64] C. M. de Melo, J. Gratch, P. Carnevale, and S. J. Read. Reverse appraisal: The importance of appraisals for the effect of emotion displays on people’s decision-making in social dilemma. In *Proceedings of the 34th Annual Meeting of the Cognitive Science Society (CogSci)*, 2012.
- [65] E. de Sevin and D. Thalmann. A motivational model of action selection for virtual humans. In *Proceedings of the Computer Graphics International 2005*, CGI ’05, pages 213–220, Washington, DC, USA, 2005. IEEE Computer Society.
- [66] B. M. DePaulo. Nonverbal behavior and self-presentation. *Psychological Bulletin*, 111(2):203–243, 1992.

- [67] S. D'Mello and S. Franklin. Computational modeling/cognitive robotics complements functional modeling/experimental psychology. *New Ideas in Psychology*, 29(3):217–227, 2011.
- [68] P. F. Dominey and F. Warneken. The basis of shared intentions in human and robot cognition. *New Ideas in Psychology*, 29(3):260–274, 2011.
- [69] H. L. Dreyfus. *What Computers Still Can't Do: A Critique of Artificial Reason*. MIT Press, 1992.
- [70] P. Ekman. An argument for basic emotions. *Cognition & Emotion*, 6(3-4):169–200, 1992.
- [71] P. Ekman and W. V. Friesen. Measuring facial movement. *Environmental psychology and nonverbal behavior*, 1(1):56–75, 1976.
- [72] M. S. El-Nasr, J. Yen, and T. R. Ioerger. Flame: Fuzzy logic adaptive model of emotions. *Autonomous Agents and Multi-Agent Systems*, 3(3):219–257, 2000.
- [73] U. X. Eligio, S. E. Ainsworth, and C. K. Crook. Emotion understanding and performance during computer-supported collaboration. *Computers in Human Behavior*, 28(6):2046–2054, 2012.
- [74] C. D. Elliott. *The Affective Reasoner: A Process Model of Emotions in a Multi-agent System*. PhD thesis, Northwestern University Institute for the Learning Sciences, 1992.
- [75] N. Esau, L. Kleinjohann, and B. Kleinjohann. Integrating emotional competence into man-machine collaboration. In *Biologically-Inspired Collaborative Computing, September 8-9, Milano, Italy*, pages 187–198, 2008.
- [76] L. Faucher and C. Tappolet. Fear and the focus of attention. *Consciousness and Emotion*, 3(2):105–144, 2002.

- [77] S. Folkman and J. T. Moskowitz. Coping: Pitfalls and promise. *Annual Review of Psychology*, 55:745–774, 2004.
- [78] B. L. Fredrickson and C. Branigan. Positive emotions broaden the scope of attention and thoughtaction repertoires. *Cognition and Emotion*, 19(3):313–332, 2005.
- [79] N. H. Frijda. *The Emotions*. Cambridge University Press, 1986.
- [80] S. R. Fussell and R. M. Krauss. Coordination of knowledge in communication: Effects of speakers’ assumptions about what others know. *Journal of personality and social psychology*, 62(1):378–391, 1992.
- [81] S. C. Gadanho and J. Hallam. Robot learning driven by emotions. *AI Magazine*, 9(1):42–64, 2001.
- [82] O. García, J. Favela, G. Licea, and R. Machorro. Extending a collaborative architecture to support emotional awareness. In *Emotion Based Agent Architectures (ebaa99)*, pages 46–52, 1999.
- [83] P. Gebhard. A layered model of affect. In *5th International Joint Conference on Autonomous Agents and Multiagent Systems*, pages 29–36, 2005.
- [84] P. J. Gmytrasiewicz and C. L. Lisetti. Emotions and personality in agent design and modeling. In J.-J. C. Meyer and M. Tambe, editors, *ATAL*, volume 2333 of *Lecture Notes in Computer Science*, pages 21–31. Springer, 2001.
- [85] E. Goffman. *The Presentation of Self in Everyday Life*. Anchor, 1959.
- [86] A. C. Gonzalez, M. Malfaz, and M. A. Salichs. An autonomous social robot in fear. *IEEE Transactions Autonomous Mental Development*, 5(2):135–151, 2013.
- [87] S. Graham and B. Weiner. *Theories and Principles of Motivation*. Prentice Hall, 1996.

- [88] J. Grant, S. Kraus, and D. Perlis. A logic-based model of intention formation and action for multi-agent subcontracting. *Artificial Intelligence*, 163(2):163–201, 2005.
- [89] J. Gratch. True emotion vs. social intentions in nonverbal communication: Towards a synthesis for embodied conversational agents. In *ZiF Workshop*, volume 4930, pages 181–197. Springer, January 2008.
- [90] J. Gratch and S. Marsella. Evaluating a computational model of emotion. *Journal of Autonomous Agents and Multi-Agent Systems. Special Issue on the Best of AAMAS 2004*, 11(1):23–43, July 2005.
- [91] J. Gratch, S. Marsella, N. Wang, and B. Stankovic. Assessing the validity of appraisal-based models of emotion. In *International Conference on Affective Computing and Intelligent Interaction*. IEEE, 2009.
- [92] J. Gratch and S. C. Marsella. A domain-independent framework for modeling emotion. *Cognitive Systems Research*, 5(4):269–306, 2004.
- [93] S. A. Green, M. Billinghurst, X. Chen, and J. G. Chase. Human-robot collaboration: A literature review and augmented reality approach in design. *International Journal of Advanced Robotic Systems*, 5(1):1–18, 2008.
- [94] J. J. Gross, G. Sheppes, and H. L. Urry. Emotion generation and emotion regulation: A distinction we should make (carefully). *Cognition and Emotion*, 25(5):765–781, 2011.
- [95] S. Grossberg and W. E. Gutowski. Neural dynamics of decision making under risk: Affective balance and cognitive-emotional interactions. *Psychological Review*, 94(3):300–318, 1987.
- [96] B. Grosz and S. Kraus. The evolution of shared plans. In *Foundations and Theories of Rational Agency*, pages 227–262, 1998.

- [97] B. J. Grosz. AAAI-94 presidential address: Collaborative systems. *AI Magazine*, 17(2):67–85, 1996.
- [98] B. J. Grosz. Beyond mice and menus. *Proceedings of the American Philosophical Society*, 149(4):523–543, 2005.
- [99] B. J. Grosz and L. Hunsberger. The dynamics of intention in collaborative activity. *Cognitive Systems Research*, 7(2-3):259–272, 2007.
- [100] B. J. Grosz, L. Hunsberger, and S. Kraus. Planning and acting together. *AI Magazine*, 20(4):23–34, 1999.
- [101] B. J. Grosz and S. Kraus. Collaborative plans for complex group action. *Artificial Intelligence*, 86(2):269–357, 1996.
- [102] B. J. Grosz and C. L. Sidner. Attention, intentions, and the structure of discourse. *Computational Linguistics*, 12(3):175–204, July 1986.
- [103] B. J. Grosz and C. L. Sidner. Plans for discourse. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 417–444. MIT Press, Cambridge, MA, 1990.
- [104] N. Guiraud, D. Longin, E. Lorini, S. Pesty, and J. Rivière. The face of emotions: A logical formalization of expressive speech acts. In *The 10th International Conference on Autonomous Agents and Multiagent Systems - Volume 3*, pages 1031–1038. International Foundation for Autonomous Agents and Multiagent Systems, 2011.
- [105] S. Hamann. Mapping discrete and dimensional emotions onto the brain: controversies and consensus. *Trends in Cognitive Sciences*, 16(9):458–466, 2012.
- [106] S. Hareli and U. Hess. What emotional reactions can tell us about the nature of others: An appraisal perspective on person perception. *Cognition & Emotion*, 24(1):128–140, 2009.

- [107] C. Hazan and P. Shaver. Romantic love conceptualized as an attachment process. *Personality & Social Psychology*, 52(3):511–524, 1987.
- [108] P. A. Heeman. *A Computational Model of Collaboration on Referring Expressions*. PhD thesis, University of Toronto, 1991.
- [109] F. Hegel, T. Spexard, B. Wrede, G. Horstmann, and T. Vogt. Playing a different imitation game: Interaction with an empathic android robot. In *Proceedings of 2006 IEEE-RAS International Conference on Humanoid Robots (Humanoids06)*, 2006.
- [110] U. Hess and P. Thibault. Darwin and emotion expression. *American Psychologist*, 64(2):120–128, 2009.
- [111] L. M. Hiatt and J. G. Trafton. A cognitive model of theory of mind. In *In Proceedings of the International Conference on Cognitive Modeling*, pages 91–96, 2010.
- [112] R. W. Hill, Jr., J. Chen, J. Gratch, P. Rosenbloom, and M. Tambe. Intelligent agents for the synthetic battlefield: A company of rotary wing aircraft. In *Innovative Applications of Artificial Intelligence (IAAI-97)*, pages 227–262, 1997.
- [113] J. Hirth, N. Schmitz, and K. Berns. Emotional architecture for the humanoid robot head roman. In *IEEE International Conference on Robotics and Automation*. IEEE, 2007.
- [114] J. R. Hobbs, A. Sagae, and S. Wertheim. Toward a commonsense theory of microsociology: Interpersonal relationships. In *Formal Ontology in Information Systems - Proceedings of the Seventh International Conference*, pages 249–262, 2012.

- [115] E. Hudlicka. To feel or not to feel: The role of affect in human-computer interaction. *International Journal of Human-Computer Studies*, 59(1-2):1–32, 2003.
- [116] E. Hudlicka. Reasons for emotinos: Modeling emotinos in integrated cognitive systems. In W. D. Gary, editor, *Integrated Models of Cognitive Systems*, volume 59, pages 1–37. New York: Oxford University Press, 2007.
- [117] E. Hudlicka. What are we modeling when we model emotion? In *AAAI Spring Symposium: Emotion, Personality, and Social Behavior*, pages 52–59. AAAI, 2008.
- [118] E. Hudlicka. Guidelines for designing computational models of emotions. *International Journal of Synthetic Emotions*, 2(1):85–102, 2011.
- [119] L. Hunsberger and B. J. Grosz. A combinatorial auction for collaborative planning. In *In Proceedings of ICMAS*, 2000.
- [120] C. E. Izard. *Human Emotions*. NY: Plenum, 1977.
- [121] B. Jarvis, D. Jarvis, and L. Jain. Teams in multi-agent systems. In *Intelligent Information Processing III*, volume 228. Springer US, 2007.
- [122] N. R. Jennings. *Joint Intentions as a Model of Multi-Agent Cooperation in Complex Dynamic Environments*. PhD thesis, Department of Electronic Engineering, University of London, 1992.
- [123] N. R. Jennings. On being responsible. In E. Werner and Y. Demazeau, editors, *Proceedings of the Third European Workshop on Modelling Autonomous Agents in a Multi-Agent World*, pages 93–102. North-Holland, 1992.
- [124] N. R. Jennings. Controlling cooperative problem solving in industrial multi-agent systems using joint intentions. *Artificial Intelligence*, 75(2):195–240, 1995.

- [125] N. R. Jennings and E. H. Mamdani. Using joint responsibility to coordinate collaborative problem solving in dynamic environments. In *10th National Conference on Artificial Intelligence (AAAI-92)*, pages 269–275, 1992.
- [126] C. L. John Robert Anderson. *The Atomic Components of Thought*. Lawrence Erlbaum Associates, 1998.
- [127] A. Kabil, C. D. Keukelaere, and P. Chavaillier. Coordination mechanisms in human-robot collaboration. In *Proceeding of the 7th International Conference on Advances in Computer-Human Interactions*, pages 389–394, 2014.
- [128] T. D. Kelley and D. N. Cassenti. Theoretical explorations of cognitive robotics using developmental psychology. *New Ideas in Psychology*, 29(3):228–234, 2011.
- [129] D. Keltner and J. Haidt. Social functions of emotions at four levels of analysis. *Cognition and Emotion*, 13(5):505–521, 1999.
- [130] H.-R. Kim and D.-S. Kwon. Computational model of emotion generation for human-robot interaction based on the cognitive appraisal theory. *Journal of Intelligent and Robotic Systems*, 60(2):263–283, 2010.
- [131] D. Kinny, M. Ljungberg, A. Rao, G. Tidhar, E. Werner, and E. Sonenberg. Planned team activity. In *Lecture notes in artificial intelligence*. Springer-Verlag, 1992.
- [132] R. Kirby, J. Forlizzi, and R. Simmons. Affective social robots. *Robotics and Autonomous Systems*, 58:322, 2010.
- [133] K. Kiryazov, R. Lowe, C. Becker-Asano, and T. Ziemke. Modelling embodied appraisal in humanoids : Grounding pad space for augmented autonomy. In *Proceedings of the Workshop on Standards in Emotion Modeling*, 2011.

- [134] H. Kitano, M. Asada, Y. Kuniyoshi, I. Noda, E. Osawai, and H. Matsubara. Robocup: A challenge problem for AI. *AI Magazine*, 18(1):73–85, 1997.
- [135] G. A. V. Kleef. How emotions regulate social life: The emotions as social information (EASI) model. *Current Directions in Psychological Science*, 18(3):184–188, 2009.
- [136] M. Klug and A. Zell. Emotion-based human-robot-interaction. In *IEEE 9th International Conference on Computational Cybernetics (ICCC)*, 2013.
- [137] J. Laird. *The Soar Cognitive Architecture*. MIT Press, 2012.
- [138] D. Laming. *Understanding Human Motivation: What Makes People Tick*. Wiley, 2003.
- [139] P. Langley, J. E. Laird, and S. Rogers. Cognitive architectures: Research issues and challenges. *Cognitive Systems Research*, 10(2):141–160, 2009.
- [140] R. S. Lazarus. *Emotion and Adaptation*. OXFORD University Press, 1991.
- [141] R. S. Lazarus, J. R. Averill, and E. M. Opton. Toward a cognitive theory of emotions. In *Feelings and Emotions*, pages 207–232. New York: Academic Press, New York, 1970.
- [142] I. Leite, A. Pereira, S. Mascarenhas, C. Martinho, R. Prada, and A. Paiva. The influence of empathy in human-robot relations. *International Journal of Human-Computer Studies*, 71(3):250–260, 2013.
- [143] N. Lesh, C. Rich, and C. L. Sidner. Collaborating with focused and unfocused users under imperfect communication. In M. Bauer, P. J. Gmytrasiewicz, and J. Vassileva, editors, *User Modeling 2001*, volume 2109, pages 64–73. Springer Berlin Heidelberg, 2001.

- [144] R. W. Levenson and J. M. Gottman. Martial interaction: Physiological linkage and affective exchange. *Personality and Social Psychology*, 45(3):587–597, 1983.
- [145] H. Leventhal and K. Scherer. The relationship of emotion to cognition. *Cognition and Emotion*, 1:3–28, 1987.
- [146] H. J. Levesque, P. R. Cohen, and J. H. T. Nunes. On acting together. In *AAAI*, pages 94–99. AAAI Press / The MIT Press, 1990.
- [147] A. Lim and H. G. Okuno. The mei robot: Towards using motherese to develop multimodal emotional intelligence. *IEEE Transactions Autonomous Mental Development*, 6(2):126–138, 2014.
- [148] J. Lin, M. Spraragen, and M. Zyda. Computational models of emotion and cognition. *Advances in Cognitive Systems*, 2:59–76, 2012.
- [149] C. L. Lisetti. Personality, affect and emotion taxonomy for socially intelligent agents. In *Proceedings of the Fifteenth International Florida Artificial Intelligence Research Society Conference*. AAAI Press, 2002.
- [150] D. J. Litman and J. F. Allen. Discourse processing and commonsense plans. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 365–388. MIT Press, Cambridge, MA, 1990.
- [151] C. Liu and N. Sarkar. Online affect detection and robot behavior adaptation for intervention of children with autism. *IEEE TRANSACTIONS ON ROBOTICS*, 24(4):883–896, 2008.
- [152] K. E. Lochbaum. A collaborative planning model of intentional structure. *Computational Linguistics*, 24(4):525–572, 1998.

- [153] K. E. Lochbaum, B. J. Grosz, and C. L. Sidner. Models of plans to support communication: An initial report. In *Proceedings of the Eighth National Conference on Artificial Intelligence*, pages 485–490. AAAI Press, 1990.
- [154] A. Luneski and R. K. Moore. Affective computing and collaborative networks: Towards emotion-aware interaction. In *Pervasive Collaborative Networks*, volume 283, pages 315–322. Springer US, 2008.
- [155] W. Mao and J. Gratch. Evaluating a computational model of social causality and responsibility. In *Fourth International Joint Conference on Autonomous Agents and Multiagent Systems*, 2006.
- [156] C. Marinetti, P. Moore, P. Lucas, and B. Parkinson. Emotions in social interactions: Unfolding emotional experience. In *Emotion-Oriented Systems, Cognitive Technologies*, pages 31–46. Springer Berlin Heidelberg, 2011.
- [157] R. P. Marinier III and J. E. Laird. Emotion-driven reinforcement learning. In *CogSci 2008*, 2008.
- [158] R. P. Marinier III, J. E. Laird, and R. L. Lewis. A computational unification of cognitive behavior and emotion. *Cognitive System Research*, 10(1):48–69, March 2009.
- [159] S. Marsella, J. Adibi, Y. Al-Onaizan, A. Erdem, R. Hill, G. A. Kaminka, Z. Qiu, and M. Tambe. Using an explicit teamwork model and learning in robocup: An extended abstract. In *RoboCup-98: Robot Soccer World Cup II*, volume 1604, pages 237–245. Springer Berlin Heidelberg, 1999.
- [160] S. Marsella, J. Gratch, N. Wang, and B. Stankovic. Assessing the validity of a computational model of emotional coping. In *International Conference on Affective Computing and Intelligent Interaction*. IEEE, 2009.
- [161] S. Marsella, J. Gratch, and P. Petta. Computational models of emotion. In E. R. Klaus R. Scherer, Tanja Banziger, editor, *A Blueprint for Affective*

Computing: A Sourcebook and Manual, pages 21–41. Oxford University Press, 2010.

- [162] S. C. Marsella and J. Gratch. EMA: A process model of appraisal dynamics. *Cognitive Systems Research*, 10(1):70–90, March 2009.
- [163] S. C. Marsella, J. Gratch, and P. Petta. Computational models of emotion. In *Blueprint for Affective Computing (Series in Affective Science)*. Oxford University Press, Oxford, 2010.
- [164] J. Martínez-Miranda and A. Aldea. Emotions in human and artificial intelligence. *Computers in Human Behavior*, 21(2):323–341, 2005.
- [165] J. Martínez-Miranda, A. Aldea, and R. Bañares-Alcántara. Simulation of work teams using a multi-agent system. In *The Second International Joint Conference on Autonomous Agents & Multiagent Systems AAMAS, July 14-18, Melbourne, Victoria, Australia*, pages 1064–1065, 2003.
- [166] L. Matignon, A. B. Karami, and A.-I. Mouaddib. A model for verbal and non-verbal human-robot collaboration. In *AAAI Fall Symposium Series*, pages 62–67, 2010.
- [167] S. W. McQuiggan and J. C. Lester. Modeling and evaluating empathy in embodied companion agents. *International Journal of Human-Computer Studies*, 65(4):348–360, 2007.
- [168] A. Mehrabian and J. A. Russell. *An approach to environmental psychology*. The MIT Press, 1974.
- [169] K. E. Merrick and K. Shafi. Achievement, affiliation, and power: Motive profiles for artificial agents. *Adaptive Behavior*, 19(1):40–62, 2011.
- [170] M. Mirolli and D. Parisi. Towards a vygotskyan cognitive robotics: The role of language as a cognitive tool. *New Ideas in Psychology*, 29(3):298–311, 2011.

- [171] V. Montreuil, A. Clodic, M. Ransan, and R. Alami. Planning human centered robot activities. In *Proceedings of the IEEE International Conference on Systems, Man and Cybernetics*, pages 2618–2623, 2007.
- [172] L. P. Morency. Computational study of human communication dynamic. In *Proceedings of the 2011 Joint ACM Workshop on Human Gesture and Behavior Understanding*, pages 13–18. ACM, 2011.
- [173] A. F. Morse, C. Herrera, R. Clowes, A. Montebelli, and T. Ziemke. The role of robotic modelling in cognitive science. *New ideas in psychology*, 29(3):312–324, 2011.
- [174] H. A. Murray. *Explorations of Personality*. Oxford University Press, New York, 1938.
- [175] B. Mutlu, A. Terrell, and C.-M. Huang. Coordination mechanisms in human-robot collaboration. In *Proceedings of the HRI 2013 Workshop on Collaborative Manipulation*, 2013.
- [176] L. C. namero. Emotion understanding from the perspective of autonomous robots research. *Neural Networks*, 18(4):445–455, 2005.
- [177] S. Nikolaidis, P. A. Lasota, G. F. Rossano, C. Martinez, T. A. Fuhlbrigge, and J. A. Shah. Human-robot collaboration in manufacturing: Quantitative evaluation of predictable, convergent joint action. In *ISR*, pages 1–6, 2013.
- [178] A. Ortony, G. L. Clore, and A. Collins. *The Cognitive Structure of Emotions*. Cambridge University Press, 1988.
- [179] A. Paiva, J. Dias, D. Sobral, R. Aylett, P. Sobrepeerez, S. Woods, C. Zoll, and L. Hall. Caring for agents and agents that care: Building empathic relations with synthetic agents. In *Proceedings of the Third International Joint Conference on Autonomous Agents and Multiagent Systems-Volume 1*, pages 194–201, 2004.

- [180] A. Paiva, J. Diasa, D. Sobrala, R. Aylettb, S. Woodsc, L. Halld, and C. Zoll. Learning by fealing: Evoking empathy with synthetic characters. *Applied Artificial Intelligence*, 19:235–266, 2005.
- [181] A. Paiva, I. Leite, and T. Ribeiro. Emotion modeling for sociable robots. In J. G. A. K. Rafael A. Calvo, Sidney D’Mello, editor, *Handbook of Affective Computing*, pages 296–308. Oxford University Press, 2014.
- [182] J. Panskepp. *Affective Neuroscience: The Foundations of Human and Animal Emotions*. NY:Oxford University Press, 1998.
- [183] B. Parkinson. Emotions are social. *British Journal of Psychology*, 87(4):663–683, 1996.
- [184] B. Parkinson. Do facial movements express emotions or communicate motives? *Personality and Social Psychology Review*, 9(4):278–311, 2005.
- [185] B. Parkinson. What holds emotions together? meaning and response coordination. *Cognitive System Research*, 10(1):31–47, 2009.
- [186] R. Pfeifer, H. Ishiguro, Y. Anzai, and N. Miyake. Robotics and emotion. In N. Miyake, D. Peebles, and R. P. Cooper, editors, *Building Bridges Across Cognitive Sciences Around the World: Proceedings of the 34th Annual Meeting of the Cognitive Science Society; Sapporo, Japan, August 1-4, 2012*, pages 30–31. Cognitive Science Society, 2012.
- [187] R. W. Picard. *Affective Computing*. The MIT Press, 2000.
- [188] S. Planalp. *Communicating Emotion: Social, Moral, and Cultural Processes*. Cambridge University Press, 1999.
- [189] M. E. Pollack. A model of plan inference that distinguishes between the beliefs of actors and observers. In *Proceedings of the 24th Annual Meeting*

- on Association for Computational Linguistics*, pages 207–214. Association for Computational Linguistics, 1986.
- [190] M. E. Pollack. Plans as complex mental attitudes. In *Intentions in Communication*, pages 77–103. MIT Press, 1990.
- [191] M. Pontier and J. F. Hoorn. How women think robots perceive them - as if robots were men. In *International Conference on Agents and Artificial Intelligence (ICAART-2)*, pages 496–504, 2013.
- [192] J. POSNER, J. A. RUSSELL, and B. S. PETERSON. The circumplex model of affect: An integrative approach to affective neuroscience, cognitive development, and psychopathology. *Development and Psychopathology*, null(3):715–734, 2005.
- [193] H. Prendinger and M. Ishizuka. The empathic companion: a character-based interface that addresses users' affective states. *Applied Artificial Intelligence*, 19(3-4):267–285, 2005.
- [194] D. V. Pynadath, M. Si, and S. C. Marsella. Modeling theory of mind and cognitive appraisal with decision-theoretic agents. In *Appraisal*, pages 1–30. 2011.
- [195] D. V. Pynadath and M. Tambe. An automated teamwork infrastructure for heterogeneous software agents and humans. *Journal of Autonomous Agents and Multi-Agent Systems, Special Issue on Infrastructure and Requirements for Building Research Grade Multi-Agent Systems*, 7(1-2):71–100, 2003.
- [196] S. Rank and P. Petta. Appraisal for a character-based story-world. In *5th International Working Conference on Intelligent Virtual Agents*, Kos, Greece, 2005. Springer.

- [197] T. W. Rauenbusch and B. J. Grossz. A decision making procedure for collaborative planning. In *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems*, 2003.
- [198] I. Ravenscroft. *Folk Psychology as a Theory*. Stanford Encyclopedia of Philosophy, 2004.
- [199] A. K. Ray, G. Leng, T. McGinnity, S. Coleman, and L. Maguire. Development of a self sustaining cognitive architecture. *Biologically Inspired Cognitive Architectures*, 6:96–108, 2013.
- [200] W. S. N. Reilly. *Believable Social and Emotional Agents*. PhD thesis, Carnegie Mellon University, 1996.
- [201] W. S. N. Reilly. Modeling what happens between emotional antecedents and emotional consequents. In *Eighteenth European Meeting on Cybernetics and Systems Research*, 2006.
- [202] C. Rich. Building task-based user interfaces with ANSI/CEA-2018. *IEEE Computer*, 42(8):20–27, July 2009.
- [203] C. Rich and C. L. Sidner. COLLAGEN: A collaboration manager for software interface agents. *User Modeling User-Adapted Interaction*, 8(3-4):315–350, 1998.
- [204] C. Rich, C. L. Sidner, and N. Lesh. COLLAGEN: Applying collaborative discourse theory to human-computer interaction. *AI Magazine*, 22(4):15–26, 2001.
- [205] J. Rickel, N. Lesh, C. Rich, C. L. Sidner, and A. Gertner. Collaborative discourse theory as a foundation for tutorial dialogue. In *Proceedings Sixth International Conference on Intelligent Tutoring Systems*, 2002.

- [206] I. J. Roseman and C. A. Smith. Appraisal theory:overview, assumptions, varieties, controversies. In K. R. Scherer, A. Schorr, and T. Johnstone, editors, *Appraisal process in emotion*, pages 3–34. NY: Oxford University Press, 2001.
- [207] T. H. Rumbell, J. A. Barnden, S. L. Denham, and T. Wennekers. Emotions in autonomous agents: comparative analysis of mechanisms and functions. *Autonomous Agents and Multi-Agent Systems*, 25(1):1–45, 2012.
- [208] J. A. Russell. A circumplex model of affect. *Journal of Personality and Social Psychology*, 39(6):1161–1178, 1980.
- [209] J. A. Russell. Core affect and the psychological construction of emotion. *Psychological Review*, 110(1):145–172, 2003.
- [210] R. M. Ryan and E. L. Deci. Intrinsic and extrinsic motivations: Classic definitions and new directions. *Contemporary Educational Psychology*, 25(1):54–67, 2000.
- [211] A. V. Samsonovich. Emotional biologically inspired cognitive architecture. *Biologically Inspired Cognitive Architectures*, 6:109–125, 2013.
- [212] D. Sander, D. Grandjean, and K. R. Scherer. A systems approach to appraisal mechanisms in emotion. *Neural Networks*, 18(4):317–352, 2005.
- [213] B. Scassellati. Theory of mind for a humanoid robot. *Autonomous Robots*, 12(1):13–24, 2002.
- [214] P. Scerri, D. Pynadath, L. Johnson, P. Rosenbloom, M. Si, N. Schurr, and M. Tambe. A prototype infrastructure for distributed robot-agent-person teams. In *Proceedings of the Second International Joint Conference on Autonomous Agents and Multiagent Systems*, AAMAS ’03, pages 433–440, New York, NY, USA, 2003. ACM.

- [215] P. Scerri, D. Pynadath, L. Johnson, P. Rosenbloom, M. Si, N. Schurr, and M. Tambe. A prototype infrastructure for distributed robot-agent-person teams. In *The Second International Joint Conference on Autonomous Agents and Multiagent Systems*, 2003.
- [216] K. R. Scherer. On the nature and function of emotion: A component process approach. In K. R. Scherer and P. Ekman, editors, *Approaches To Emotion*, pages 293–317. Lawrence Erlbaum, Hillsdale, NJ, 1984.
- [217] K. R. Scherer. Vocal affect expression: A review and a model for future research. *Psychological Bulletin*, 99(2):143–165, 1986.
- [218] K. R. Scherer. On the sequential nature of appraisal processes: Indirect evidence from a recognition task. *Cognition & Emotion*, 13(6):763–793, 1999.
- [219] K. R. Scherer. What are emotions? and how can they be measured? *Social Science Information*, 44:695–729, 2005.
- [220] K. R. Scherer. The dynamic architecture of emotion: Evidence for the component process model. *Cognition and Emotion*, 23(7):1307–1351, 2009.
- [221] K. R. Scherer. Emotions are emergent processes: they require a dynamic computational architecture. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 364(1535):3459–3474, 2009.
- [222] K. R. Scherer and H. Elgiring. Are facial expressions of emotion produced by categorical affect programs or dynamically driven by appraisal? *Emotion*, 7(1):113–130, 2007.
- [223] K. R. Scherer, A. Schorr, and T. Johnstone. *Appraisal Processes in Emotion: Theory, Methods, Research*. Oxford University Press, 2001.

- [224] M. Scheutz. Affective action selection and behavior arbitration for autonomous robots. In H. Arabnia, editor, *Proceedings of the 2002 International Conference on Artificial Intelligence*, pages 334–340. CSREA Press, 2002.
- [225] M. Scheutz. Useful roles of emotions in artificial agents: A case study from artificial life. In *Proceedings of AAAI*, pages 31–40. AAAI Press, 2004.
- [226] M. Scheutz and V. Andronache. Architectural mechanisms for dynamic changes of behavior selection strategies in behavior-based systems. *IEEE Transactions on Systems, Man, and Cybernetics, Part B*, 34(6):2377–2395, 2004.
- [227] M. Scheutz and A. Sloman. Affect and agent control: Experiments with simple affective states. In *Intelligent Agent Technology*, pages 200–209. World Scientific, 2001.
- [228] M. Scheutz, A. Sloman, and B. Logan. Emotional states and realistic agent behaviour. In *SOCIETY FOR COMPUTER SIMULATION*, pages 81–88, 2000.
- [229] O. Schultheiss and J. Brunstein. *Implicit Motives*. Oxford University Press, 2010.
- [230] J. R. Searle. Collective intentionality. In *Intentions in Communication*, pages 401–415. MIT Press, 1990.
- [231] M. Sellers. Toward a comprehensive theory of emotion for biological and artificial agents. *Biologically Inspired Cognitive Architectures*, 4:3–26, 2013.
- [232] M. Shayganfar, C. Rich, and C. L. Sidner. A design methodology for expressing emotion on robot faces. In *IROS*, pages 4577–4583. IEEE, 2012.
- [233] S. Shott. Emotion and social life: A symbolic interactionist analysis. *The American Journal of Sociology*, 84(6):1317–1334, 1979.

- [234] M. Si, S. C. Marsella, and D. V. Pynadath. Modeling appraisal in theory of mind reasoning. *Autonomous Agents and Multi-Agent Systems*, 20(1):14–31, 2010.
- [235] M. Si, S. C. Marsella, and D. V. Pynadath. Modeling appraisal in theory of mind reasoning. *Journal of Autonomous Agents and Multi-Agent Systems*, 20(1):14–31, 2010.
- [236] C. Sidner. An artificial discourse language for collaborative negotiation. In *Proceedings of the Twelfth National Conference on Artificial Intelligence*, pages 814–819. MIT Press, 1994.
- [237] C. L. Sidner and M. Dzikovska. A first experiment in engagement for human-robot interaction in hosting activities. In *Advances in Natural Multimodal Dialogue Systems*, volume 30 of *Cognitive Technologies*, pages 55–76. Springer Netherlands, 2005.
- [238] H. A. Simon. Motivational and emotional controls of cognition. *Psychology Review*, 74(1):29–39, 1967.
- [239] H. Skubch. *Modelling and Controlling of Behaviour for Autonomous Mobile Robots*. Springer Science Business Media, 2012.
- [240] A. Sloman. Motives, mechanisms, and emotions. *Cognition & Emotion*, 1(3):217–233, 1987.
- [241] A. Sloman. Beyond shallow models of emotion. *Cognitive Processing: International Quarterly of Cognitive Science*, 2(1):177–198, 2001.
- [242] A. Sloman and M. Croucher. Why robots will have emotions. In *In Proceeding 7th International Joint Conference on AI*, pages 197–202, 1981.
- [243] C. A. Smith and H. S. Scott. A componential approach to the meaning of facial expressions. In J. A. Russell and J. M. Fernández-Dols, editors, *The psychology*

- of facial expression. Studies in emotion and social interaction*, pages 229–254. NY: Cambridge University Press, 1997.
- [244] D. Sofge, M. D. Bugajska, J. G. Trafton, D. Perzanowski, S. Thomas, M. Skubic, S. Blisard, N. Cassimatis, D. P. Brock, W. Adams, and A. C. Schultz. Collaborating with humanoid robots in space. *International Journal of Humanoid Robotics*, 2(2):181–201, 2005.
- [245] A. Staller and P. Petta. Introducing emotions into the computational study of social norms: A first evaluation. *Journal of Artificial Societies and Social Simulation*, 4:615–625, 2001.
- [246] R. A. Subramanian, S. Kumar, and P. Cohen. Integrating joint intention theory, belief reasoning, and communicative action for generating team-oriented dialogue. In *AAAI*, pages 1501–1507. AAAI Press, 2006.
- [247] Suppressed for Anonymity. 2016.
- [248] M. Tambe. Agent architecture for flexible, practical teamwork. In *Proceedings of the National Conference on Artificial Intelligence*, pages 22–28, 1997.
- [249] M. Tambe. Towards flexible teamwork. *Journal of Artificial Intelligence Research*, 7:83–124, 1997.
- [250] M. Tambe. Establishing and maintaining long-term human-computer relationships. *ACM Transactions on Computer-Human Interaction (TOCHI)*, 12(2):293–327, 2005.
- [251] P. Thagard. Why wasn't o. j. convicted: emotional coherence in legal inference. *Cognition and Emotion*, 17(3):361–383, 2003.
- [252] P. A. Thoits. The sociology of emotion. *Annual Review of Sociology*, 15:317–342, 1989.

- [253] K. R. Thórisson and H. P. Helgason. Cognitive architectures and autonomy: A comparative review. *Journal of Artificial General Intelligence*, 3(2):1–30, 2012.
- [254] L. Z. Tiedens and C. W. Leach. *The Social Life of Emotions (Studies in Emotion and Social Interaction)*. Cambridge University Press, 2004.
- [255] S. S. Tomkins. *Affect, Imagery, Consciousness*. NY: Springer, 1962.
- [256] D. Traum, J. Rickel, onathan Gratch, and S. Marsella. Negotiation over tasks in hybrid human-agent teams for simulation-based training. In *International Conference on Autonomous Agents and Multiagent Systems*, 2003.
- [257] C. Urban. Pecs: A reference model for human-like agents. In *Deformable Avatars*. Netherlands: Kluwer Academic Publishers, 2001.
- [258] S. van Hooft. Scheler on sharing emotions. *Philosophy Today*, 38(1):18–28, 1994.
- [259] J. D. Velàsquez. Modeling emotions and other motivations in synthetic agents. In *Proceedings of the 14th Nnatiional Conference on Artificial Intelligence AAAI-97*, pages 10–15, 1997.
- [260] D. Vogiatzis, C. Spyropoulos, V. Karkaletsis, Z. Kasap, C. Matheson, and O. Deroo. An affective robot guide to museums. In *Proceedings of the 4th International Workshop on Human-Computer Conversation*, 2008.
- [261] D. Watson and A. Tellegen. Toward a consensual structure of mood. *Psychological Bulletin*, 98(2):219–235, 1985.
- [262] T. Wehrle. Motivations behind modeling emotional agents: Whose emotion does your robot have?, 1998.
- [263] B. Wilsker. A study of multi-agent collaboration theories. In *ISI Research Report*, pages 396–449, 1996.

- [264] A. K. Wisecup, D. T. Robinson, and L. Smith-Lovin. *The Sociology of Emotions*. SAGE Publications, Inc., 2007.
- [265] I. Wright and A. Sloman. *MINDER1: An Implementation of a Protoemotional Agent Architecture*. Cognitive science research papers. University of Birmingham, Cognitive Science Research Centre, 1997.
- [266] J. Yen, J. Yin, T. R. Ioerger, M. S. Miller, D. Xu, and R. A. Volz. Cast: Collaborative agents for simulating teamwork. In *Proceedings of IJCAI2001*, pages 1135–1142, 2001.
- [267] J. Yin, M. S. Miller, T. R. Ioerger, J. Yen, and R. A. Volz. A knowledge-based approach for designing intelligent team training systems. In *Proceedings of the Fourth International Conference on Autonomous Agents*, pages 427–434. ACM, 2000.
- [268] R. B. Zajonc. Social facilitation. *Science*, 149(3681):269–274, 1965.
- [269] Zamfirescu and Candea. On integrating agents into gdss. In *Preprints of the 9th IFAC / IFORS / IMACS / IFIP/ Symposium on Large Scale Systems: Theory and Applications*, 2001.
- [270] Z. Zeng, M. Pantic, G. I. Roisman, and T. S. Huang. A survey of affect recognition methods: Audio, visual and spontaneous expressions. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 31(1):39–58, 2009.
- [271] T. Zhang, D. B. Kaber, B. Zhu, M. Swangnetr, P. Mosaly, and L. Hodge. Service robot feature design effects on user perceptions and emotional responses. *Intelligent Service Robotics*, 3(2):73–88, 2010.
- [272] J. Zhu and P. Thagard. Emotion and action. *Journal of Philosophical Psychology*, 15(1):19–36, 2002.

- [273] E. L. Zurbriggen and T. S. Sturman. Linking motives and emotions: A test of mcclelland's hypotheses. *Personality and Social Psychology*, 28(4):521–535, 2002.

APPENDIX A

Instructions & Questionnaires for Crowd-Sourcing Study:

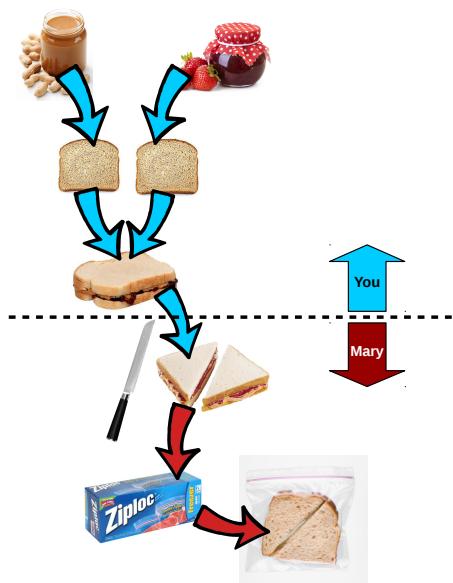
You are preparing for hiking. You are supposed to make two different snacks, a peanut butter and jelly sandwich and a simple boiled egg sandwich, together with Mary. The peanut butter and jelly sandwich MUST be prepared using peanut butter and strawberry jam. The egg sandwich MUST be prepared using two boiled eggs, salt and some pickles. All required ingredients are available to you and Mary, including:

- a) a jar of peanut butter,
- b) a jar of strawberry jam,
- c) whole wheat bread,
- d) two eggs,
- e) pickles, and
- f) salt.

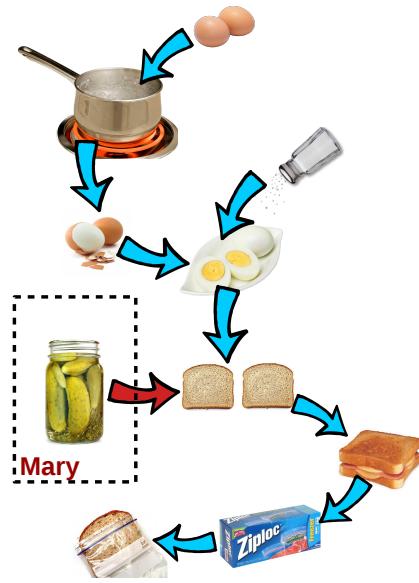
First Sandwich:

To prepare the peanut butter and jelly sandwich the following steps MUST be taken in the following order (as shown in the picture):

Spread some peanut butter on one slice of bread. Spread some strawberry jam on another slice of bread. Press the two slices together, and pass the sandwich to Mary. Mary uses a knife to cut the sandwich in half. Mary takes a zip lock bag, and puts the sandwich in the bag.



(a) Peanut butter sandwich instruction



(b) Boiled egg sandwich instruction

Second Sandwich:

To prepare the boiled egg sandwich the following steps MUST be taken in the following order (as shown in the picture):

Boil the water in a pot, and put two eggs in the pot when the water is boiling. Remove the eggs from the pot after 5 minutes. Peel the boiled eggs and slice them into a few pieces. Shake some salt on the eggs. Put the sliced eggs on bread. Wait for Mary to put some pickles on top of the eggs. Press another slice of bread on top. Grab a zip lock bag and put your sandwich inside of the zip lock bag.

NOTE: The instructions provided above were the same for all four questionnaires, and are only presented once. The four separate questionnaires are provided in the following pages.

Relevance Questionnaire:

In this questionnaire, there are 19 questions about different situations while you are preparing the two sandwiches together with Mary (as mentioned above). Each question provides three possible answers regarding the relevance of various situations to the goal of preparing the sandwiches.

A situation or an event is **relevant** if it requires your attention because of its positive or negative effect, since it can impact your state or what you want. For example, if you need Mary to take the peanut butter sandwich you made, cut it in half and put it in a zip lock for you (see the picture on the previous page), You will still find it relevant if she cuts your sandwich in half and puts it in a small paper bag for you.

Now, please answer the following questions:

1. Imagine you have just pressed two slices of bread together after spreading peanut butter and jelly on them. Which of the following two actions is more relevant?
 - A. Mary takes your sandwich and cuts it in half.
 - B. Mary looks for a jar of pickles for your second sandwich.
 - C. Equally relevant.
2. Imagine you want to make the peanut butter sandwich. Which of the following two actions is more relevant?
 - A. You trying to find the jar of peanut butter.
 - B. Mary wants to figure out which knife to use to cut the sandwich you make.
 - C. Equally relevant.
3. Imagine you have made the peanut butter sandwich and passed it to Mary to

cut it in half. You begin to boil some water in the pot for the eggs. Which of the following two actions is more relevant?

- A. Mary has cut the peanut butter sandwich and put it in a zip lock bag.
 - B. Mary asks you whether you have found the jar of peanut butter.
 - C. Equally relevant.
4. Imagine you have spread peanut butter on one slice of bread and you have just spread the strawberry jam on the second slice of bread. Now you want to press them together and pass the sandwich to Mary to cut it in half. Which of the following two actions is more relevant?
- A. Mary asks you to pass the sandwich to her.
 - B. Mary asks you whether the eggs are done.
 - C. Equally relevant.
5. Imagine you are about to start making a hard boiled egg sandwich. Which of the following two actions is more relevant?
- A. At the beginning, you discover that you have run out of eggs.
 - B. You discover that Mary will not be able to find the zip lock bags when the sandwich is ready at the end.
 - C. Equally relevant.
6. Imagine you and Mary want to make the hard boiled egg sandwich. Which of the following two actions is more relevant?
- A. Mary tells you that you can not boil the eggs since the stove is broken.
 - B. Mary tells you that she is going to grab a knife to cut the sandwich in half.
 - C. Equally relevant.
7. Imagine you want to make the peanut butter sandwich. You open the lid on both jars of peanut butter and strawberry jam. Which of the following two actions

is more relevant?

- A. You pick a slice of bread to spread peanut butter onto it.
- B. You ask Mary whether she has a zip lock bag to put the sandwich inside.
- C. Equally relevant.

8. Imagine you and Mary are going to make a peanut butter and a hard boiled egg sandwich for your hiking trip. Which of the following two actions is more relevant?

- A. You take two eggs and put them in a pot with boiling water.
- B. You think of making a snack for your hiking trip.
- C. Equally relevant.

9. Which of the following two actions is more relevant?

- A. Pressing two slices of bread after spreading peanut butter and jelly on them.
- B. Pressing two slices of bread after putting slices of boiled egg, and adding some salt and pickles.
- C. Equally relevant.

10. Imagine you want to make the peanut butter sandwich. Which of the following two actions is more relevant?

- A. Spreading peanut butter on one slice of bread.
- B. Spreading strawberry jam on another slice of bread.
- C. Equally relevant.

11. Imagine you have made the peanut butter sandwich and passed it to Mary to cut it in half. You begin to boil some water in the pot for the eggs. Which of the following two actions is more relevant?

- A. Mary has cut the peanut butter sandwich and put it in a zip lock bag.
- B. You put the eggs in the boiling water.

C. Equally relevant.

12. Which of the following two actions is more relevant?

- A. You want to press two slices of bread together after spreading peanut butter and strawberry jam on them.
- B. You want to spread strawberry jam on another slice of bread after spreading peanut butter on one slice of the bread.
- C. Equally relevant.

13. Which of the following two actions is more relevant?

- A. Mary tells you that you have run out of eggs when you wanted to make a hard boiled egg sandwich.
- B. Mary tells you that you have run out of peanut butter when you wanted to make a peanut butter sandwich.
- C. Equally relevant.

14. Which of the following two actions is more relevant?

- A. You try to find a pot to boil the eggs since you want to make the hard boiled egg sandwich.
- B. You look for the jar of peanut butter since you want to make peanut butter sandwich .
- C. Equally relevant.

15. Which of the following two actions is more relevant?

- A. You open the lid on both jars of peanut butter and strawberry jam since you want to make the peanut butter sandwich.
- B. Mary opens the lid on jar of pickles since she wants to add some pickles to the top of your sliced boiled eggs.
- C. Equally relevant.

16. Which of the following two actions is more relevant?

- A. You take two eggs and put them in a pot with boiling water, since you want to make a hard boiled egg sandwich.
- B. You open the lids on the jars of both the peanut butter and strawberry jam since you want to make the peanut butter sandwich.
- C. Equally relevant.

17. Imagine you have made the peanut butter sandwich and passed it to Mary to cut it in half. Which of the following two actions is more relevant?

- A. Mary starts crying since she cut her finger with a knife.
- B. You begin to boil the water to boil the eggs for your second sandwich.
- C. Equally relevant.

18. Imagine you have spread peanut butter on one slice of bread and you just spread the strawberry jam on the second slice of bread. Now you want to press them together and pass the sandwich to Mary to cut it in half. Which of the following two actions is more relevant?

- A. Mary calmly looks at you.
- B. Mary begins waving her hands and yelling at you!
- C. Equally relevant.

19. Imagine you want to make the peanut butter sandwich. You open the lids on the jars of both the peanut butter and strawberry jam. Which of the following two actions is more relevant?

- A. You know what to do, but Mary begins to laugh so hard. She can't say anything, but is pointing at the jar of peanut butter. You want to see what happened!
- B. You know what to do, and Mary is not in the kitchen at the moment.
- C. Equally relevant.

Desirability Questionnaire:

In this questionnaire, there are 12 questions about different situations while you are preparing two sandwiches together with Mary (as mentioned above). Each question provides three possible answers regarding the desirability of certain circumstances.

Desirability is a measure of the value of an event or situation. Desirability can be positive if an event facilitates a positive value for you, and it can be negative if it inhibits a positive value for you. For example, you know Mary is supposed to add some pickles to the top of your sliced boiled eggs (see the picture above). It will not be desirable if she adds another garnish to your sandwich that you do not like.

Now, please answer the following questions:

1. Imagine you want to make one peanut butter and one hard boiled egg sandwich for your hike. Which of the following two actions is more desirable?
 - A. Mary tells you the peanut butter and eggs you used in your sandwiches are spoiled and out dated.
 - B. Your peanut butter sandwich is done and ready. Mary tells you to wait for her to cut the egg sandwich in half and put it in a zip lock bag for you.
 - C. Equally desirable.

2. Imagine you want to make one peanut butter and one hard boiled egg sandwich for your hike. Which of the following two actions is more desirable?
 - A. You have both sandwiches ready to go.
 - B. Your peanut butter sandwich is done and ready. Mary tells you to wait for her to cut the egg sandwich in half and put it in a zip lock bag for you.
 - C. Equally desirable.

3. Imagine you want to make the egg sandwich. You have sliced the eggs, put them on one slice of bread, salted them, and are waiting for Mary to put some pickles on your eggs. Which of the following two actions is more desirable?

- A. Mary tells you there are no more pickles left in the jar to put on your sandwich.
- B. Mary finds the pickle jar to put some pickles on your eggs.
- C. Equally desirable.

4. Imagine you pressed two slices of bread together with peanut butter and strawberry jam on them, and passed them to Mary. Which of the following two actions is more desirable?

- A. Mary cuts the peanut butter sandwich in half and put it in a zip lock bag.
- B. Mary can not find a knife to cut the sandwich in half.
- C. Equally desirable.

5. Imagine you have put two eggs in a boiling pot. Which of the following two actions is more desirable?

- A. You go back to the pot to remove the eggs from the pot, but the egg shells are broken and the eggs are mixed with water.
- B. You decide to wait for two more minutes for the eggs to boil.
- C. Equally desirable.

6. Imagine you want to make the egg sandwich. You have put eggs in the pot to boil. Which of the following two actions is more desirable?

- A. You come back to the pot and you see that you have forgotten to turn on the stove.
- B. You come back to the pot and eggs are boiling, and ready to be removed from the pot.
- C. Equally desirable.

7. Imagine you have pressed the two slices of bread together (one covered with strawberry jam and one covered with peanut butter) and passed it to Mary. Which of the following two actions is more desirable?

- A. Mary puts two pieces of the sandwich into a zip lock bag after cutting the sandwich in half.
- B. You go ahead and begin to boil the water to make the boiled egg sandwich.
- C. Equally desirable.

8. Imagine you want to make one peanut butter and one hard boiled egg sandwich for your hike. Which of the following two actions is more desirable?

- A. Mary tells you the peanut butter and eggs you used in your sandwiches are spoiled and out dated.
- B. Mary finds some mold on the bread in both sandwiches.
- C. Equally desirable.

9. Imagine you want to make the egg sandwich. You have sliced the eggs, put them on one slice of bread, salted them, and are waiting for Mary to put some pickles on your eggs. Which of the following two actions is more desirable?

- A. Mary tells you there are no more pickles left in the jar to put on your sandwich.
- B. Mary finds the pickles are out dated and not edible.
- C. Equally desirable.

10. Which of the following two actions is more desirable?

- A. Imagine you pressed two slices of bread together with peanut butter and strawberry jam on them, and passed them to Mary. Mary cuts the peanut butter sandwich in half and puts them in the zip lock bag.
- B. Imagine you want to make the egg sandwich. You have sliced the eggs, put them on one slice of bread, salted them, and waiting for Mary to put some pickles on your eggs. Mary puts some pickles on your eggs.

C. Equally desirable.

11. Imagine you have put two eggs in a pot to boil. Which of the following two actions is more desirable?

A. You go back to the pot to remove the eggs from the pot, but the egg shells are broken and eggs are mixed with water and are not edible any more.

B. You go back to the pot to remove the eggs from the pot, but you find out somebody has removed the eggs before you and eaten them.

C. Equally desirable.

12. Which of the following two actions is more desirable?

A. Imagine you want to make the egg sandwich, but you cannot find the eggs.

B. Imagine you want to make the peanut butter sandwich, but you find the peanut butter jar empty.

C. Equally desirable.

Expectedness Questionnaire:

In this questionnaire, there are 10 questions about different situations while you are preparing two sandwiches together with Mary (as mentioned above). Each question provides three possible answers regarding how expected the possible situations are.

Expectedness is the extent to which a situation or an event could be predicted from your past knowledge. For example, if you know Mary is supposed to add some pickles to the top of your sliced boiled eggs (see the picture above), you will not find it expected if she does something completely different in comparison with what you and she planned for.

Now, please answer the following questions:

1. Imagine you have just pressed two slices of bread together after spreading peanut butter and strawberry jam on them. Which of the following two actions is more expected?
 - A. Mary takes your sandwich and cuts it in half.
 - B. You begin to put the water in the pot without passing the peanut butter sandwich to Mary.
 - C. Equally expected.

2. Imagine you have just pressed the two slices of bread together after spreading peanut butter and strawberry jam on them, and passed it to Mary. Which of the following two actions is more expected?
 - A. Mary takes your sandwich and cuts it in half.
 - B. Mary takes a bite of the given sandwich.
 - C. Equally expected.

3. Imagine you are spreading peanut butter onto one slice of bread. Which of

the following two actions is more expected?

- A. Mary waits for you to prepare the sandwich and pass it to her.
- B. Mary puts some pickles on another slice of bread.
- C. Equally expected.

4. Imagine Mary cuts your given peanut butter sandwich in half. Which of the following two actions is more expected?

- A. Mary takes a zip lock bag and put the sandwich inside of the bag.
- B. Mary leaves the apartment to buy an energy drink.
- C. Equally expected.

5. Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Which of the following two actions is more expected?

- A. Mary puts the given sandwich into a zip lock bag after cutting it in half.
- B. Mary puts some pickles on another slice of bread.
- C. Equally expected.

6. Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Which of the following two actions is more expected?

- A. Mary puts some pickles on another slice of bread.
- B. Mary begins to cook some chicken for you.
- C. Equally expected.

7. Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Which of the following two actions is more expected?

- A. Mary puts the given sandwich into a zip lock bag after cutting it in half.

B. You go ahead and begin to boil the water in a pot, since Mary doesn't need your help anymore.

C. Equally expected.

8. Imagine you are spreading strawberry jam onto the second slice of bread. Which of the following two actions is more expected?

- A. Mary wants to cut the one slice of bread which is already covered by peanut butter.
- B. Mary puts some pickles on another slice of bread.
- C. Equally expected.

9. Which of the following two actions is more expected?

- A. Imagine you have pressed the two slices of bread (one covered with strawberry jam and one covered with peanut butter) together and passed it to Mary. Mary puts the given sandwich into a zip lock bag after cutting it in half.
- B. Imagine you have shaken salt on the slices of the boiled eggs on the bread. Mary puts some pickles on another slice of bread.
- C. Equally expected.

10. Imagine you have peeled the boiled eggs and sliced them into some pieces. Which of the following two actions is more expected?

- A. Mary leaves the apartment to buy an energy drink.
- B. Mary begins to cook some chicken for you.
- C. Equally expected.

Controllability Questionnaire:

In this questionnaire, there are 10 questions about different situations while you are preparing two sandwiches together with Mary (as mentioned above). Each question provides three possible answers regarding how controllable the situations are.

Controllability is about whether the outcome of an event can be changed under your or your collaborator's control. For example, boiling the eggs is more controllable for you since you have the eggs and the pot and you know how to turn on the stove, but putting pickles on the sliced eggs is not controllable for you since your hands are dirty and you should not open the lid on the pickle jar with your dirty hands.

Now, please answer the following questions:

1. Imagine you want to quickly make a peanut butter sandwich. You know that you are supposed to press two slices of bread together that are covered by peanut butter and strawberry jam, and pass it to Mary to cut it in half. Which of the following two actions is more controllable?
 - A. You spread the peanut butter on one slice of bread and strawberry jam on the other slice of bread to prepare the sandwich.
 - B. Mary asks you whether you know how to make your own homemade bread for your sandwich.
 - C. Equally controllable.

2. Imagine you want to quickly make a peanut butter sandwich. Which of the following two actions is more controllable?
 - A. You spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread. You do not need Mary to cut the sandwich in half. You press the two slices together and put it in a zip lock bag.

B. You spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread. You need to pass the sandwich to Mary to cut it in half and put it in a zip lock bag.

C. Equally controllable.

3. Imagine you want to quickly make a peanut butter sandwich. Which of the following two actions is more controllable?

A. You can spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread.

B. You can spread the peanut butter on one slice of bread and you need Mary to spread strawberry jam on the second slice of bread.

C. Equally controllable.

4. Imagine you want to quickly make a hard boiled egg sandwich. Which of the following two actions is more controllable?

A. You want to put the sandwich into a zip lock bag. You check the sandwich. It is made the way you wanted.

B. You put the sandwich into a zip lock bag. You check the sandwich. You find out the eggs are not boiled enough in a way you wanted.

C. Equally controllable.

5. Imagine you want to quickly make a peanut butter sandwich. Which of the following two actions is more controllable?

A. You have a jar of peanut butter and a jar of strawberry jam opened and ready to use.

B. You cannot find the jar of peanut butter, and Mary tells you that you ran out of bread.

C. Equally controllable.

6. Which of the following two actions is more controllable?

- A. You spread the peanut butter on one slice of bread and strawberry jam on the other slice of bread to prepare the peanut butter sandwich.
- B. You have boiled, peeled, and sliced the eggs. You put the sliced eggs on a bread and add some salt to prepare the hard boiled egg sandwich.
- C. Equally controllable.

7. Which of the following two actions is more controllable?

- A. Imagine you want to quickly make a peanut butter sandwich. You spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread. You do not need Mary to cut the sandwich in half. You press two slices together and put it in a zip lock bag.
- B. Imagine you want to quickly make a hard boiled egg sandwich. You have boiled and peeled the eggs. You do not need Mary to put some pickles on your eggs.
- C. Equally controllable.

8. Which of the following two actions is more controllable?

- A. Imagine you want to quickly make a peanut butter sandwich. You can spread the peanut butter on one slice of bread and strawberry jam on the second slice of bread.
- B. Imagine you want to quickly make a hard boiled egg sandwich. You can boil the eggs, peel them, put them on bread and add some pickles on top without getting any help from Mary.
- C. Equally controllable.

9. Which of the following two actions is more controllable?

- A. Imagine you want to quickly make a hard boiled egg sandwich. You want to put the sandwich into a zip lock bag. You check the sandwich. It is made the way you wanted.

B. Imagine you want to quickly make a peanut butter sandwich. You want to put the sandwich into a zip lock bag. You check the sandwich. It has extra amount of strawberry jam as the way you wanted.

C. Equally controllable.

10. Which of the following two actions is more controllable?

A. Imagine you want to quickly make a peanut butter sandwich. You have a jar of peanut butter and a jar of strawberry jam opened ready to use.

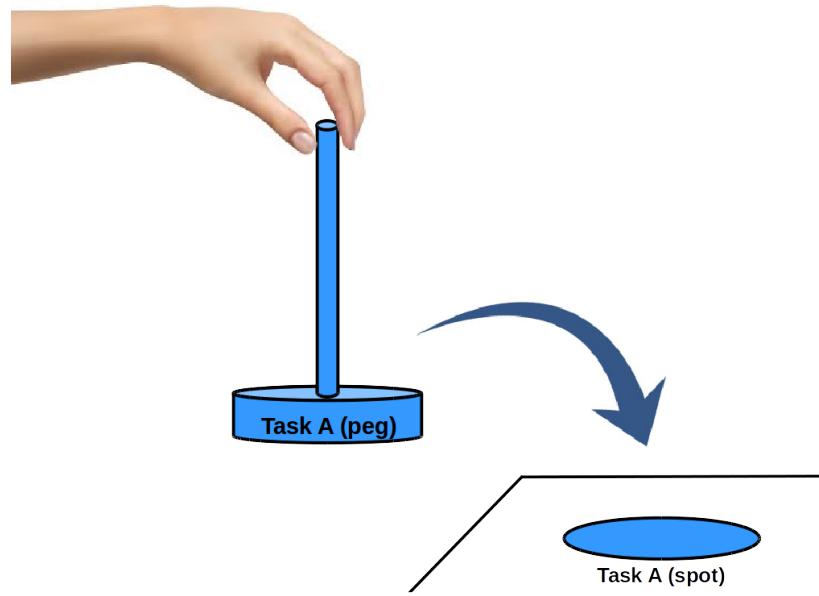
B. Imagine you want to quickly make a hard boiled egg sandwich. You have two eggs, and a boiling pot on the stove.

C. Equally controllable.

APPENDIX B


Instructions:

- You are going to collaborate with a robot to install a solar panel in a simulated environment. The robot is going to operate as your collaborative partner. Therefore, some of the tasks are going to be done by the robot, and some other tasks will be done by you.
- All the tasks are implemented with labeled pegs. You and the robot are going to pick up and place these pegs on the predefined spots on the shared table-top in front of you.
- Pegs that are properly placed in the predefined spots are considered to be successfully completed tasks.
- Pegs are provided in two colors. Red pegs represent the robot's tasks and blue pegs represent your tasks.
- The predefined spots on the board are labeled with tasks and are colored to match the corresponding pegs.
- The robot is supposed to place its own pegs on the red spots.
- You should place your own blue pegs only in the corresponding blue spots, i.e., if you have a blue peg labeled A (representing a task called A that you are responsible for), there will be a blue spot on the shared board labeled A (see the figure below).



Success and Failure of Achieving a Task:

- Picking up and placing a peg on the right spot of the shared board means that particular task has been successfully achieved.
- Task failures are simulated using magnets: a task fails when the magnet in the board does not let the peg's settle fully into its spot (see the following figure).



Order of Doing Tasks:

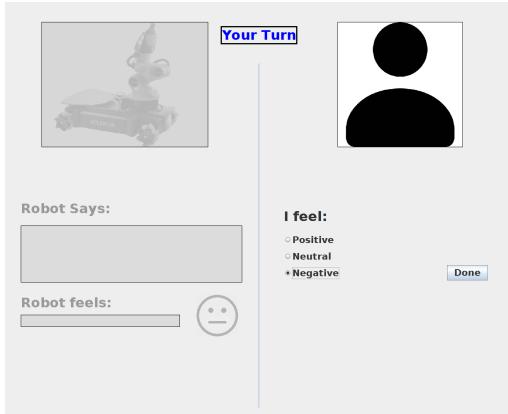
The numerical labels on the shared board under each spot provide the correct order of the task completion.

Your Objectives:


- a) Successful completion of installing the solar panel,
- b) Expressing yourself to the robot through the interface provided on the screen.
- c) Timely completion of the overall task (e.g., if there is a failure that requires the robot's supervisor, the extra time will be counted against you).
- d) Ensuring overall satisfaction of the robot during collaboration.
- e) Avoiding impasse; where there is a failure and the robot supervisor has to come in, you will lose points unless you are already working on a different task suggested by the robot.

Note 1: The winner's prize will be sent by e-mail the week after the study finishes.

Note 2: Your overall score will be calculated based on these objectives.




(f) Human turn



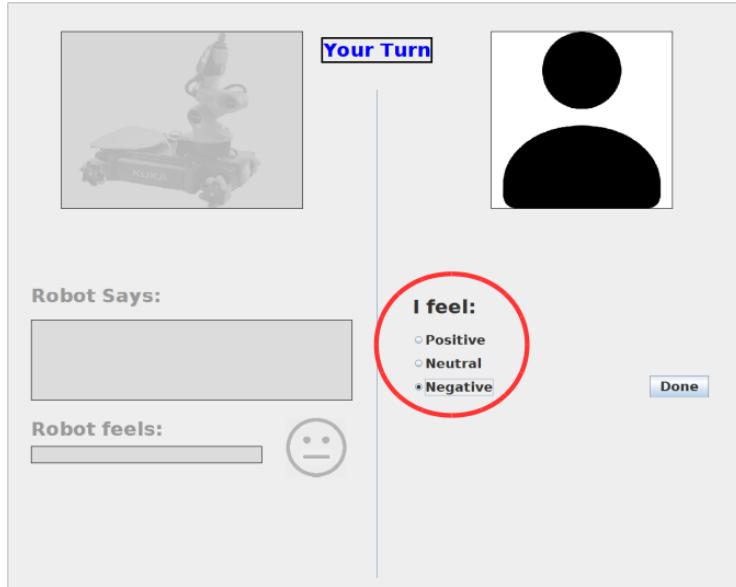
(g) Robot turn

Playing Rules:

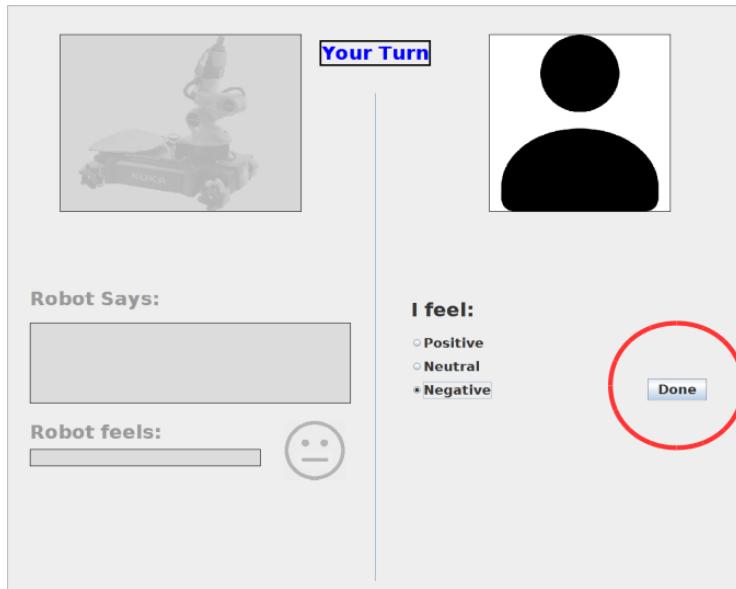
1. You should not touch the robot under any circumstance.
2. You are responsible for removing a failed task peg (the robot's or your own) from the shared board.
3. You should wait for the robot to inform you that it is your turn through voice and/or the visual interface (see the figure below the visual interface will change from what you see in the left to the right condition).
4. When it is your turn to do your own task, you should strictly obey the following:
 - 4.1) Move ONLY one of your pegs according to the current numerical label of your task, e.g., if the robot's last move was task 8, you should do task 9 if you are responsible for it (see the following figure).



4.2) Based on the failure or success of your task, choose how you feel (i.e., Positive, Neutral, Negative) about the outcome of your task on the interface provided on the screen (see the following figure).



4.3) Press the Done button to give the turn back to the robot (see the following figure).



5. You must not remove any peg from the board unless it is a clear case of a failure caused by the magnetic field.

6. It is okay to change the order of the tasks ONLY if the robot makes such a decision and announces it to you.
7. Do not touch your pegs until it is your turn again.
8. The robot might ask its supervisor to come and help if there is a task failure that causes an impasse during your collaboration.
9. The robot might want to help you with its own pegs; if so, the robot will place a peg in the Handoff Area for use during your next turn.